

Lecture 7: Deadlock


Yinqian Zhang @ 2021, Spring

copyright@Bo Tang

Deadlock

I won't give you the paper unless you come out

I won't come out unless you give me the paper




waiting

holding

waiting




Starvation vs Deadlock



◆ Starvation vs. Deadlock

- ◆ Starvation: thread waits indefinitely
 - ◆ Low-priority thread waiting for resources constantly in use by high-priority threads
- ◆ Deadlock: circular waiting for resources
 - ◆ Thread A owns Res 1 and is waiting for Res 2
Thread B owns Res 2 and is waiting for Res 1



- ◆ Deadlock \Rightarrow Starvation but not vice versa
 - ◆ Starvation can end (but does not have to)
 - ◆ Deadlock cannot end without external intervention

Conditions for Deadlock

◆ Deadlock not always deterministic

<u>Thread A</u>	<u>Thread B</u>
<code>sem_wait(x);</code>	<code>sem_wait(y);</code>
<code>sem_wait(y);</code>	<code>sem_wait(x);</code>
<code>sem_post(y);</code>	<code>sem_wait(x);</code>
<code>sem_post(x);</code>	<code>sem_wait(y);</code>

- ◆ Deadlock will not always happen with this code
 - ◆ Have to have exactly the right timing
 - ◆ So you release a piece of software, and you tested it, and there it is, controlling a nuclear power plant...

◆ Deadlocks occur with multiple resources

- ◆ Means you cannot decompose the problem
- ◆ Cannot solve deadlock for each resource independently
- ◆ System with 2 disk drives and two threads
 - ◆ Each thread needs 2 disk drives to function
 - ◆ Each thread gets one disk and waits for another one

Four requirements for Deadlock

◆ Mutual exclusion

- ◆ Only one thread at a time can use a resource.

◆ Hold and wait

- ◆ Thread holding at least one resource is waiting to acquire additional resources held by other threads

◆ No preemption

- ◆ Resources are released only voluntarily by the thread holding the resource, after thread is finished with it


◆ Circular wait

- ◆ There exists a set $\{T_1, \dots, T_n\}$ of waiting threads
 - ◆ T_1 is waiting for a resource that is held by T_2
 - ◆ T_2 is waiting for a resource that is held by T_3
 - ◆ ...
 - ◆ T_n is waiting for a resource that is held by T_1

Resource-Allocation Graph

◆ System Model

- ◆ A set of Threads T_1, T_2, \dots, T_n
- ◆ Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- ◆ Each resource type R_i has W_i instances
- ◆ Each thread utilizes a resource as follows:
 - ◆ Request () / Use () / Release ()




◆ Resource-Allocation Graph:

- ◆ V is partitioned into two types:
 - ◆ $T = \{T_1, T_2, \dots, T_n\}$, the set threads in the system.
 - ◆ $R = \{R_1, R_2, \dots, R_m\}$, the set of resource types in system
- ◆ request edge – directed edge $T_1 \rightarrow R_j$
- ◆ assignment edge – directed edge $R_j \rightarrow T_i$


Resource Allocation Graph Examples

◆ Recall:


- ◆ request edge – directed edge $T_1 \rightarrow R_j$
- ◆ assignment edge – directed edge $R_j \rightarrow T_i$



Simple Resource
Allocation Graph




Allocation Graph
With Deadlock



Allocation Graph with
Cycle, but No Deadlock


Methods for Handling Deadlocks



- ◆ Allow system to enter deadlock and then recover
 - ◆ Requires deadlock detection algorithm
 - ◆ Some technique for forcibly preempting resources and/or terminating tasks
- ◆ Ensure that system will *never* enter a deadlock
 - ◆ Need to monitor all resource acquisitions
 - ◆ Selectively deny those that *might* lead to deadlock
- ◆ Ignore the problem and pretend that deadlocks never occur in the system
 - ◆ Used by most operating systems, including UNIX

Deadlock Detection Algorithm

- ◆ Only one of each type of resource \Rightarrow look for cycles



- ◆ More than one resource of each type
 - ◆ More deadlock detection algorithm
 - ◆ Next page

Several Instances of a Resource Type

- ◆ **Available:** A vector of length m indicates the number of available resources of each type.
- ◆ **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process.
- ◆ **Request:** An $n \times m$ matrix indicates the current request of each process. If $Request[i_j] = k$, then process P_i is requesting k more instances of resource type. R_j .

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively Initialize:
 - (a) $Work = Available$
 - (b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then
 $Finish[i] = false$; otherwise, $Finish[i] = true$
2. Find an index i such that both:
 - (a) $Finish[i] == false$
 - (b) $Request_i \leq Work$

If no such i exists, go to step 4

Detection Algorithm (Cont.)

3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2
4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $Finish[i] == false$, then P_i is deadlocked

Example of Detection Algorithm

- ◆ Five processes P_0 through P_4 ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- ◆ Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- ◆ Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i

Example (Cont.)

- ◆ P_2 requests an additional instance of type C

	<u>Request</u>		
	A	B	C
P_0	0	0	0
P_1	2	0	2
P_2	0	0	1
P_3	1	0	0
P_4	0	0	2

- ◆ State of system?
 - ◆ Can reclaim resources held by process P_0 (not deadlocked), but insufficient resources to fulfill other processes; requests
 - ◆ Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4

What to do when detect deadlock?

- ◆ Terminate thread, force it to give up resources
 - ◆ In Bridge example, Godzilla picks up a car, hurls it into the river. Deadlock solved!
 - ◆ Shoot a dining philosopher
 - ◆ But, not always possible
- ◆ Preempt resources without killing off thread
 - ◆ Take away resources from thread temporarily
 - ◆ Does not always fit with semantics of computation
- ◆ Roll back actions of deadlocked threads
 - ◆ For bridge example, make one car roll backwards (may require others behind him)
 - ◆ Common technique in databases (transactions)
 - ◆ Of course, if you restart in exactly the same way, may reenter deadlock once again
- ◆ Many operating systems use other options

Techniques for Preventing Deadlock

◆ Infinite resources

- ◆ Include enough resources so that no one ever runs out of resources. Examples:
 - ◆ Bay bridge with 12,000 lanes. Never wait!
 - ◆ Infinite disk space (not realistic yet?)

◆ No sharing of resources (totally independent threads)

- ◆ Not very realistic

◆ Do not allow waiting

- ◆ Technique used in Ethernet/some multiprocessor nets
 - ◆ Everyone speaks at once. On collision, back off and retry
- ◆ Inefficient, since have to keep retrying
 - ◆ Consider: driving to SUSTech; when hit traffic jam, suddenly you are transported back home and told to retry!

Techniques for Preventing Deadlock

- ◆ Make all threads request everything they will need at the beginning.
 - ◆ Problem: Predicting future is hard, tend to over-estimate resources. Example:
 - ◆ If need 2 chopsticks, request both at same time
 - ◆ Don not leave home until we know no one is using any intersection between home and SUSTech; only one car on the Bay Bridge at a time
- ◆ Force all threads to request resources in a particular order preventing any cyclic use of resources
 - ◆ Thus, preventing deadlock
 - ◆ Example (x.P, y.P, z.P,...)
 - ◆ Make tasks request disk, then memory, then...
 - ◆ Keep from deadlock on freeways around SF by requiring everyone to go clockwise

Banker's Algorithm

- ◆ Multiple instances of each resource type
- ◆ Each process must a priori claim maximum use
- ◆ When a process requests a resource it may have to wait
- ◆ When a process gets all its resources it must return them in a finite amount of time

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- ◆ **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available
- ◆ **Max:** $n \times m$ matrix. If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- ◆ **Allocation:** $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j
- ◆ **Need:** $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:

$Work = Available$

$Finish[i] = false$ for $i = 0, 1, \dots, n-1$

2. Find an index i such that both:

(a) $Finish[i] = false$

(b) $Need_i \leq Work$ (i.e., for all k , $Need_i[k] \leq Work[k]$)

If no such i exists, go to step 4

3. $Work = Work + Allocation_i$

$Finish[i] = true$

go to step 2

4. If $Finish[i] == true$ for all i , then the system is in a safe state

Resource-Request Algorithm for Process P_i

Request = request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

Available = *Available* - *Request*;

*Allocation*_{*i*} = *Allocation*_{*i*} + *Request*_{*i*};

*Need*_{*i*} = *Need*_{*i*} - *Request*_{*i*};

If safe \Rightarrow the resources are allocated to P_i

If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example of Banker's Algorithm

◆ 5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			

Example (Cont.)

- ◆ The content of the matrix *Need* is defined to be *Max – Allocation*

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

- ◆ The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria

Example: P_1 Request (1,0,2)

- Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow$ true

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?

Thank You!