

Final report

SID: 11812318 Name: 张嘉兮

Task 1 Efficient Alarm Clock

Dara structures and functions

- *thread.h*:

- Add field `blocking_tick` into *struct thread*. The field will record the ticks this thread will be blocking, and then unblock.

```
1 | int64_t blocking_tick;           /* The ticks this thread  
   | will be blocking. */
```

- Add function ***thread_blocking_check*** to check `blocking_tick` for each thread and check whether unblock it.

```
1 | void thread_blocking_check (struct thread *, void *);
```

Algorithms

This task is to avoid the 'busy waiting' in `timer_sleep`, where the thread will constantly check whether sleep time is finish.

So I use a 'sleep-being woken up' mechanism to avoid the loop for checking.

1. Assume thread ***t*** calls the ***timer_sleep(n)*** with ***n*** time ticks.
2. Then I will block the thread ***t*** and set ***n*** to its field `blocking_tick`.
3. In order to wake up ***t***, I modify the function ***timer_interrupt***, where I call ***thread_blocking_check*** for all thread.
4. And in the ***thread_blocking_check***, if the `status` of thread input is `BLOCKING` and the `blocking_tick` is greater than 0, then the `blocking_tick` will be minus by 1.
5. And if the `blocking_tick` reaches 0, then I will unblock the thread input.
6. So the `blocking_tick` of ***t*** will continuous minus by 1 until reach 0 (for ***n*** time ticks), when is the time for waking it up. And ***t*** will be unblocked.

Synchronization

Potential concurrent accesses to shared resources

- The accesses to field `blocking_tick`, which will be accessed by function ***timer_sleep*** and function ***thread_blocking_check***. And the ***thread_blocking_check*** will be called every ***timer_interrupt***.
- The call of ***thread_block*** and ***thread_unblock***, which will access some shared resources including `status` of thread, the `ready_list` and etc.

Strategy

- I use the code described below to guarantee the access of field `blocking_tick` is atomic, unable to be interrupted.

```
1 enum intr_level old_level = intr_disable ();
2 /* Codes */
3 intr_set_level (old_level);
```

Then the "codes" inside will exclusive access the resources.

- Function ***thread_block*** and ***thread_unblock*** has been guaranteed the correct synchronization by pintos using the same mechanism.

Rationale

- Shortcomings: If there are too many threads, then the execution of ***timer_interrupt*** will be time-consuming.
- Time complexity: $O(n)$ for each execution of ***timer_interrupt*** where n is the number of threads in `all_list`. As each execution of ***timer_interrupt*** will check the threads one by one, and the operation of check for one thread is $O(1)$.
- Space complexity: $O(n)$ where n is the number of threads in `all_list`. As we use a field `blocking_tick` for each thread to save the number of ticks it may be blocking.

Task 2 Priority Scheduler

Data structures and functions

- *thread.h*:
 - Add field `lock_waiting_for` in *struct thread*, for saving the lock that thread acquire for but occupied.

```
1 struct lock * lock_waiting_for; /* The lock this thread
  is waiting for. */
```

- Add field `locks_holding` in *struct thread*, for saving a list of locks that thread hold.

```
1 | struct list locks_holding;          /* The locks this thread
   | is holding. */
```

- Add field `undonated_priority` in *struct thread*, for saving the priority set.

```
1 | int undonated_priority;            /* The origin priority of
   | the thread. */
```

- Add function ***thread_donate_priority*** for a thread donating its priority for a hold lock recursively.

```
1 | void thread_donate_priority (struct thread *, struct thread*);
```

- Add function ***thread_check_priority*** for a thread to calculate the real priority it should have after donation.

```
1 | void thread_check_priority (struct thread *);
```

- Add function ***thread_priority_comparator*** as a *list_less_func* that compare two *threads* by the `priority` of the thread.

```
1 | bool thread_priority_comparator (const struct list_elem *,
   | const struct list_elem *, void *);
```

- *synch.h*

- Add field `elem` into *struct lock* let *struct lock* becoming a type of list element.

```
1 | struct list_elem elem;            /* For used in list as lock can
   | only held by one thread. */
```

- Add function ***lock_priority_comparator*** as a *list_less_func* that compare two *locks* by the max priority of threads waiting for the lock.

```
1 | bool lock_priority_comparator (const struct list_elem *,
   | const struct list_elem *, void *);
```

- Add function ***locks_max_priority*** to find the maximal priority of *threads* in the waiting list of any *lock* which is in the list of locks.

```
1 | int locks_max_priority (const struct list *);
```

- Add function ***cond_sema_priority_comparator*** as a *list_less_func* that compare two *conditions* by the max priority of threads waiting for the *semaphore* that *condition* has.

```
1 | bool cond_sema_priority_comparator(const struct list_elem
    *, const struct list_elem *, void *);
```

Algorithms

This task is to implement the priority schedule and priority donation.

In this part, I will separately introduce my design for priority schedule and priority donation.

Priority schedule:

1. For **next_thread_to_run**, I modify **next_thread_to_run** as using **list_max** with **thread_priority_comparator** to get the **thread** that has maximal **priority** in the **ready_list**, and return this **thread**. As **schedule** will call **next_thread_to_run** to find the next thread should run, I have guarantee **schedule** will choose the **thread** has maximal **priority**.
2. For **semaphore**, I modify **sema_up** as using **list_max** with **thread_priority_comparator** to get the **thread** that has maximal **priority** in the **waiters** of the **semaphore**, and **unblock** this **thread**.
3. For **condition**, I modify **cond_signal** as using **list_max** with **cond_sema_priority_comparator** to get the **semaphore_elem** which has the **thread** with maximal **priority** in **all semaphore's waiters** from the **waiters** of the **condition**, and then **sema_up** the **semaphore** I selected.
4. For **lock**, as I have modified the **sema_up**, and **lock_release** call **sema_up**, so I have implemented this.

Priority donation:

1. When **acquire lock**, I need to check whether there is a holder for this lock, if so, then I need try to donate priority recursively along the lock requirement chain. I implement this by **thread_donate_priority**. Otherwise, it will execute the **sema_down**.
2. **thread_donate_priority**:
 1. Let we call the **lock's holder** as **l_holder**. I will call **thread_check_priority** on **l_holder** to find the correct priority after donation.
 2. **thread_check_priority**: I will search all the **waiters** of the **semaphore** of the **locks_holding** (a list containing locks holding by **l_holder**) of the **l_holder**, and find the maximal priority all the threads in **waiters** have. Let we call it **max_pri**. Then I will let the **priority** of the **l_holder** become the maximal one between **l_holder's priority** and **max_pri**.
 3. Then **l_holder's priority** has been donated. Finally, if **l_holder** is requiring a **lock l2**, then I will recursively call **thread_donate_priority** from **l_holder** to **l2's holder**. Recusing like this until a **lock's holder** no longer has the requested lock.
3. After priority donation, it will call **sema_down**.

4. After **sema_down**, it get the *lock* and become the *holder* of the *lock*, then I need to add the *lock* into the *locks_holding* of the current thread.
5. When **release lock**, I need to remove the *lock* from *locks_holding* of the current thread. Then I call **thread_check_priority** to let the *undonated_priority* become the *priority* of the current thread.
6. When **changing thread's priority**, I will change the *undonated_priority* of the thread, then call **thread_check_priority** to evaluate the correct *priority* after donation. Finally, call the **thread_yield** to do a schedule.

Synchronization

Potential concurrent accesses to shared resources

- Operation of the *locks_holding* of *struct thread* as a *lock* may be required concurrently.
- Modification of the *priority* of a thread (can be modified by itself and other thread by donating, even other threads may concurrently donate to the same thread).
- Original operations about *semaphore*, *condition*, *lock* and other sync objects.
- *lock_waiting_for*, *undonated_priority* are safe because they can only modified by the thread itself.

Strategy

- I use the code described below to guarantee the access of the fields in 'Codes' is atomic, unable to be interrupted.

```
1  enum intr_level old_level = intr_disable ();
2  /* Codes */
3  intr_set_level (old_level);
```

Then the "codes" inside will exclusive access the resources.

- I assume the original operations about sync objects are safe implementation by pintos. In fact, the implementation is the same as I mentioned above.

Rationale

Advantages

My implementation is easy to maintain the *priority* of a *thread*, because it can be calculated easily by call **thread_check_priority**.

Shortcomings

The dynamic computing mode for priority donation may take more time in especial the number of threads is large.

The priority schedule can be improved by a heap, which will let the complexity reduced to $O(\log n)$.

Time complexity

- $O(n)$ for next thread to run, semaphore, condition, lock finding the next thread to switch to or to take over the sync object, where n is the number of threads in `all_list`.
- $O(n)$ for priority donation, as one thread can only require for one lock in a time. Then the complexity to find the maximal priority of donators is $O(n)$, where n is the number of threads in `all_list`.

Space complexity

- $O(n+m)$ for the field `locks_holding` and `undonated_priority` in *struct thread* and field `elem` in *struct lock*, where n is the number of threads in `all_list` as each thread has a `locks_holding` and m is the number of exist locks as each *lock* has a `elem`.

Task 3 Multi-level Feedback Queue Scheduler (MLFQS)

Data structures and functions

- *thread.h*:
 - Add field `nice` into *struct thread* to save the value of 'nice' in the formula.

```
1 | int nice;                                /* Nice value    for  
   | mlfqs. */
```

- Add field `recent_cpu` into *struct thread* to save the value of 'recent_cpu' in the formula.

```
1 | fixed_t recent_cpu;                      /* The recent    cpu for  
   | mlfqs. */
```

- Add function *threads_update_mlfqs* for *timer_interrupt* to call to deal with the updating of values about mlfqs.

```
1 | void threads_update_mlfqs (int, int);
```

- Add function *thread_update_recent_cpu_pri* for updating the `recent_cpu` and `priority` for a thread.

```
1 | void thread_update_recent_cpu_pri (struct thread *, void *);
```

- Add function *thread_update_pri* for updating the `priority` for a thread.

```
1 | void thread_update_pri (struct thread *, void *);
```

- *thread.c*:
 - Add static variable `load_avg` for 'load_average' in formula.

```
1 | /* The load_average for mlfqs. */  
2 | static fixed_t load_avg = FP_CONST (0);
```

Algorithms

In this task, I just implement the formulas given to me.

1. I call *thread_update_mlfqs* in *timer_interrupt*.
2. In *thread_update_mlfqs*, I check that the `thread_mlfqs` is enabled, and then I increase the `recent_cpu` of the current thread by 1.
3. Then if the time ticks is 4's multiples, I update the priority of all thread by
$$priority = PRI_MAX - (recent_cpu/4) - (nice \times 2)$$
4. Then if the time ticks is `TIMER_FREQ`'s multiples, I update the `load_avg` by
$$load_avg = (59/60) \times load_avg + (1/60) \times ready_threads$$

And update the `recent_cpu` by

$$recent_cpu = (2 \times load_avg) / (2 \times load_avg + 1) \times recent_cpu + nice$$

The priority of threads changed in time increasing. The running thread having the highest priority will get increasing `recent_cpu` more, then its `priority` will decrease. Therefore, MLFQS guarantees the bounded-waiting and no starvation.

Synchronization

Potential concurrent accesses to shared resources

- the `priority` of a thread, as it can be modified by both *timer_interrupt* and *thread_set_nice*, where may occur synchronization problem.

Strategy

- I use the code described below to guarantee the access of the fields in 'Codes' is atomic, unable to be interrupted.

```
1 | enum intr_level old_level = intr_disable ();  
2 | /* Codes */  
3 | intr_set_level (old_level);
```

Then the "codes" inside will exclusive access the resources.

Rationale

Time complexity

- $O(n)$ where n is the number of threads in `all_list`. As each thread should be update.

Space complexity

- $O(n)$ where n is the number of threads in `all_list`. As each thread should save the field `recent_cpu` and `nice`.

3.1.2

timer ticks	R(A)	R(B)	R(C)	P(A)	P(B)	P(C)	thread to run
0	0	0	0	63	61	59	A
4	4	0	0	62	61	59	A
8	8	0	0	61	61	59	B
12	8	4	0	61	60	59	A
16	12	4	0	60	60	59	B
20	12	8	0	60	59	59	A
24	16	8	0	59	59	59	C
28	16	8	4	59	59	58	B
32	16	12	4	59	58	58	A
36	20	12	4	58	58	58	C

Yes, there are ambiguities make value uncertain as we can implement the selecting of thread among threads having same priority in different way.

According to my codes, I choose the thread enqueued firstly when there are some threads have the same priority.

And if I change the strategy for choosing the thread among threads having same priority, the result will be different.