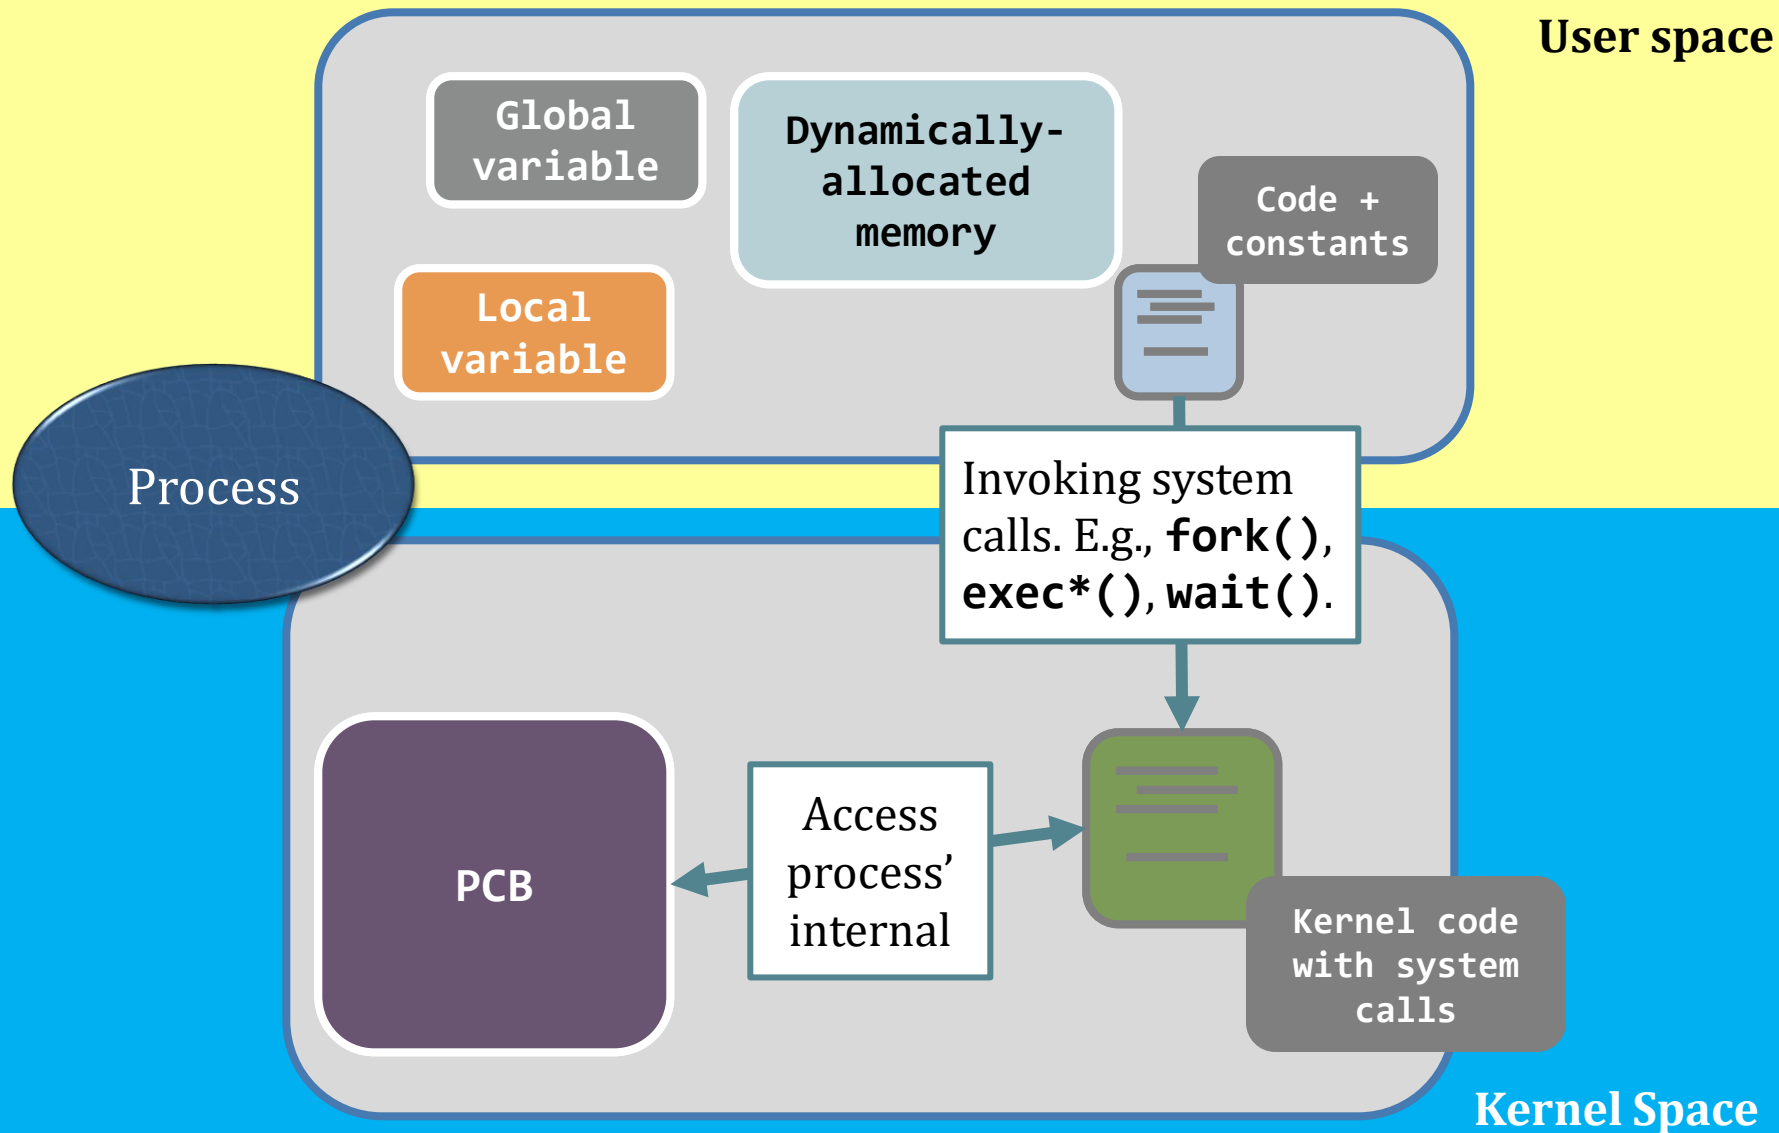


Lecture 4: Process II


Bo Tang @ 2020, Spring

The story so far...




When invoking a system call (memory view)

- ◆ When running a program code of a user process.
- ◆ As the code is in user-space memory, so the program counter is pointing to that region.



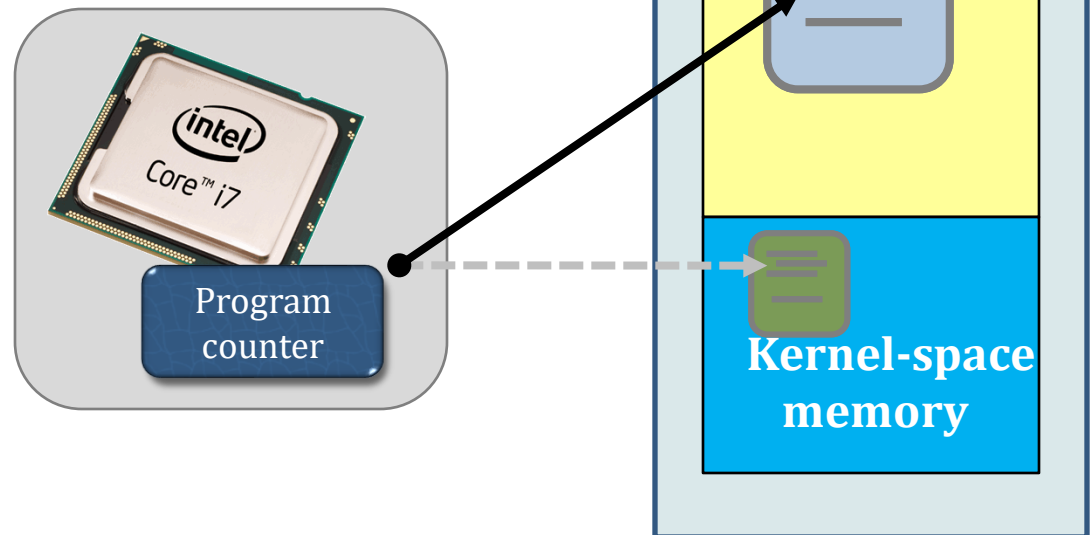
When invoking a system call (memory view)

- ◆ When the process is calling the system call “**getpid()**”.
- ◆ Then, the CPU switches from the user-space to the kernel-space, and reads the PID of the process from the kernel.

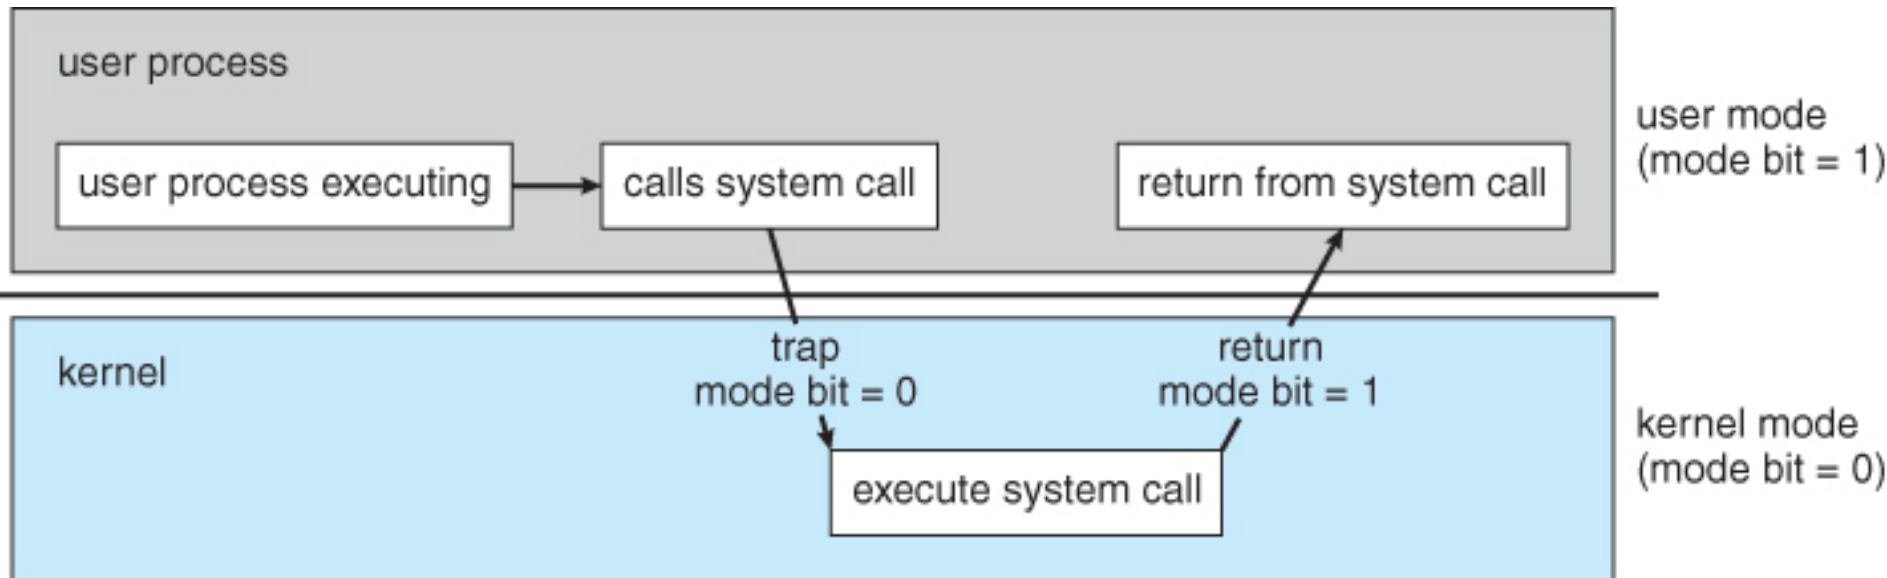


When invoking a system call (memory view)

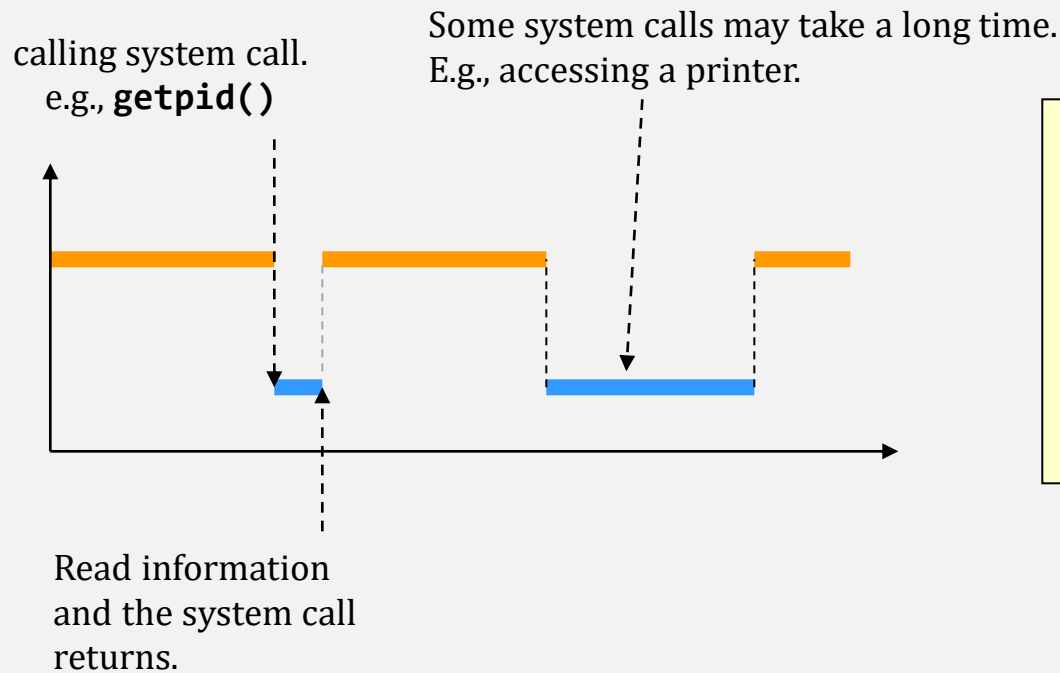
- ◆ When the CPU has finished executing the “**getpid()**” system call
 - ◆ it switches back to the user-space memory, and continues running that program code.



When invoking a system call (CPU view)



Process real time cost (wall-clock time)



- User time** – CPU time spent on codes in user-space memory.
- Sys time** – CPU time spent on codes in kernel-space memory.

User time VS System time – example 1

- Let's tell the difference...with the tool “**time**”.

```
$ time ./time_example
```

```
real    0m0.001s
user    0m0.000s
sys     0m0.000s
$ _
```

The Real-time elapsed when “./time_example” terminates.

The user time of “./time_example”.

The sys time of “./time_example”.

It's possible:
real > user + sys
real < user + sys

Why?

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```


User time VS System time – example 1

- ◆ Let's tell the difference...with the tool “**time**”.

```
$ time ./time_example
```

```
real    0m0.001s
user    0m0.000s
sys     0m0.000s
$ _
```

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

Commented on purpose.

```
$ time ./time_example
```

```
real 0m2.795s
user 0m0.084s
sys  0m0.124s
$ _
```

See? Accessing hardware costs the process more time.

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        printf("x = %d\n", x);
    }
    return 0;
}
```

Comment released.

User time VS Sys time – example 2

- ◆ The user time and the sys time together **define the performance of an application.**
- ◆ When writing a program, you must consider both the user time and the sys time.
 - ◆ E.g., the output of the following two programs are exactly the same. But, their running time is not.

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX / 5 ; i++)
        printf("x\nx\nx\nx\nx\n");
    return 0;
}
```

User time VS Sys time – example 2

```
#define MAX 1000000
```

```
int main(void) {  
    int i;  
    for(i = 0; i < MAX; i++)  
        printf("x\n");  
    return 0;  
}
```

```
$ time ./time_example_slow
```

```
real 0m1.562s  
user 0m0.024s  
sys  0m0.108s  
$ _
```

```
#define MAX 1000000
```

```
int main(void) {  
    int i;  
    for(i = 0; i < MAX / 5 ; i++)  
        printf("x\nx\nx\nx\nx\n");  
    return 0;  
}
```

```
$ time ./time_example_fast
```

```
real 0m1.293s  
user 0m0.012s  
sys  0m0.084s  
$ _
```

User time VS Sys time

- ◆ Function calls cause overhead
 - ◆ Stack pushing (will see later)
- ◆ Sys calls may cause even more
 - ➔ Sys call is from another “process” (the kernel)
 - ➔ Switching to another “process” ➔ context switch (will see later)


<https://www.quora.com/Is-an-OS-kernel-itself-a-process>

Working of system calls

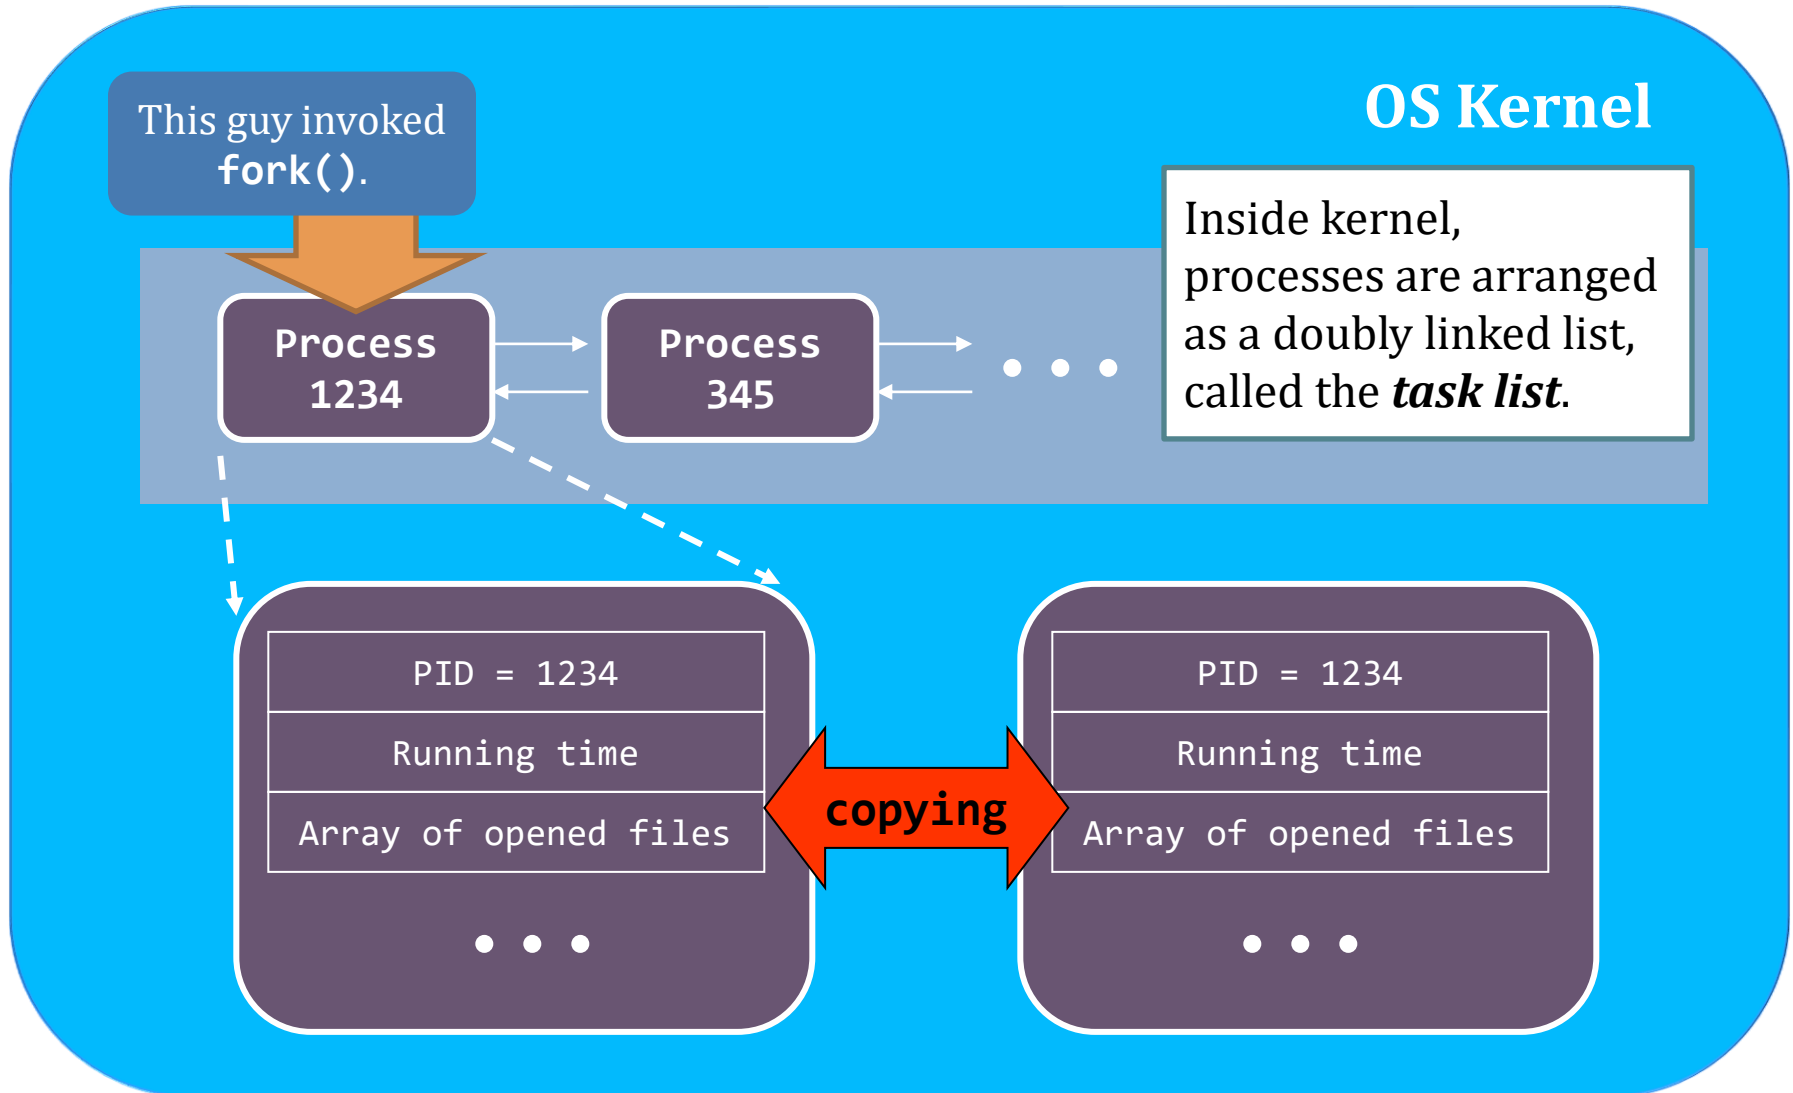
- `fork()`;



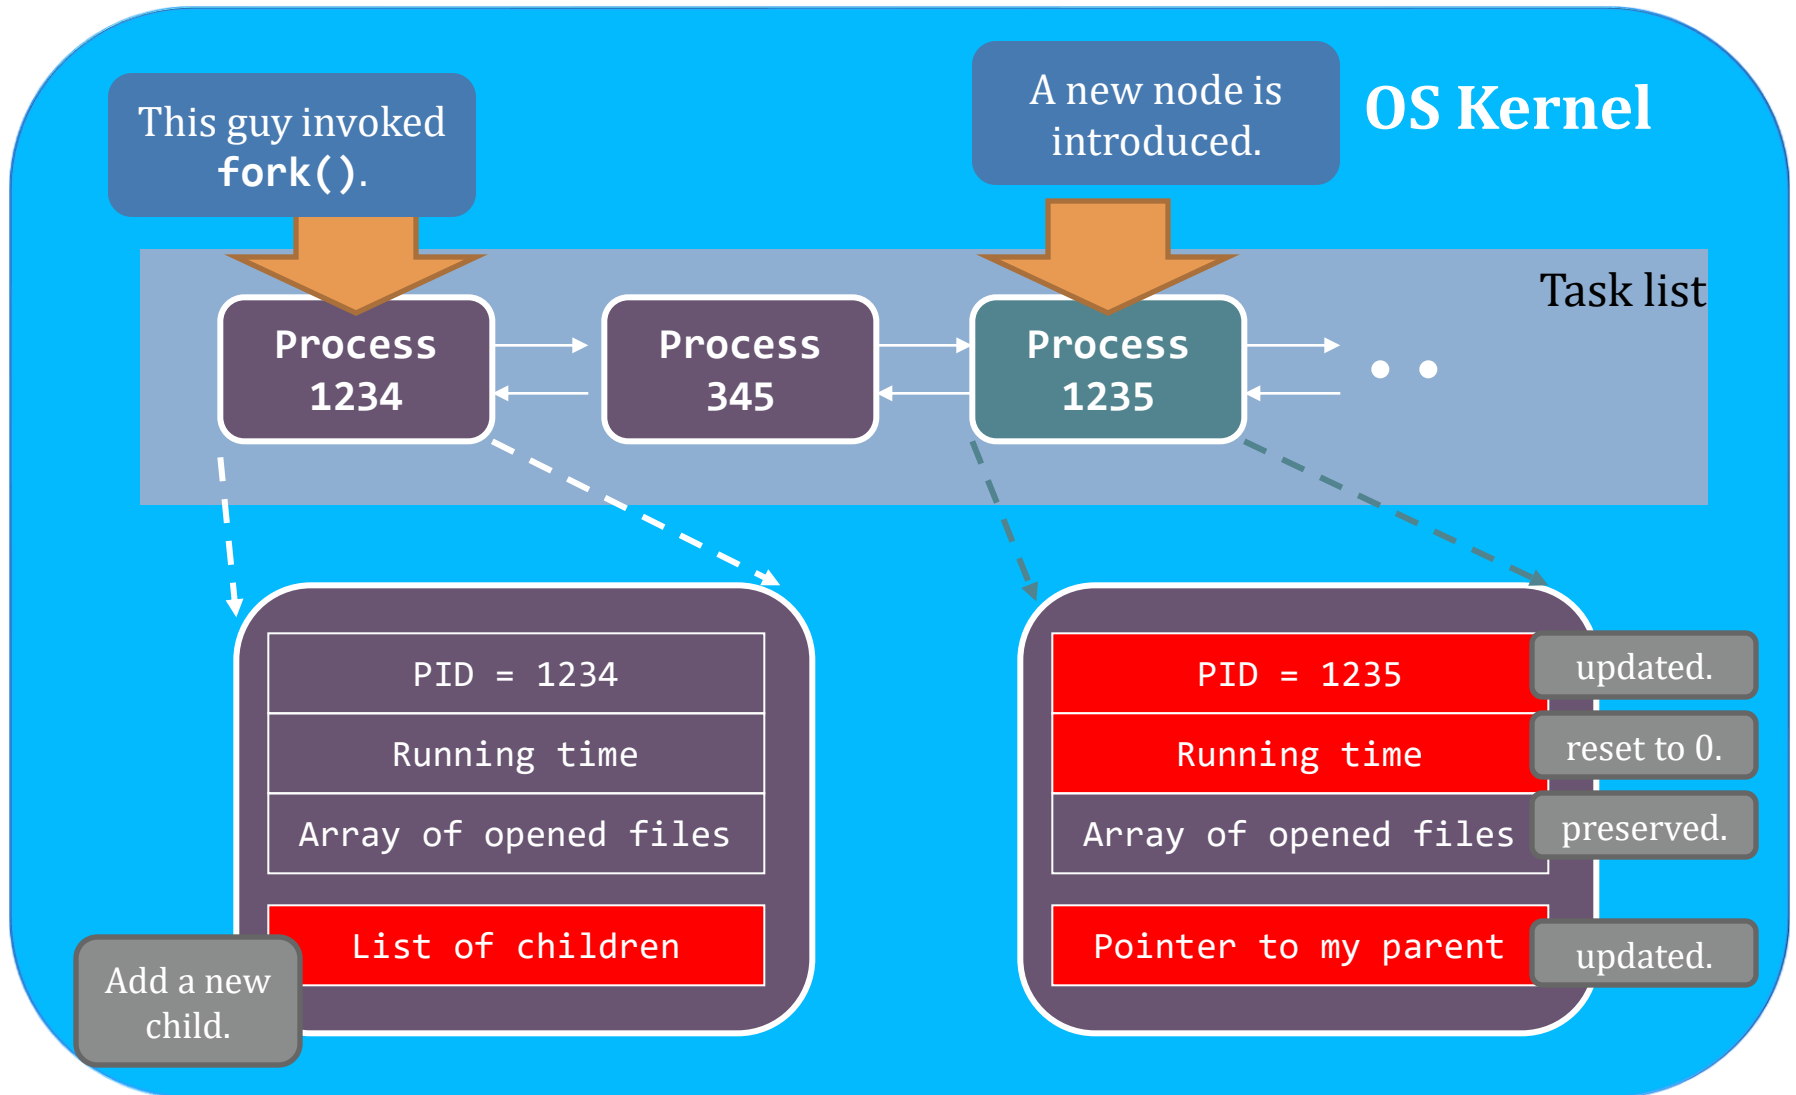
Programmer view of fork()



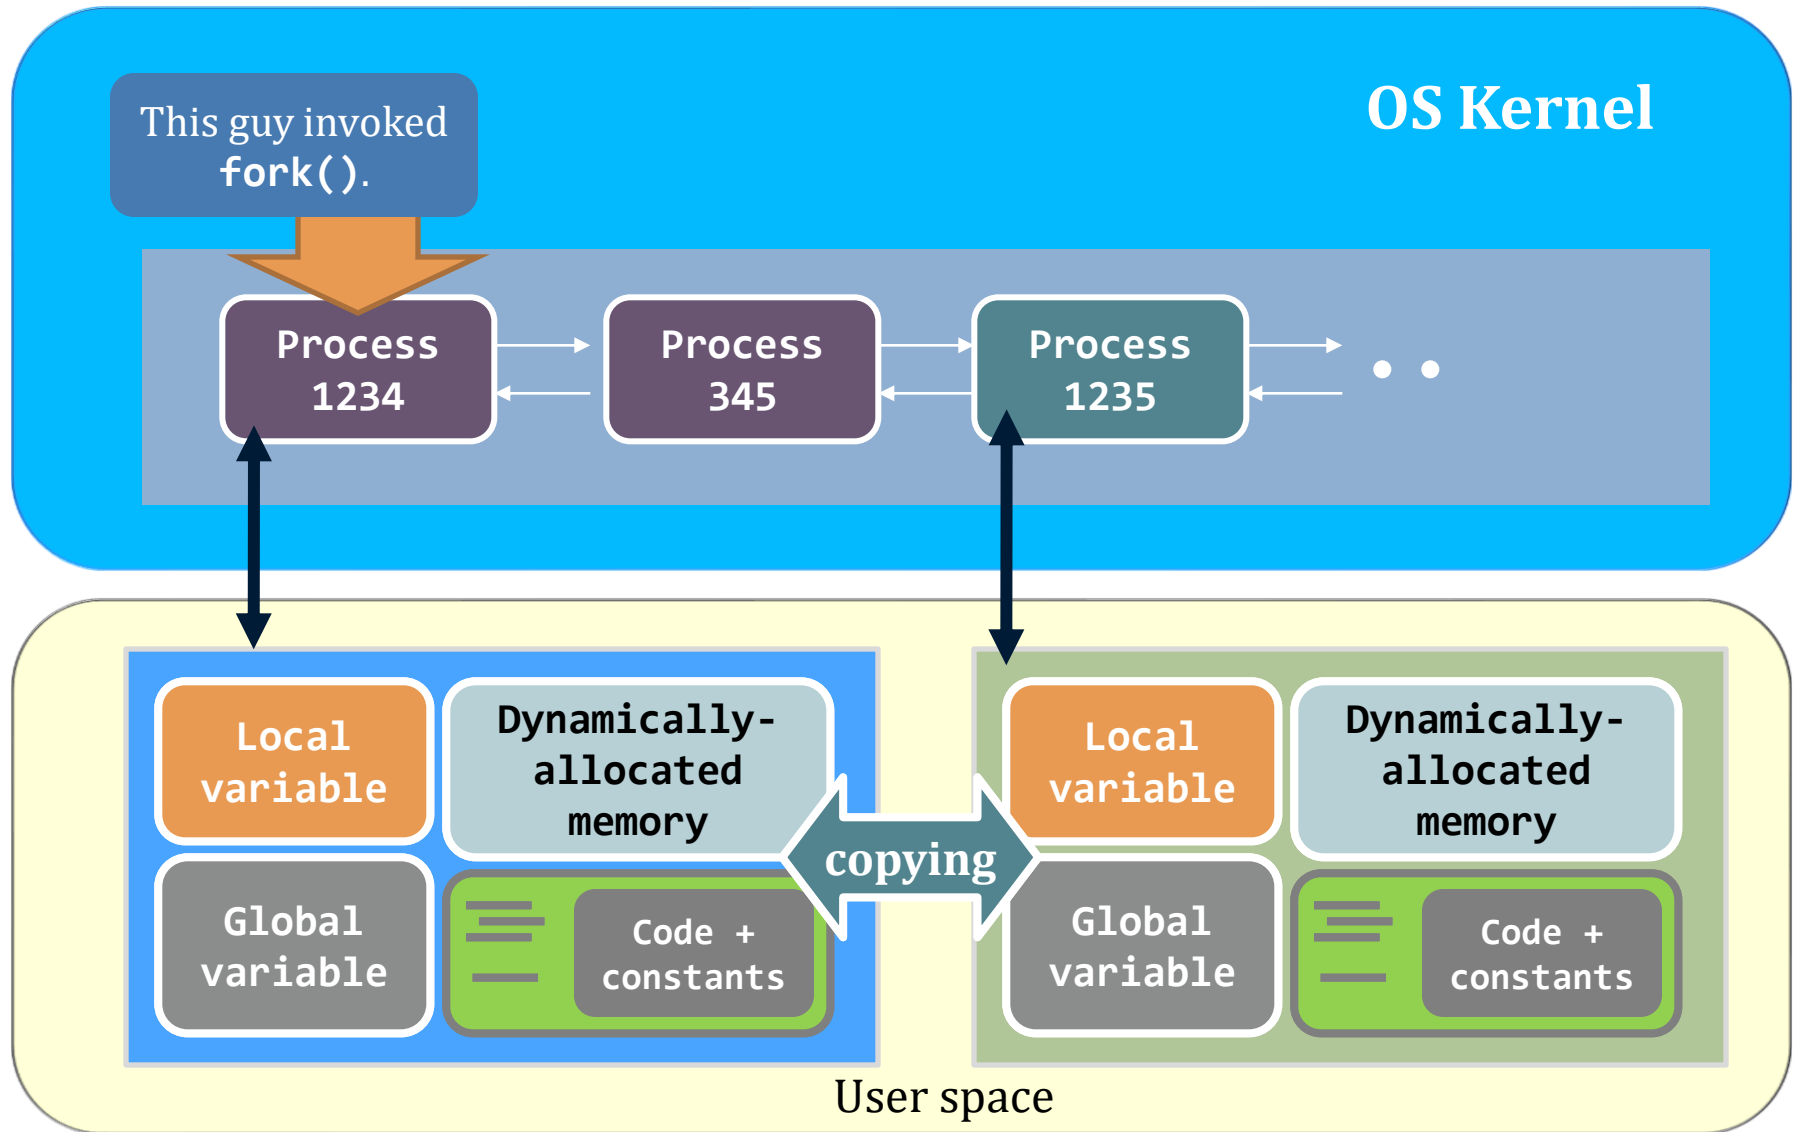
fork() inside the kernel



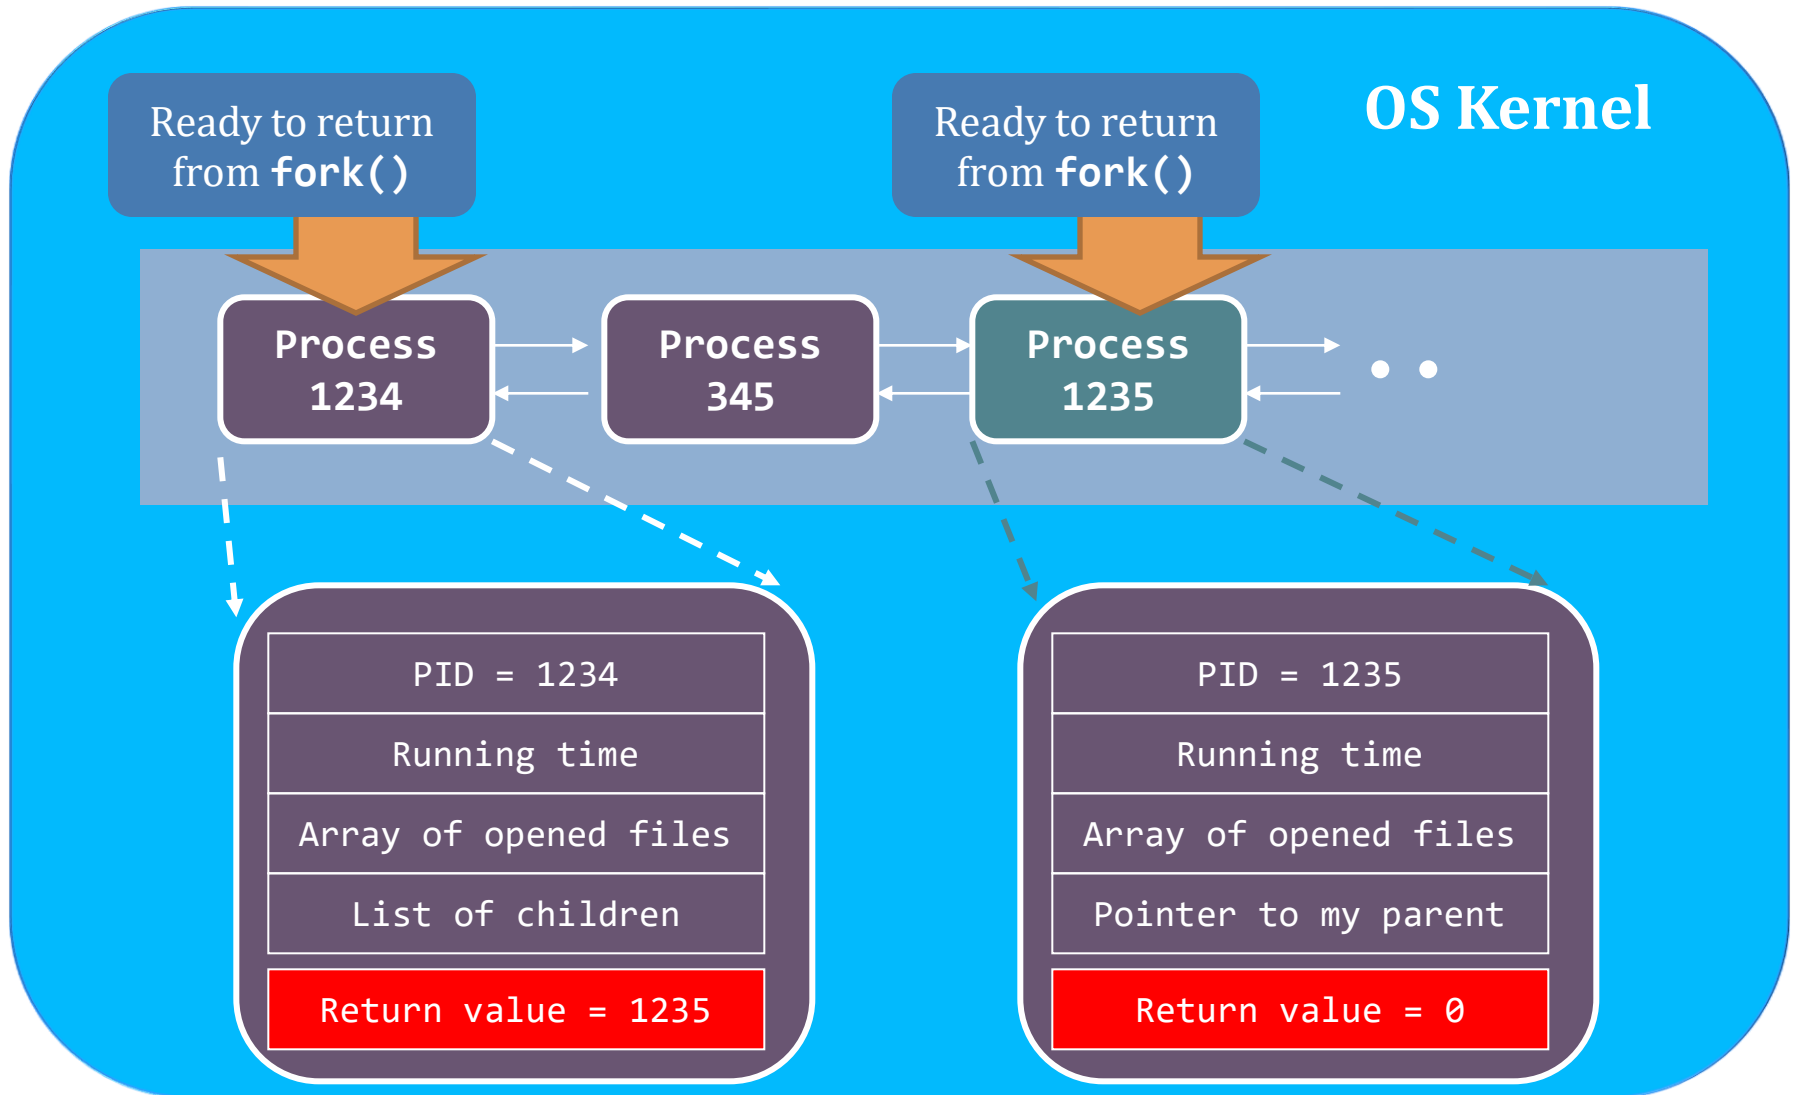
fork() in action – kernel-space update



fork() in action – user-space update



fork() in action – finish



fork() in action – array of opened files?

- ◆ Array of opened files contains:

| Array Index | Description |
|-------------|---|
| 0 | Standard Input Stream; FILE *stdin; |
| 1 | Standard Output Stream; FILE *stdout; |
| 2 | Standard Error Stream; FILE *stderr; |
| 3 or beyond | Storing the files you opened, e.g., fopen() , open() , etc. |


- ◆ That's why a parent process **shares the same terminal output stream** as the child process.

Stream is just a logical object for you to read as a sequence of bytes

So, how can you random access the middle of a file? Read Stream → Array → Array[mid-point]₉

fork() in action – sharing opened files?


- What if two processes, **sharing the same opened file**, write to that file together?



Let's see what will happen when the program finishes running!

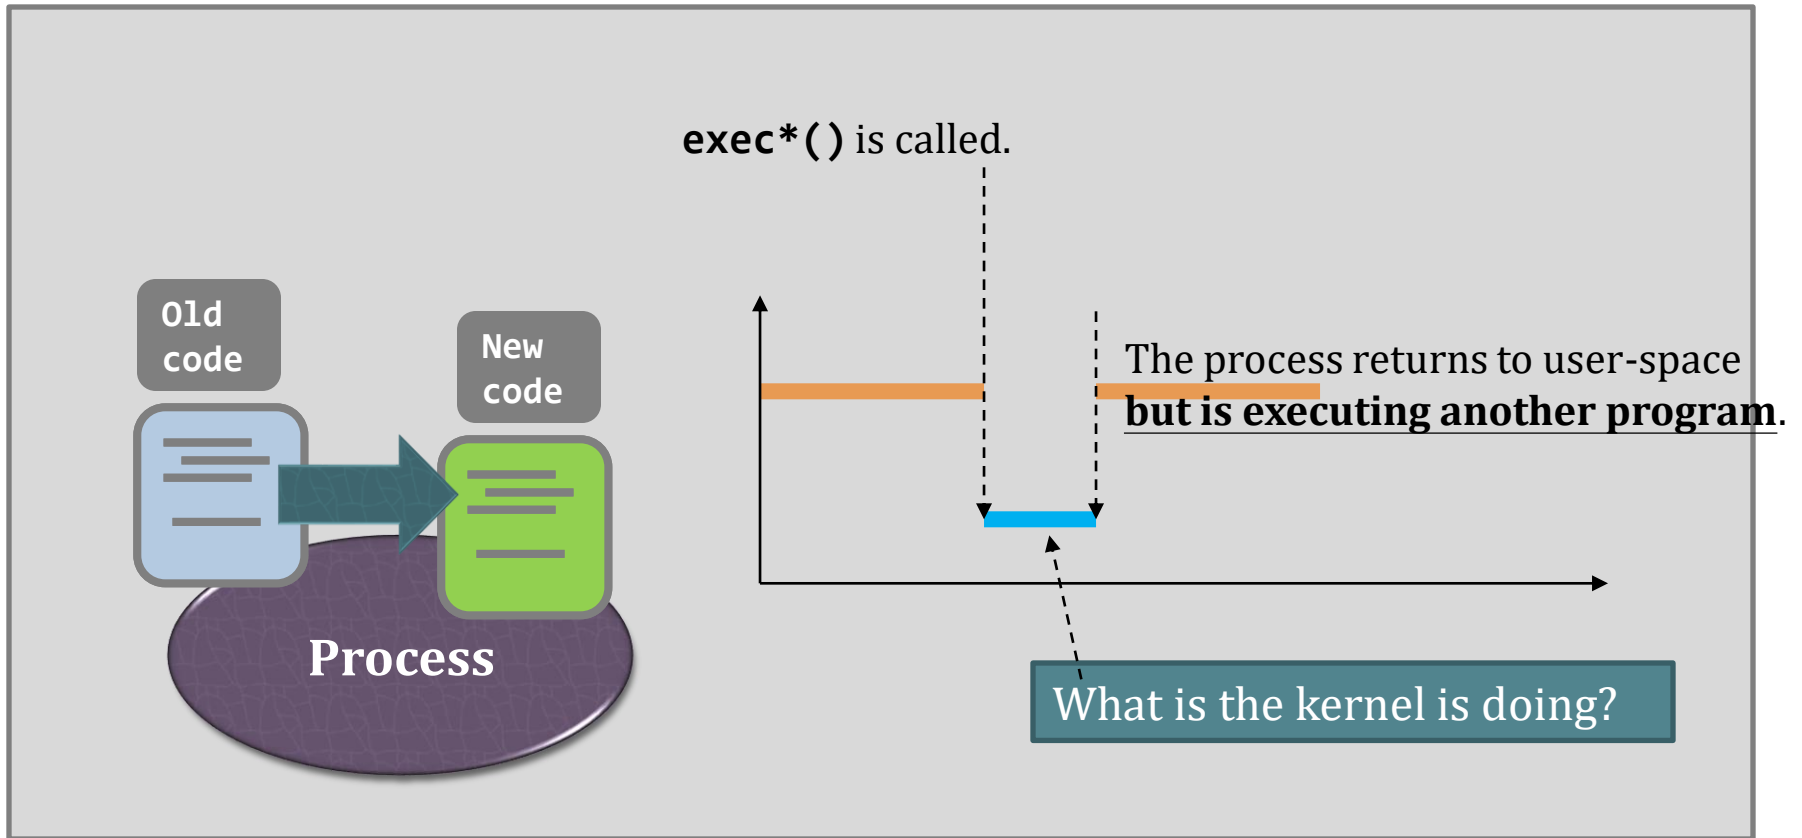
Working of system calls

- `fork()`;
- `exec*()`;

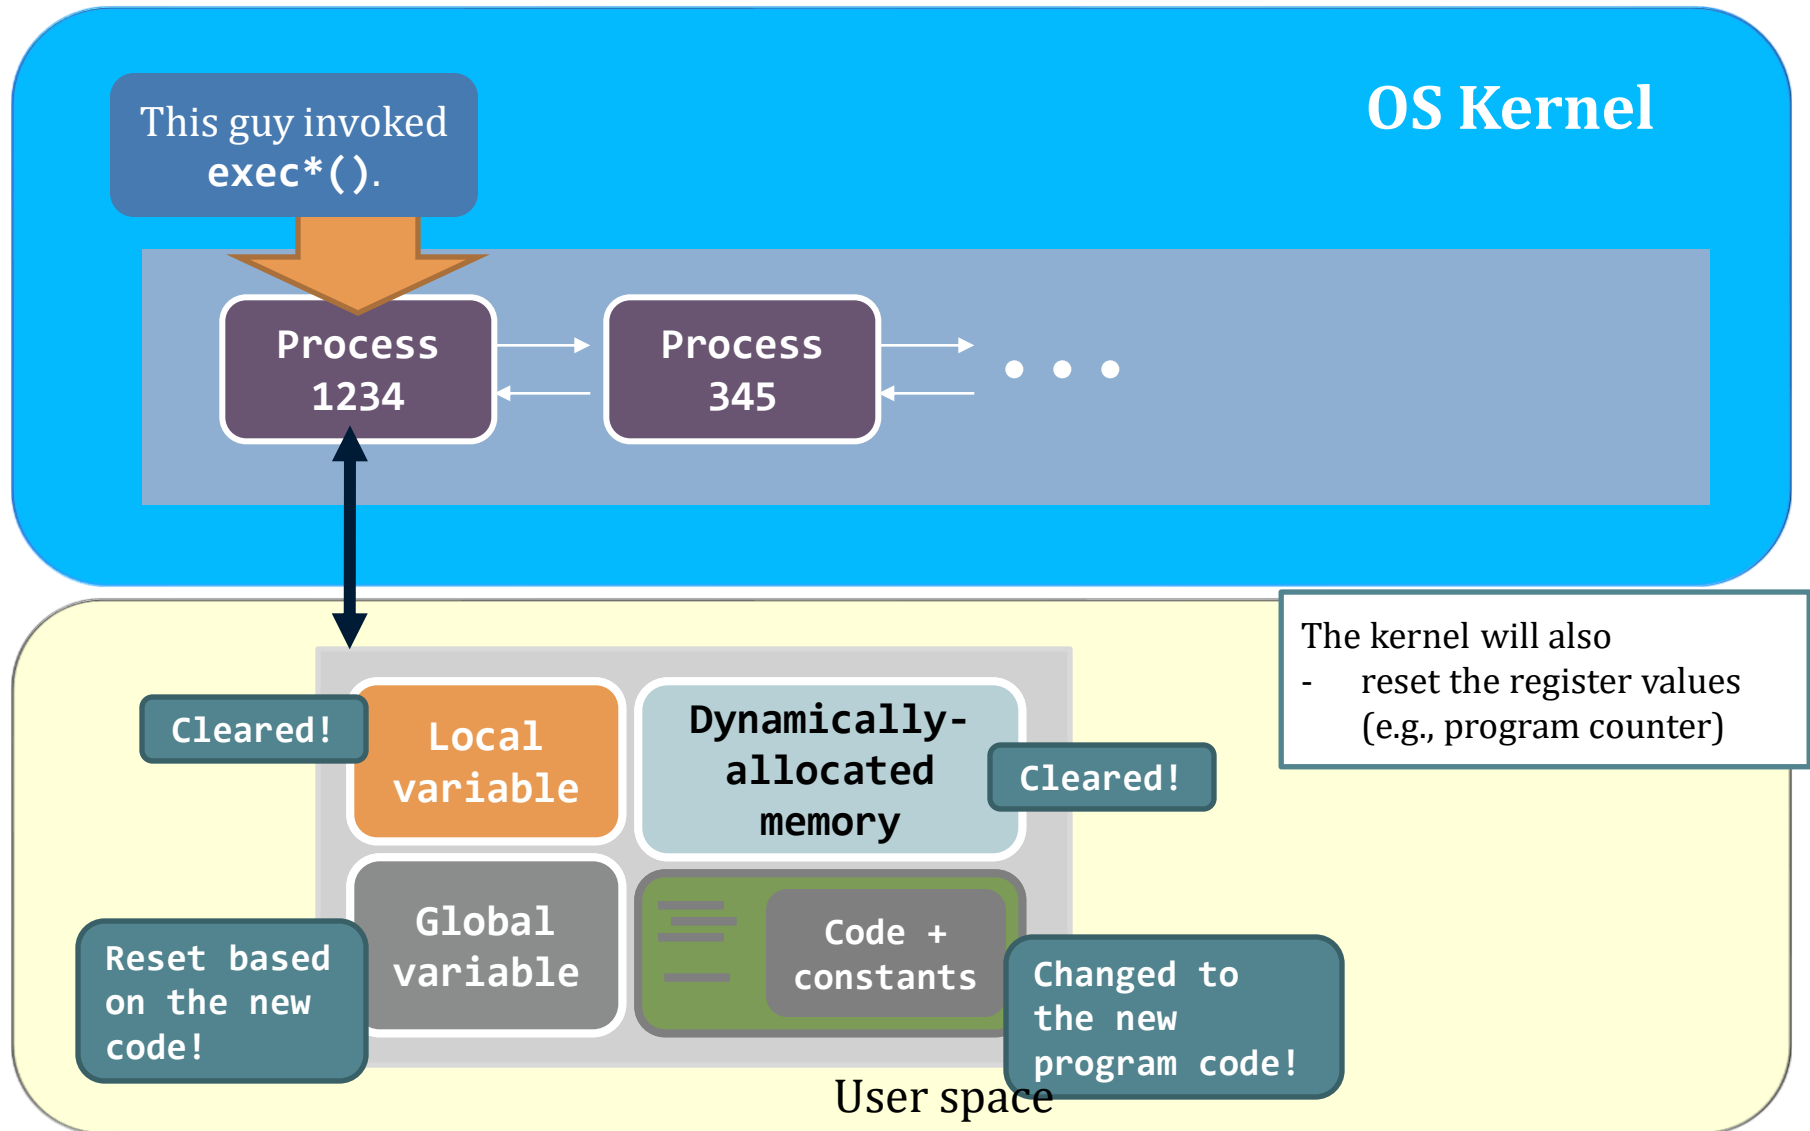


`exec*()` that you've learnt...

- ◆ How about the `exec*()` call family?




exec*() in action

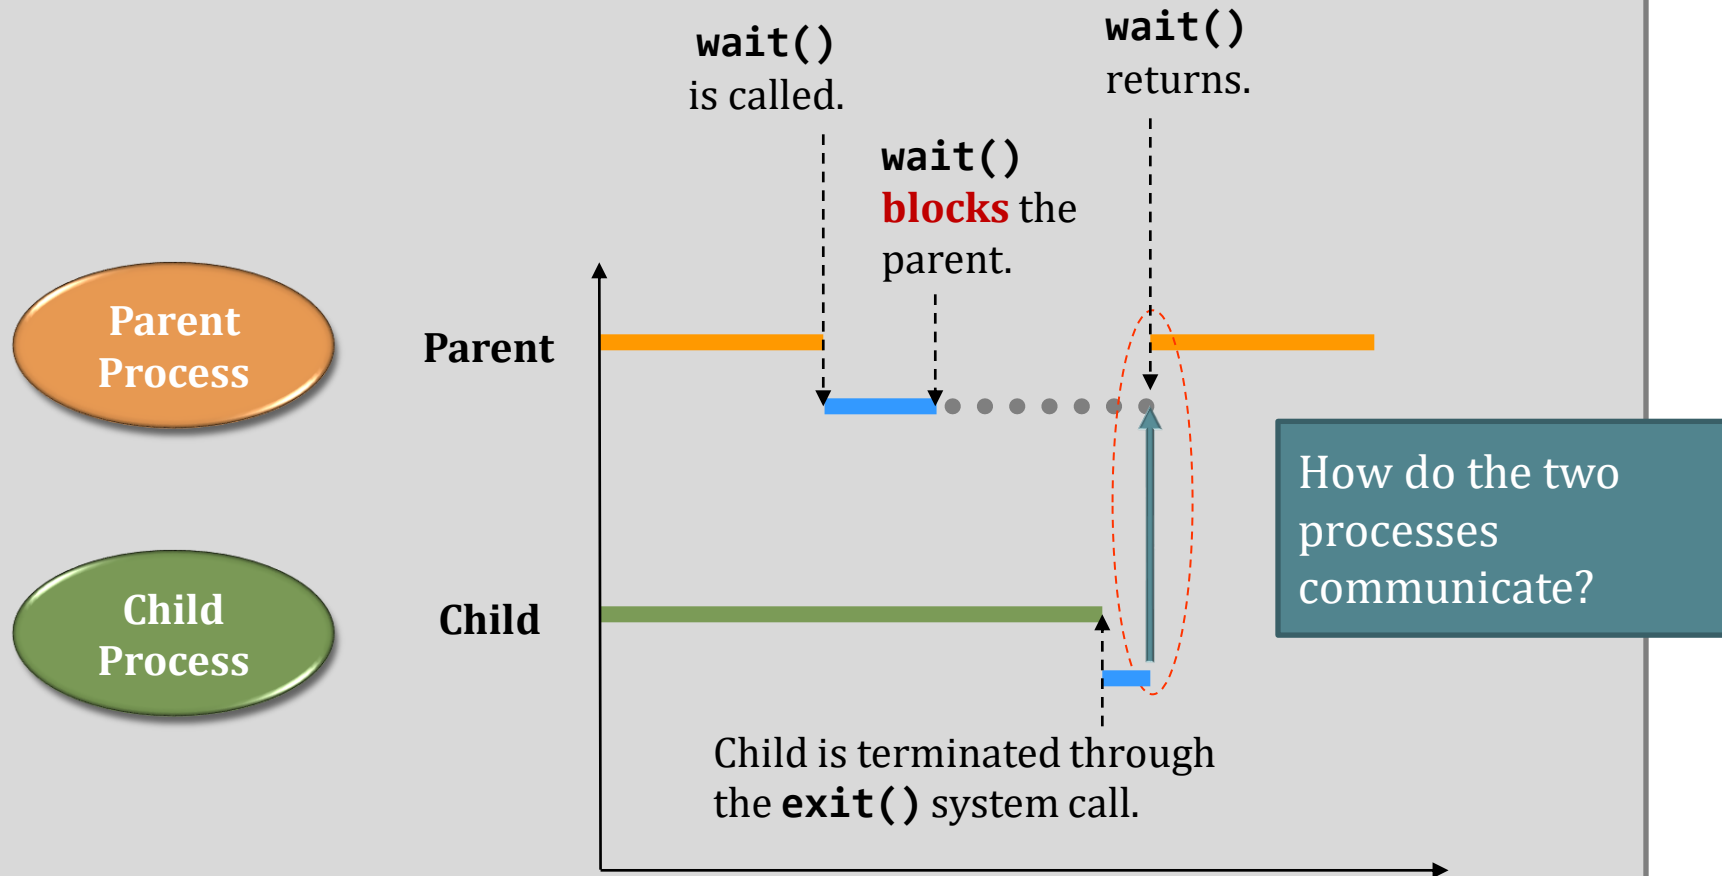


Working of system calls

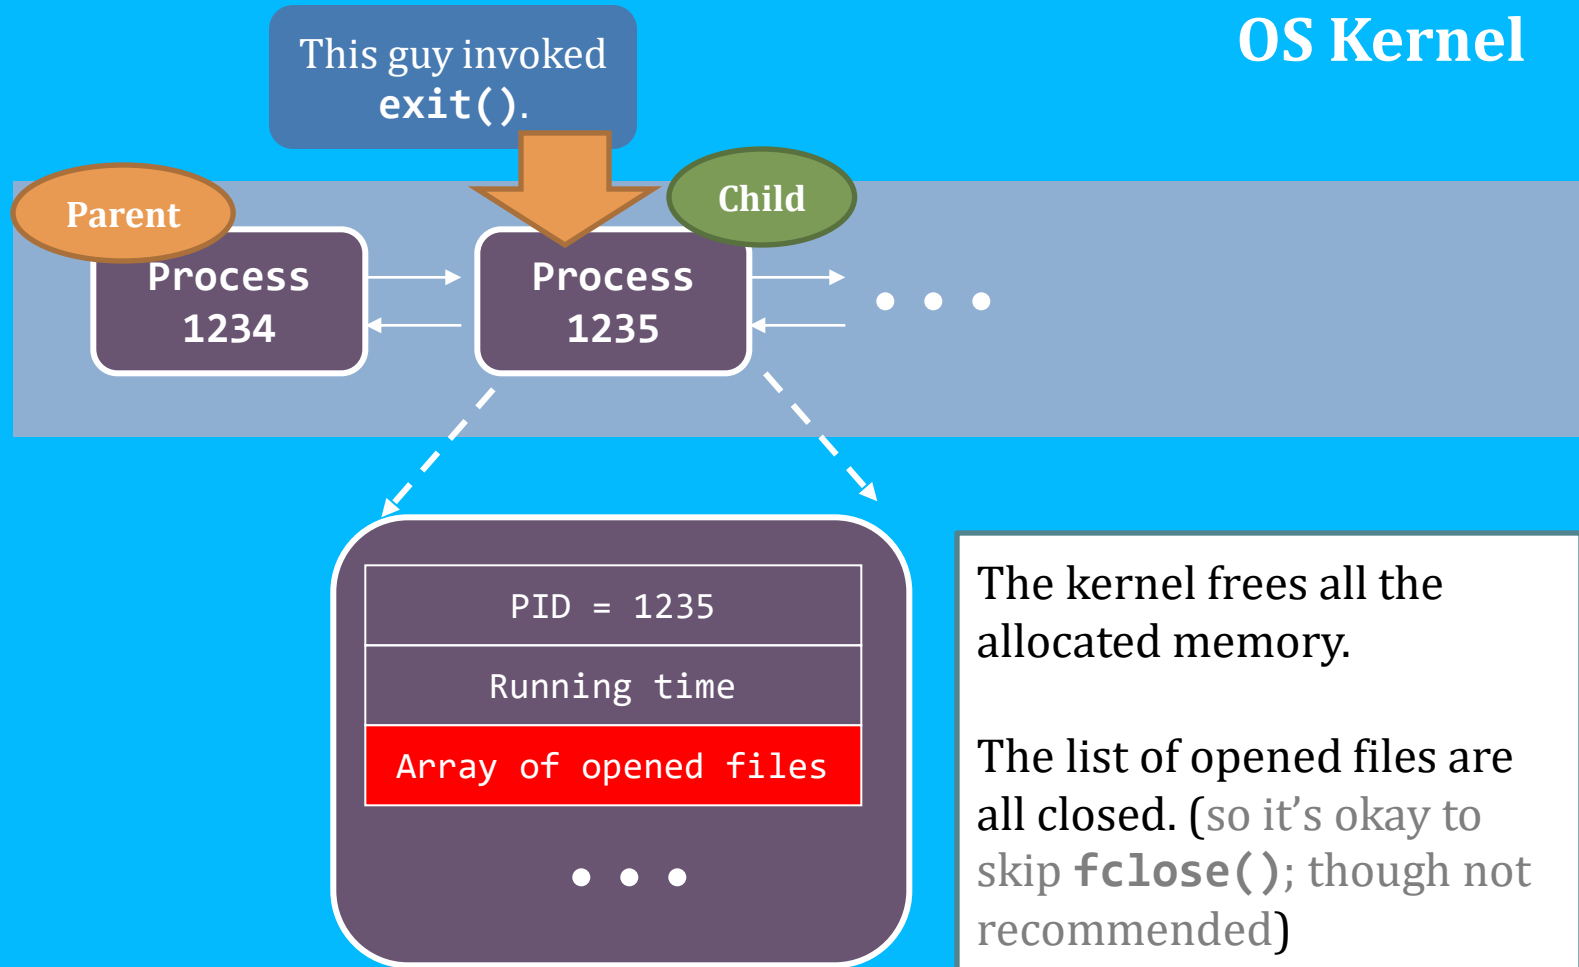
- `fork()`;
- `exec*()`;
- `wait() + exit()`;



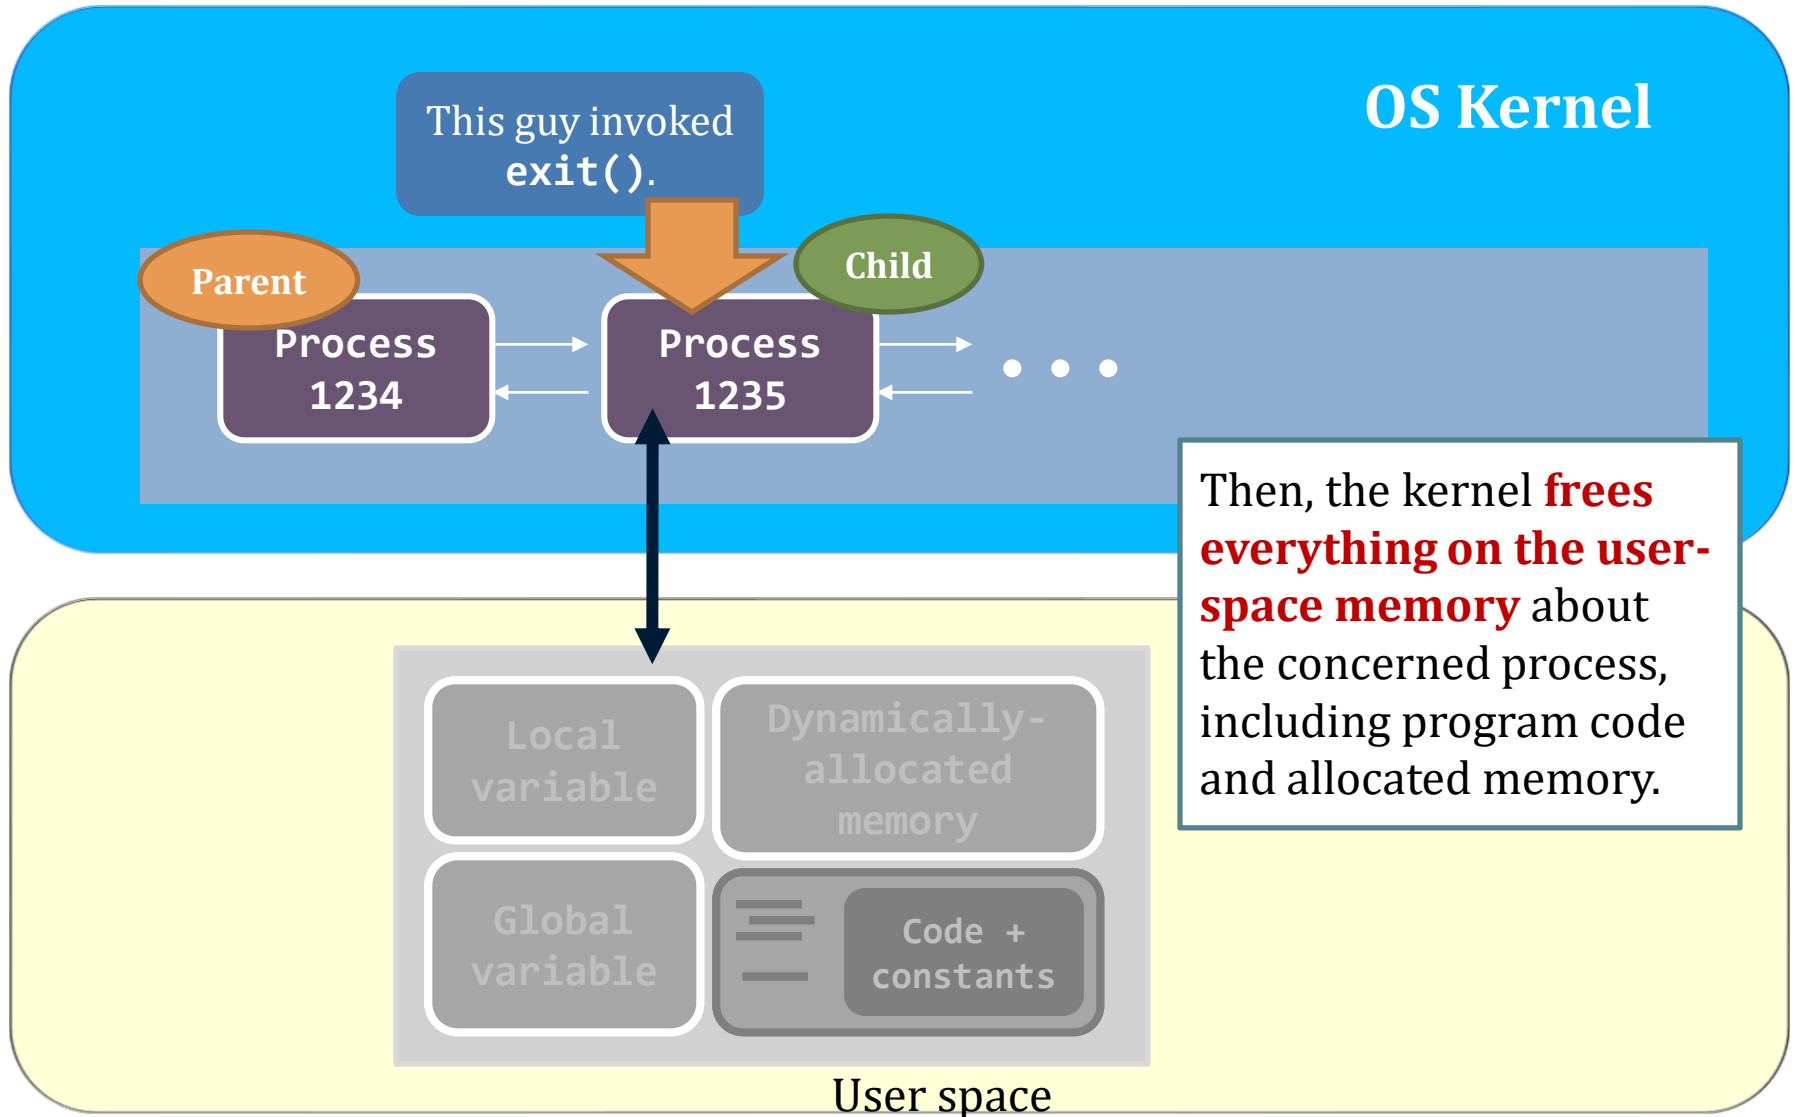
wait() and exit()



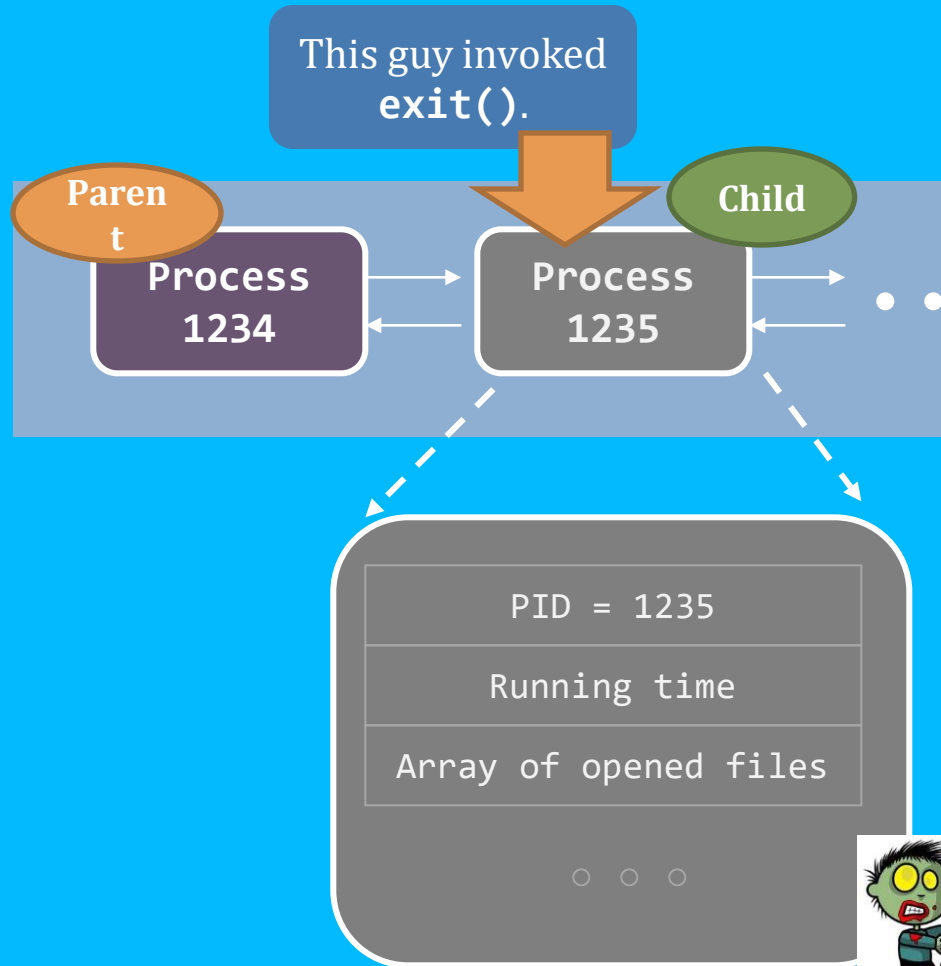
exit() (kernel-view)



exit() (kernel-view)



exit() (kernel-view)



OS Kernel

Process ID stills in the kernel's process table

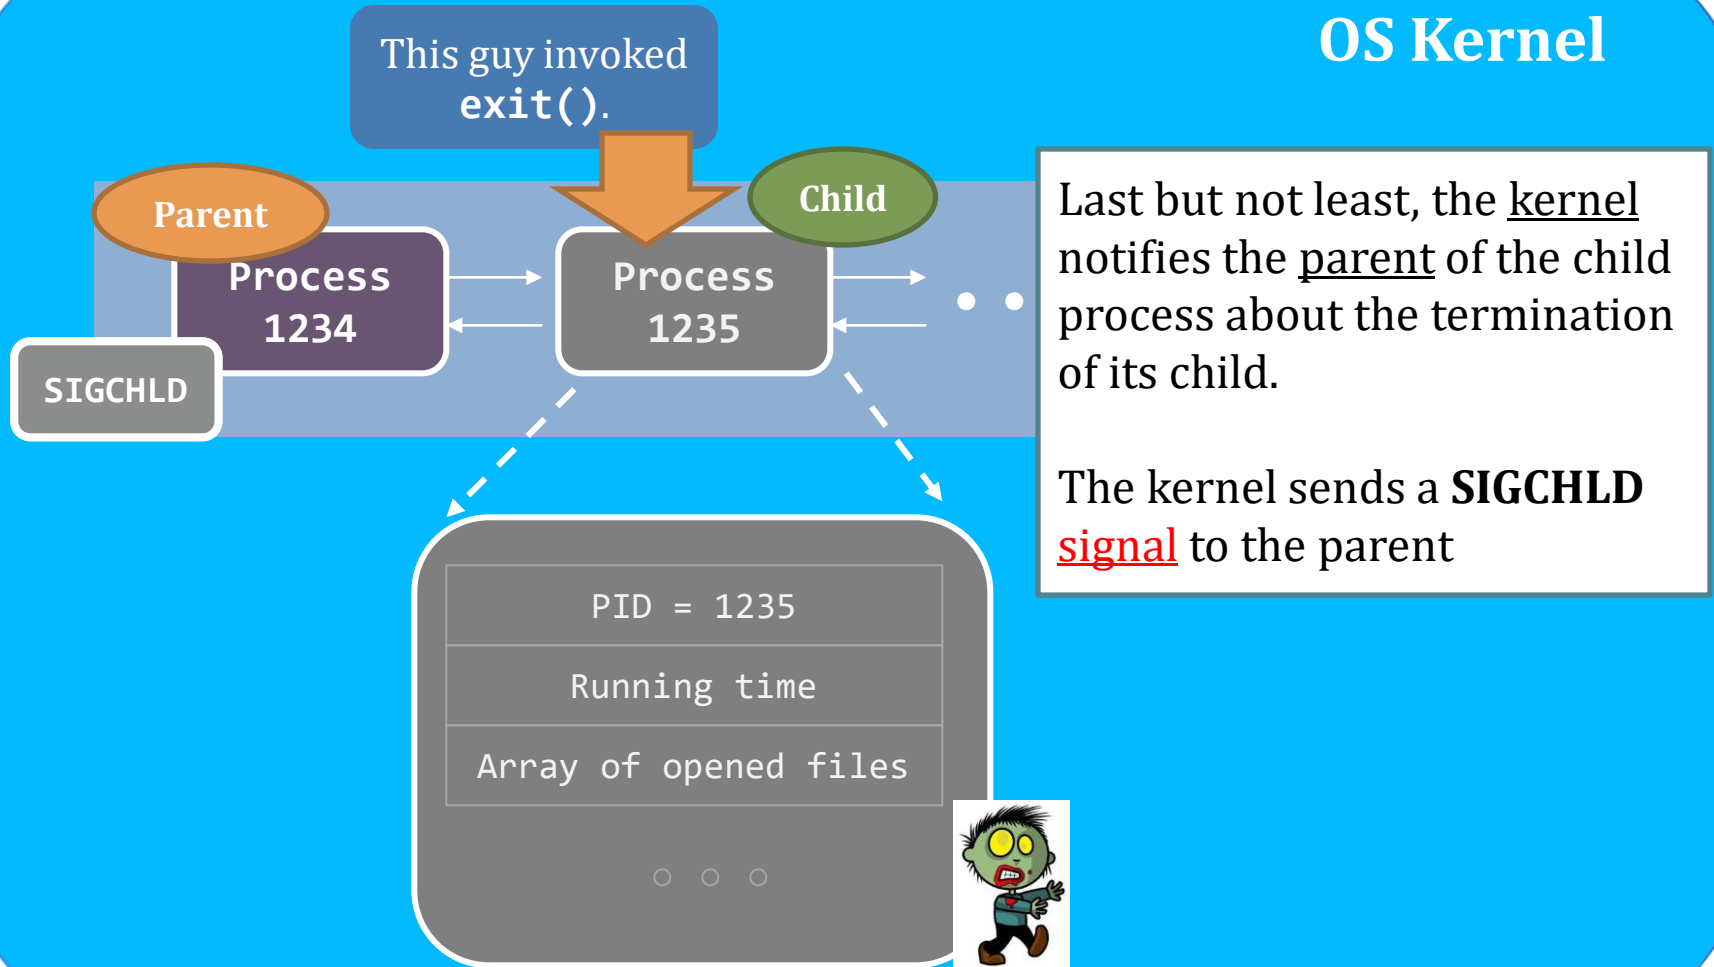
- Why?

[Wiki] *This entry is still needed to allow the process that started the (now zombie) process to read its exit status.*

The status of the child is now called **zombie** ("terminated").



exit() (kernel-view)




Summary -- what the kernel does for `exit()`

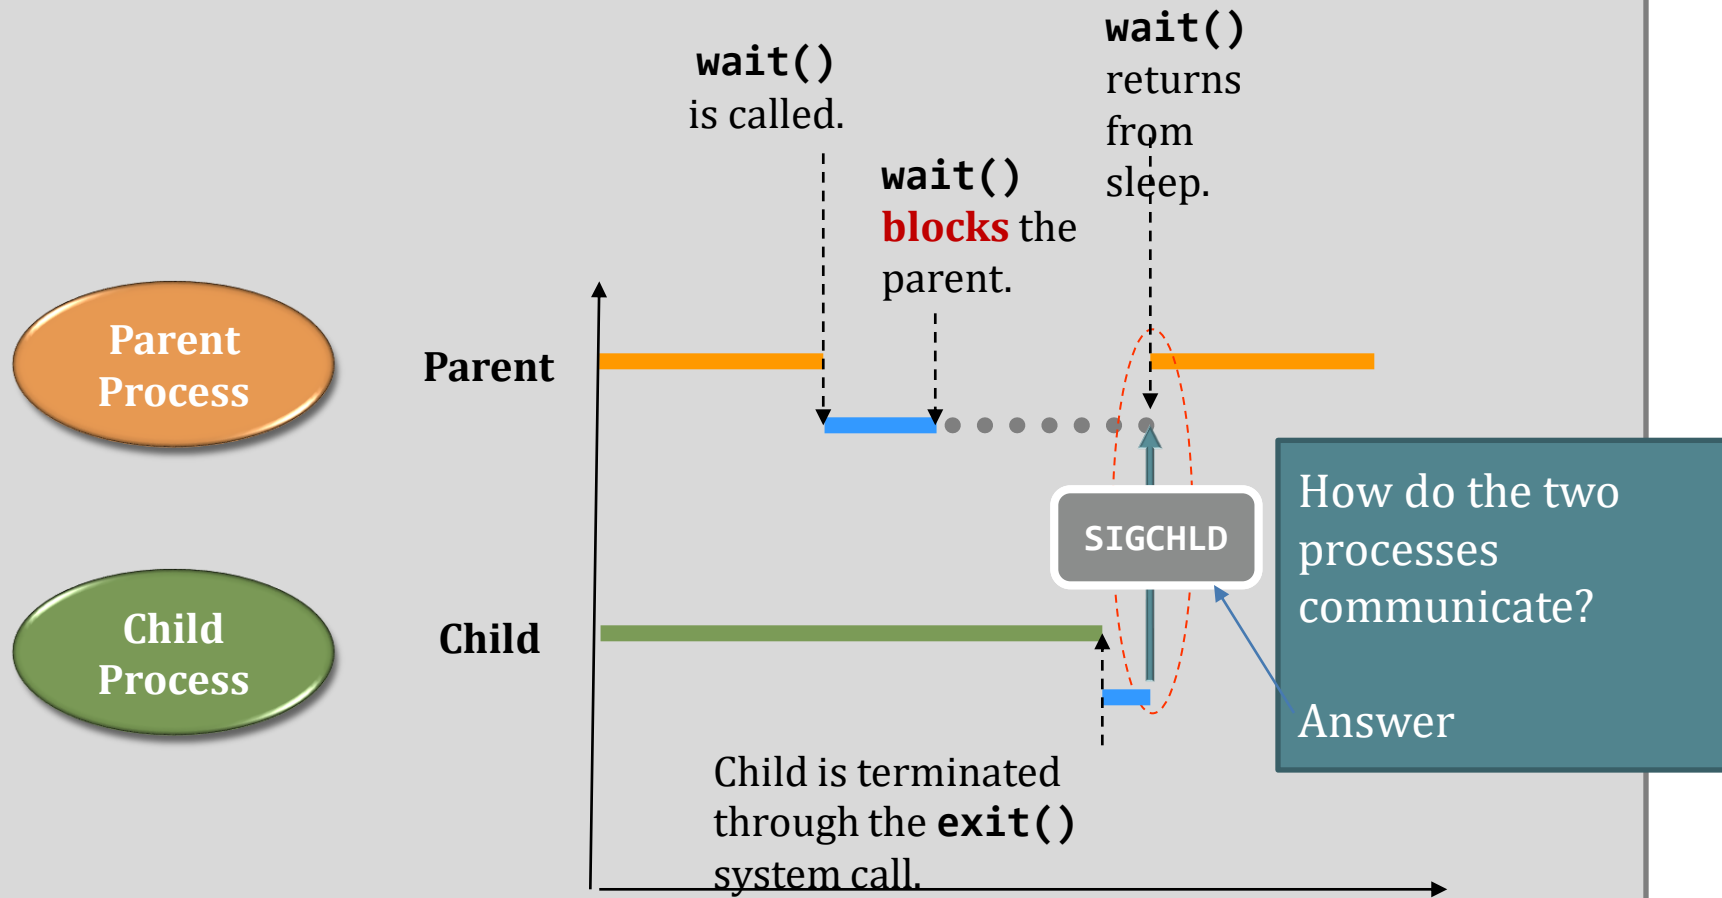
Step (1) Clean up most of the allocated kernel-space memory (e.g., process's running time info).

Step (2) Clean up the exit process's user-space memory.

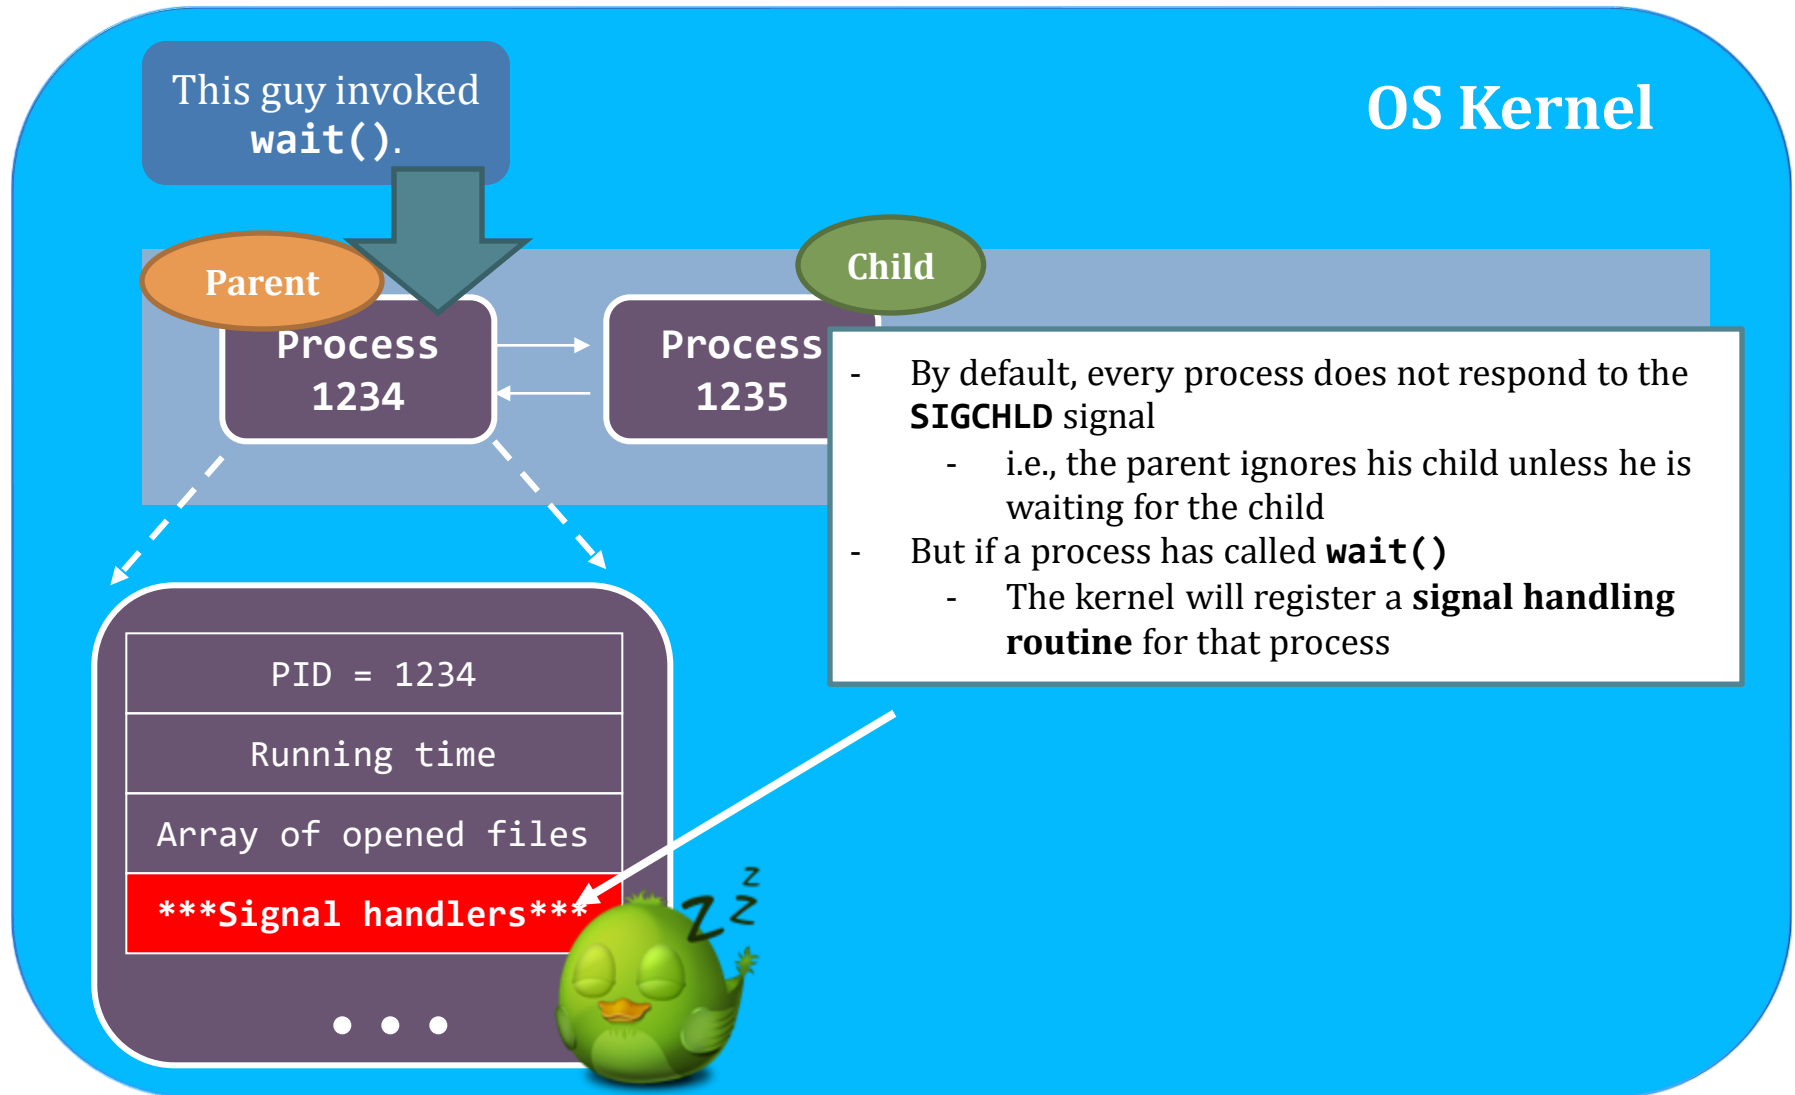
Step (3) Notify the parent with SIGCHLD.



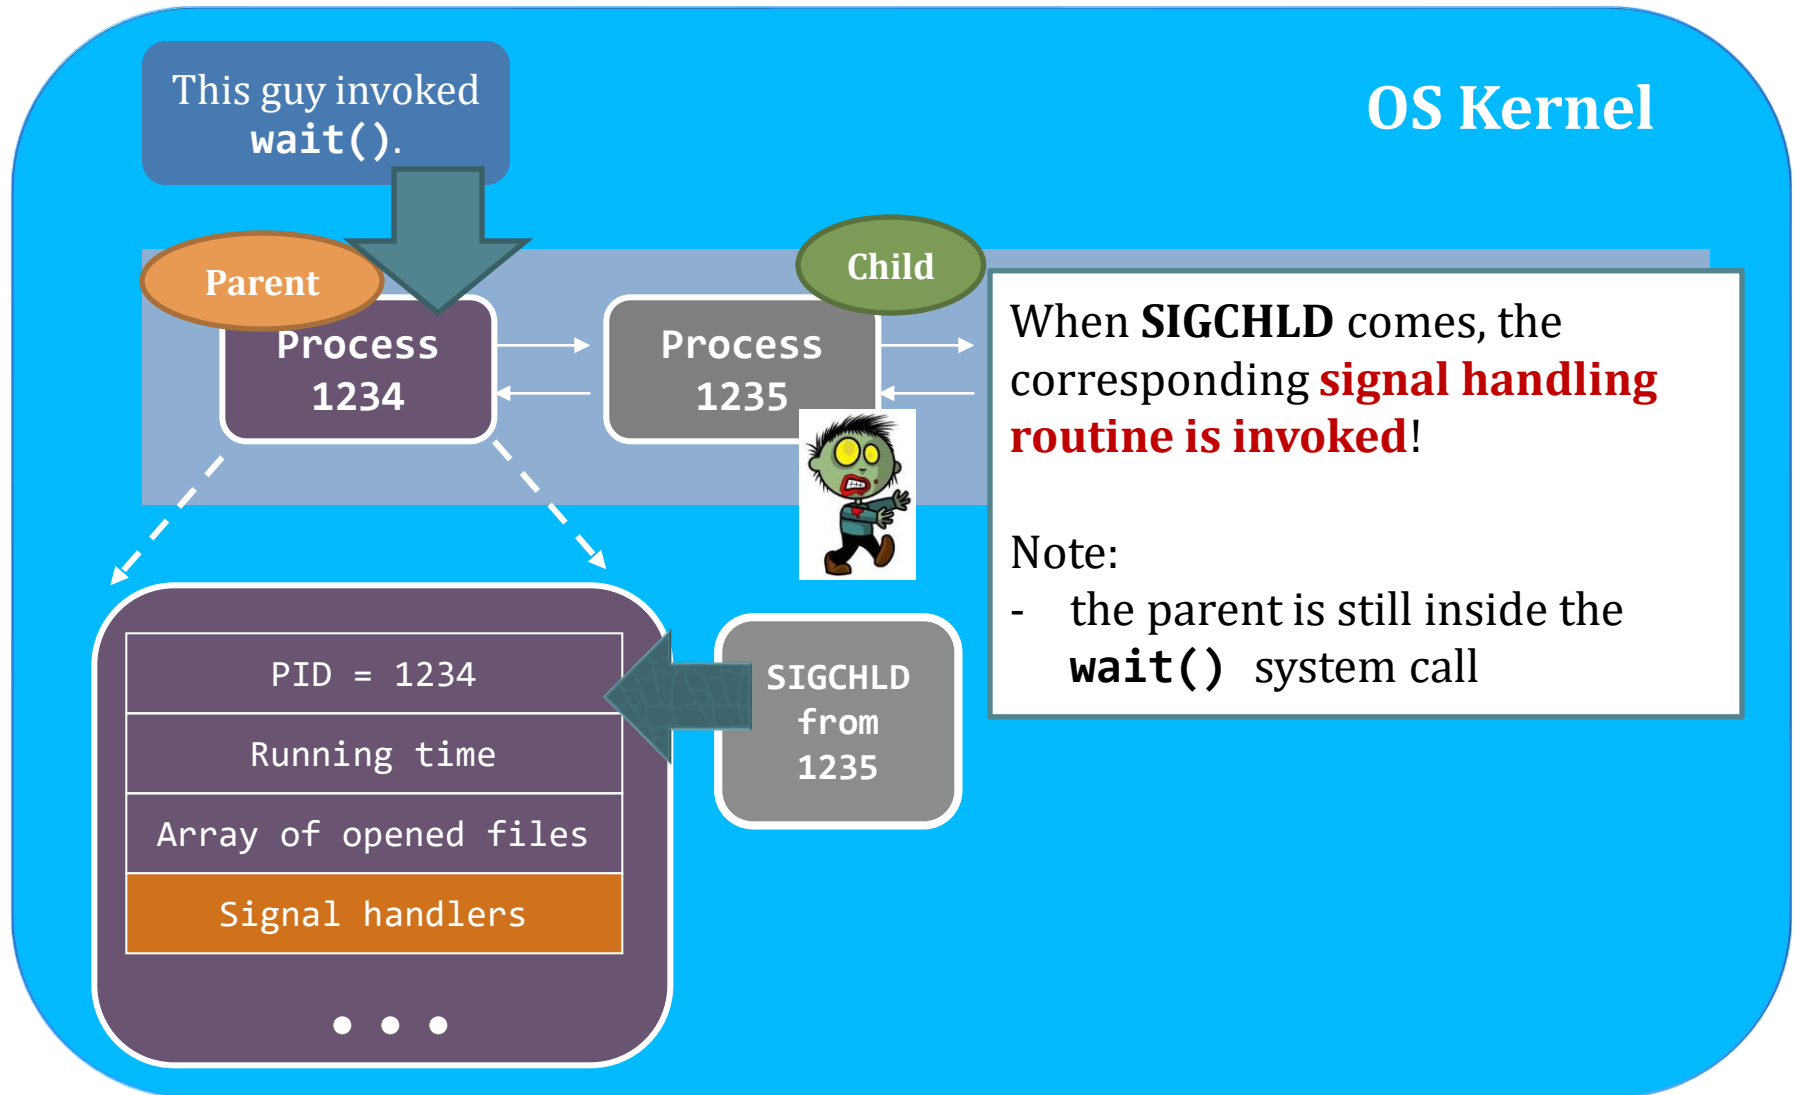
wait() and exit()



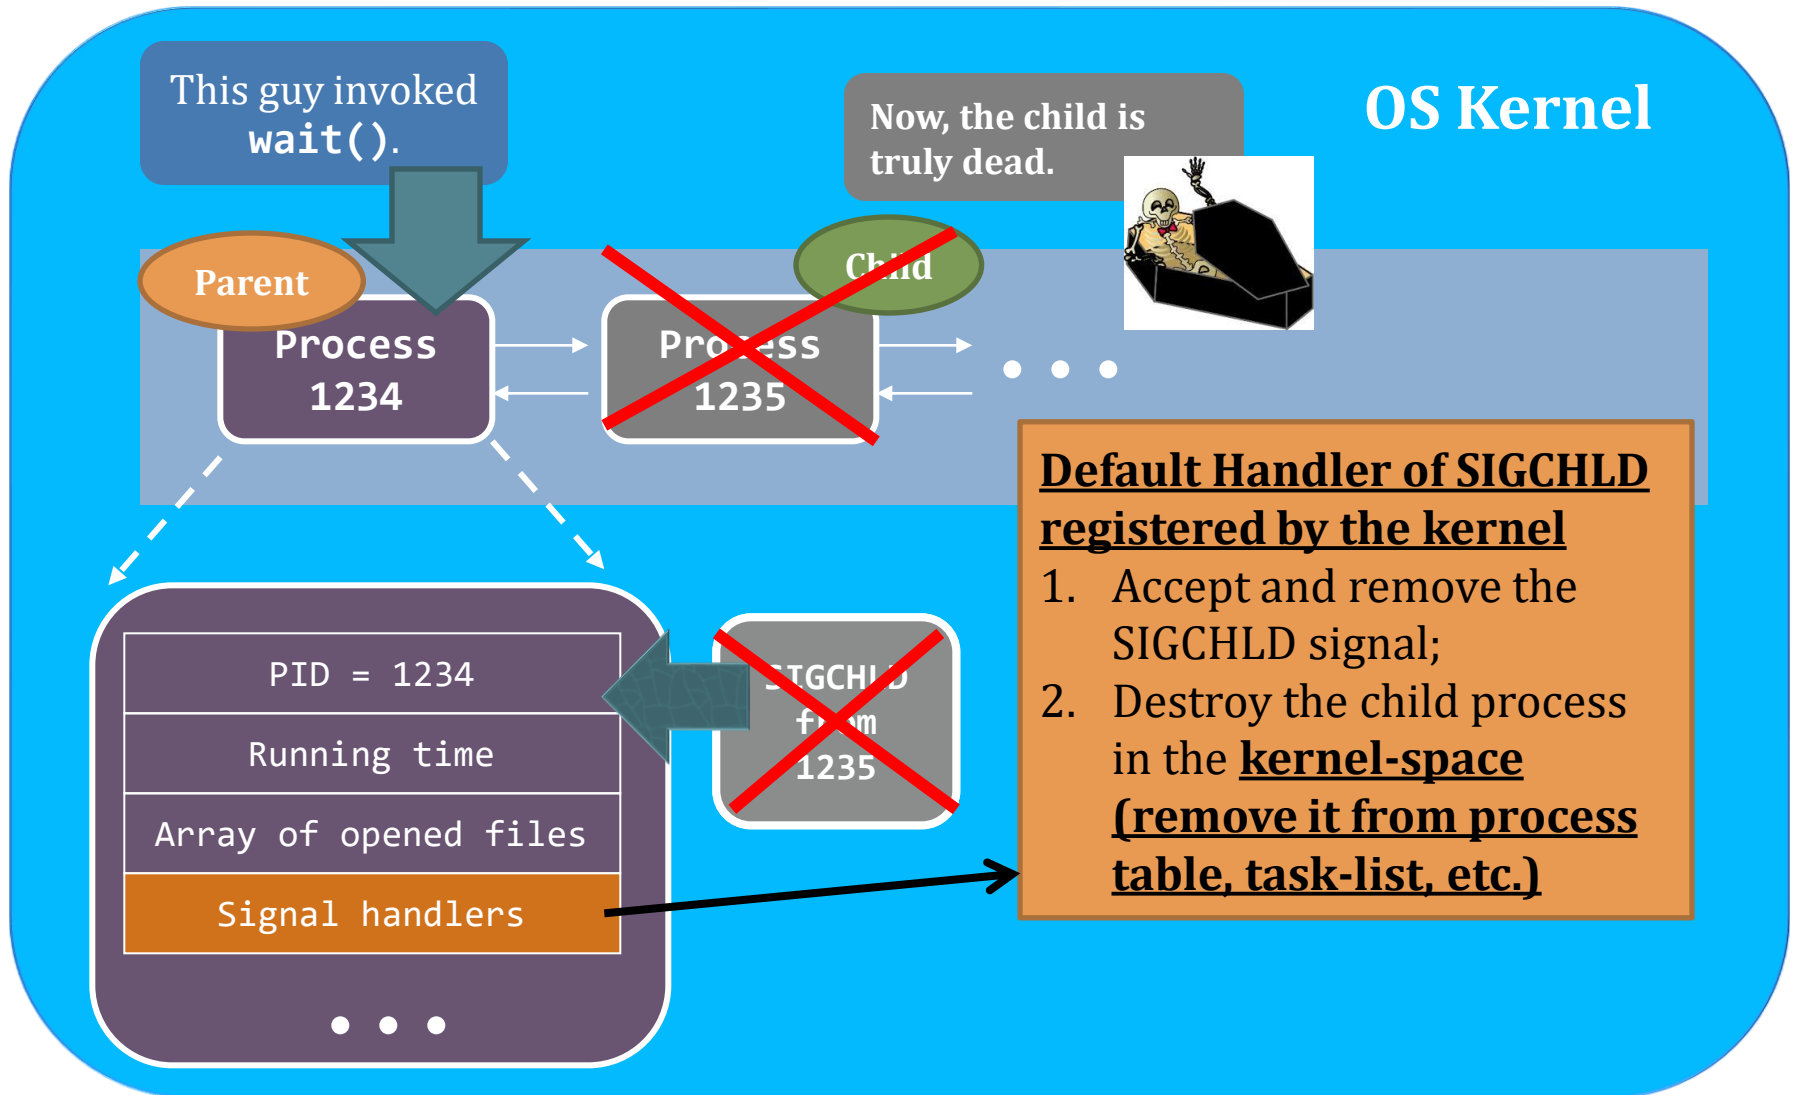
`wait()` kernel view's – registering signal handling routine



wait() kernel's view



wait() kernel's view



wait() kernel's view

OS Kernel

Ready to return
from `wait()`.

Parent

Process
1234

The kernel

- deregisters the **signal handling routine** for the parent
- returns the PID of the terminated child as the return value of `wait()`

The parent is ignoring **SIGCHLD** again.

PID = 1234

Running time

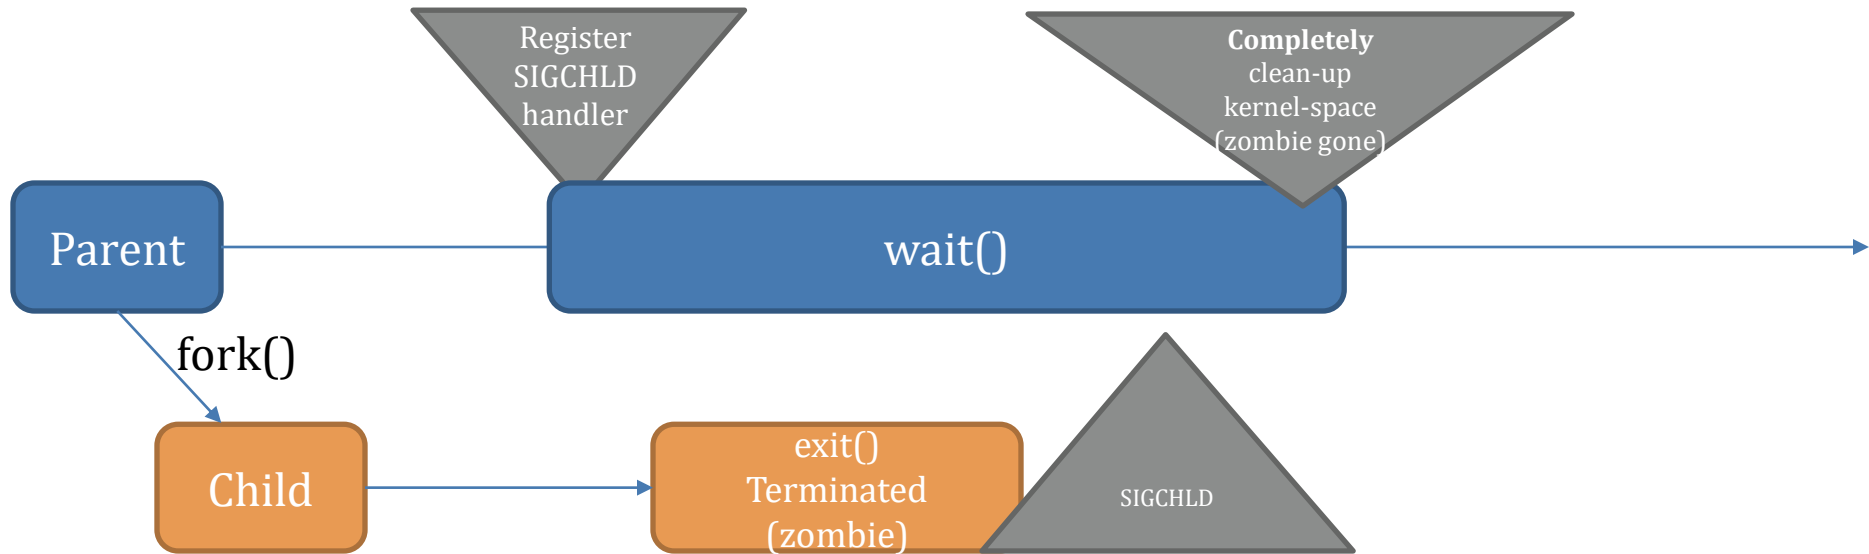
Array of opened files

~~Signal handlers~~

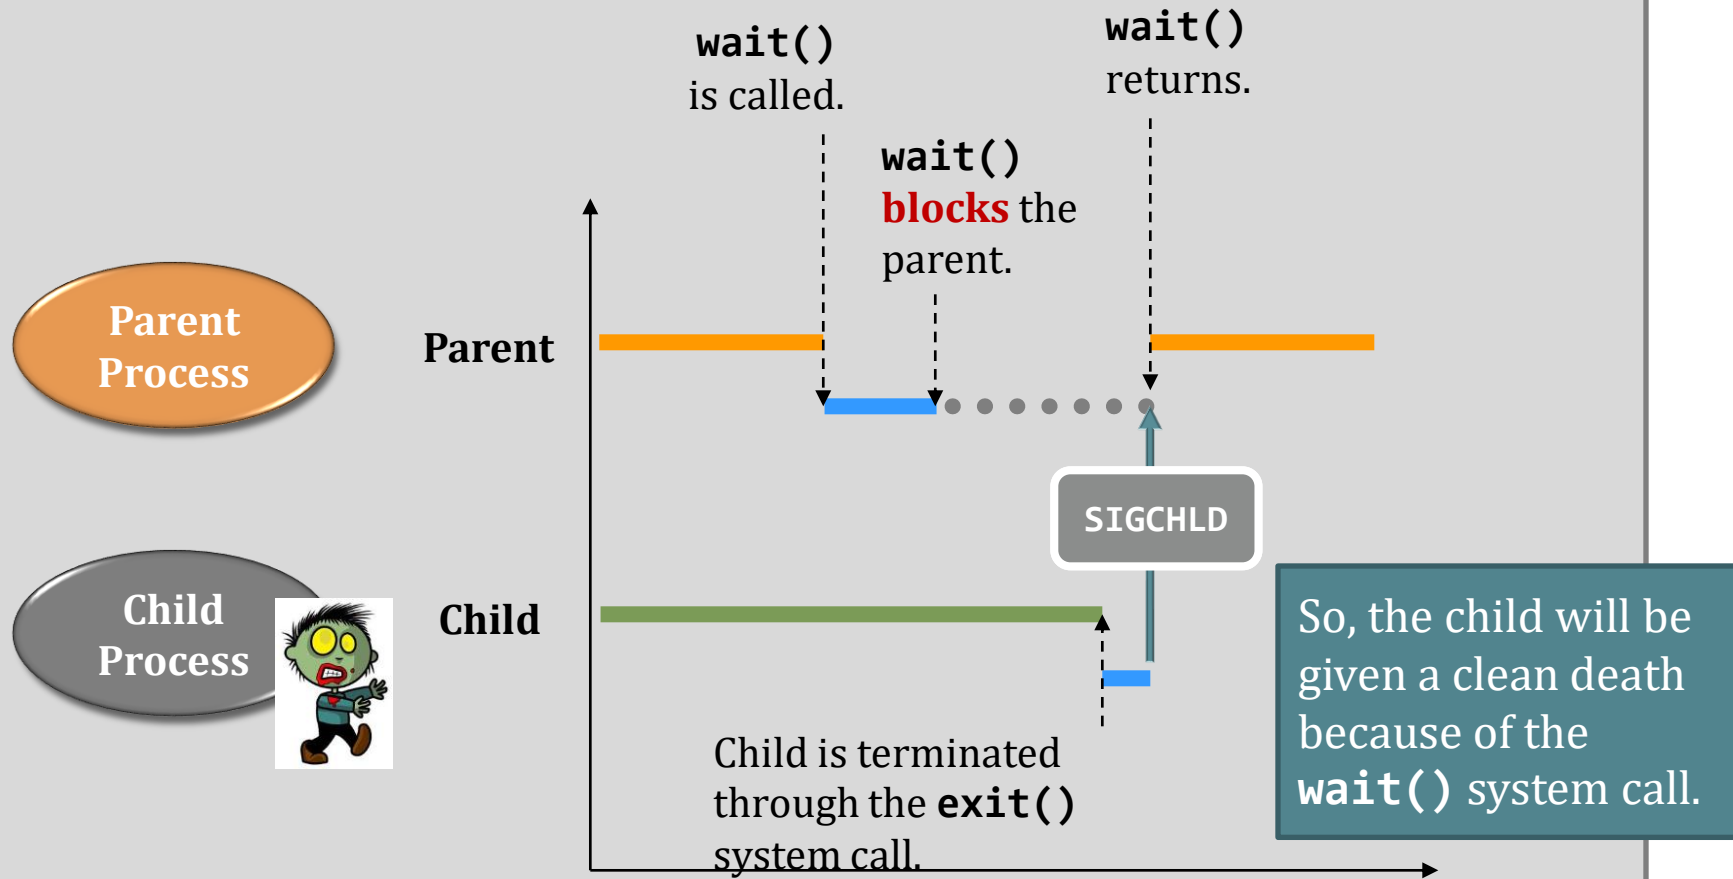
Return value = 1235



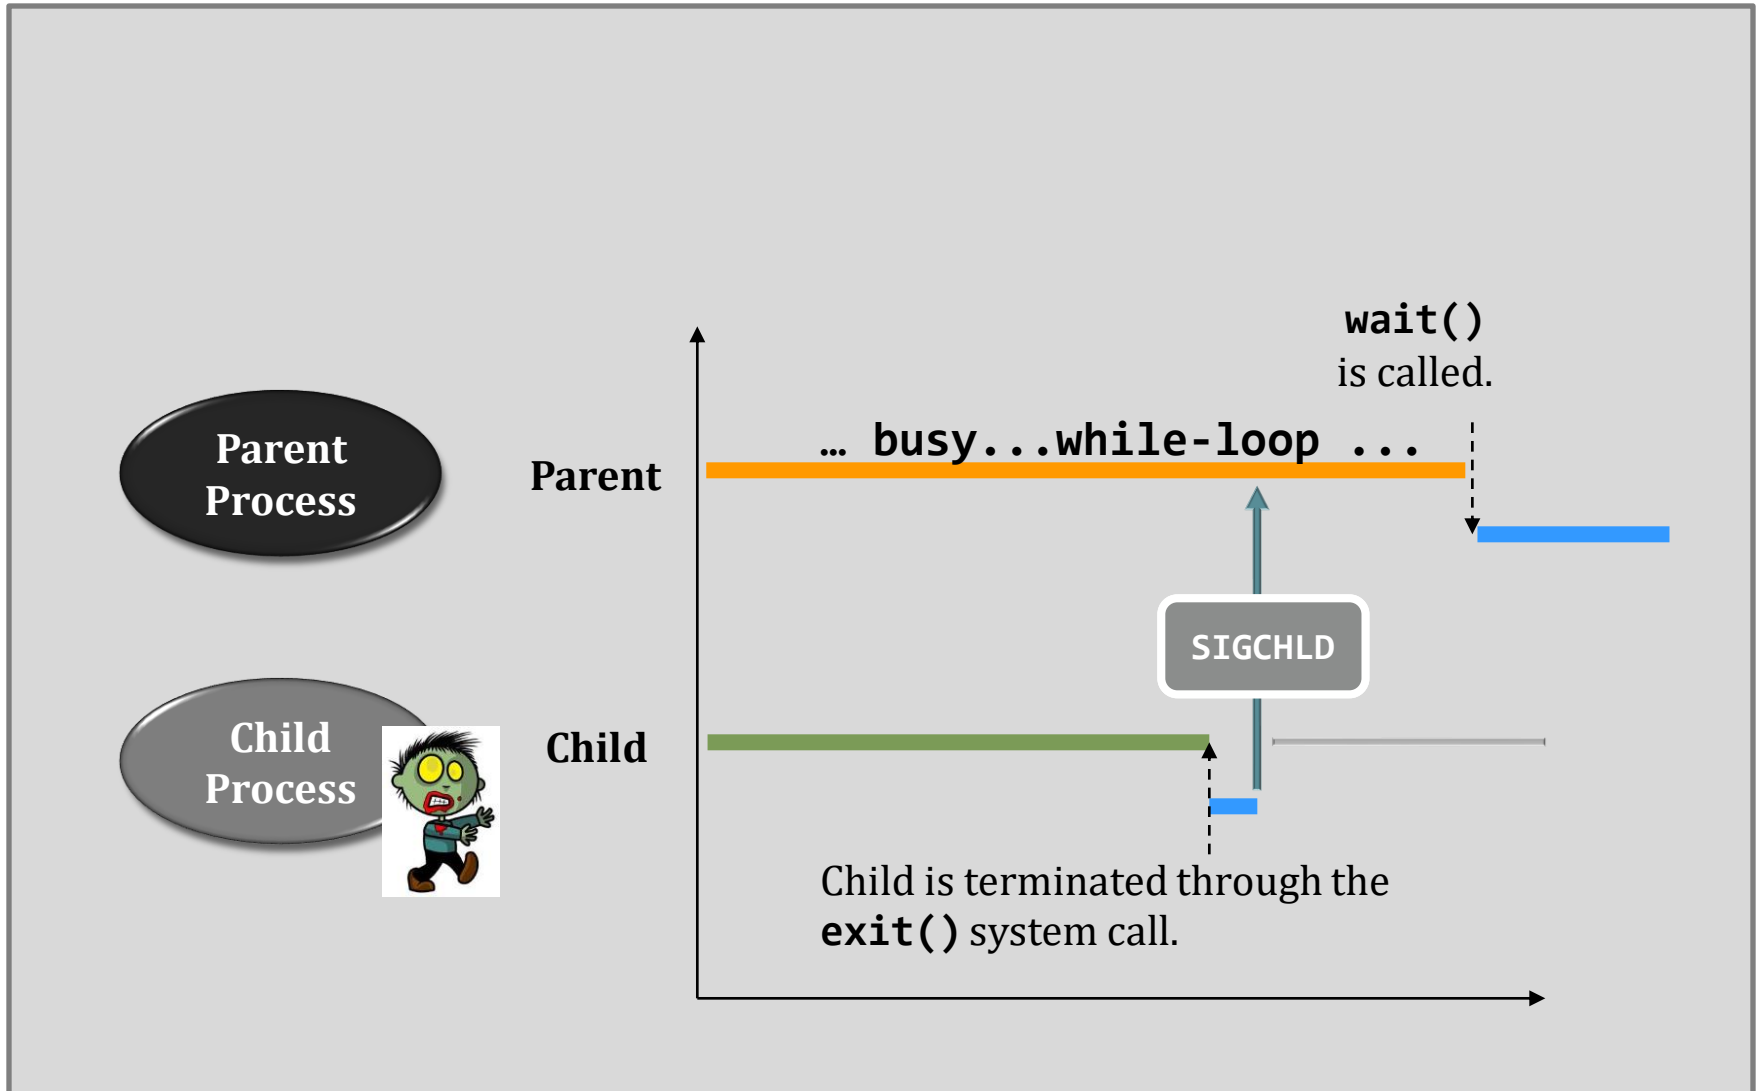
Overall – normal case



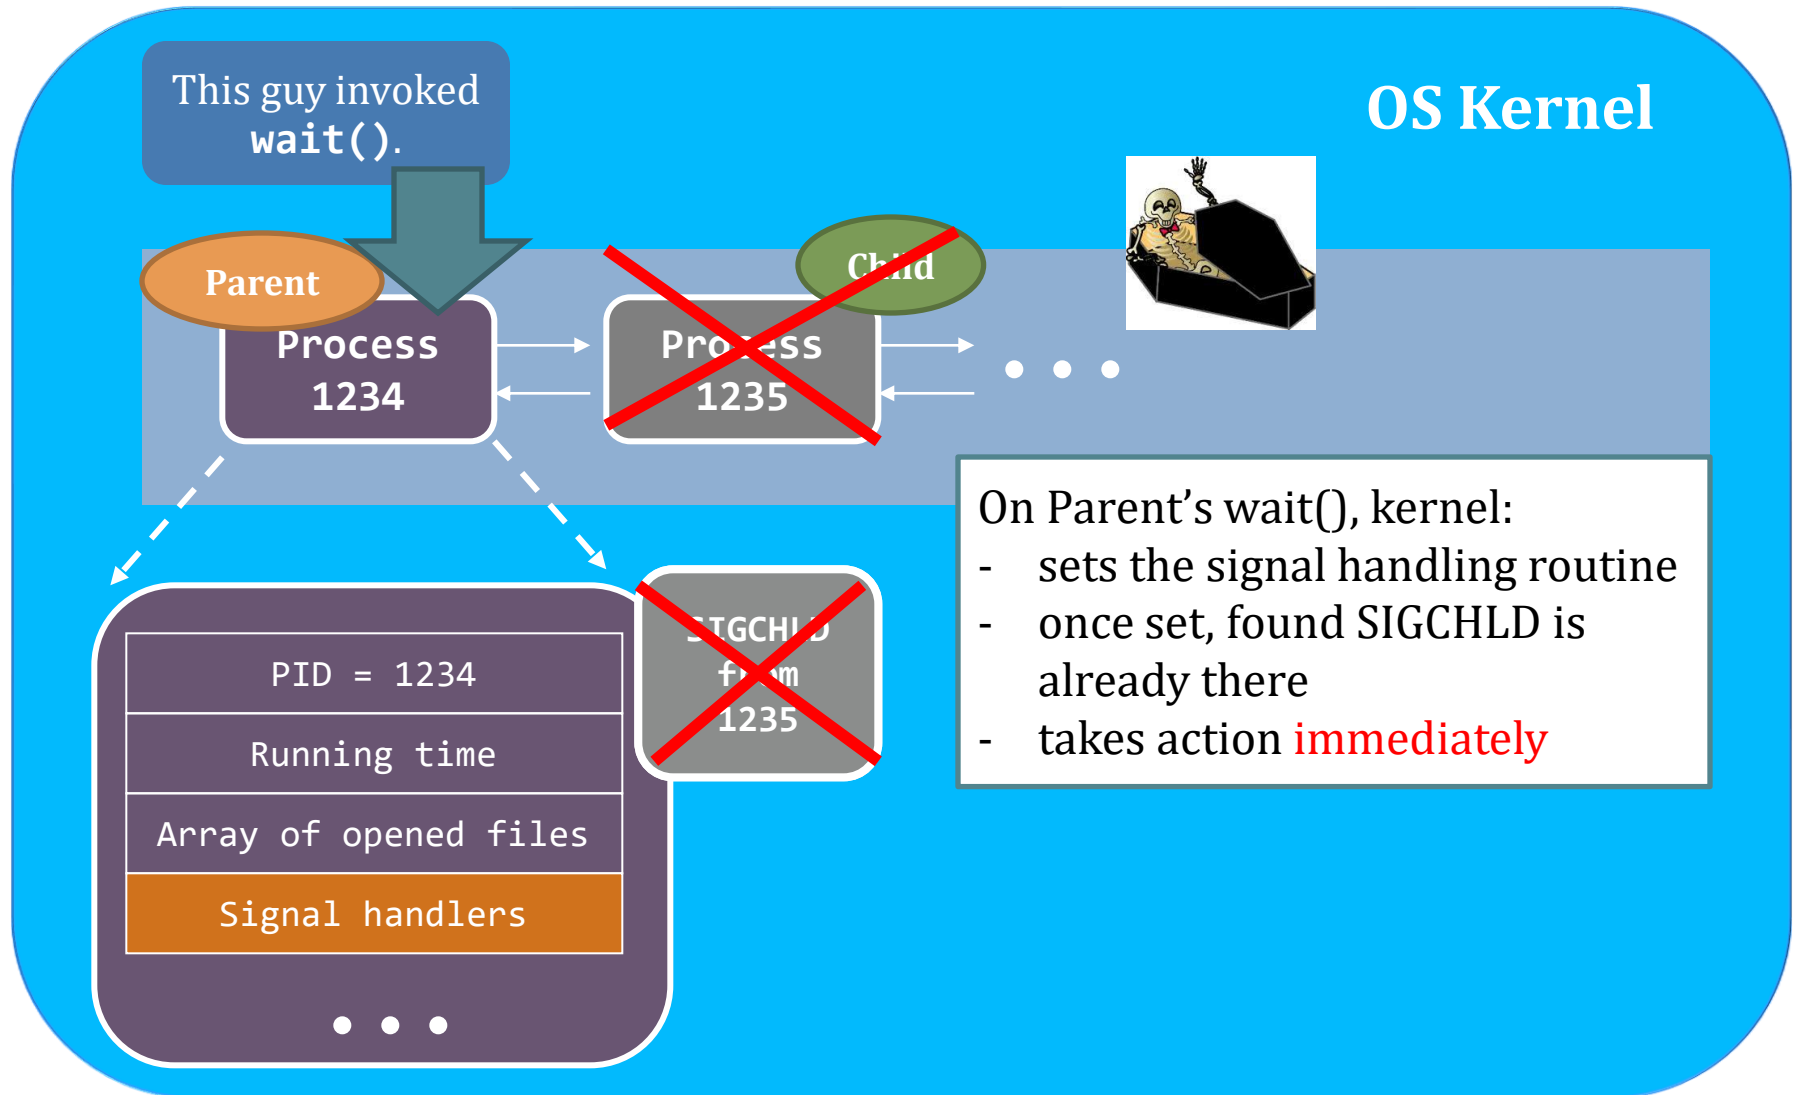
Normal Case



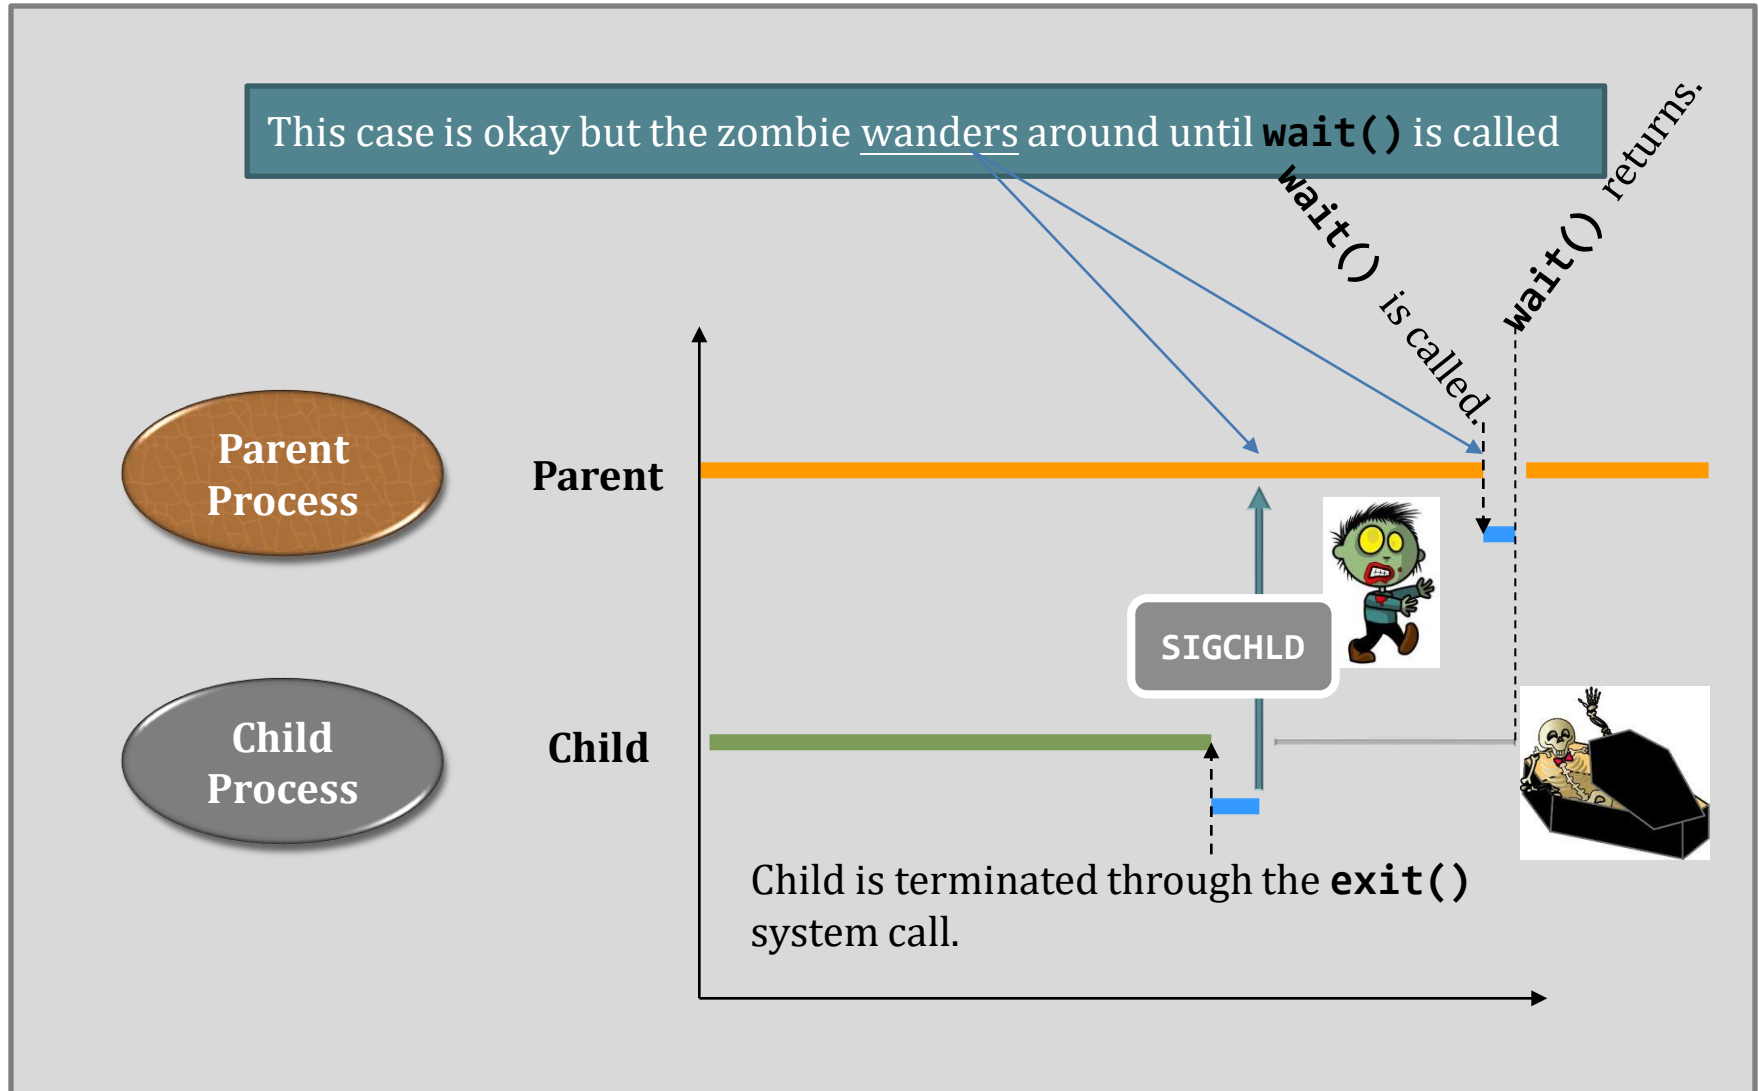
Parent's wait() after Child's exit()



Parent's Wait() after Child's exit()



Parent's Wait() after Child's exit()



wait() and **exit()** – short summary

- ◆ **exit()** system call turns a process into a zombie when...
 - ◆ The process calls **exit()**.
 - ◆ The process returns from **main()**.
 - ◆ The process terminates abnormally.
 - ◆ The kernel knows that the process is terminated abnormally. Hence, the kernel invokes **exit()** for it.

`wait()` and `exit()` – short summary

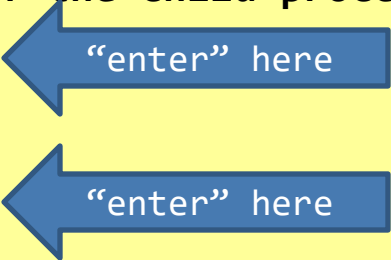
- ◆ `wait()` & `waitpid()` are to reap zombie child processes.
 - ◆ It is a must that you should never leave any zombies in the system.
 - ◆ `wait()` & `waitpid()` pause the caller until
 - ◆ A child terminates/stops, OR
 - ◆ The caller receives a signal (i.e., the signal interrupted the `wait()`)
- ◆ Linux will label zombie processes as “<**defunct**>”.
 - ◆ To look for them:

```
$ ps aux | grep defunct
..... 3150 ... [ls] <defunct>
$ _
```

PID of the
process

wait() and exit() – short summary

```
1 int main(void)
2 {
3     int pid;
4     if( (pid = fork()) !=0 ) {
5         printf("Look at the status of the child process %d\n", pid);
6         while( getchar() != '\n' );
7         wait(NULL);
8         printf("Look again!\n");
9         while( getchar() != '\n' );
10    }
11    return 0;
12 }
```



This program requires you to type “enter” twice before the process terminates.

You are expected to see **the status of the child process changes (ps aux [PID])** between the 1st and the 2nd “enter”.

Working of system calls

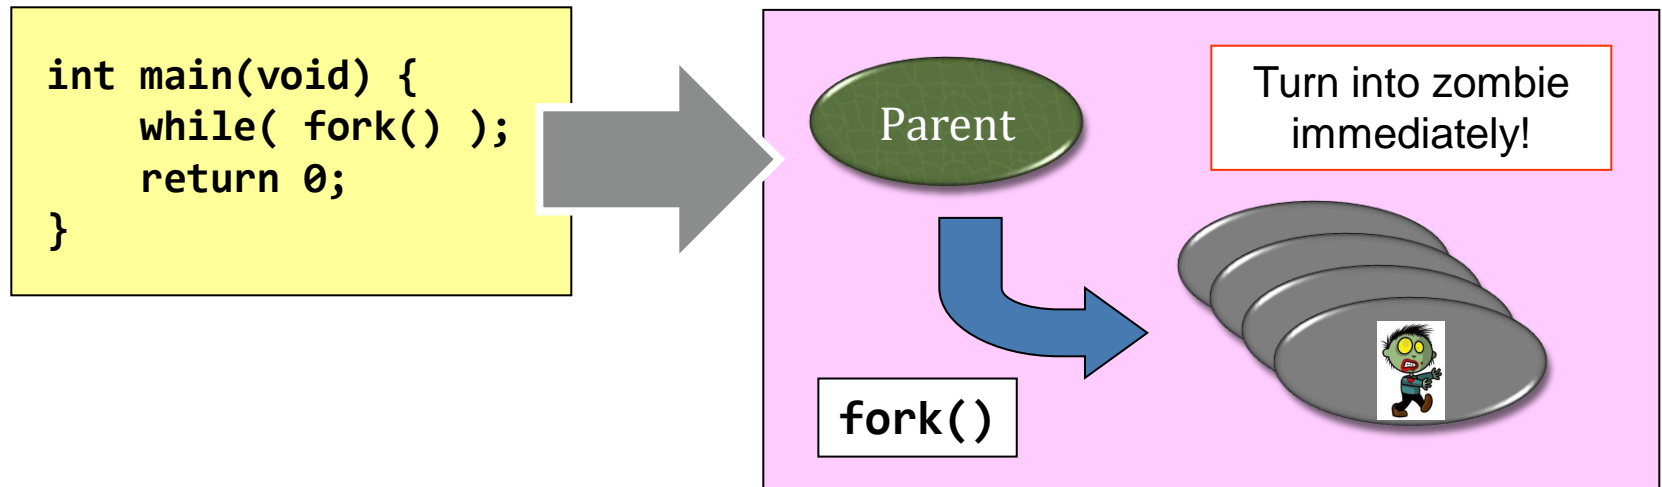
- `fork()`;
- `exec*()`;
- `wait()` + `exit()`;
- **importance/fun in knowing the above things?**

Calling `wait()` is important.

- ◆ It is not only about process execution / suspension...
- ◆ It is about **system resource management**.
 - ◆ A zombie takes up a PID;
 - ◆ The total number of PIDs are limited;
 - ◆ Read the limit: `cat /proc/sys/kernel/pid_max`
 - ◆ It is 32,768.
 - ◆ What will happen if we don't clean up the zombies?

The fork bomb


- ❖ Deliberately missing wait()
- ❖ Do not try this on department's machines...



An infinite, zombie factory!

When `wait()` is absent...

```
int main(void) {  
    while( fork() );  
    return 0;  
}
```



\$./interesting

—

Terminal A

\$ **ls**

No process left.

\$ **poweroff**

No process left.

\$ **=__=**

No process left.

\$ —

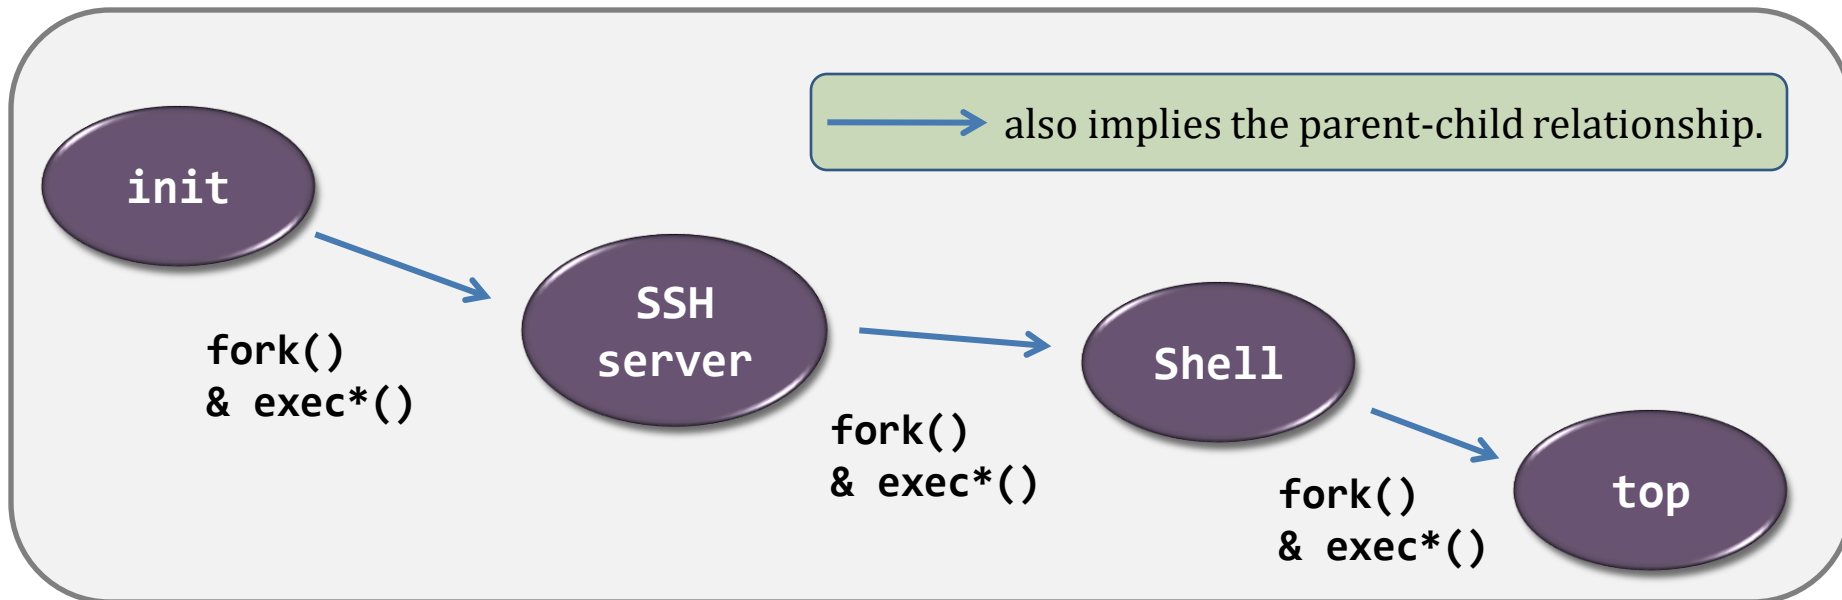
Terminal B

The first process

- ◆ We now focus on the process-related events.
 - ◆ The kernel, while it is booting up, creates the first process – **init**.
- ◆ The “**init**” process:
 - ◆ has **PID = 1**, and
 - ◆ is running the program code “**/sbin/init**”.
- ◆ Its first task is to **create more processes...**
 - ◆ Using **fork()** and **exec*()**.

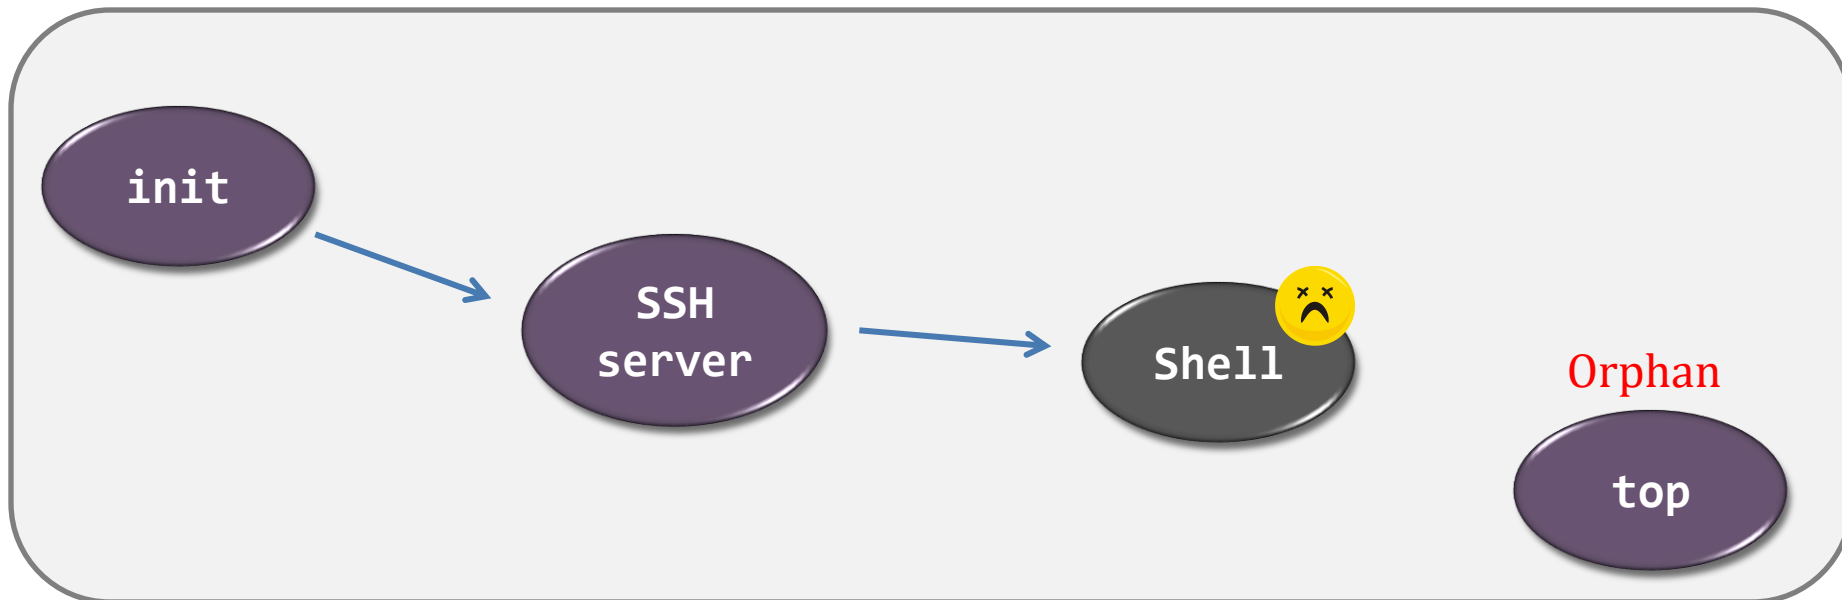
Process blossoming

- ◆ You can view the tree with the command:
 - ◆ “**pstree**”; or
 - ◆ “**pstree -A**” for ASCII-character-only display.



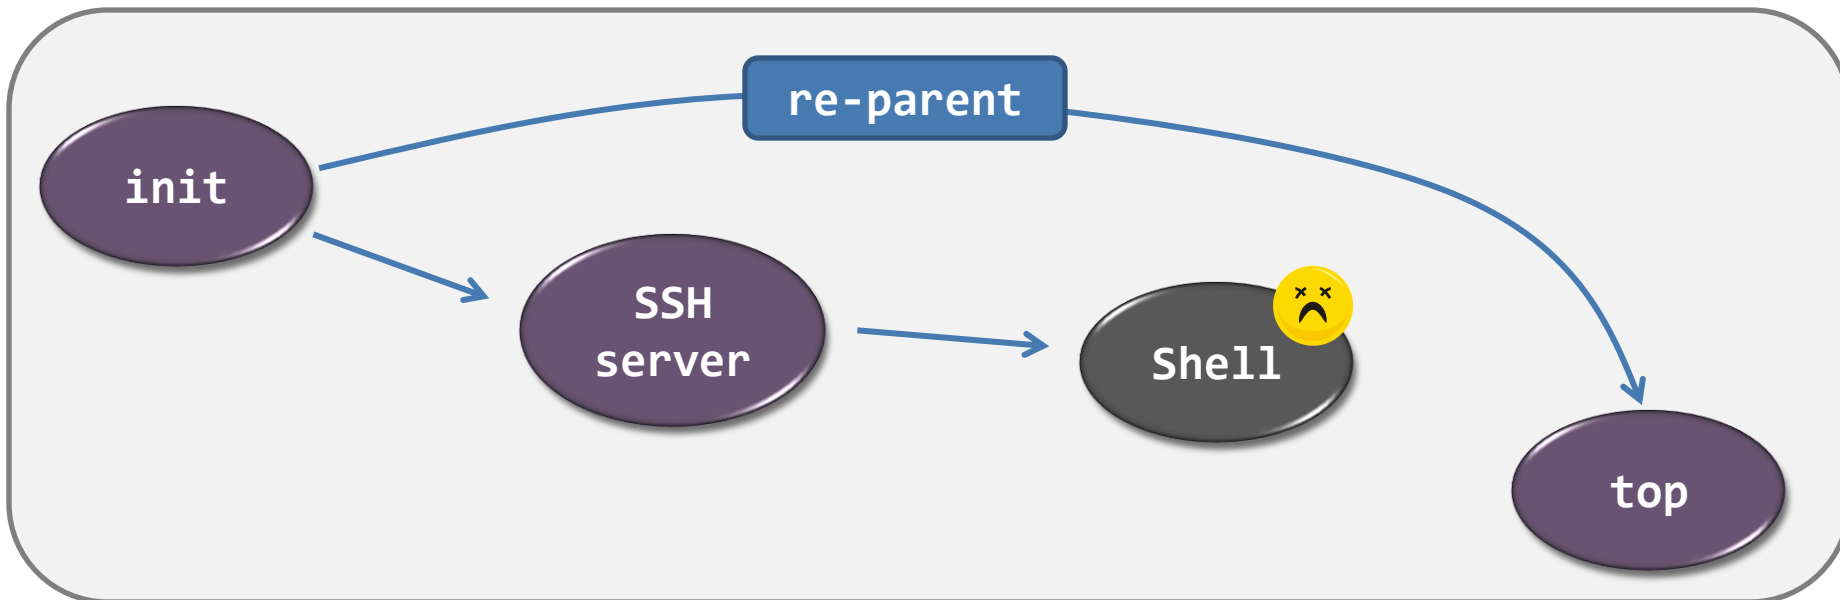
Process blossoming...with orphans?

- ◆ However, termination can happen, at any time and in any place...
 - ◆ This is no good because an orphan turns the hierarchy from a **tree** into a **forest**!
 - ◆ Plus, no one would know the termination of the orphan.



Process blossoming...with re-parent!

- ◆ In Linux
 - ◆ The “**init**” process will become the step-mother of all orphans
 - ◆ It's called **re-parenting**
- ◆ In Windows
 - ◆ It maintains a *forest-like process hierarchy*.....



*New Linux kernels may choose someone else (e.g., the grandparent, user-level init)

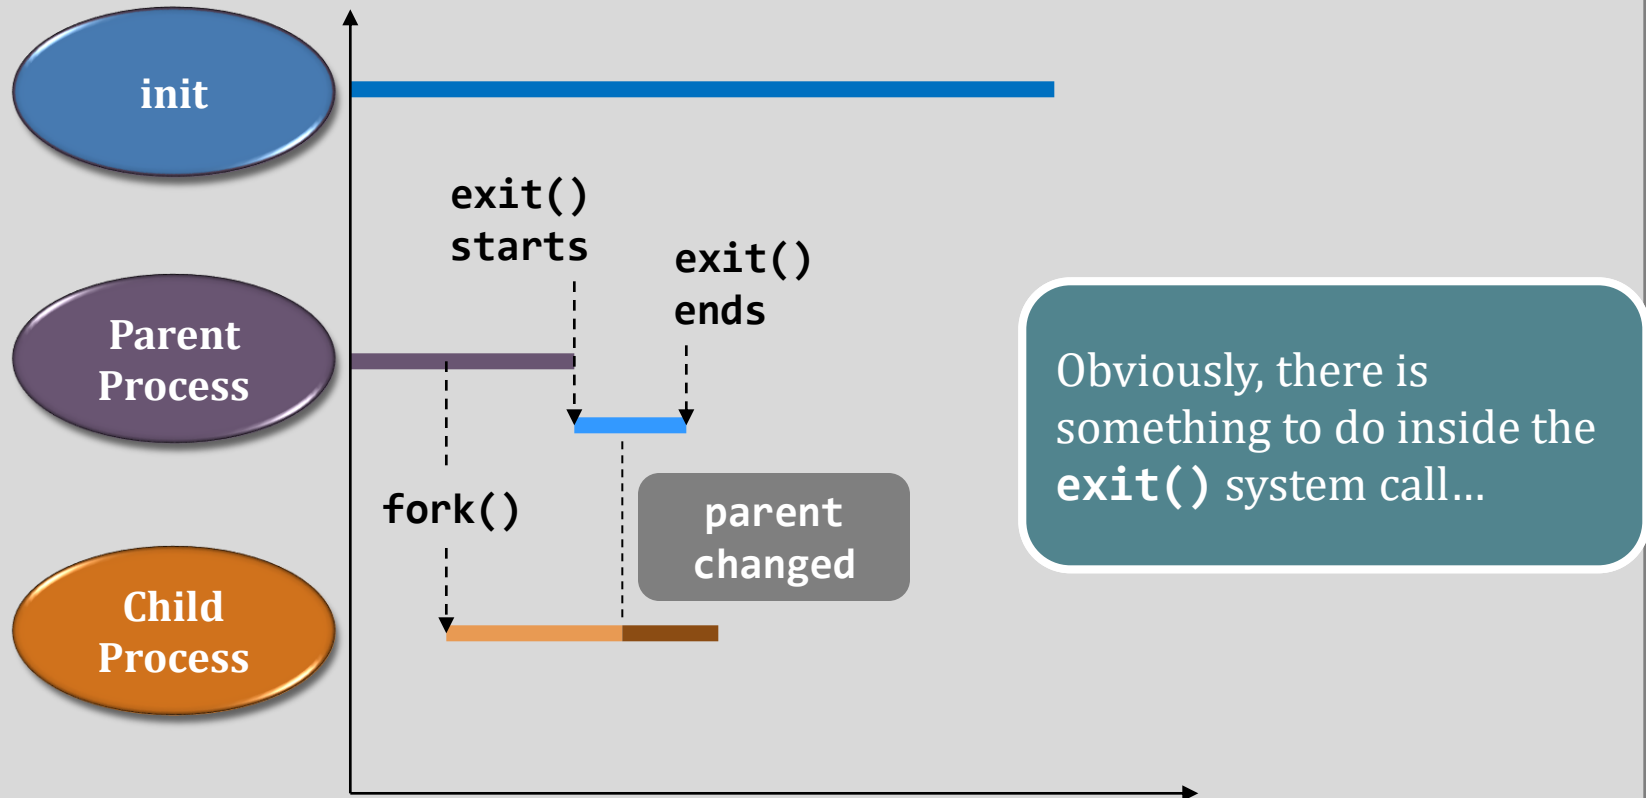
Re-parenting example

```
1  int main(void) {
2      int i;
3      if(fork() == 0) {
4          for(i = 0; i < 5; i++) {
5              printf("(%d) parent's PID = %d\n",
6                  getpid(), getppid() );
7              sleep(1);
8          }
9      }
10     else
11         sleep(1);
12     printf("(%d) bye.\n", getpid());
13 }
```

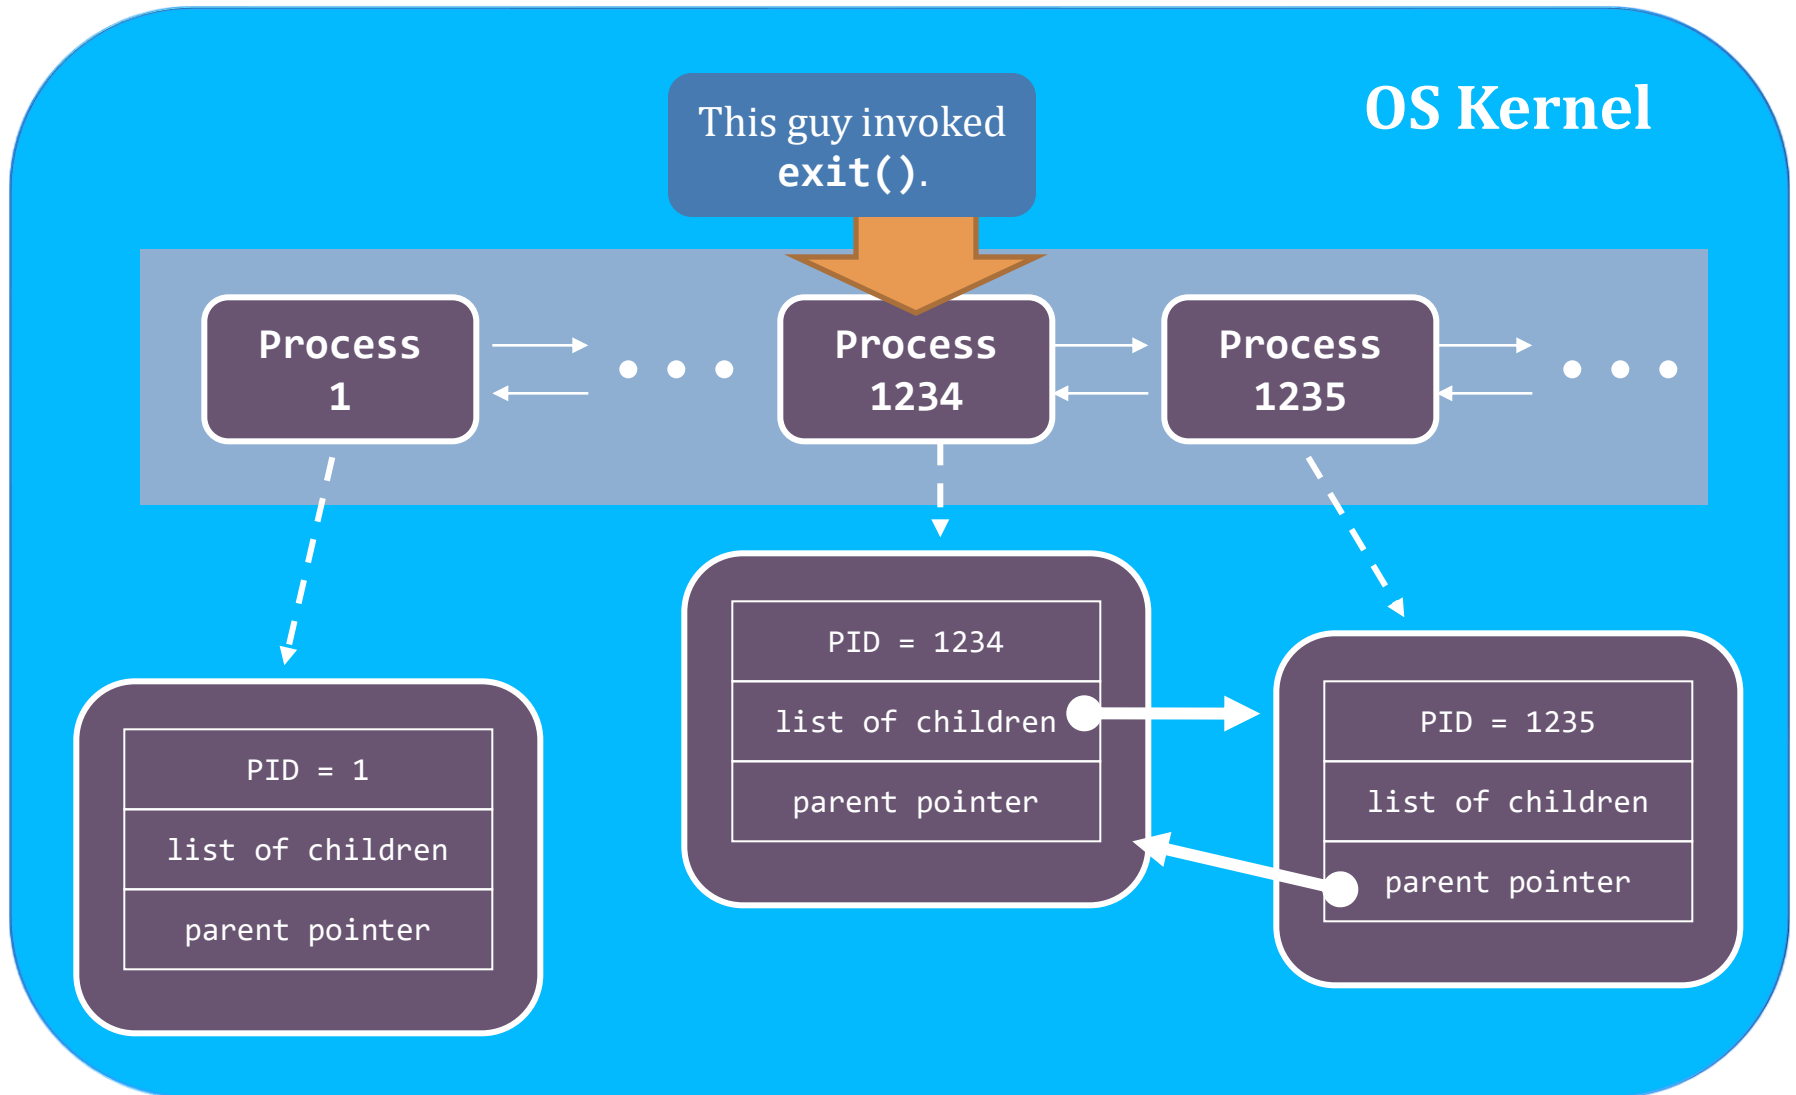
`getppid()` is the system call that returns the parent's PID of the calling process.

```
$ ./reparent
(1235) parent's PID = 1234
(1235) parent's PID = 1234
(1234) bye.
$ (1235) parent's PID = 1
(1235) parent's PID = 1
(1235) parent's PID = 1
(1235) bye.
$ _
```

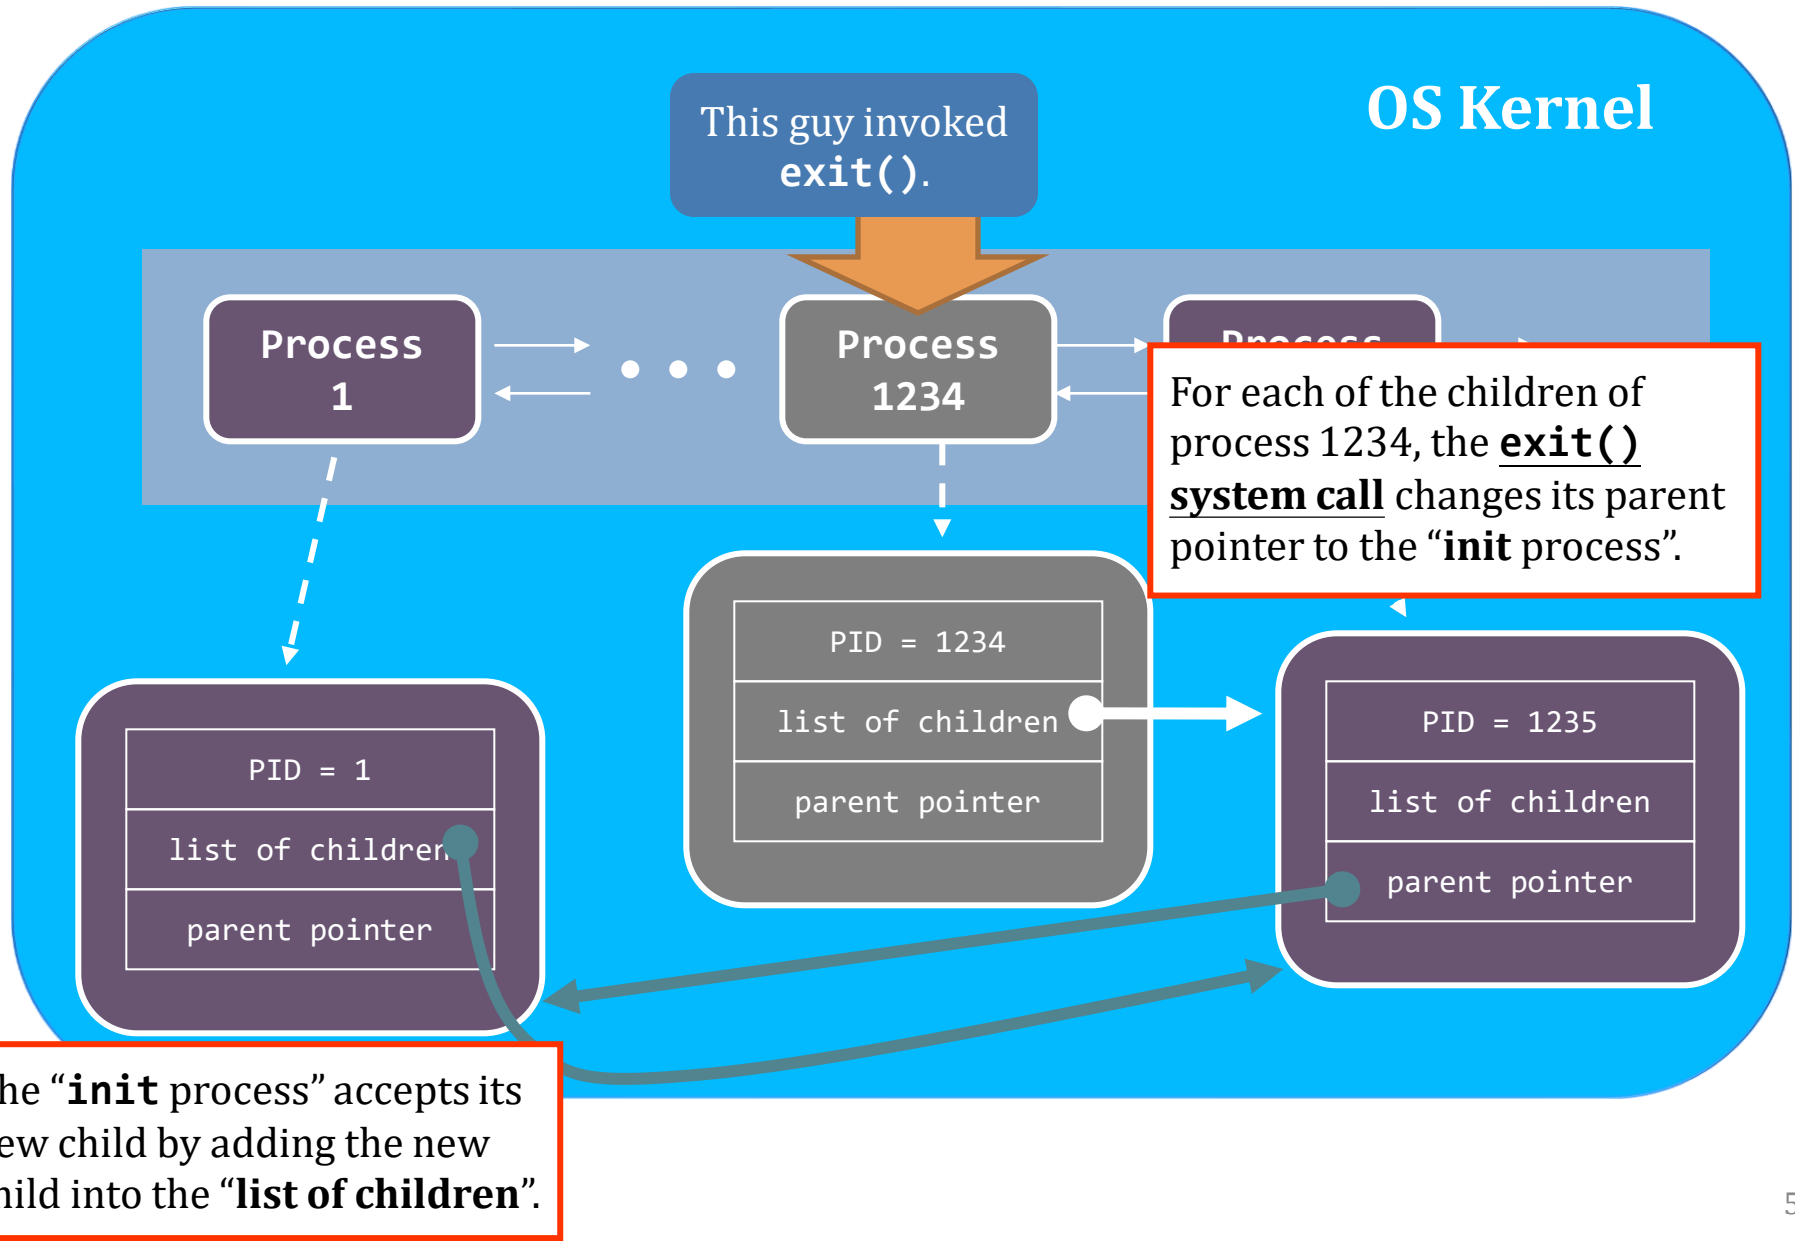
What had happened during re-parenting?



What had happened during re-parenting?



What had happened during re-parenting?



Background jobs

- ◆ The re-parenting operation enables something called **background jobs** in Linux
 - ◆ It allows a process runs **without a parent terminal/shell**

[Back to home](#)

```
$ ./infinite_loop &  
$ exit  
  
[ The shell is gone ]
```

```
$ ps -C infinite_loop  
PID  TTY  
1234  ... ./infinite_loop  
$ _
```


Process lifecycle

The birth of a process.

Except the first process “**init**”, every process is created using **fork()**.

Just
fork()-ed

Zombie
(or terminated)

Ready

Running

Interruptible

Un-interruptible

Blocked/Waiting

Process lifecycle - Ready

Just
fork()-ed

Ready

Interruptible

Blocked

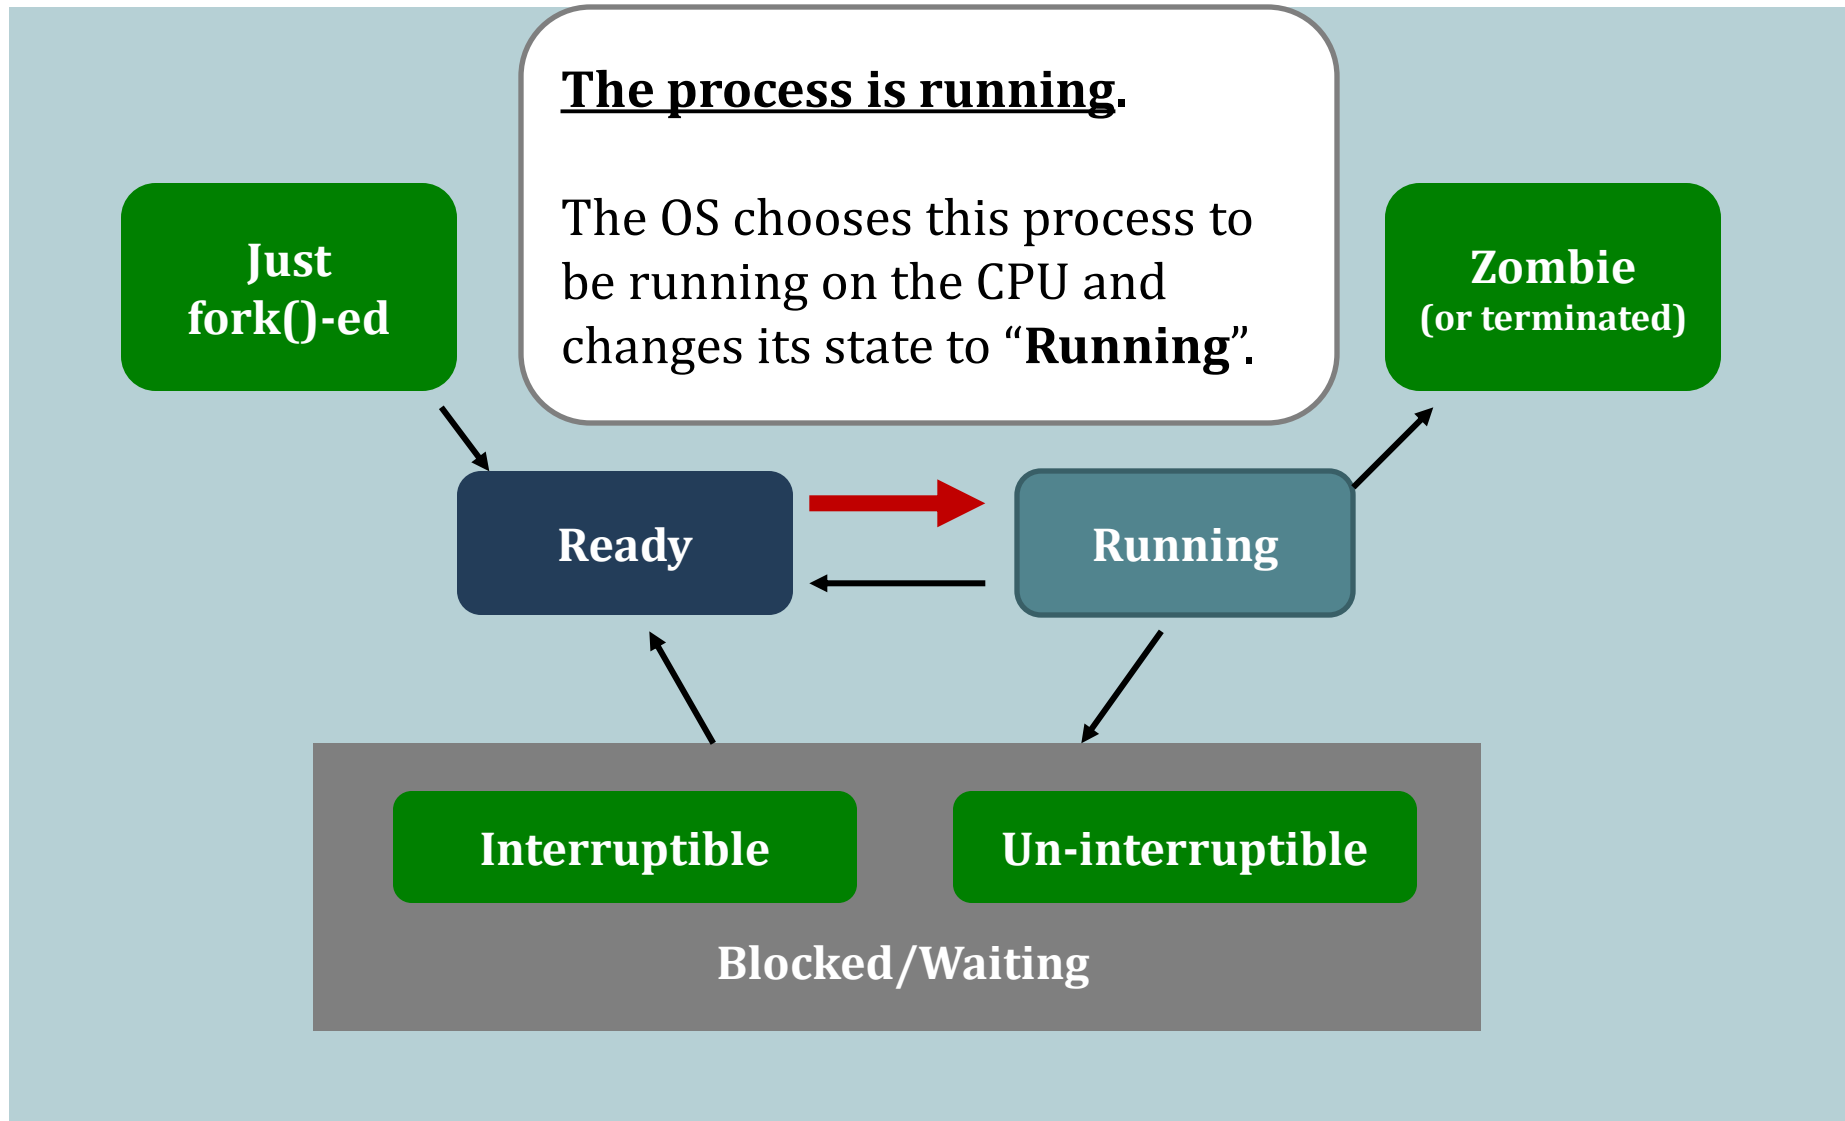
The process is ready.

It means it is **ready to run but is not running.**

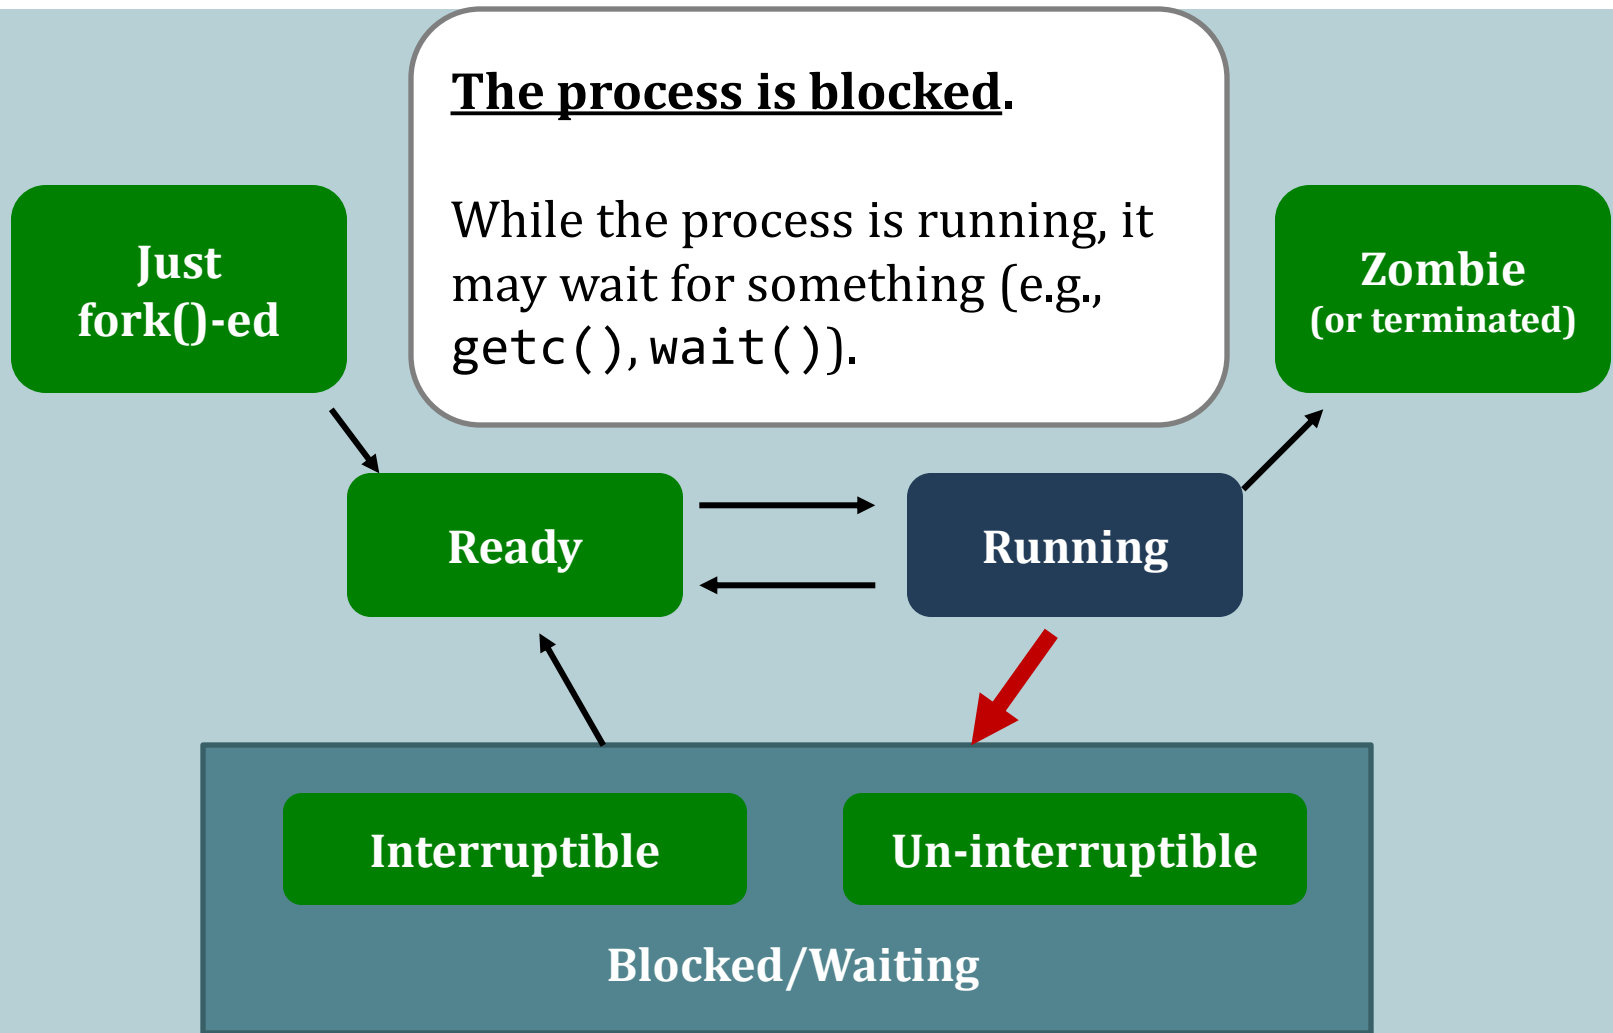
A process may become “ready” (*runnable*) after...

- it is just created by **fork()**;
- it has been running on the CPU for some time and the OS chooses another process to run (scheduled context switch)
- returning from blocked states.

Process lifecycle - Running



Process lifecycle - Blocking



Process lifecycle – Interruptible wait

Example. Reading a file.

Sometimes, the process has to wait for the response from the device and, therefore, it is **blocked**

- this blocking state is **interruptible**
 - E.g., “**Ctrl + C**” can get the process out of the waiting state (but goes to termination state instead).



Interruptible

Un-interruptible

Blocked/Waiting

Process lifecycle – Un-Interruptible wait

Sometimes, a process needs to wait for a resource until it really gets what it wants

- Doesn't want to be "Ctrl-C" interruptible
- **Un-interruptible** status
 - No way to signal it to wake up unless it returns itself
 - Check online! The only solution is ...

Who set this?

- E.g., syscall call (http://man7.org/linux/man-pages/man2/delete_module.2.html)

Why set this?

- Easier programming for lazy programmer (e.g., a driver program for a DVD drive)
- The programmer "thinks" the wait is very short and robust
 - This is one the top reasons that hang your machine / process today!
- ...



Interruptible

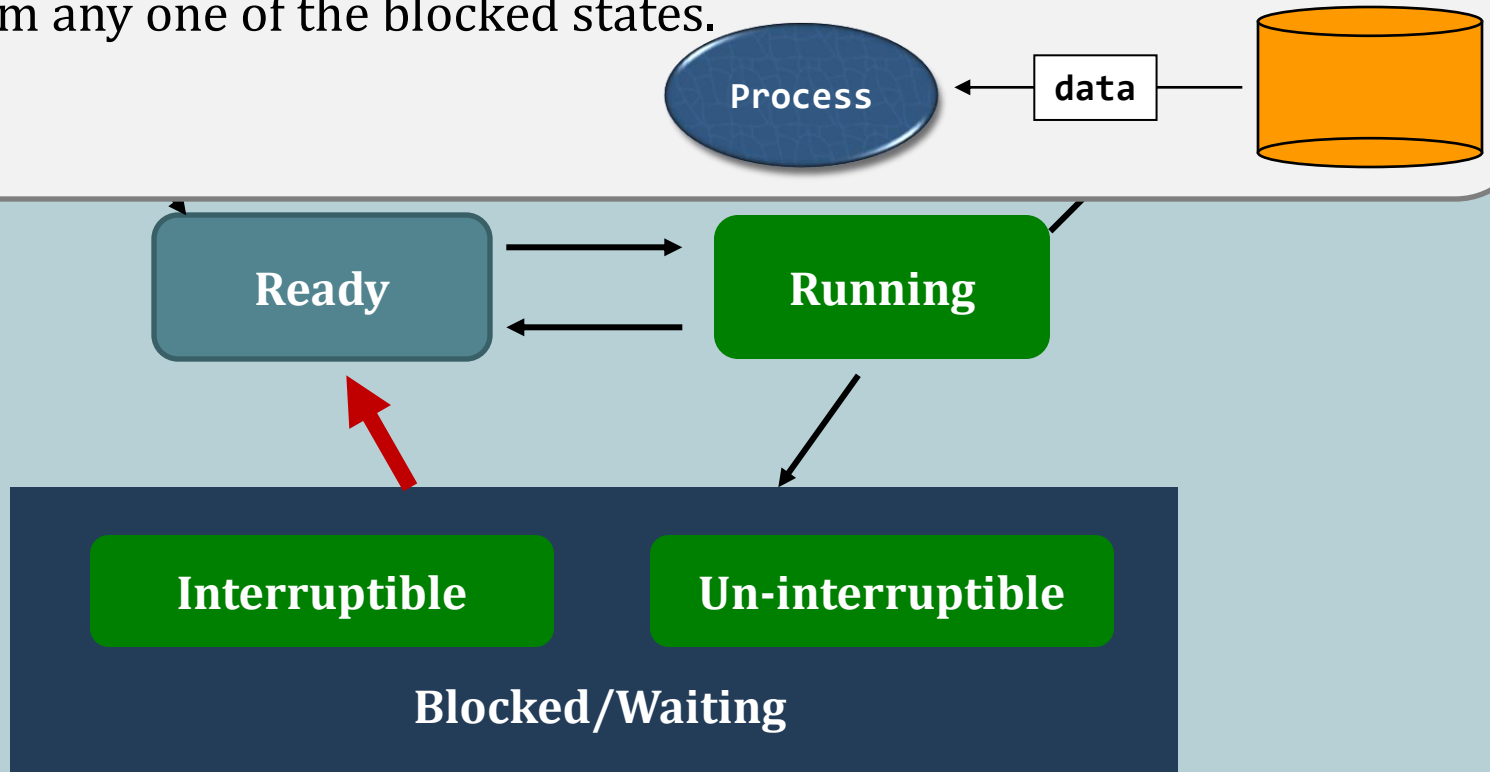
Un-interruptible

Blocked/Waiting

Process lifecycle

Return back to ready.

When response arrives, the status of the process changes back to **Ready**. from any one of the blocked states.



Process lifecycle

The process is going to die.

The process may

- choose to terminate itself; or
- force to be terminated.

Running

Zombie
(or terminated)

Interruptible

Un-interruptible

Blocking / Waiting

Thank You!