

Lecture 6: Synchronization

Bo Tang @ 2020, Spring

Inter-process Communication (IPC)


◆ Pipe

- ◆ Unidirectional
- ◆ Between processes with a common ancestor (e.g., `ls | less`; ancestor=shell)

◆ Signal

- ◆ More kernel-level
- ◆ Limited (SIGCHLD, SIG...)

Pipe is a **shared object** between two processes.

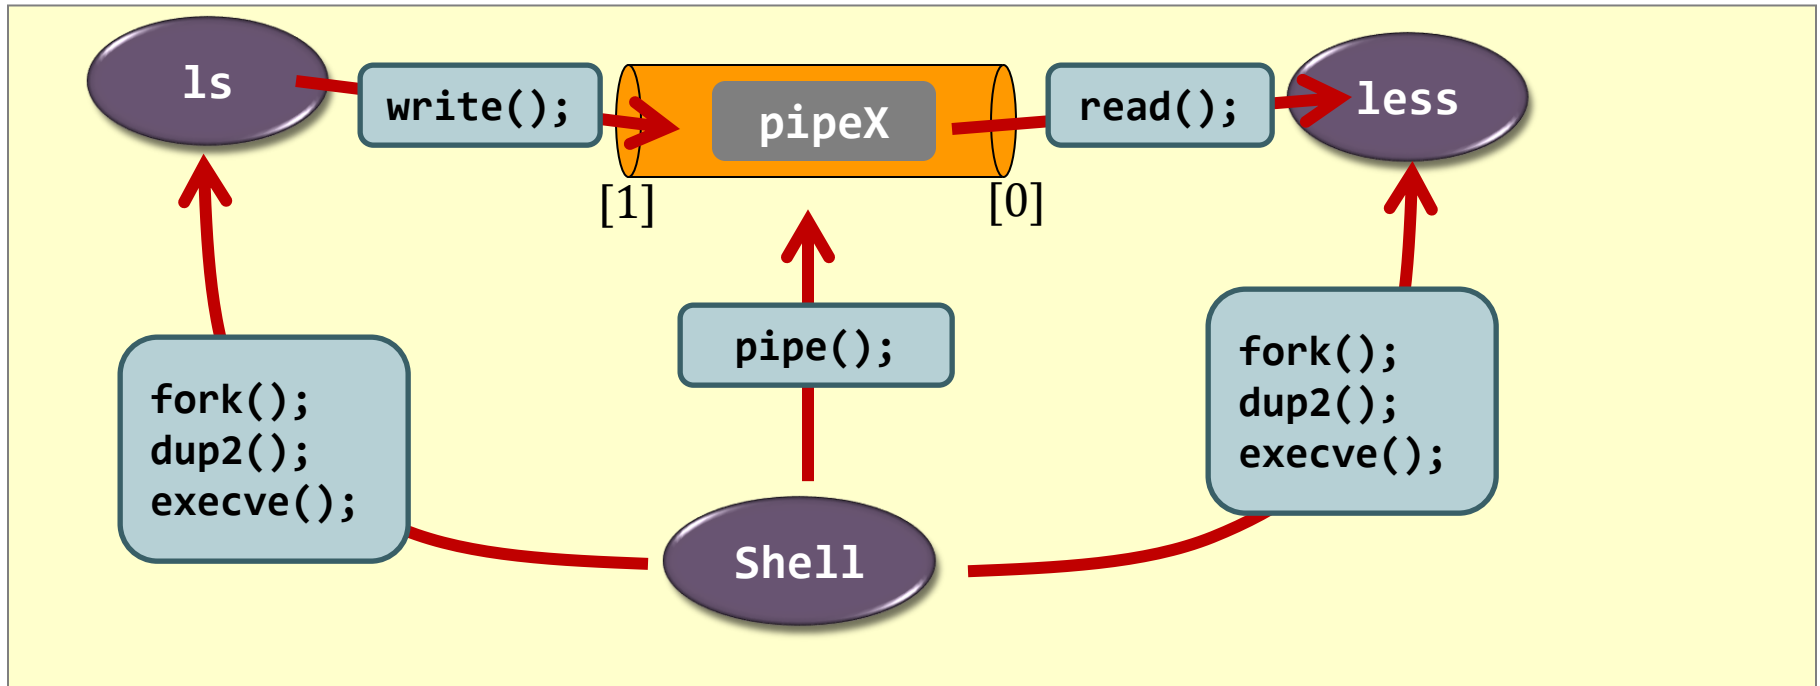


Programming “ls | less”

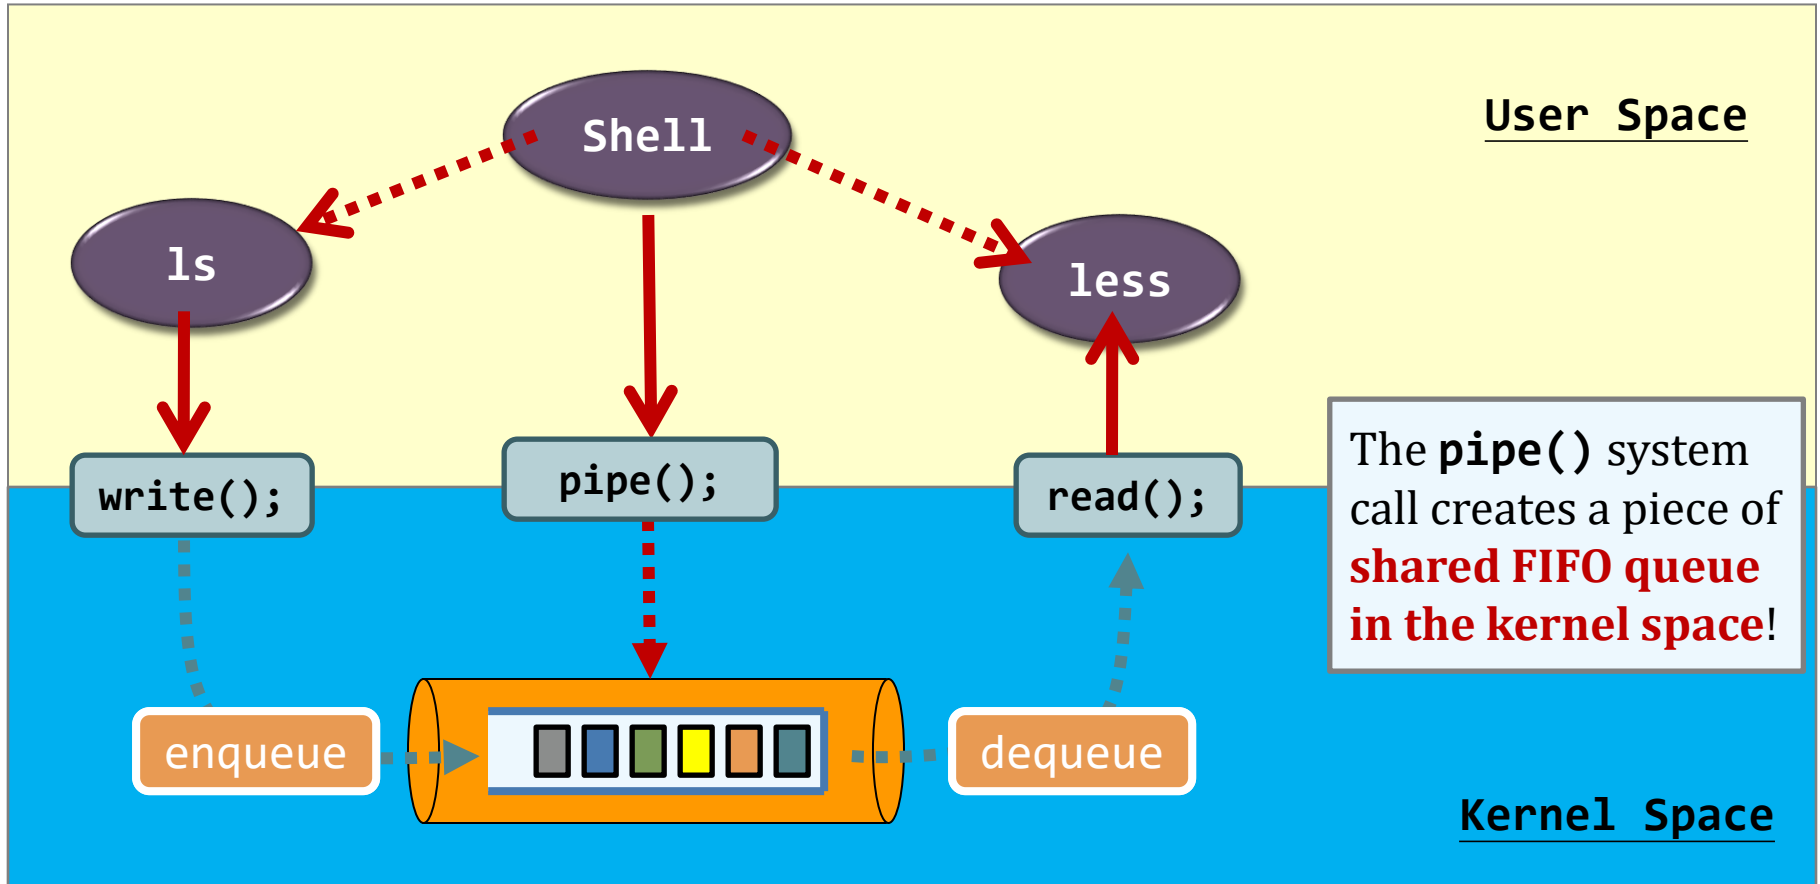
```
fork();  
if (pid==0) { // child; “ls”  
    //dup2: replace “ls” default stdout  
    //by the write end of the pipe  
    dup2(pipeX[1], STDOUT_FILENO);  
    execlp(“ls”, “ls”, NULL);  
} else ... //parent; “less”
```

In UNIX*, “everything is a file”

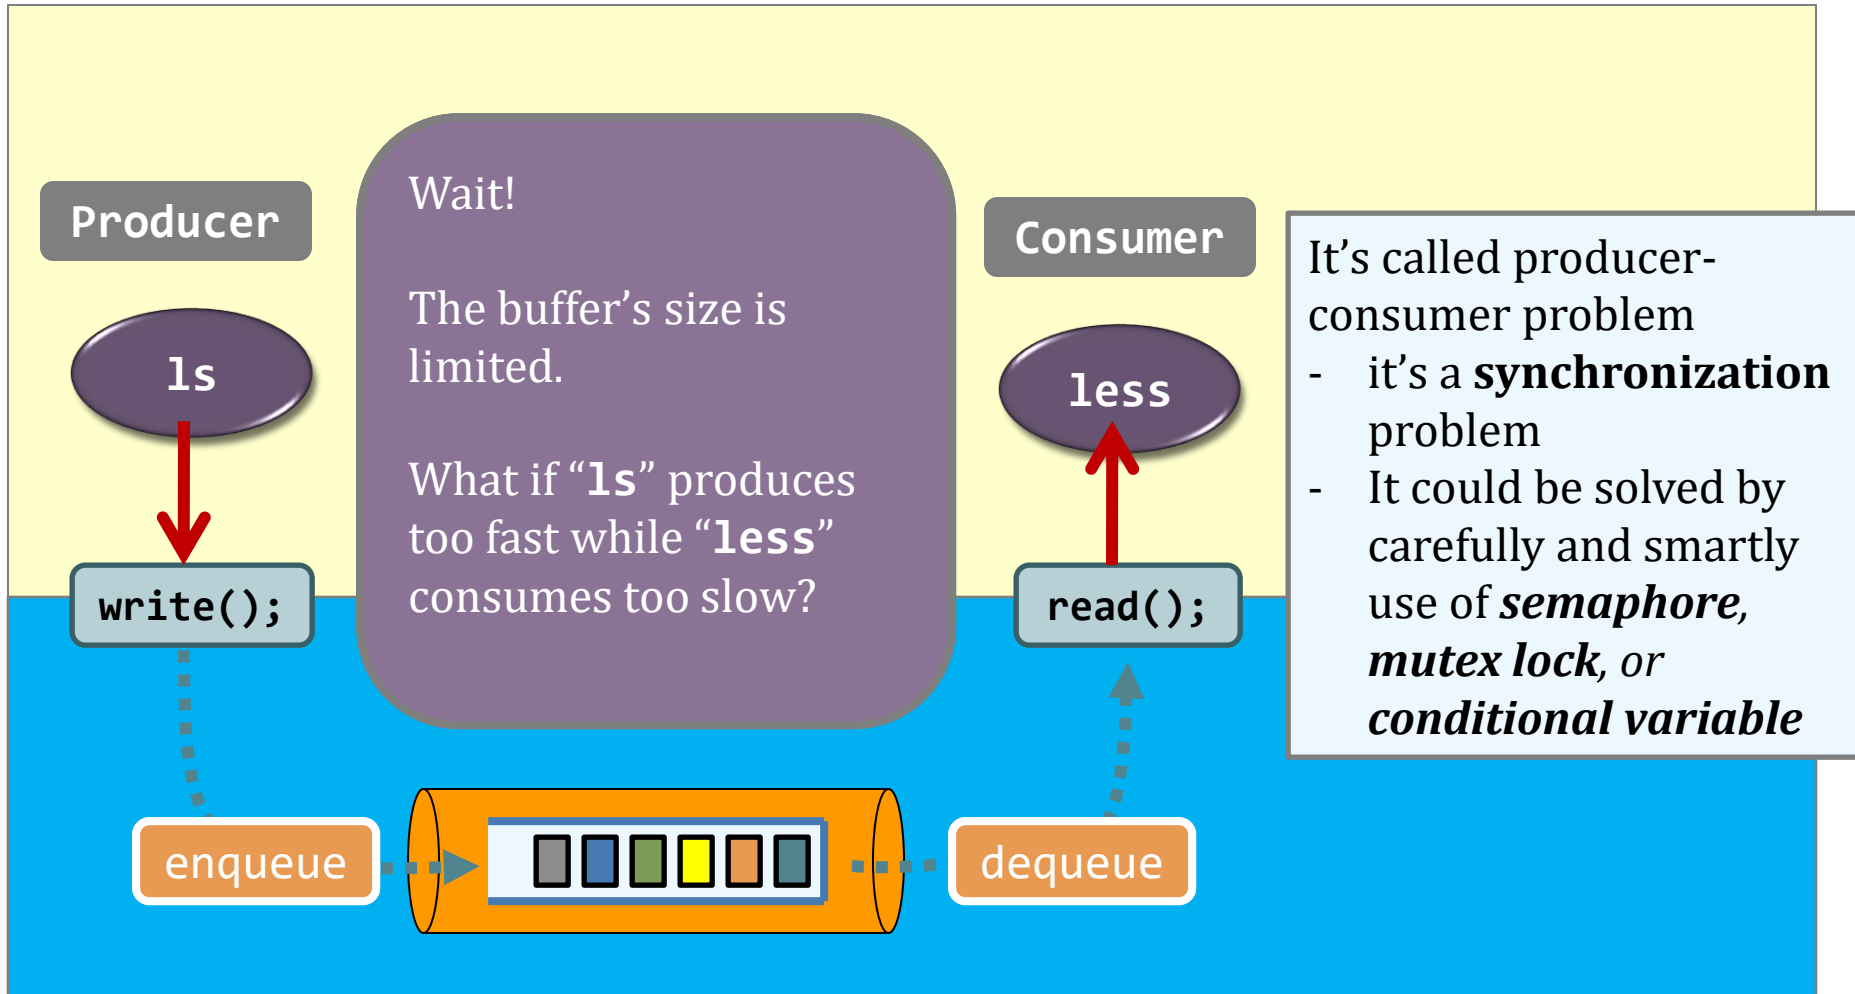
- Every resource that can read/write is represented as a file. E.g.,
 - Network, Disk, Keyboard
- A “file” is indexed by a number called *file descriptor*



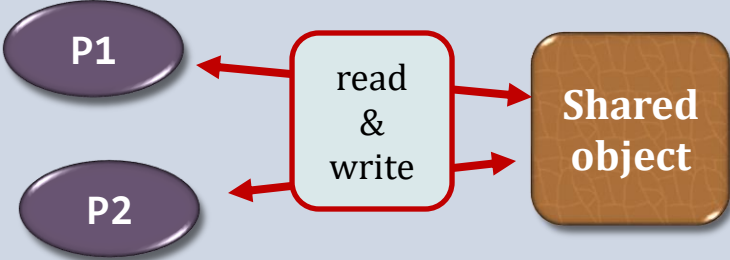
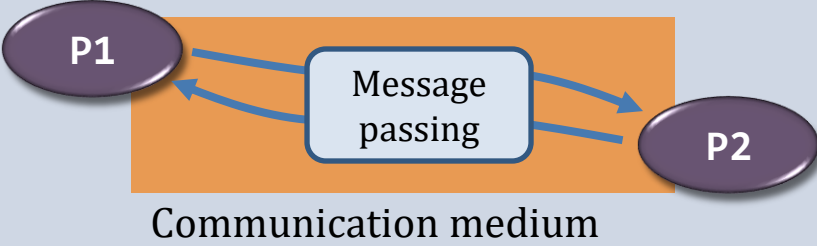
“ls | less” in kernel



Synchronization problems



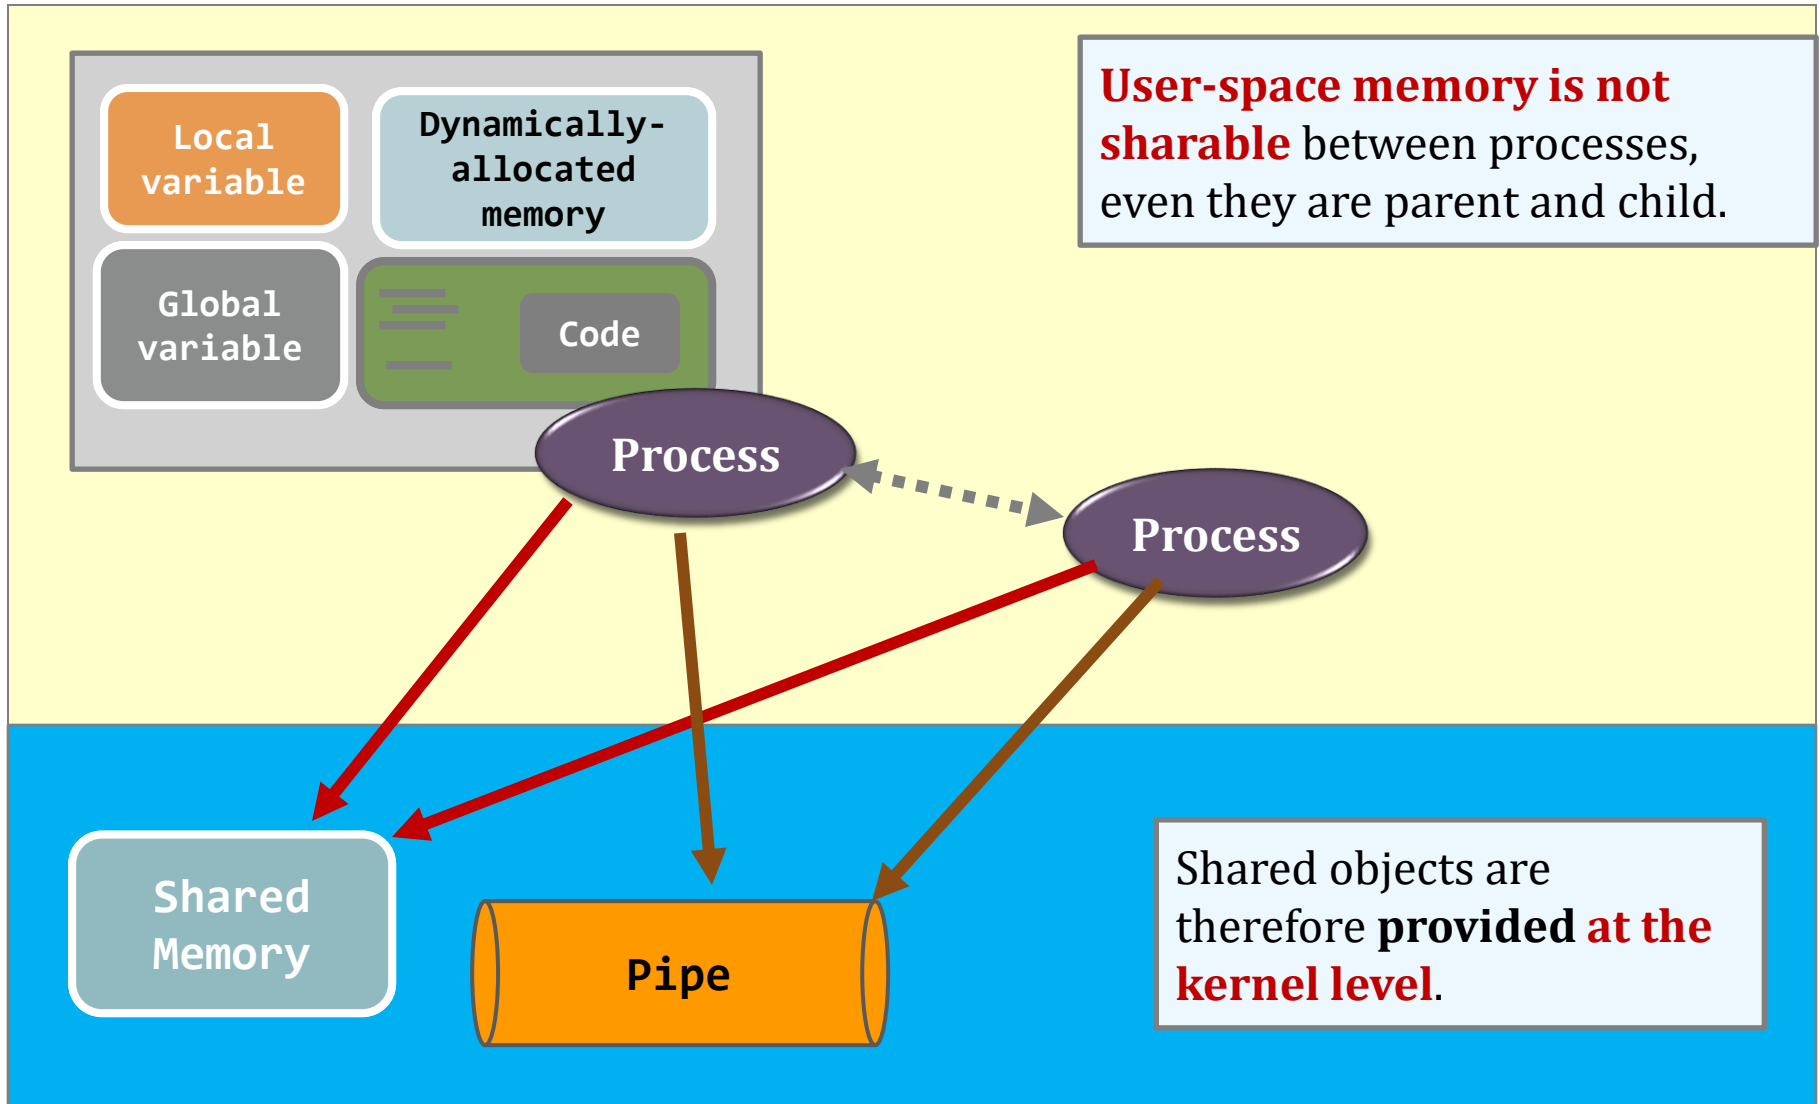
Summary of IPC models

Shared Objects	Message Passing
 <p>The diagram illustrates the Shared Objects IPC model. Two processes, P1 and P2, are shown as purple ovals on the left. A central light blue box labeled 'read & write' is connected to a brown rounded rectangle labeled 'Shared object' on the right. Red arrows point from the 'read & write' box to both P1 and P2, and from both P1 and P2 to the 'Shared object' box, indicating bidirectional communication between the processes and the shared object.</p>	 <p>The diagram illustrates the Message Passing IPC model. Two processes, P1 and P2, are shown as purple ovals. A light blue box labeled 'Message passing' is situated between them, all within an orange rectangular area labeled 'Communication medium' below. Blue arrows show bidirectional communication between P1 and the 'Message passing' box, and between P2 and the 'Message passing' box.</p>
<ul style="list-style-type: none">• shared files (on disk; slow)• pipes (restricted, but OS takes care of synchronization for you)• shared memory (primitive, general, but synchronization is on you)• shared address space (threading)	<ul style="list-style-type: none">• socket programming• message passing interface (MPI) library for computing clusters.
<ul style="list-style-type: none">- Usually single-node communication- More efficient- Need to take great care of synchronization because of sharing the same object	<ul style="list-style-type: none">- Usually multi-node communication- Less efficient- Less troublesome in synchronization- But need to care of other faults (e.g., what if a network link is broken?)

Inter-process communication (IPC)

- What, why, and how?
- The problem: race condition

Evil source: the shared objects.



Evil source: the shared objects.

- ❖ Kernel provides you “pipe” to do **one-way** data flow between **2 processes** from the **same ancestor**
 - ❖ Super restrictive
- ❖ Other IPC problems beyond pipe?
 - ❖ You have to use shared memory directly
 - ◆ **concurrent access may yield unpredictable outcomes!**
 - ◆ Kernel will not take care of that for you
 - ◆ You take care of that

Race Condition: Understanding the problem


High-level language for Program A

```
1 attach to the shared memory X;  
2 add 10 to X;  
3 exit;
```

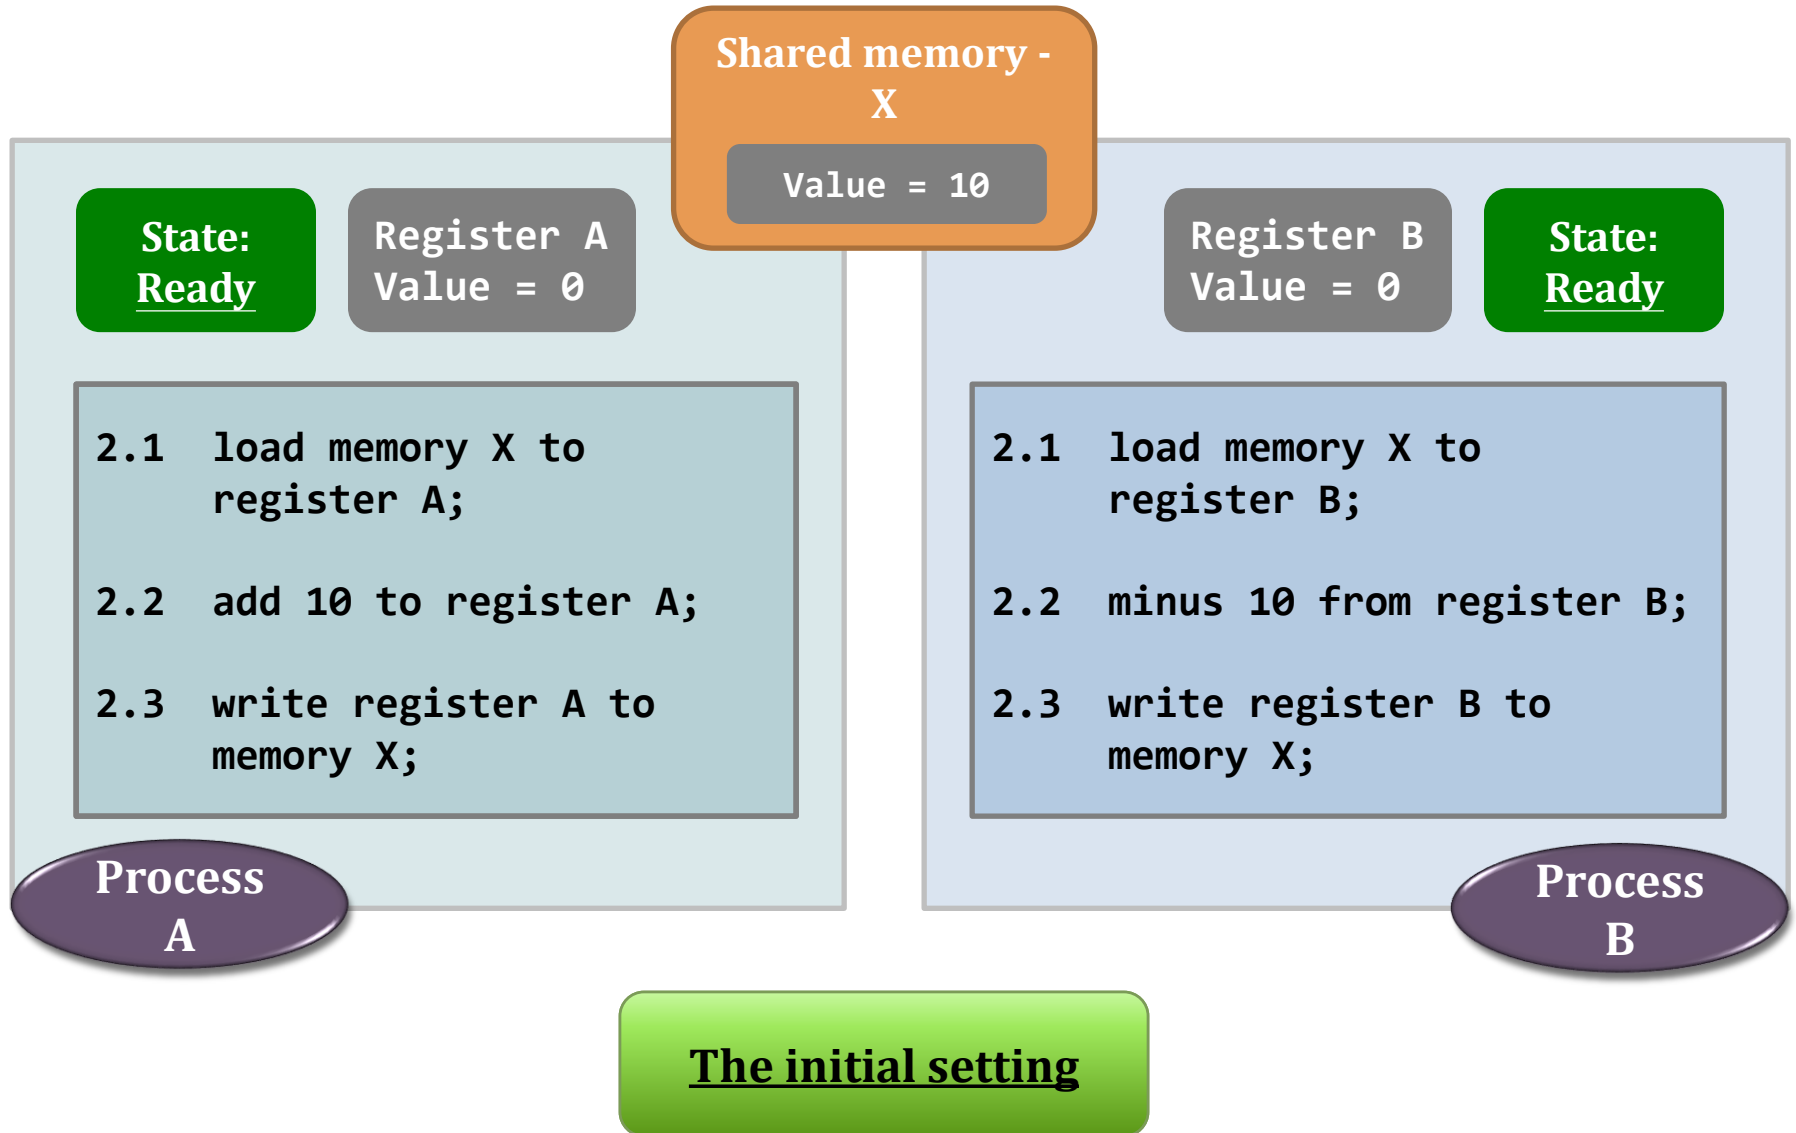
Partial low-level language for Program A

```
1 attach to the shared memory X;  
.....  
2.1 load memory X to register A;  
2.2 add 10 to register A;  
2.3 write register A to memory X;  
.....  
3 exit;
```

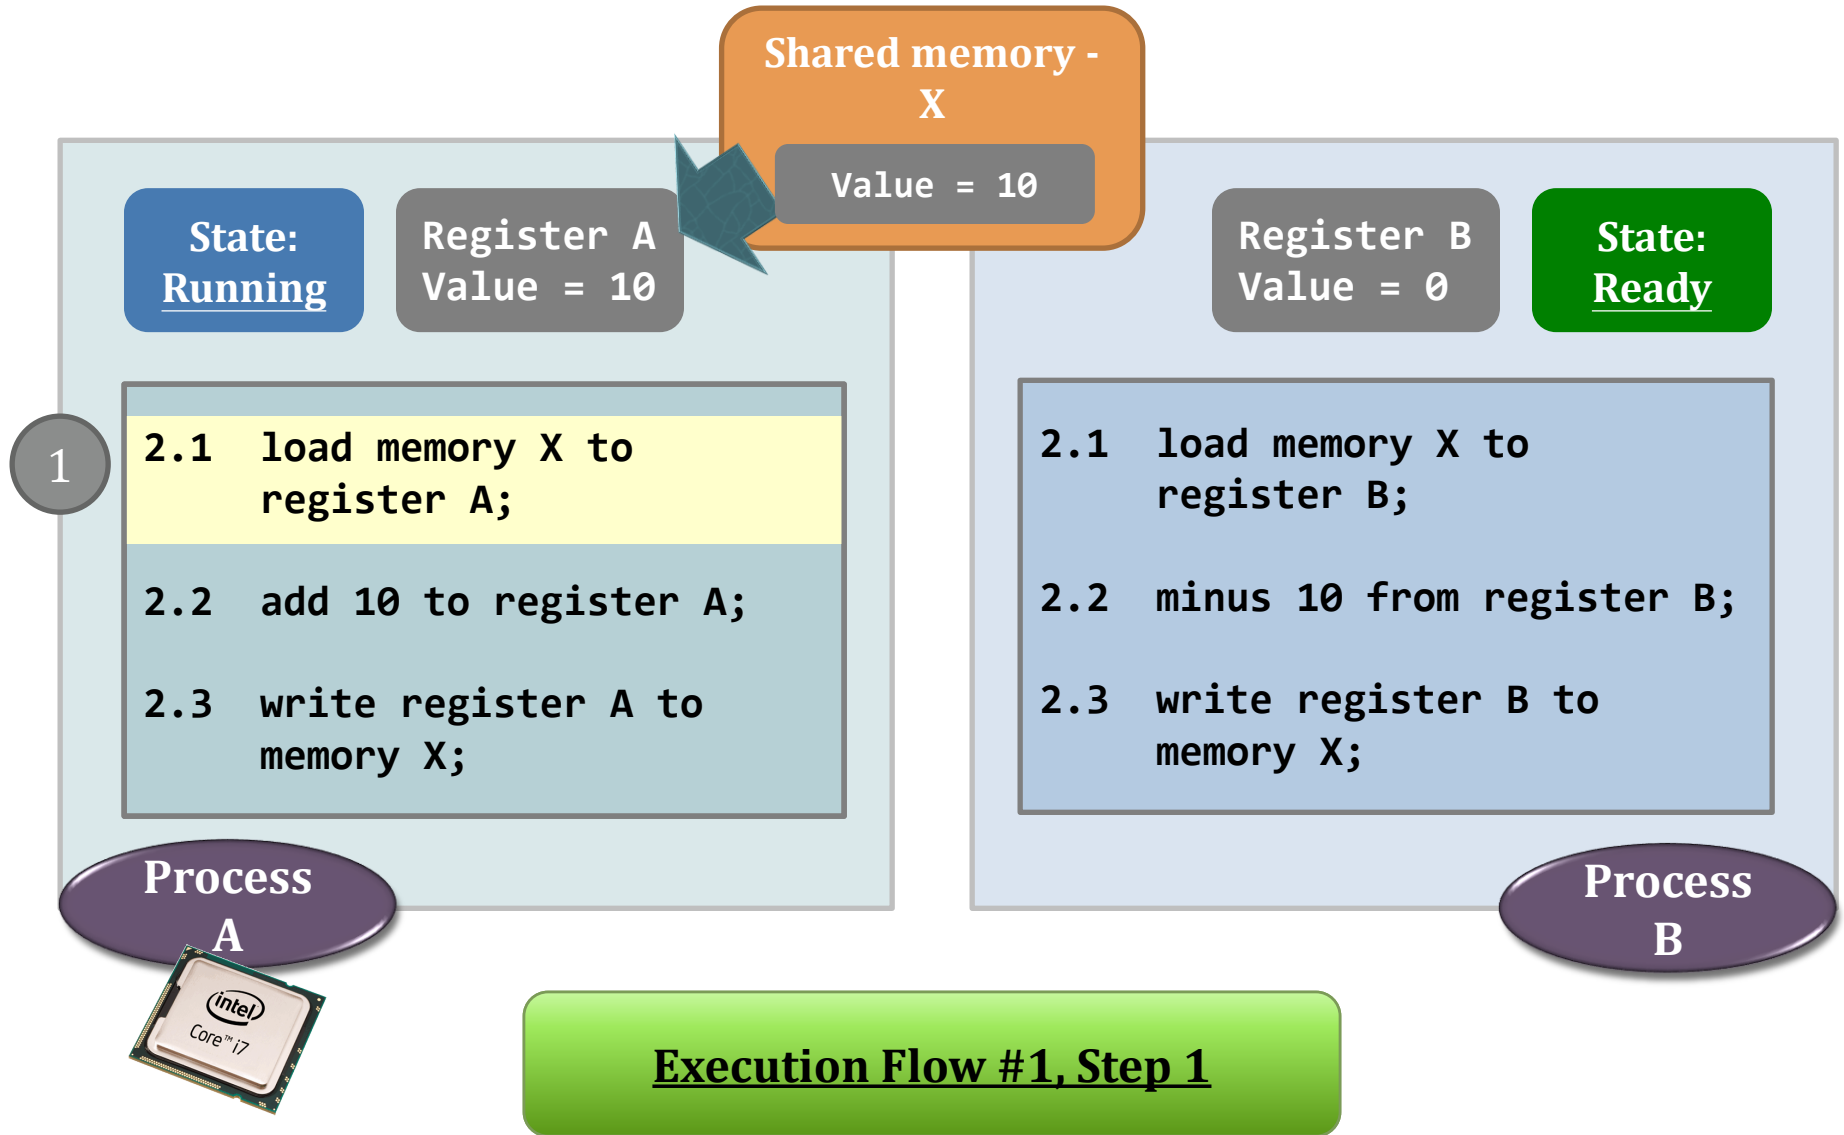
The Scenario



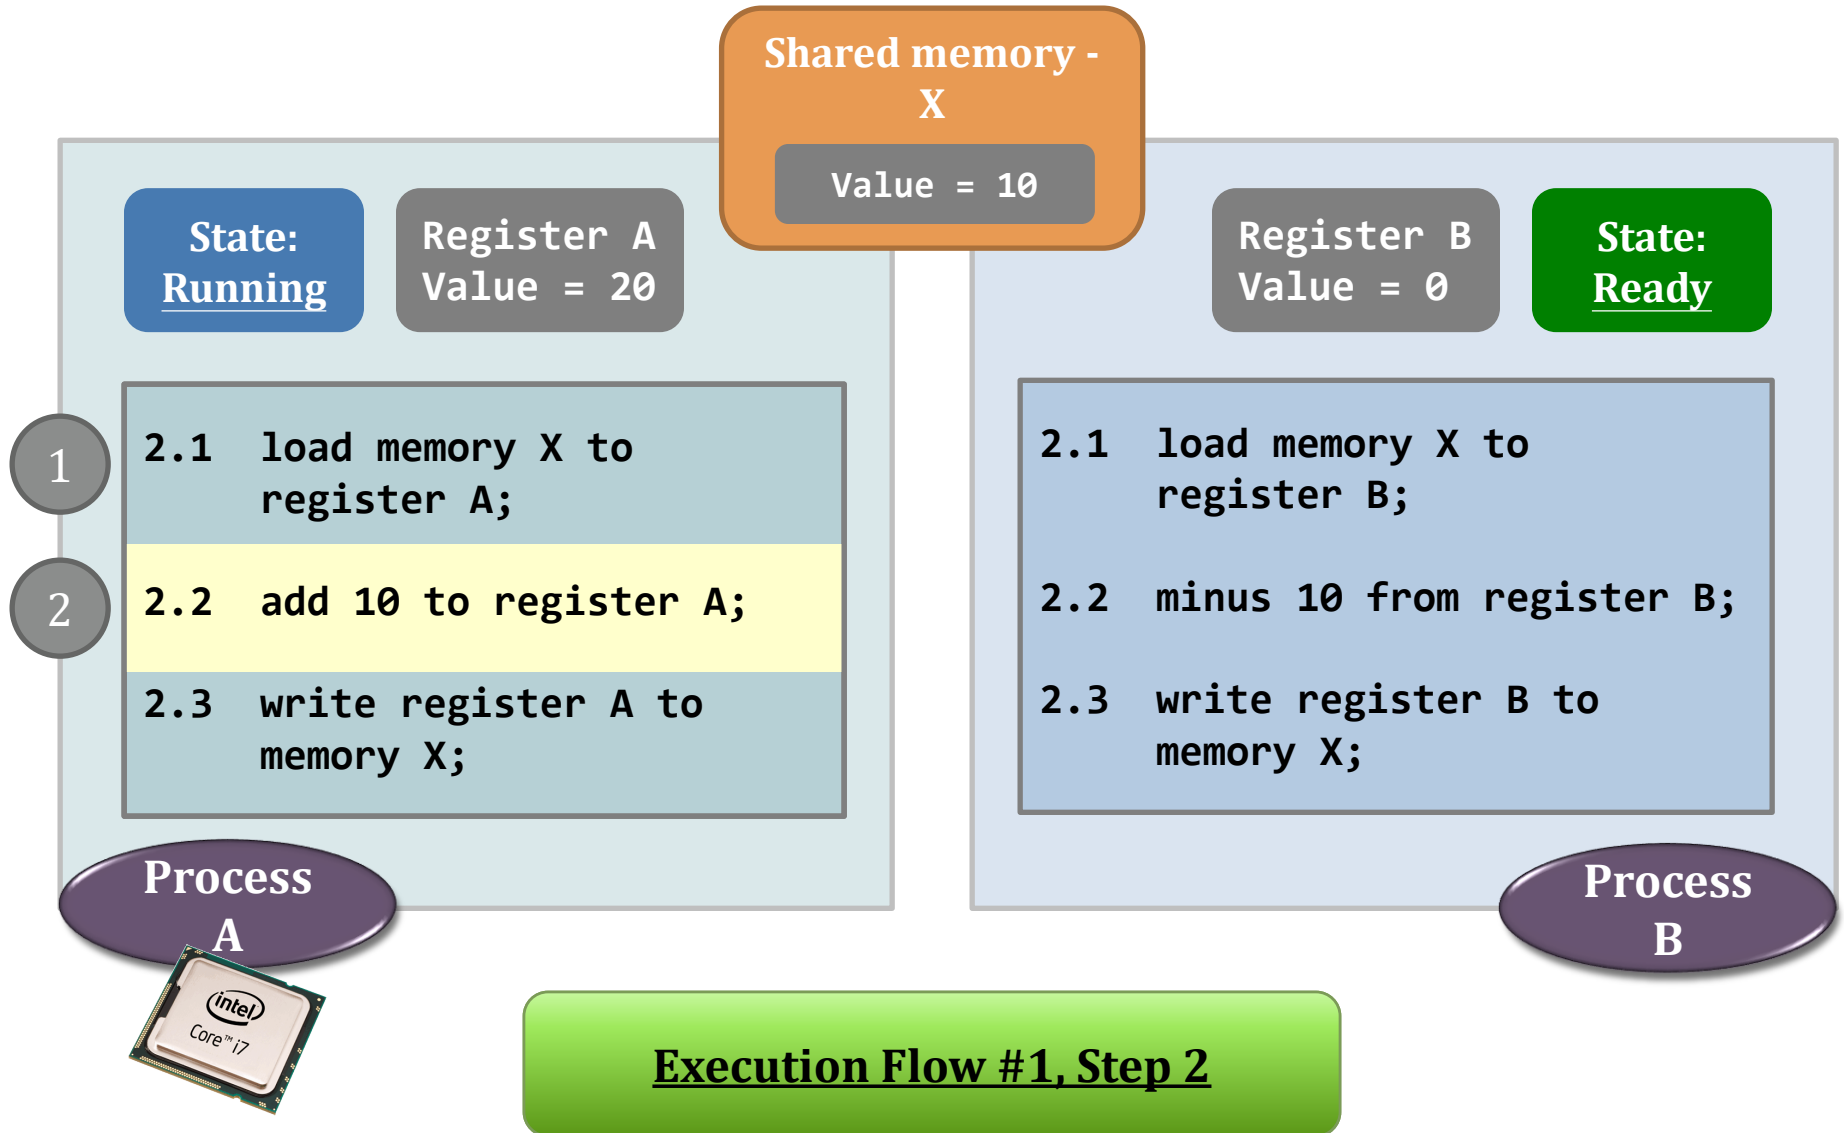
Race Condition



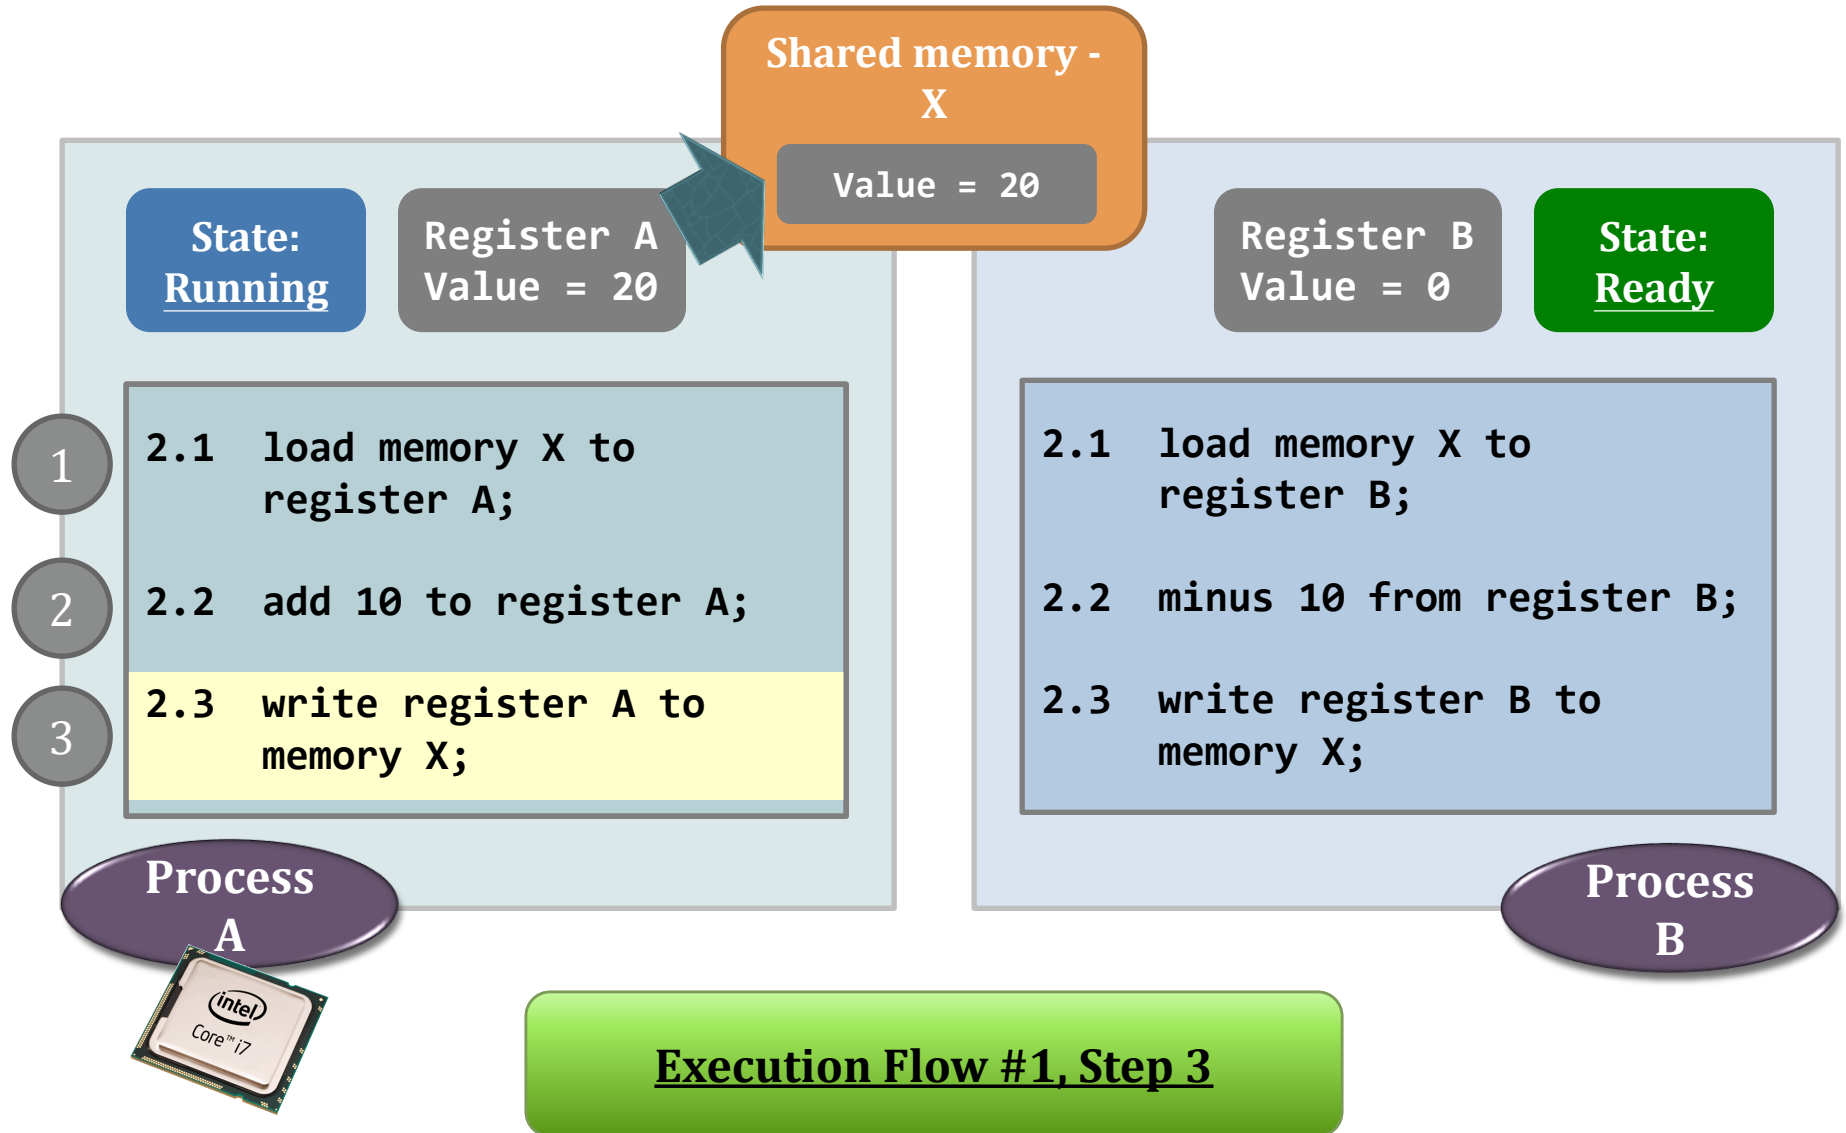
Problem not yet arise...



Problem not yet arise...

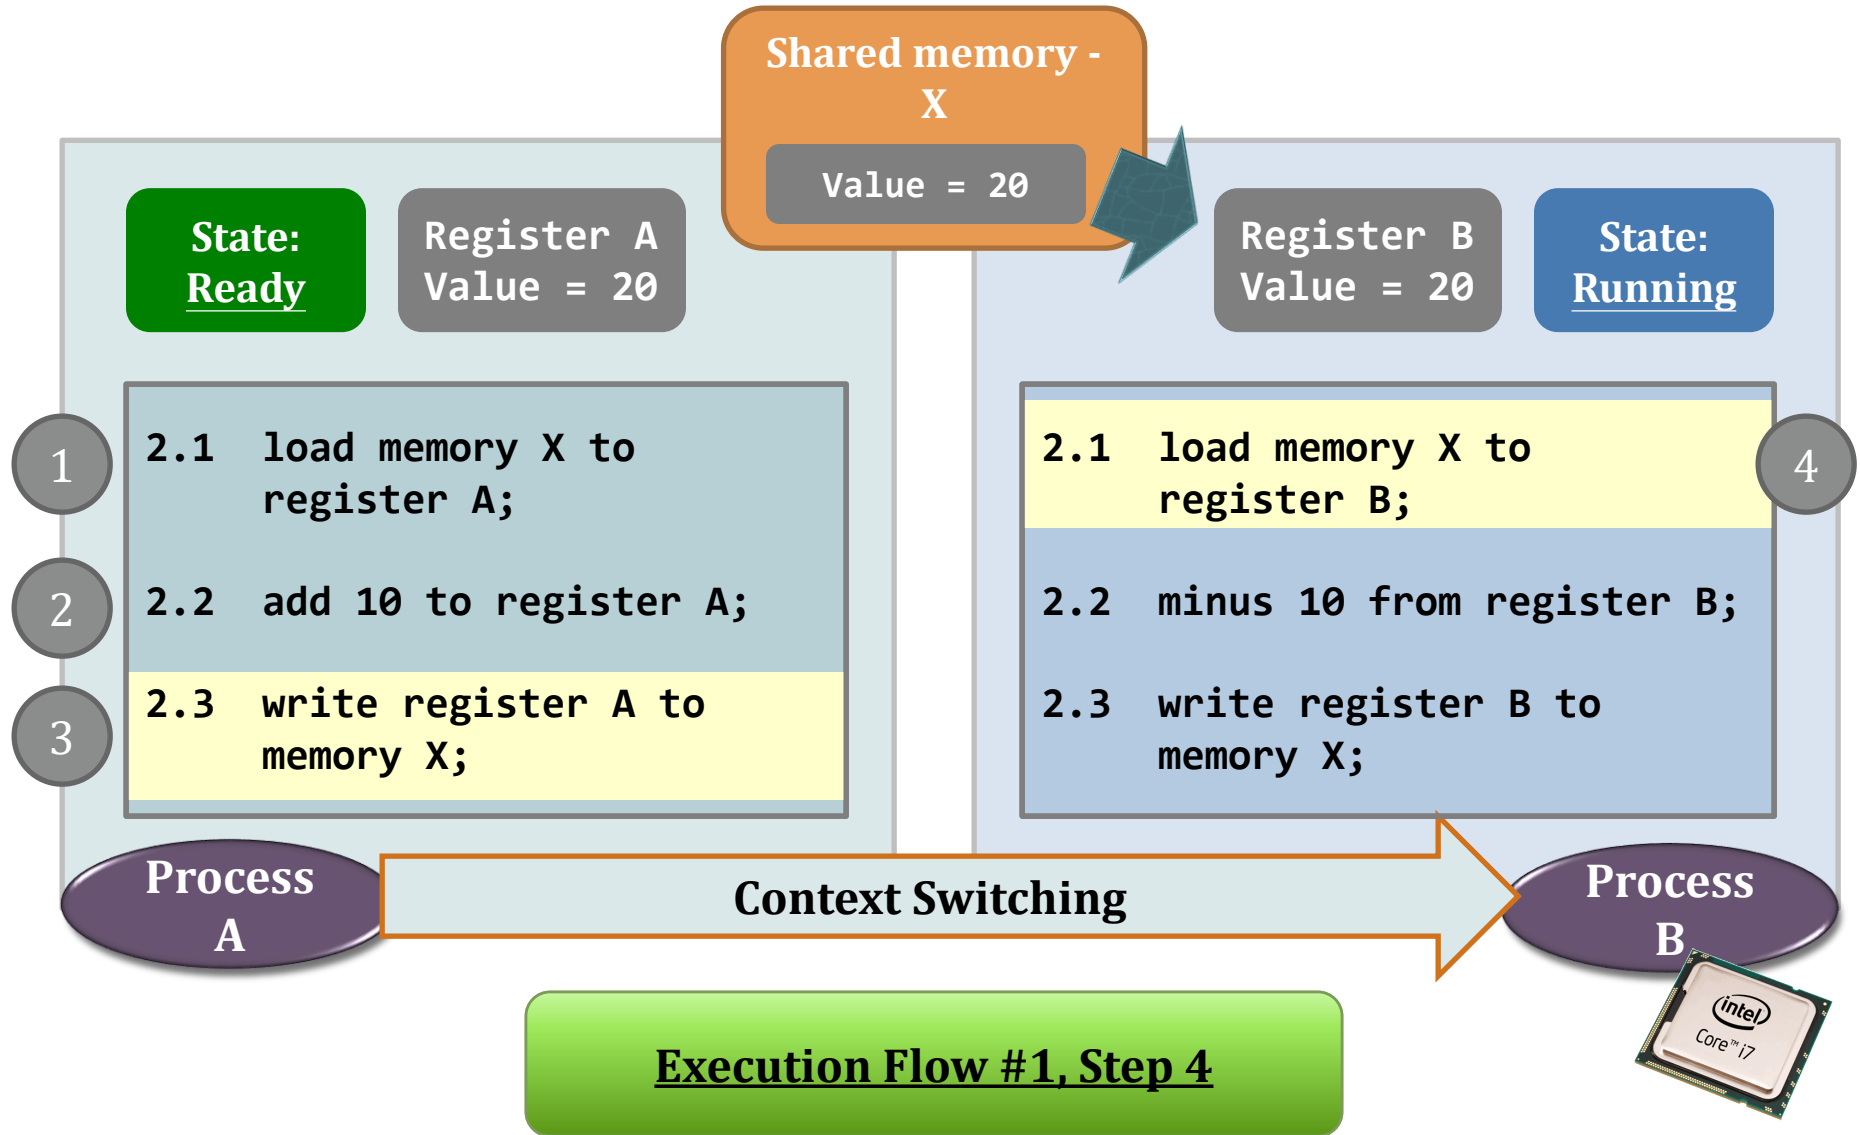


Problem not yet arise...



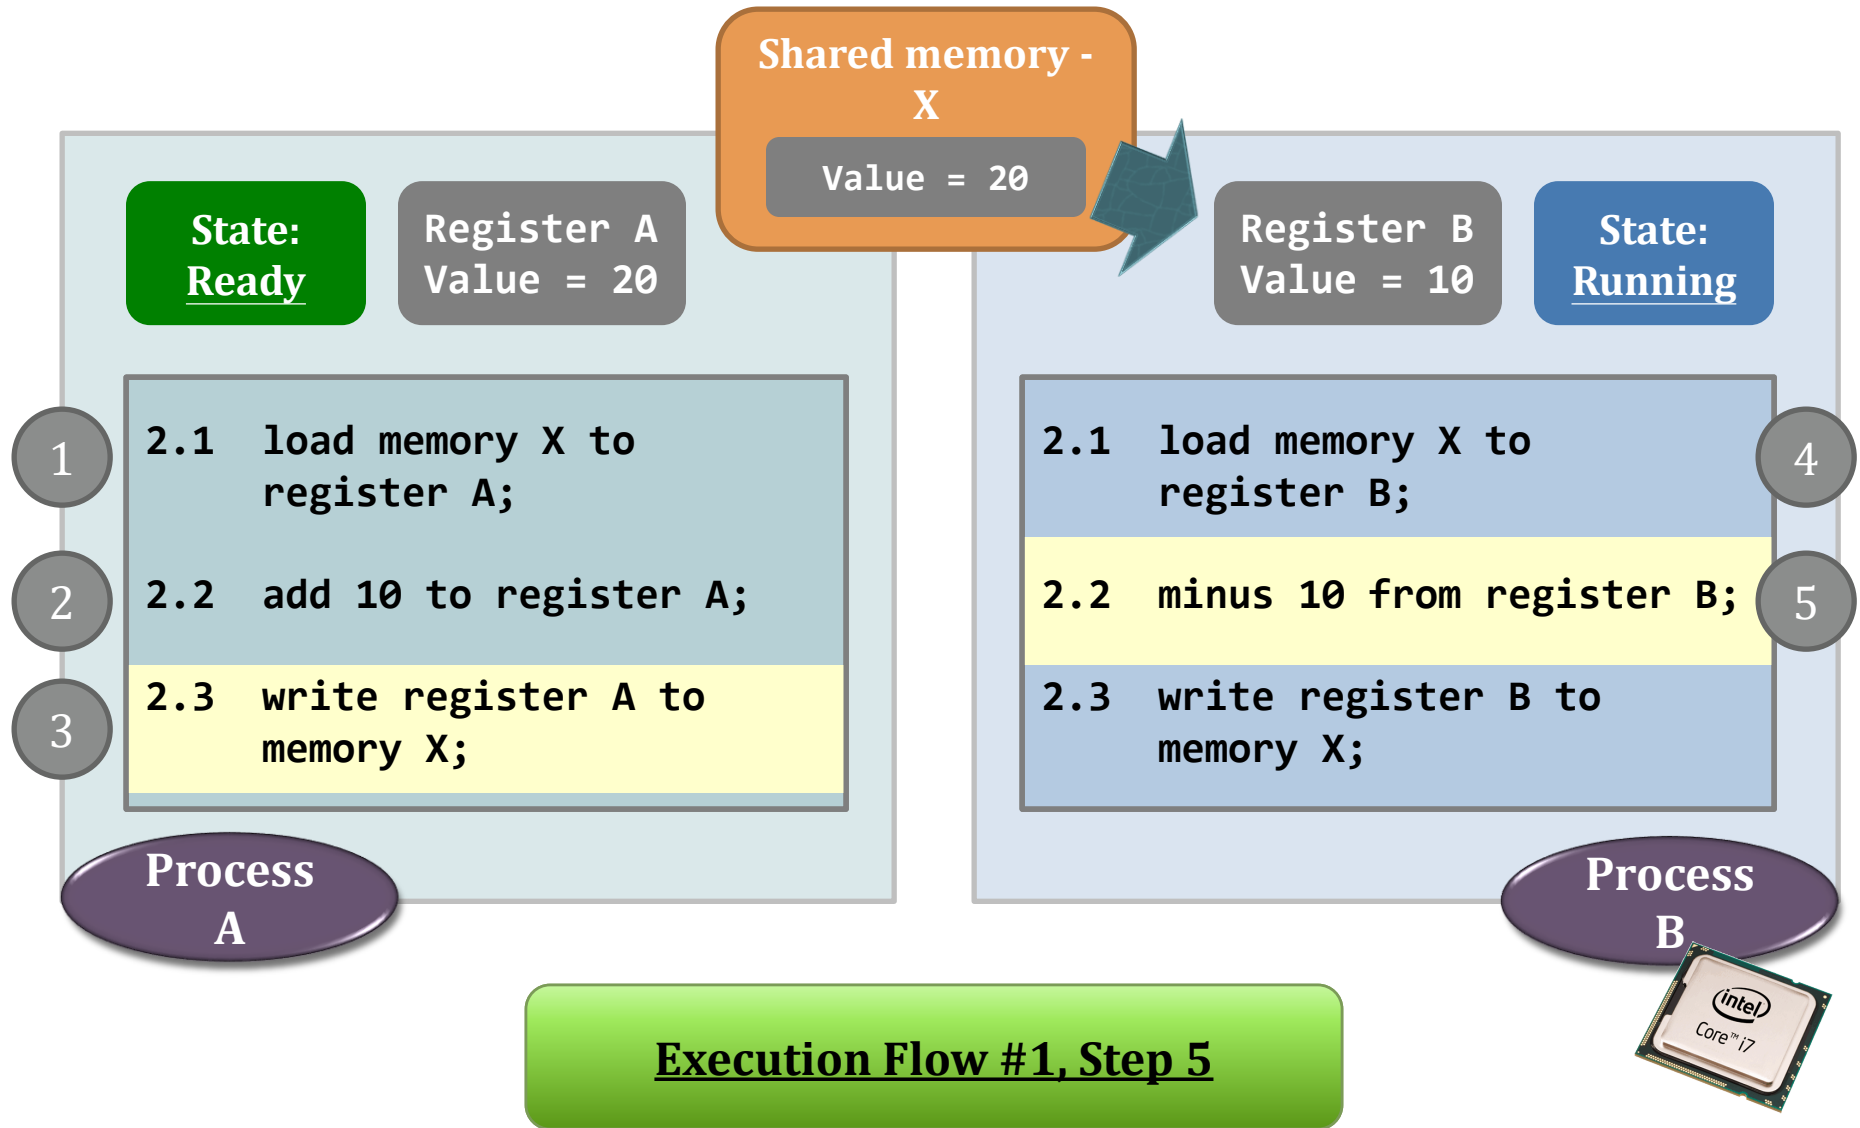
Problem not yet arise...

Don't
print

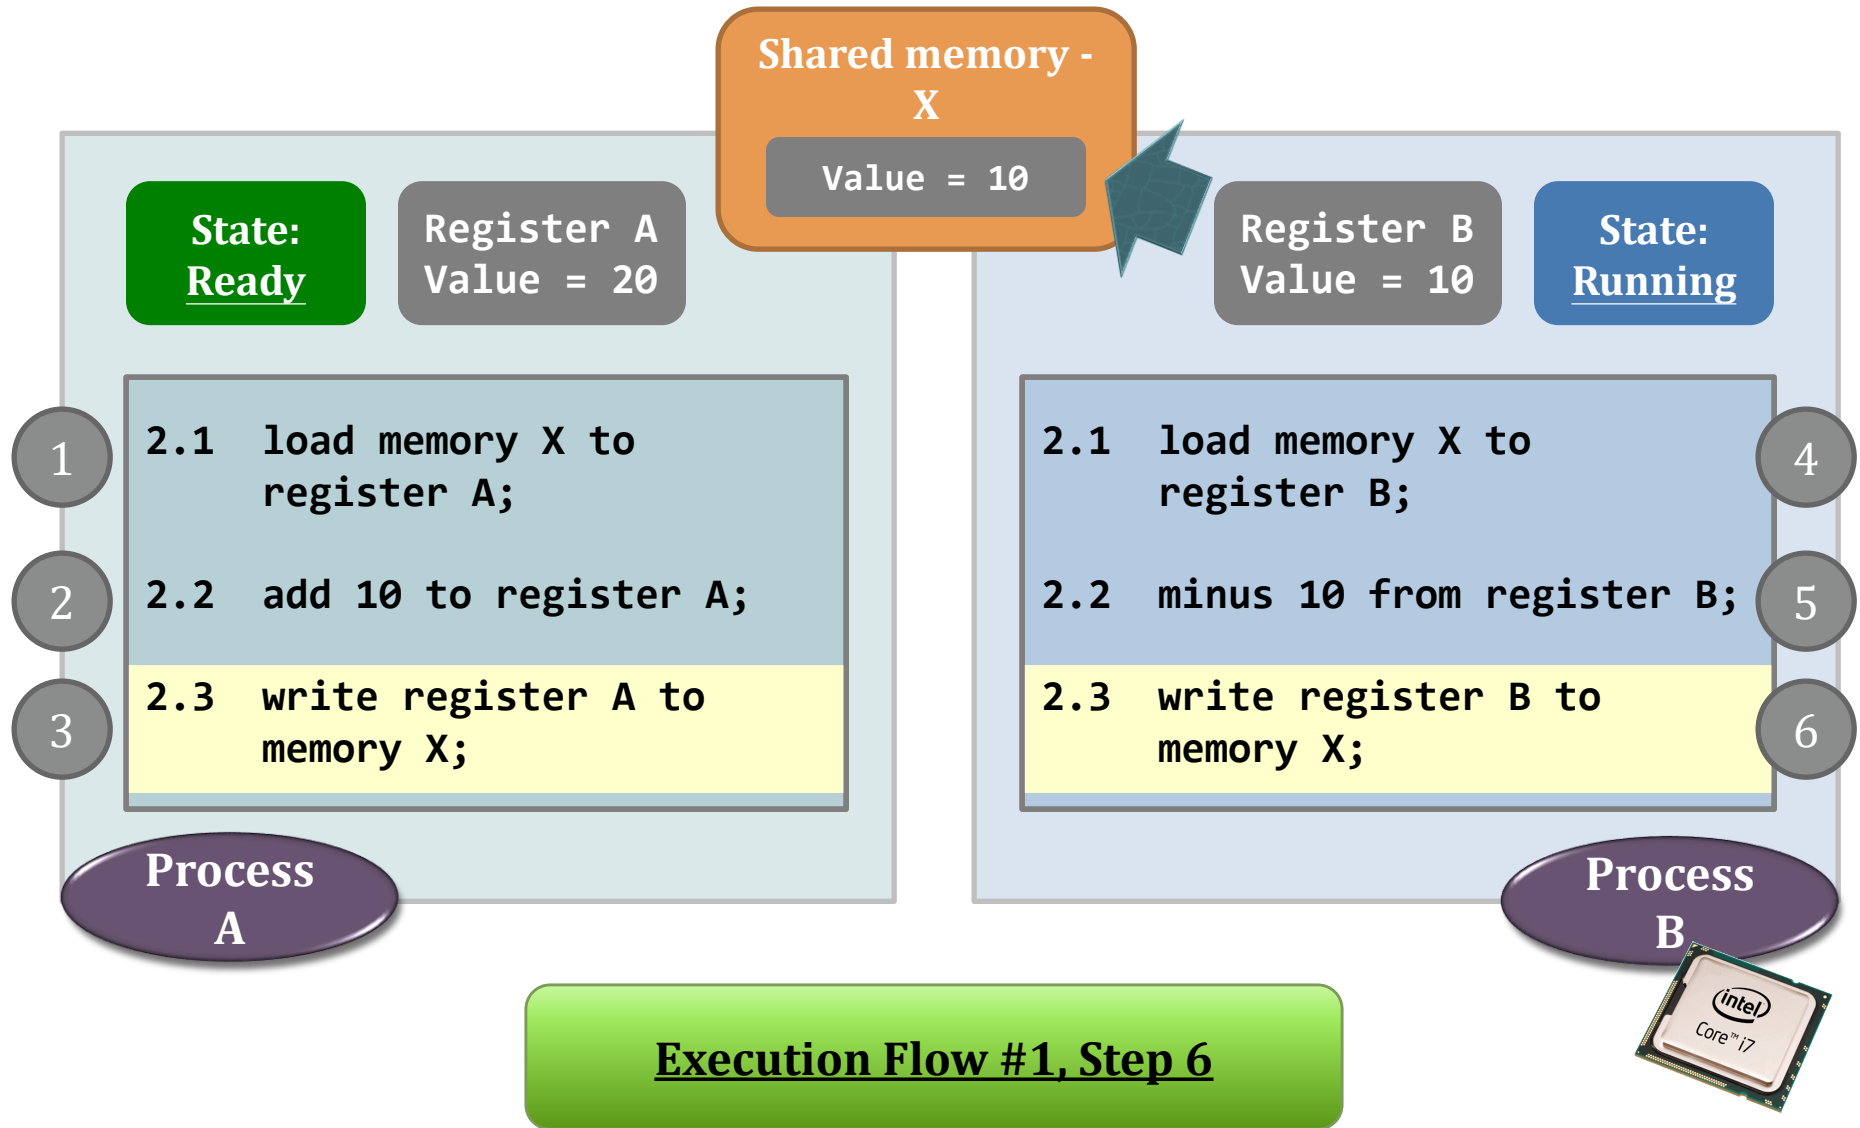


Problem not yet arise...

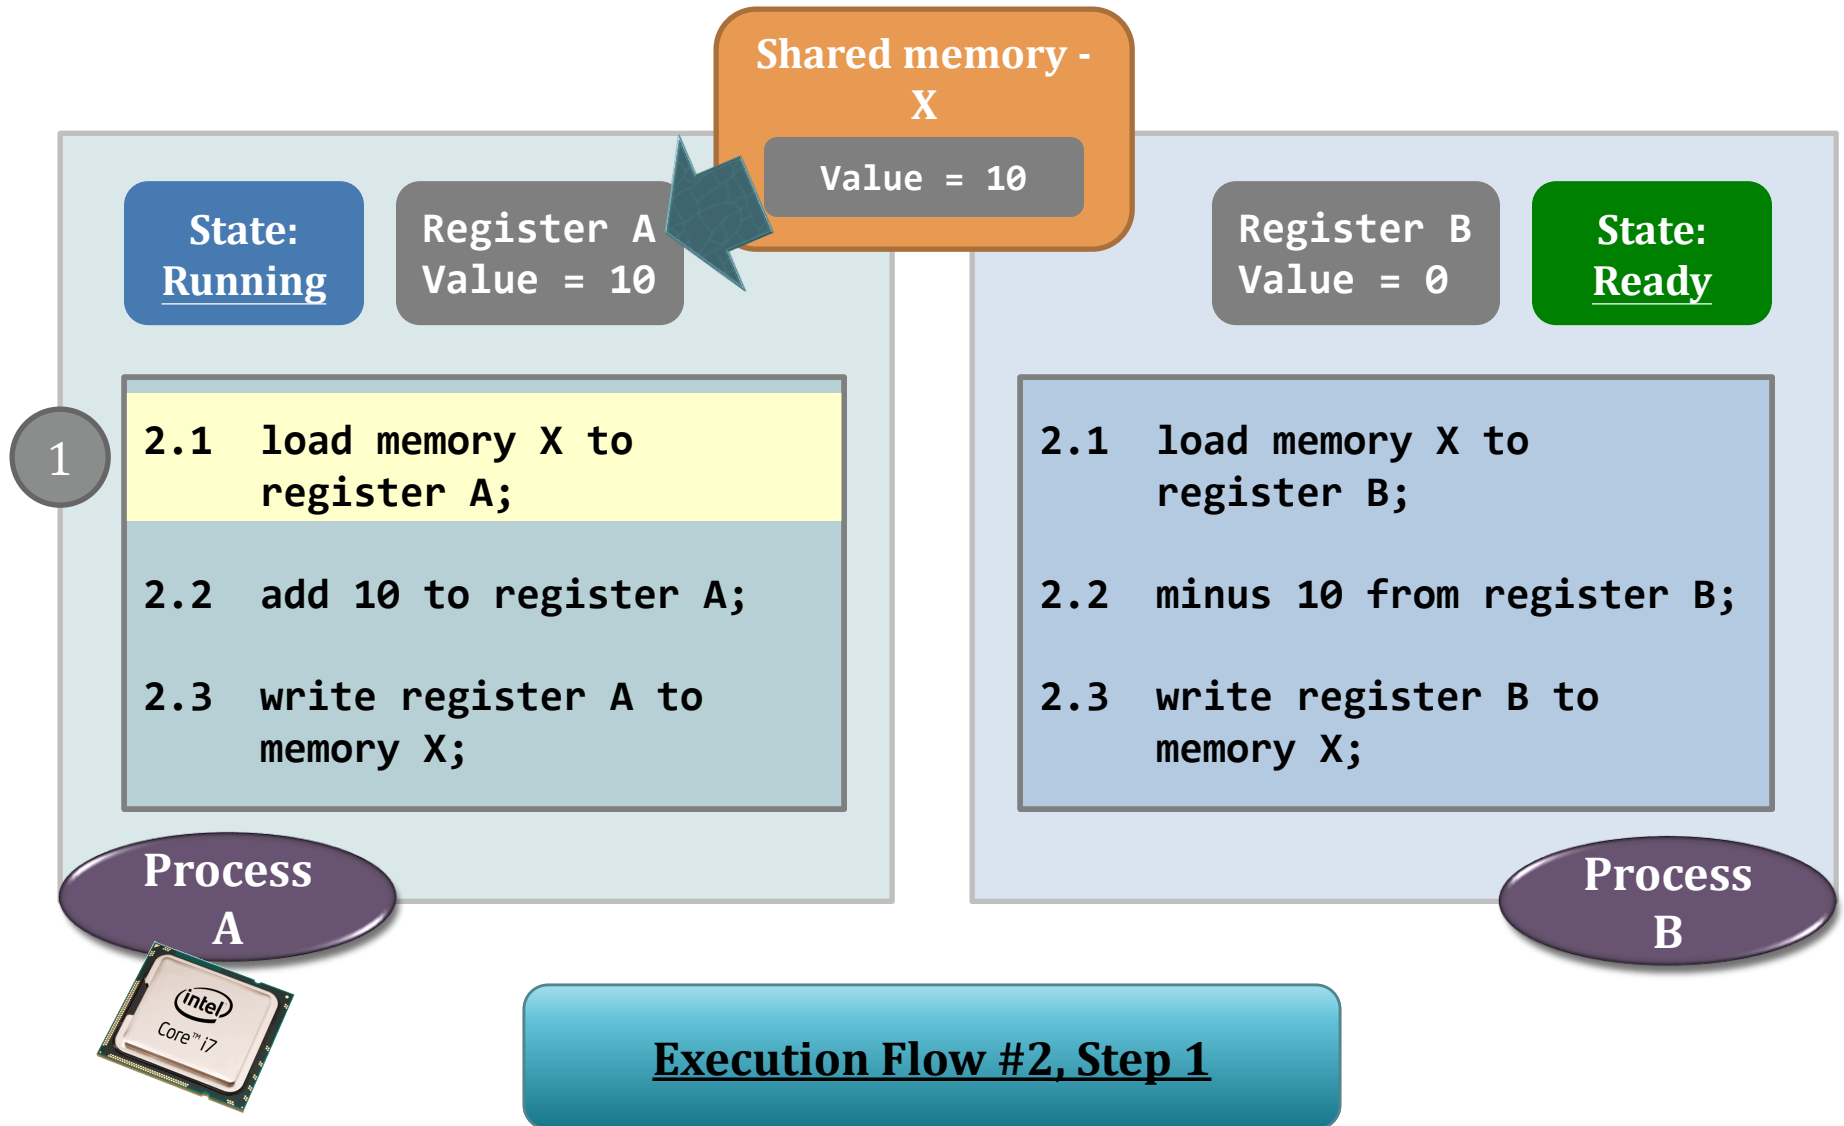
Don't
print



Problem not yet arise...

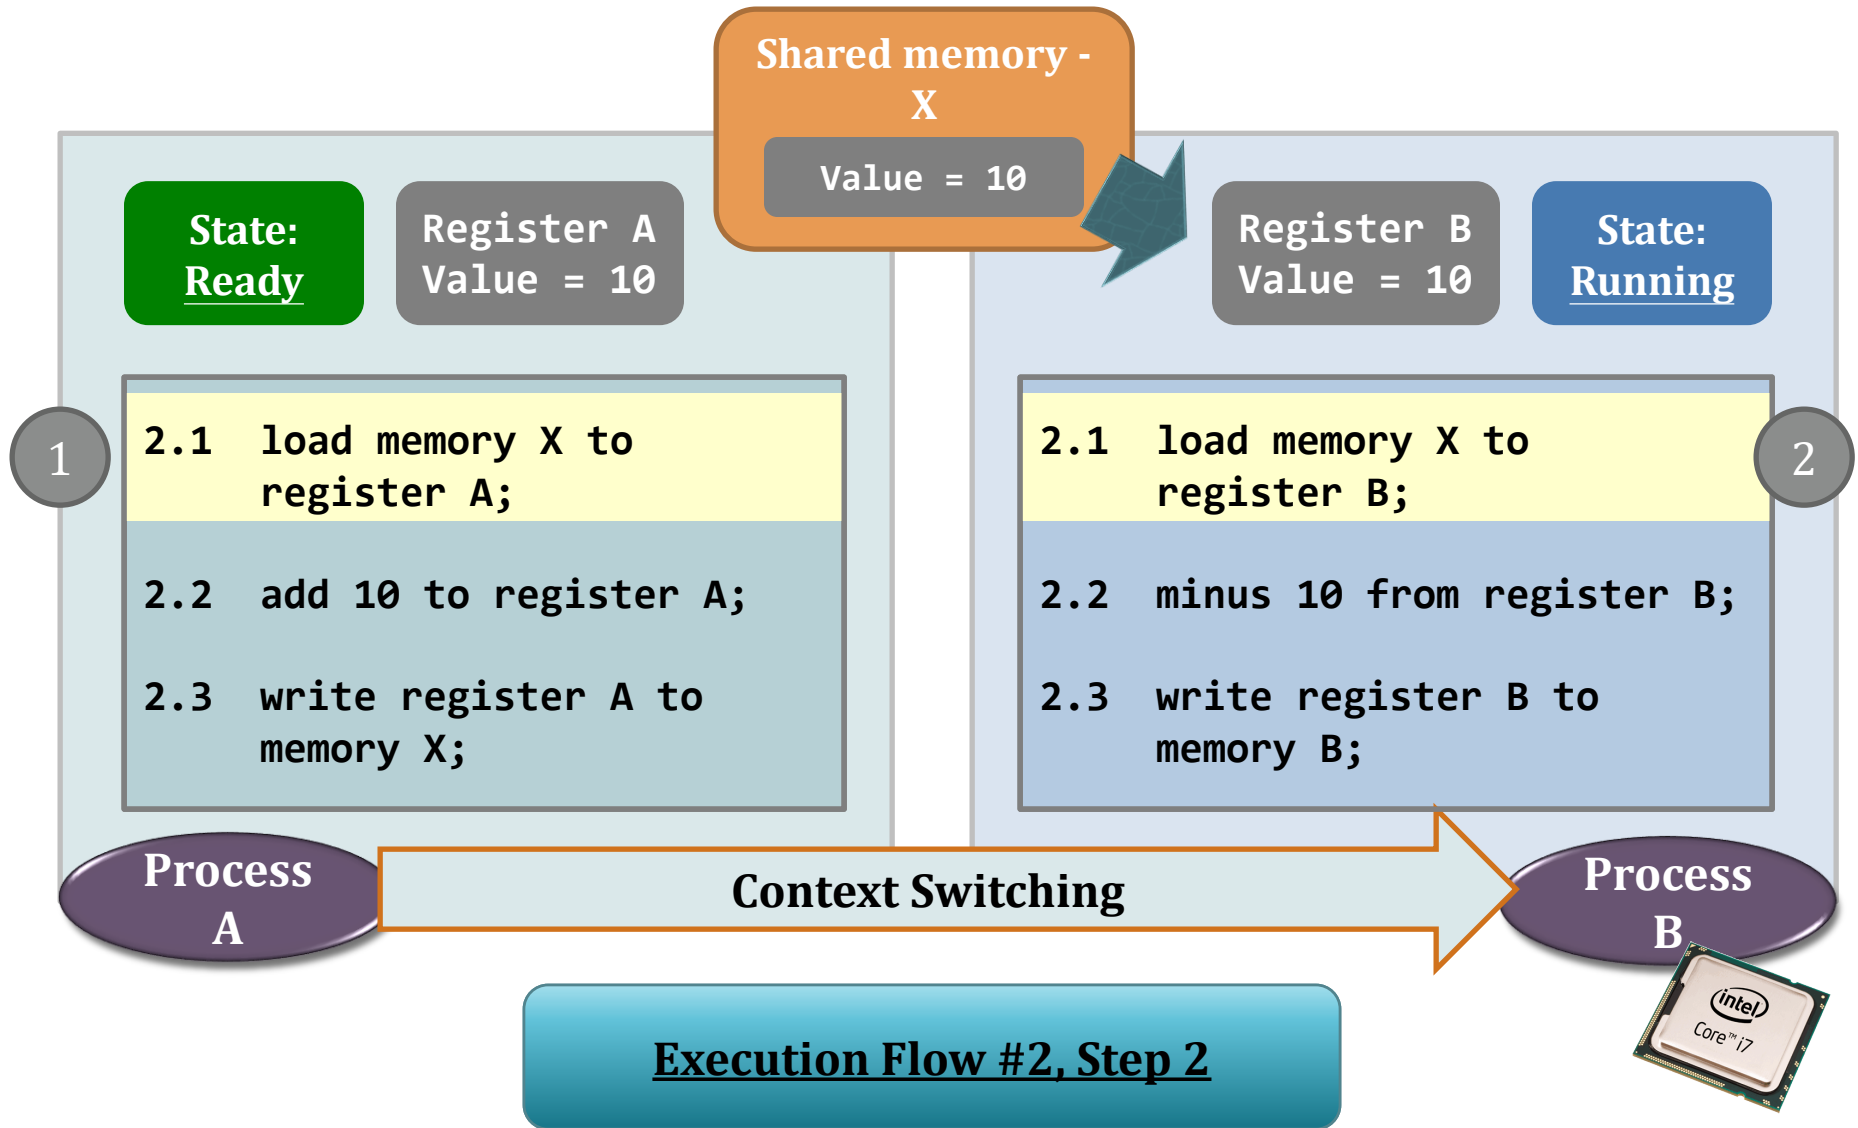


Problem arise...

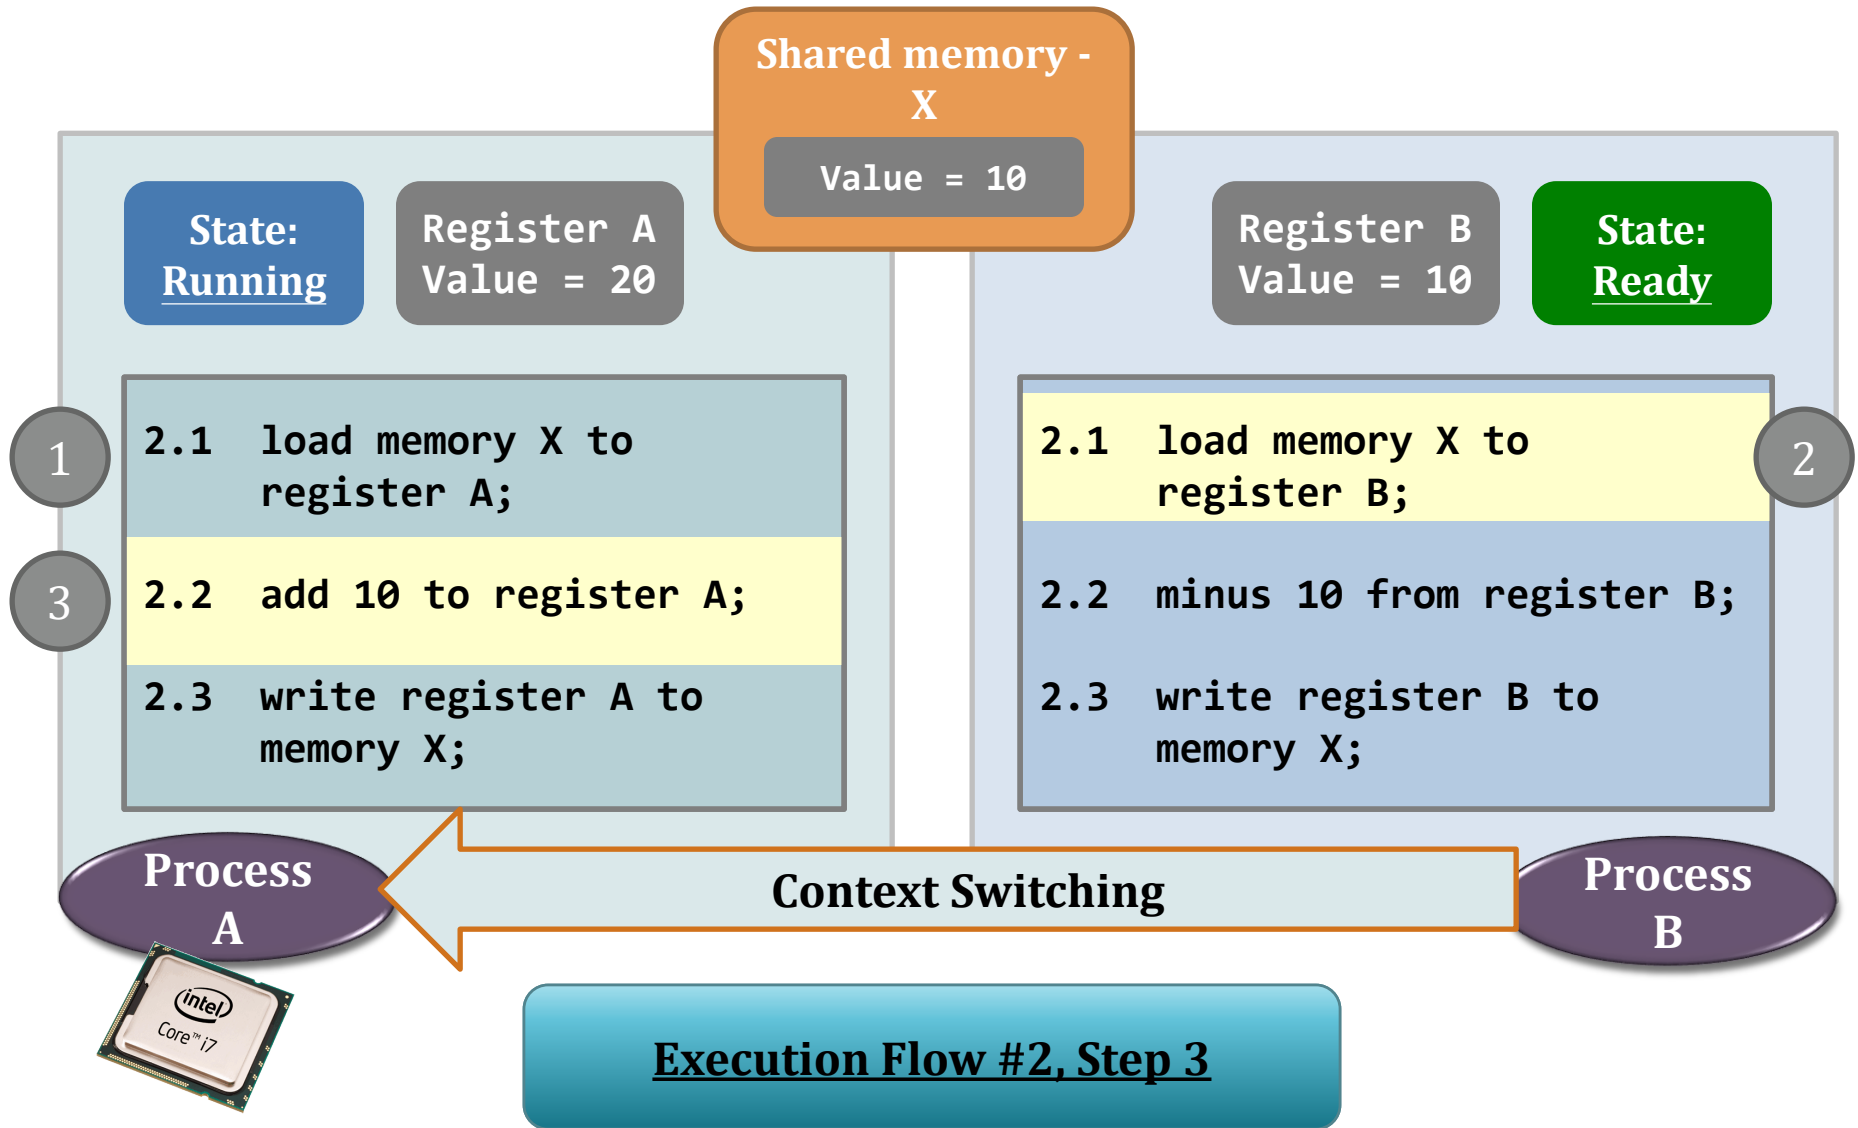


Problem arise...

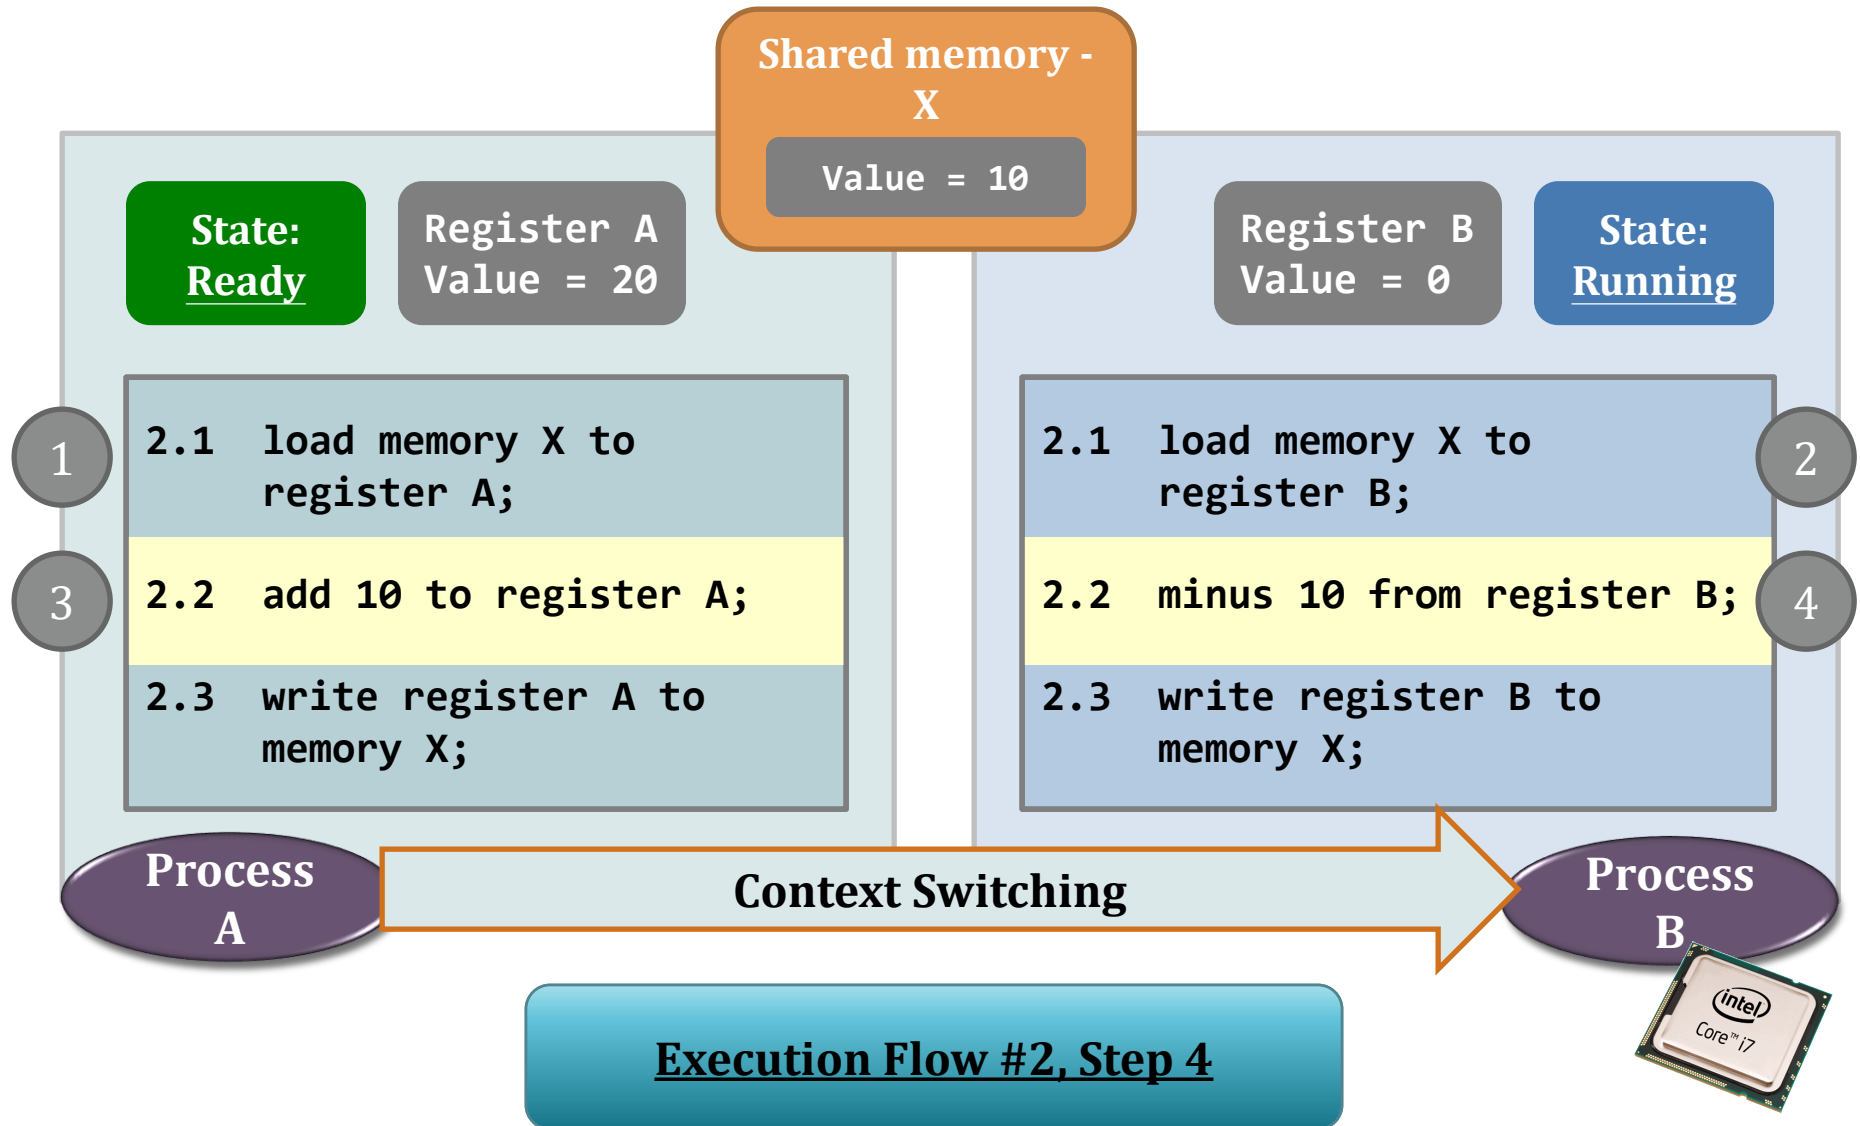
Don't
print



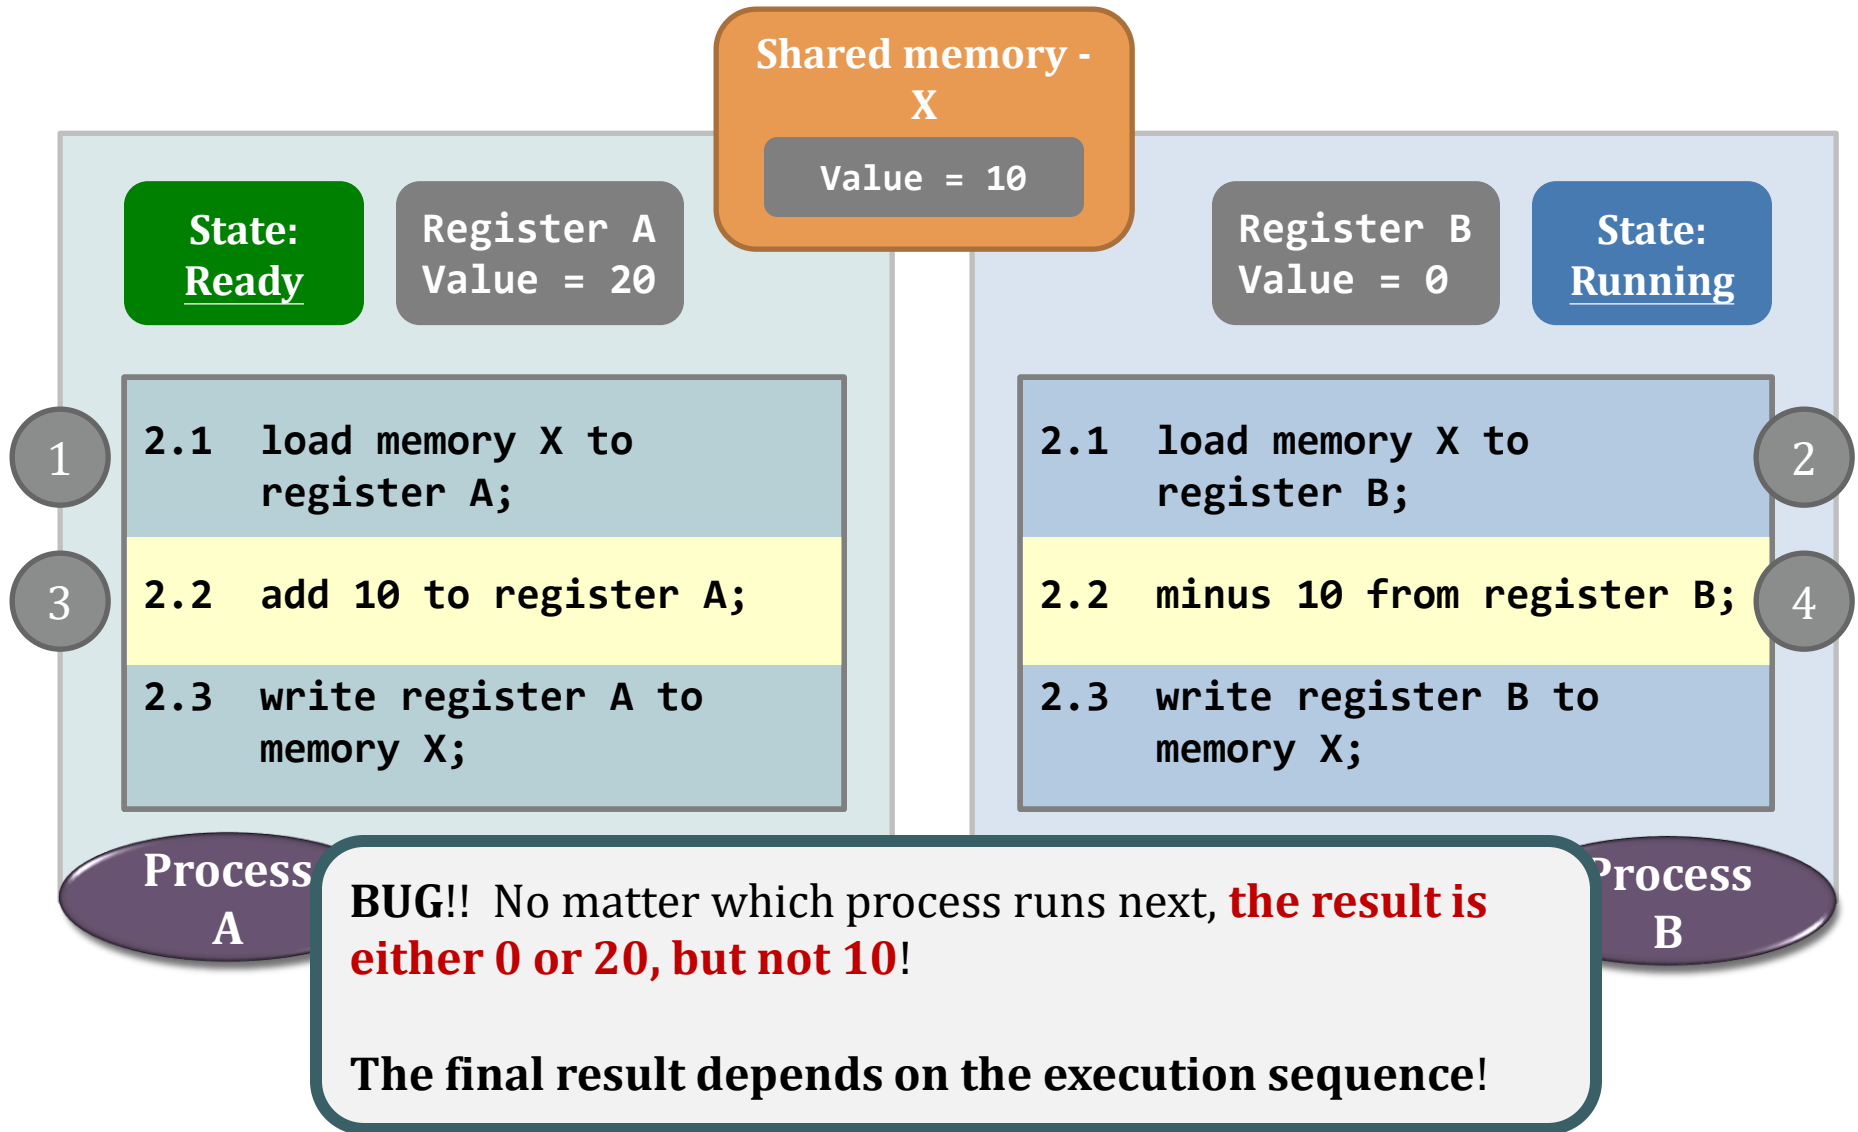
Problem arise...



Problem arise...



Problem arise...




Race condition

- ◆ The above scenario is called the **race condition**.
 - ◆ May happen whenever “**shared object**” + “**multiple processes**” + “**concurrently**”
- ◆ A **race condition** means
 - ◆ the outcome of an execution depends on a particular order in which the shared resource is accessed.
- ◆ Remember: race condition is always a bad thing and debugging race condition is a **nightmare**!
 - ◆ It may end up ...
 - ◆ 99% of the executions are fine.
 - ◆ 1% of the executions are problematic.

Inter-process communication (IPC)

- What, why, and how?
- The problem: race condition.
- **How to resolve race condition on a shared object?**
 - **Mutual Exclusion**

Mutual Exclusion – the cure




How to resolve race condition?

Solution: **mutual exclusion**

When I'm playing with the shared memory, no one could touch it.

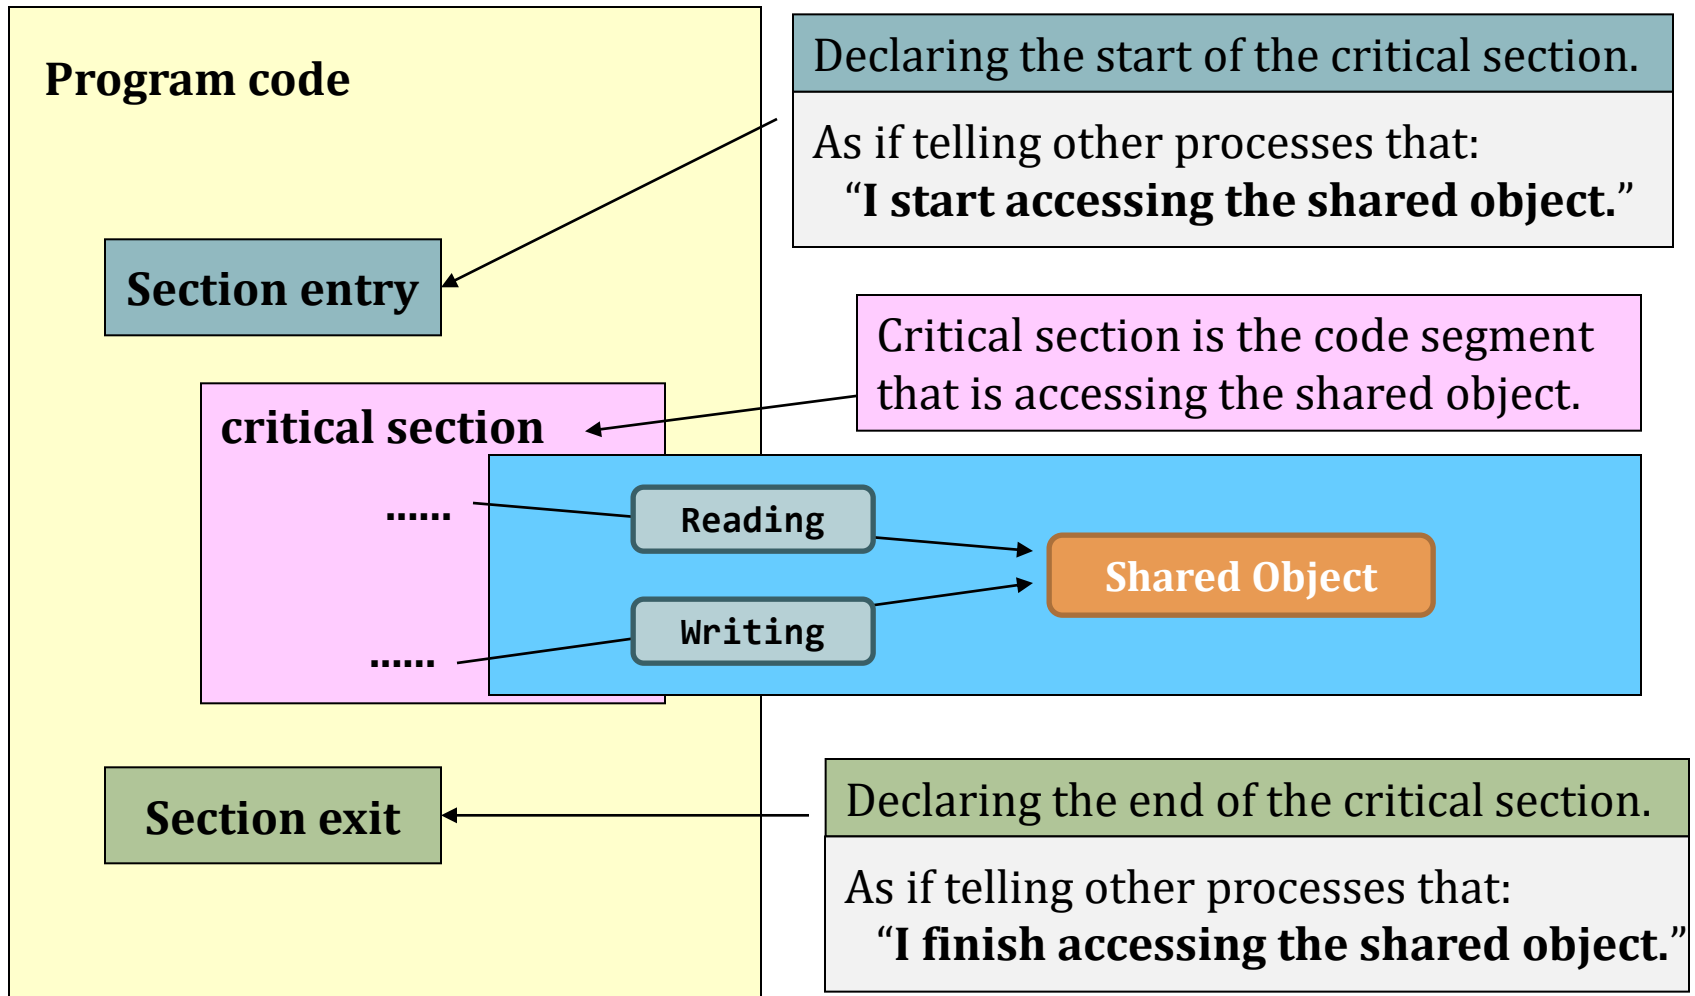
A set of processes would not have the problem of race condition *if **mutual exclusion is guaranteed***.



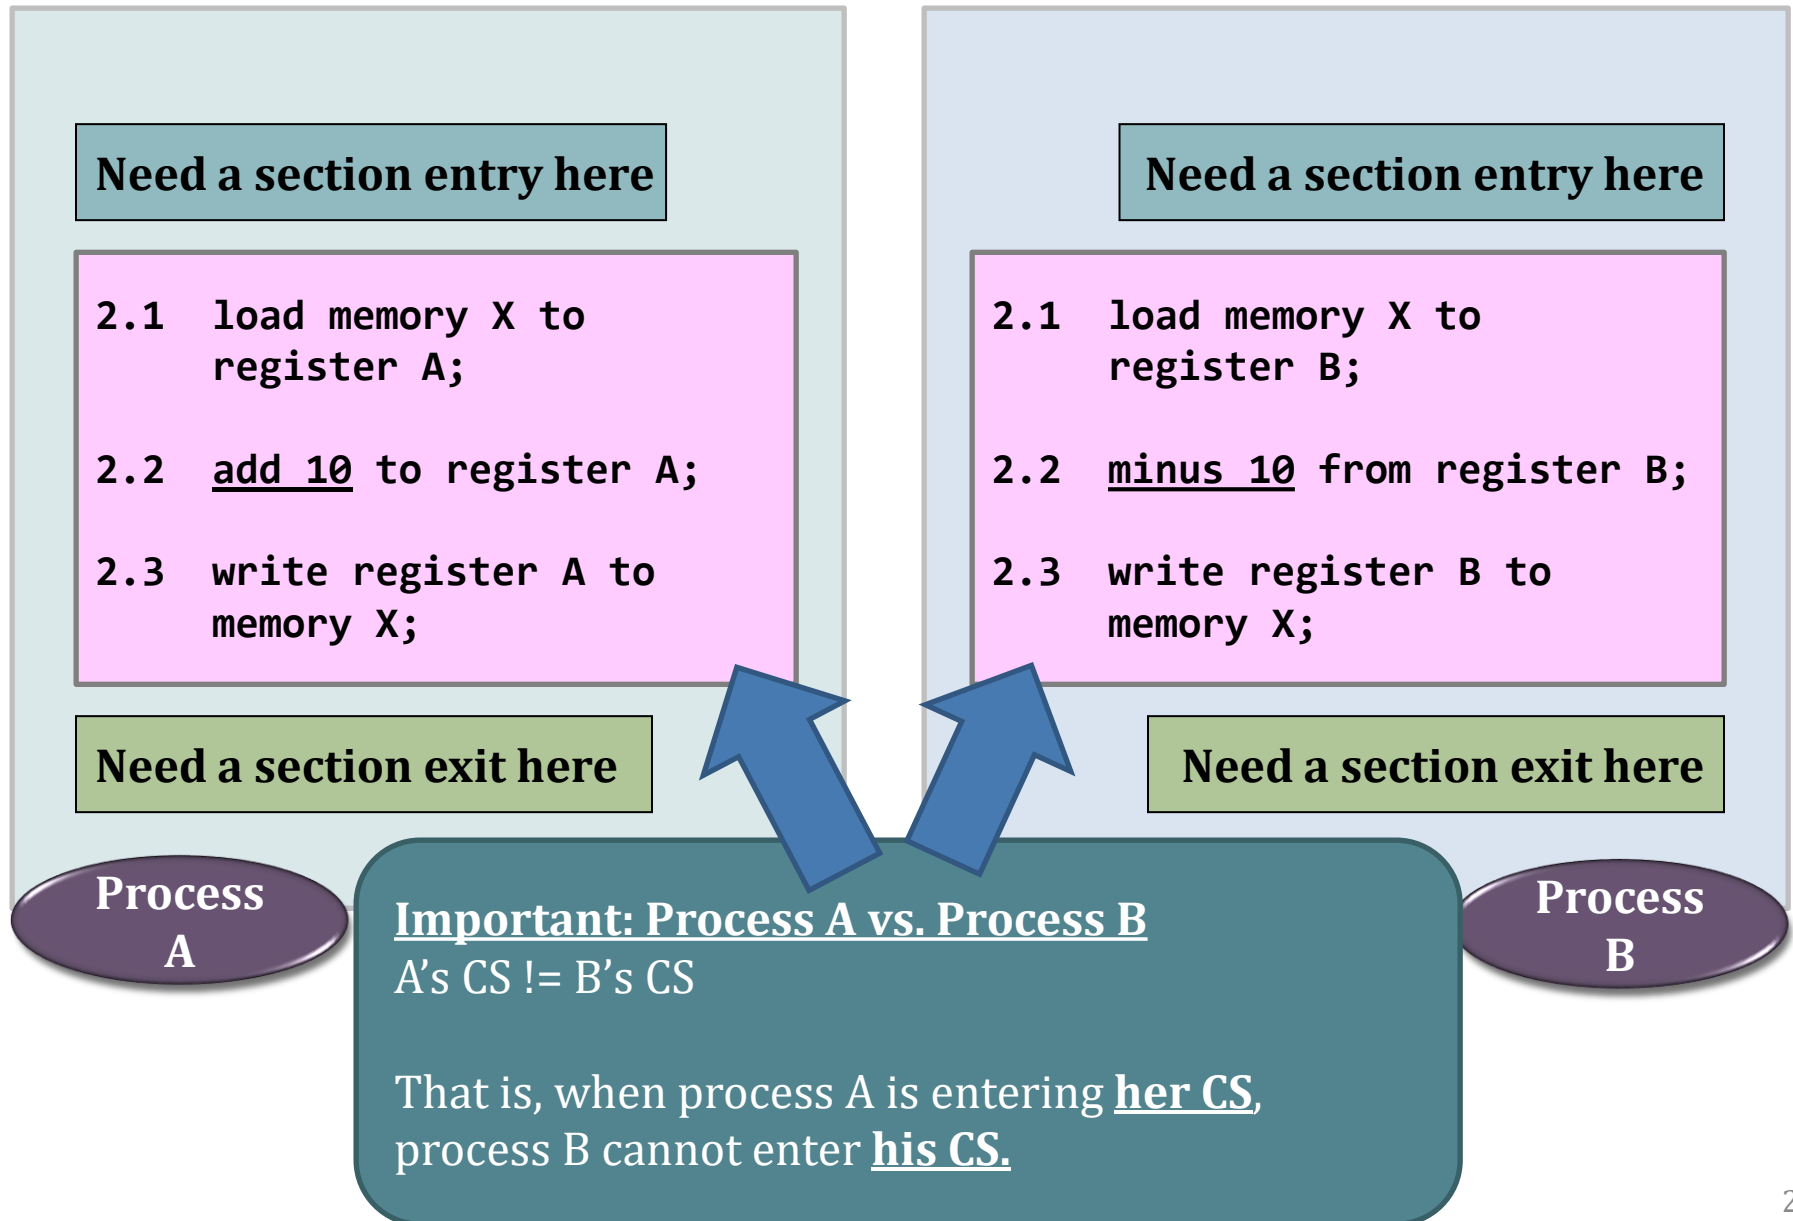
Solution – Mutual exclusion

- ◆ Shared object is still sharable, but
- ◆ Not to share the “shared object” at the same time
- ◆ Share the “shared object” one by one

Critical Section – the realization



Critical Section (CS) – the realization



What's really matter is the section entry/exit



Summary

◆ **Race condition**

- ◆ happens when programs accessing a shared object
- ◆ The outcome of the computation **totally depends on the execution sequences** of the processes involved.

◆ **Mutual exclusion** is a requirement.

- ◆ If it could be achieved, then the problem of the race condition would be gone.

Summary

- ◆ **A critical section** is the code segment that access shared objects.
- ◆ Critical section should be **as tight as possible**.
 - ◆ Well, you can set the entire code of a program to be a big critical section.
 - ◆ But, the program will have a very high chance to block other processes or to be blocked by other processes.
- ◆ Note that **one critical section** can be designed for **accessing more than one shared objects**.

Summary

- ◆ **Implementing section entry and exit** is a challenge.
 - ◆ The entry and the exit are **the core parts that guarantee mutual exclusion**, but not the critical section.
 - ◆ Unless they are correctly implemented, race condition would appear.
- ◆ **Mutual exclusion hinders the performance of parallel computations.**

Entry and exit implementation - requirements

◆ Requirement #1. **Mutual Exclusion**

- ◆ No two processes could be simultaneously go inside their **own** critical sections.

◆ Requirement #2. **Bounded Waiting**

- ◆ One a process starts trying to enter her CS, there is a bound on the number of times other processes can enter theirs.

Entry and exit implementation - requirements

◆ Requirement #3. Progress

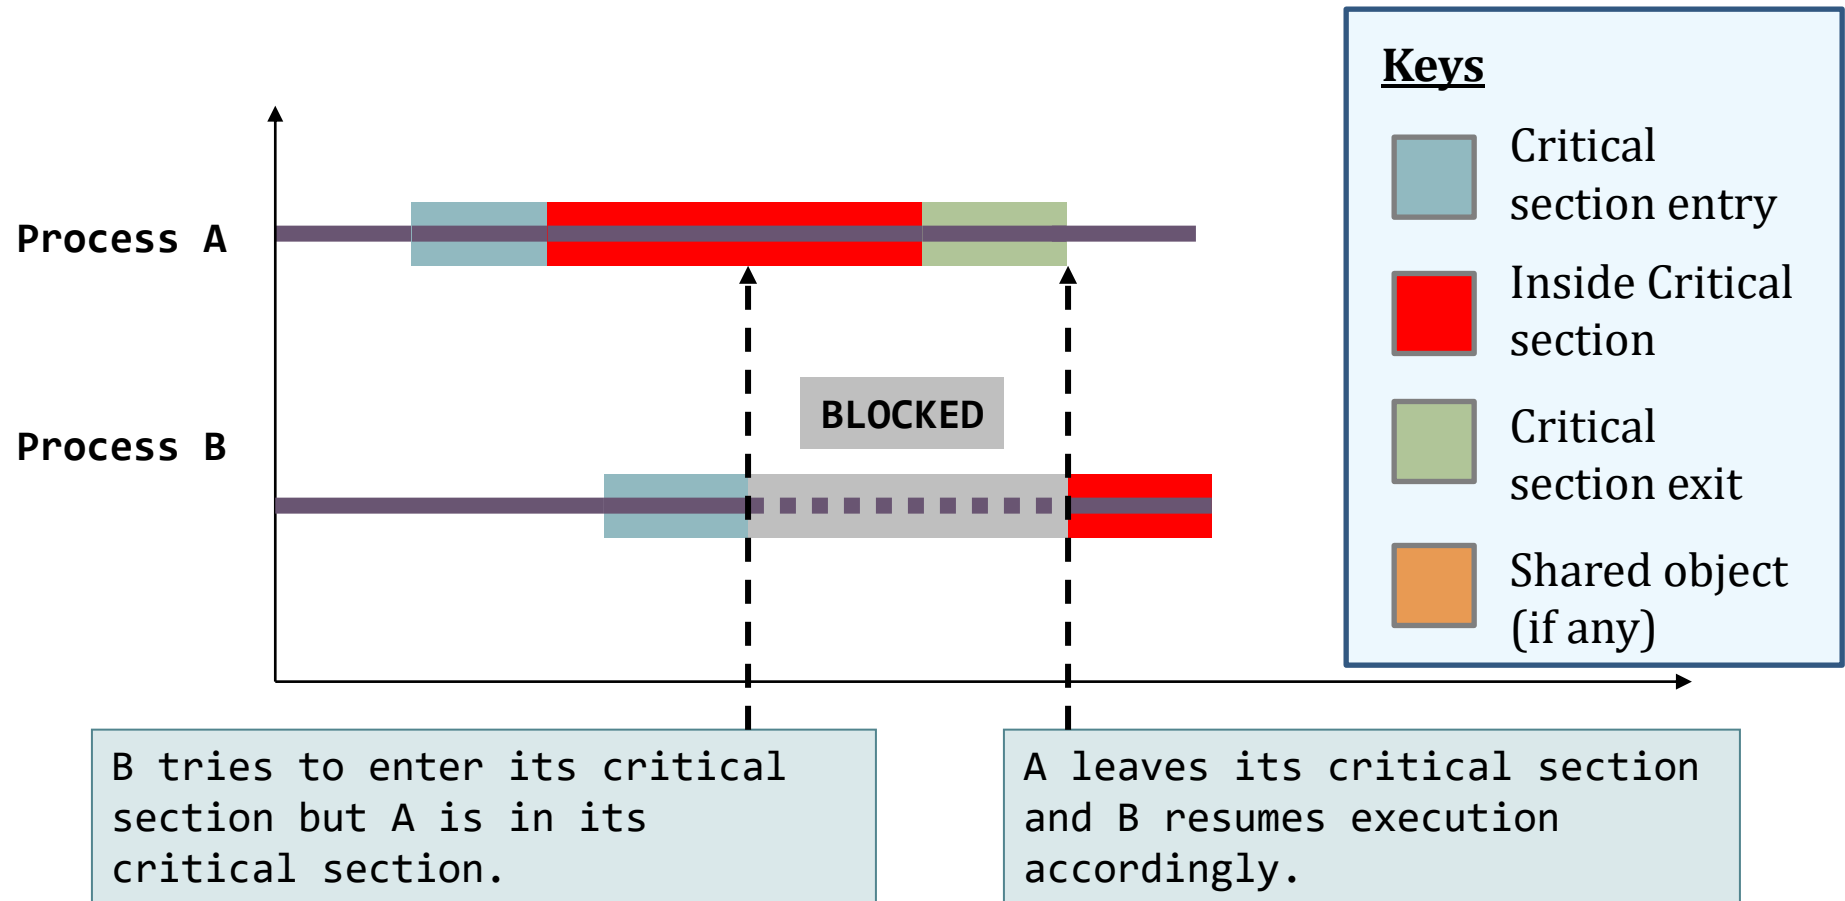
- ◆ Say no process currently in C.S.
- ◆ One of the processes trying to enter will eventually get in



Progress vs. bounded waiting

- ◆ If no process can enter C.S, do not have *progress*
- ◆ If A waiting to enter its C.S, while B repeated leaves and re-enters its C.S *and infinitum*
- ◆ A does not have *bounded waiting (but B is having progress)*

A typical mutual exclusion scenario



Achieving Mutual Exclusion

◆ Lock-based

◆ Spin-based lock

- ◆ E.g., use of “pthread_spin_lock”

- ◆ What is inside?

- ◆ **Basic spinning using 1 shared variable**
- ◆ **Spin using 2 shared variables + good algorithm (=Peterson's solution)**
- ◆ Spin using atomic instructions + smart algorithm (=Ticket, MCS algorithm, etc.)

◆ Sleep-based lock

- ◆ E.g., **POSIX semaphore**, pthread_mutex_lock

- ◆ What is inside?

- ◆ wait, yield(), atomic instructions + smart algorithm

◆ Lock-free

User-level synchronization

- ◆ You can treat that as:
 - ◆ You add some more global variables and write some while-loop smartly
 - ◆ Proper use of some library functions (e.g., pthread library)
 - ◆ You follow some synchronization “algorithm” to work on your case
 - ◆ ...

#0 – disabling interrupt for the whole CS

◆ Aim

- ◆ To **disable context switching** when the process is inside the critical section.

◆ Effect

- ◆ When a process is in its critical section, no other processes could be able to run.

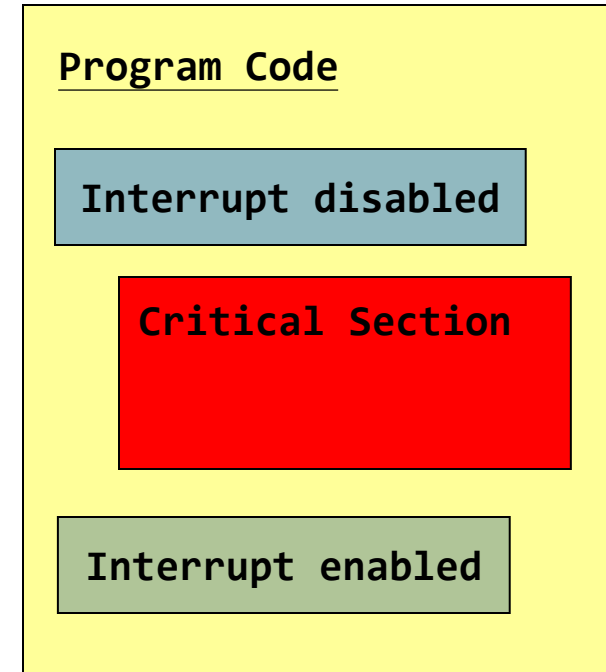
◆ Correctness?

◆ **Uni-core: Correct but not permissible**

- ◆ at userspace: what if one writes a CS that loops infinitely and the other process (e.g., the shell) never gets the context switch back to kill it?
- ◆ At kernel level: yes, correct and permissible

◆ **Multi-core: Incorrect**

- ◆ if there is another core modifying the shared object in the memory (unless you disable interrupts on all cores!!!!)



#1: Basic Spin lock (busy waiting)

◆ Aim.

- ◆ Loop on **yet another shared object**, **turn**, to detect the status of other processes

Shared object "turn"

initial value = 0


```
1 while (TRUE) {  
2   while( turn != 0 )  
3     ; /* busy waiting */  
4   critical_section();  
5   turn = 1;  
6   remainder_section();  
7 }
```

Process 0

```
1 while (TRUE) {  
2   while( turn != 1 )  
3     ; /* busy waiting */  
4   critical_section();  
5   turn = 0;  
6   remainder_section();  
7 }
```

Process1

#1: Basic Spin lock (busy waiting)



Shared object "turn"

initial Value = 0

```
1 while (TRUE) {  
2   while( turn != 0 )  
3     ; /* busy waiting */  
4   critical_section();  
5   turn = 1;  
6   remainder_section();  
7 }
```

Process 0

```
1 while (TRUE) {  
2   while( turn != 1 )  
3     ; /* busy waiting */  
4   critical_section();  
5   turn = 0;  
6   remainder_section();  
7 }
```

Process1

#1: Basic Spin lock (busy waiting)

- ◆ Correct
 - ◆ but it wastes CPU resources
 - ◆ OK for short waiting
 - ◆ Especially these days we have multi-core
 - ◆ Will not block other irrelevant processes a lot
 - ◆ Ok when spin-time < context-switch-overhead
- ◆ Impose a “strict alternation” order
 - ◆ Sometimes you give me my turn but I’m not ready to enter CS yet
 - ◆ Then you have to wait long

#1: Basic Spin lock (busy waiting)

- ◆ You can wrap them as lock() and unlock() functions
- ◆ In fact, some nice people wrap their super efficient implementation of the spinlock concept as pthread_mutex_lock() and pthread_mutex_unlock() functions

```
1 while (TRUE) {  
2     while( turn != 0 )  
3         ; /* busy waiting */  
4     critical_section();  
5     turn = 1;  
6     remainder_section();  
7 }
```

```
1 while (TRUE) {  
2     lock();  
4     critical_section();  
5     unlock();  
6     remainder_section();  
7 }
```

#1: Basic Spin lock violates *progress*

- ◆ Consider the following sequence:
 - ◆ Process0 leaves `cs()`, set `turn=1`
 - ◆ Process1 enters `cs()`, leaves `cs()`,
 - ◆ set `turn=0`, work on `remainder_section-slow()`
 - ◆ Process0 loops back and enters `cs()` again, leaves `cs()`, set `turn=1`
 - ◆ Process0 finishes its `remainder_section()`, go back to top of the loop
 - ◆ It can't enter its `cs()` (as `turn=1`)
 - ◆ That is, process0 gets blocked, but Process1 is outside its `cs()`, it is at its `remainder_section-slow()`

```
1 while (TRUE) {  
2   while( turn != 0 )  
3     ; /* busy waiting */  
4   cs();  
5   turn = 1;  
6   remainder_section();  
7 }
```

Process 0

```
1 while (TRUE) {  
2   while( turn != 1 )  
3     ; /* busy waiting */  
4   cs();  
5   turn = 0;  
6   remainder_section_slow ();  
7 }
```

Process 1

#2: Spin Smarter (by Peterson's solution)

◆ Highlight:

- ◆ Use one more extra shared object: **interested**
 - ◆ If I don't show interest
 - ◇ I let you **all** go
 - ◆ If we both show interest
 - ◇ Take turns

Shared objects:

- turn &
- "interested[2]"

#2: Spin Smarter (by Peterson's solution)

```
1  int turn;                                /* who is last enter cs */
2  int interested[2] = {FALSE,FALSE}; /* express interest to enter cs*/
3
4  void lock( int process ) { /* process is 0 or 1 */
5      int other;                /* number of the other process */
6      other = 1-process;        /* other is 1 or 0 */
7      interested[process] = TRUE; /* express interest */
8      turn = process;
9      while ( turn == process &&
              interested[other] == TRUE )
10         ; /* busy waiting */
11 }
12
13 void unlock( int process ) { /* process: who is leaving */
14     interested[process] = FALSE; /* I just left critical region */
15 }
```


#2: Spin Smarter (by Peterson's solution)

```
1  int turn;
2  int interested[2] = {FALSE,FALSE};
3
4  void lock( int process ) {
5      int other;
6      other = 1-process;
7      interested[process] = TRUE;
8      turn = process;
9      while ( turn == process &&
              interested[other] == TRUE )
10         ;    /* busy waiting */
11 }
12
13 void unlock( int process ) {
14     interested[process] = FALSE;
15 }
```

Express interest to enter CS



If others not show interest,
I can always go ahead

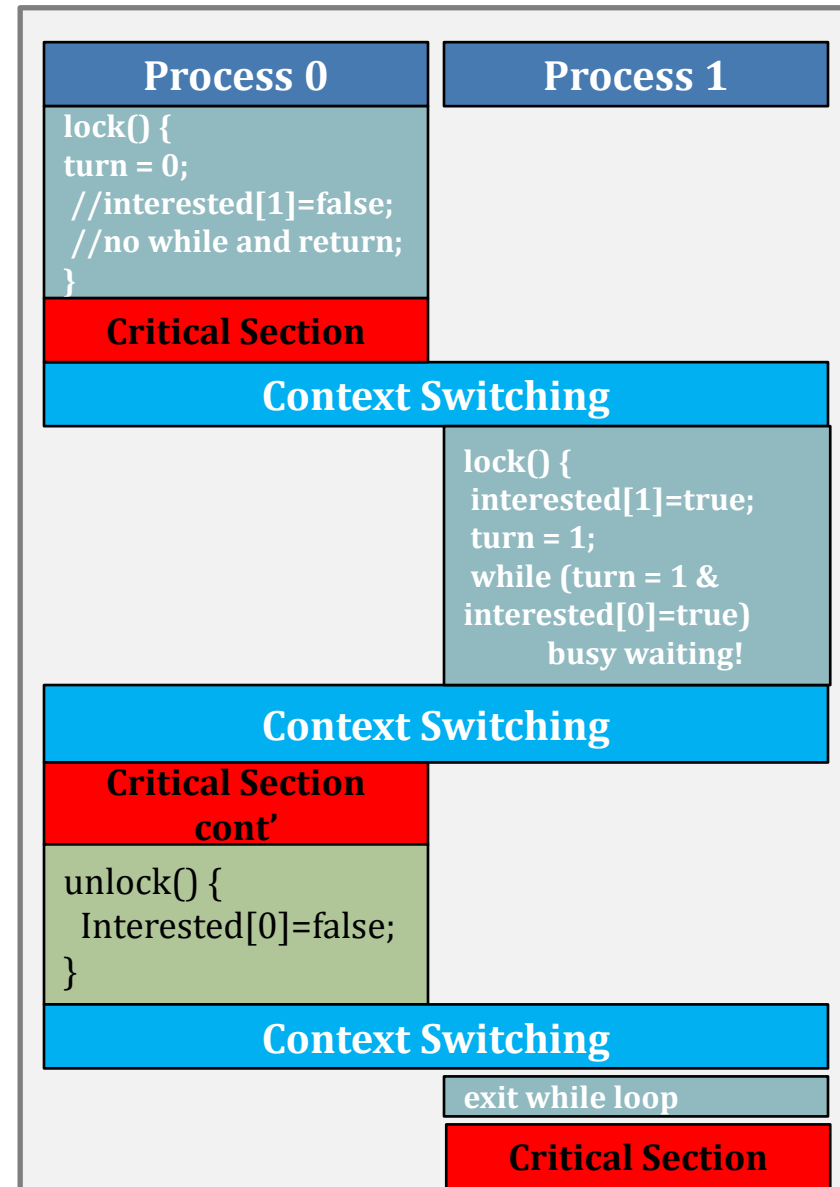


#2: Spin Smarter (by Peterson's solution)

```

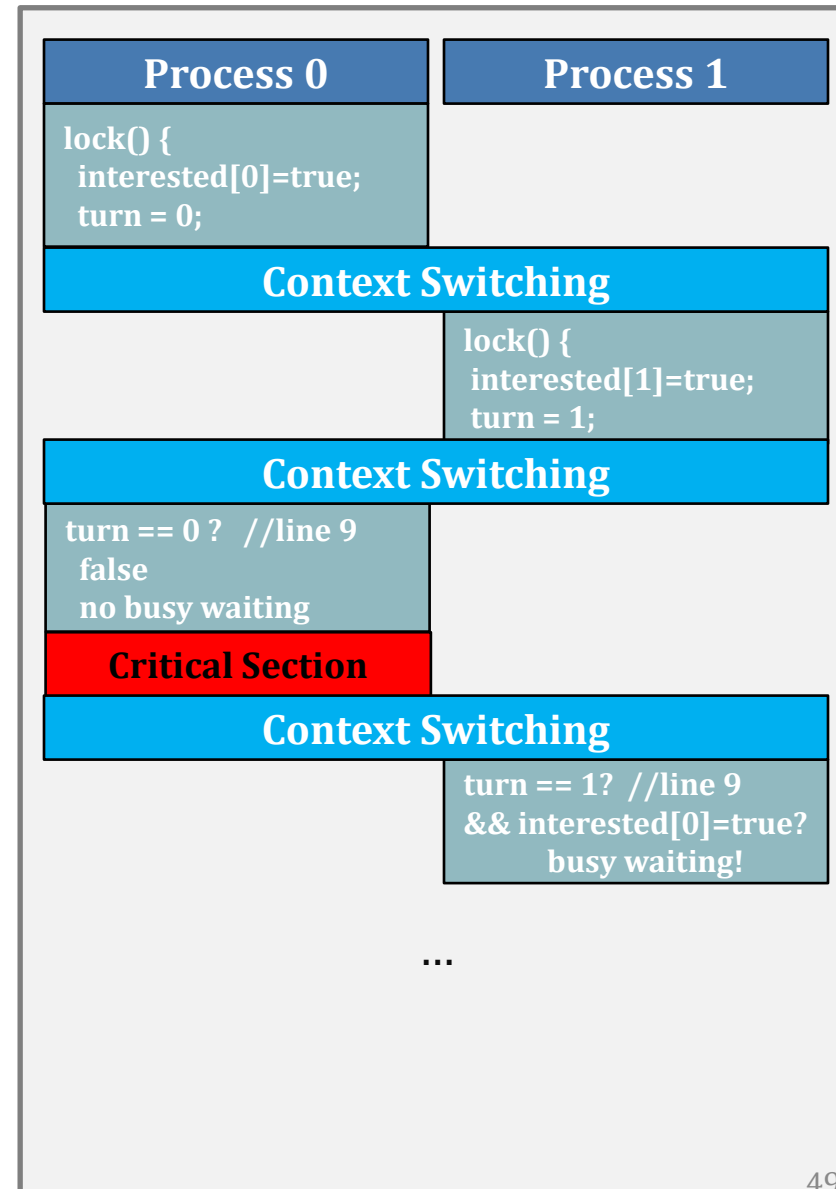
1  int turn;
2  int interested[2] = {FALSE,FALSE};
3
4  void lock( int process ) {
5      int other;
6      other = 1-process;
7      interested[process] = TRUE;
8      turn = process;
9      while ( turn == process &&
              interested[other] == TRUE )
10         ;    /* busy waiting */
11 }
12
13 void unlock( int process ) {
14     interested[process] = FALSE;
15 }

```



#2: Spin Smarter (by Peterson's solution) (another case)

```
1  int turn;
2  int interested[2] = {FALSE,FALSE};
3
4  void lock( int process ) {
5      int other;
6      other = 1-process;
7      interested[process] = TRUE;
8      turn = process;
9      while ( turn == process &&
              interested[other] == TRUE )
10         ;    /* busy waiting */
11 }
12
13 void unlock( int process ) {
14     interested[process] = FALSE;
15 }
```



Spin Smarter (by Peterson's solution)

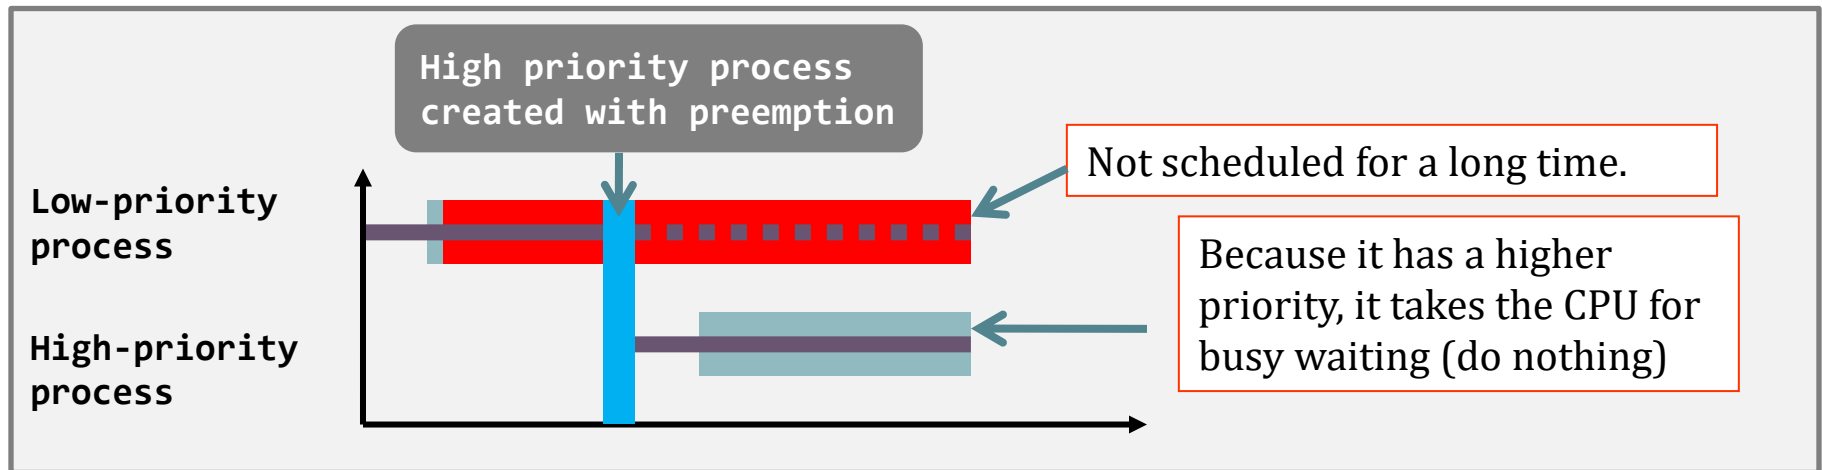
- ◆ = Busy waiting
 - + shared variable **turn** for mutual exclusion
 - + shared variables **interest** to resolve strict alternation
- ◆ Wikipedia:
 - ◆ *"It satisfies the three essential criteria to solve the critical section problem, provided that changes to the variables turn, interest[0], and interest[1] propagate immediately and atomically."*
- ◆ Suffer from **priority inversion problem**

Does it work for >2 processes?

https://en.wikipedia.org/wiki/Peterson's_algorithm

Peterson spinlock suffers from Priority Inversion

- ◆ Priority/Preemptive Scheduling (Linux, Windows... all OS...)
 - ◆ A low priority process **L** is inside the critical region, but ...
 - ◆ A high priority process **H** gets the CPU and wants to enter the critical region.
 - ◆ But **H** can not **lock** (because **L** has not **unlock**)
 - ◆ So, **H** gets the CPU to do nothing but spinning



#3: Sleep-based lock: Semaphore

- ◆ Semaphore is just a struct
 - ◆ Include
 - ◆ an integer that counts the # of resources available
 - ◇ Can do more than solving mutual exclusion
 - ◆ a wait-list
- ◆ The trick is still the section entry/exit function implementation
 - ◆ **Need to interact with scheduler (must involve kernel, e.g., syscall)**
 - ◆ **Implement uninterruptable section entry/exit**
 - ◆ Section entry/exit function are **short**
 - ◇ **Compared with Implementation #0 (uninterruptable throughout the whole CS)**



Semaphore **logical** view

```
typedef struct {  
    int value;  
    list process_id;  
} semaphore;
```

Section Entry: sem_wait()

```
1 void sem_wait(semaphore *s) {  
2     disable_interrupt();  
3     *s = *s - 1;  
4     if ( *s < 0 ) {  
5         enable_interrupt();  
6         sleep();  
7         disable_interrupt();  
8     }  
9     enable_interrupt();  
10 }
```

Initialize $s = 1$

“sem_wait(s)”

- I wait until I get **an s**
(i.e., wait(s) only returns when I get **an s**)
- Implementation:
of $s--$;
sleep if # of $s < 0$;

Important 1
 s can be a plural

Important 2

This wait is different from parent's folk wait(child). When programming, it is sem_wait()

“sem_post(s)”

- I notify the others that one s is added
- Implementation:
of $s++$;
If someone is waiting s , wakeup one of them

Section Exit: sem_post()

```
1 void sem_post(semaphore *s) {  
2     disable_interrupt();  
3     *s = *s + 1;  
4     if ( *s <= 0 )  
5         wakeup();  
6     enable_interrupt();  
7 }
```

Example

Process
1234

Sem_wait(X)

Assuming someone else (process **1357**) has already taken the only one resource:

$X = 1$ (initial) $\Rightarrow X = 0$

Now, process **1234** arrives

Section Entry: sem_wait()

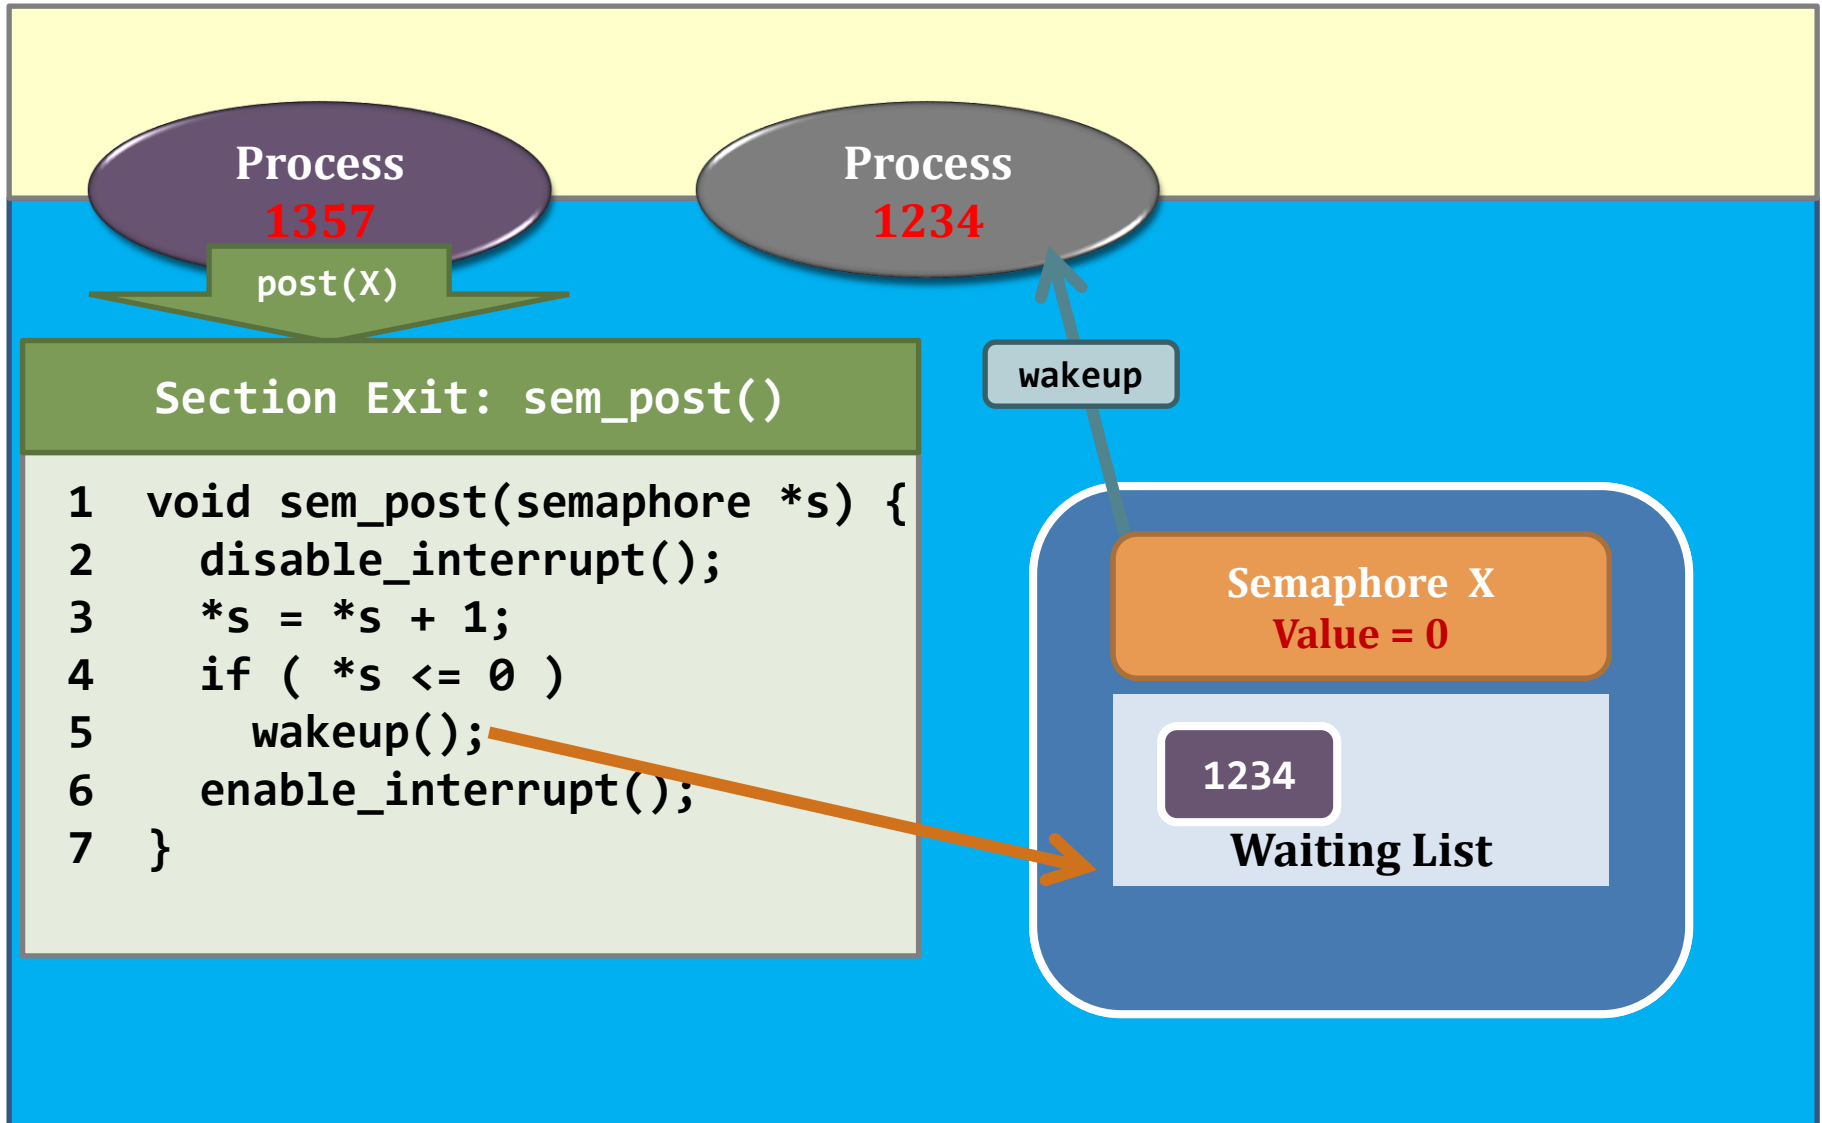
```
1 void sem_wait(semaphore *s){
2     disable_interrupt();
3     *s = *s - 1;
4     if ( *s < 0 ) {
5         enable_interrupt();
6         sleep();
7         disable_interrupt();
8     }
9     enable_interrupt();
10 }
```

Semaphore X
Value = -1

1234

Waiting List

Example



Example


Process
1234

Section Entry: sem_wait()

return

```
void sem_wait(sem. *s) {  
2   disable_interrupt();  
3   *s = *s - 1;  
4   if ( *s < 0 ) {  
5       enable_interrupt();  
6       sleep();  
7       disable_interrupt();  
8   }  
9   enable_interrupt();  
10 }
```


Using Semaphore (user-level)



```
semaphore *s; /* from kernel */  
*s = 1;      /* initial value */
```

```
1 while(TRUE) {  
2     sem_wait(s);  
3     critical_section();  
4     sem_post(s);  
5 }
```

entry

exit

Using Semaphore beyond mutual exclusion

- ◆ Can also be used as a synchronization primitive to solve IPC problems
 - ◆ E.g., make sure “ls” waits until “less” consumes things from the pipe (otherwise the pipe will overflow)



Which one is the shared object in this picture?

Achieving Mutual Exclusion

◆ Lock-based

◆ Spin-based lock

- ◆ E.g., use of “pthread_spin_lock”

- ◆ What is it inside?

- ◆ **Basic spinning using 1 shared variable**
- ◆ **Spin using 2 shared variables + good algorithm (=Peterson's solution)**
- ◆ Spin using atomic instructions + smart algorithm (=Ticket, MCS algorithm, etc.)

◆ Sleep-based lock

- ◆ E.g., **POSIX semaphore**, pthread_mutex_lock

- ◆ What is it inside?

- ◆ wait, yield(), atomic instructions + smart algorithm

◆ Lock-free

IPC / Synchronization problems

	Properties	Examples
Producer-Consumer Problem	Two classes of processes: <u>producer</u> and <u>consumer</u> ; At least one producer and one consumer. [Single-Object Synchronization]	Pipe, Named Pipe
Dining Philosopher Problem	They are all running the same program; At least two processes. [Multi-Object Synchronization]	Cross-road traffic control
Reader Writer Problem	Multiple reads, 1 write	...
...	<p>Named Pipe (a.k.a. FIFO in Linux)</p> <ul style="list-style-type: none"> - Like Shared File (so <u>multiple</u> producers and consumers; unlike pipe) - Like Pipe (unidirectional) - In-memory (unlike file) - Use like pipe (more restrictive; unlike shared memory) 	

Inter-process communication (IPC)

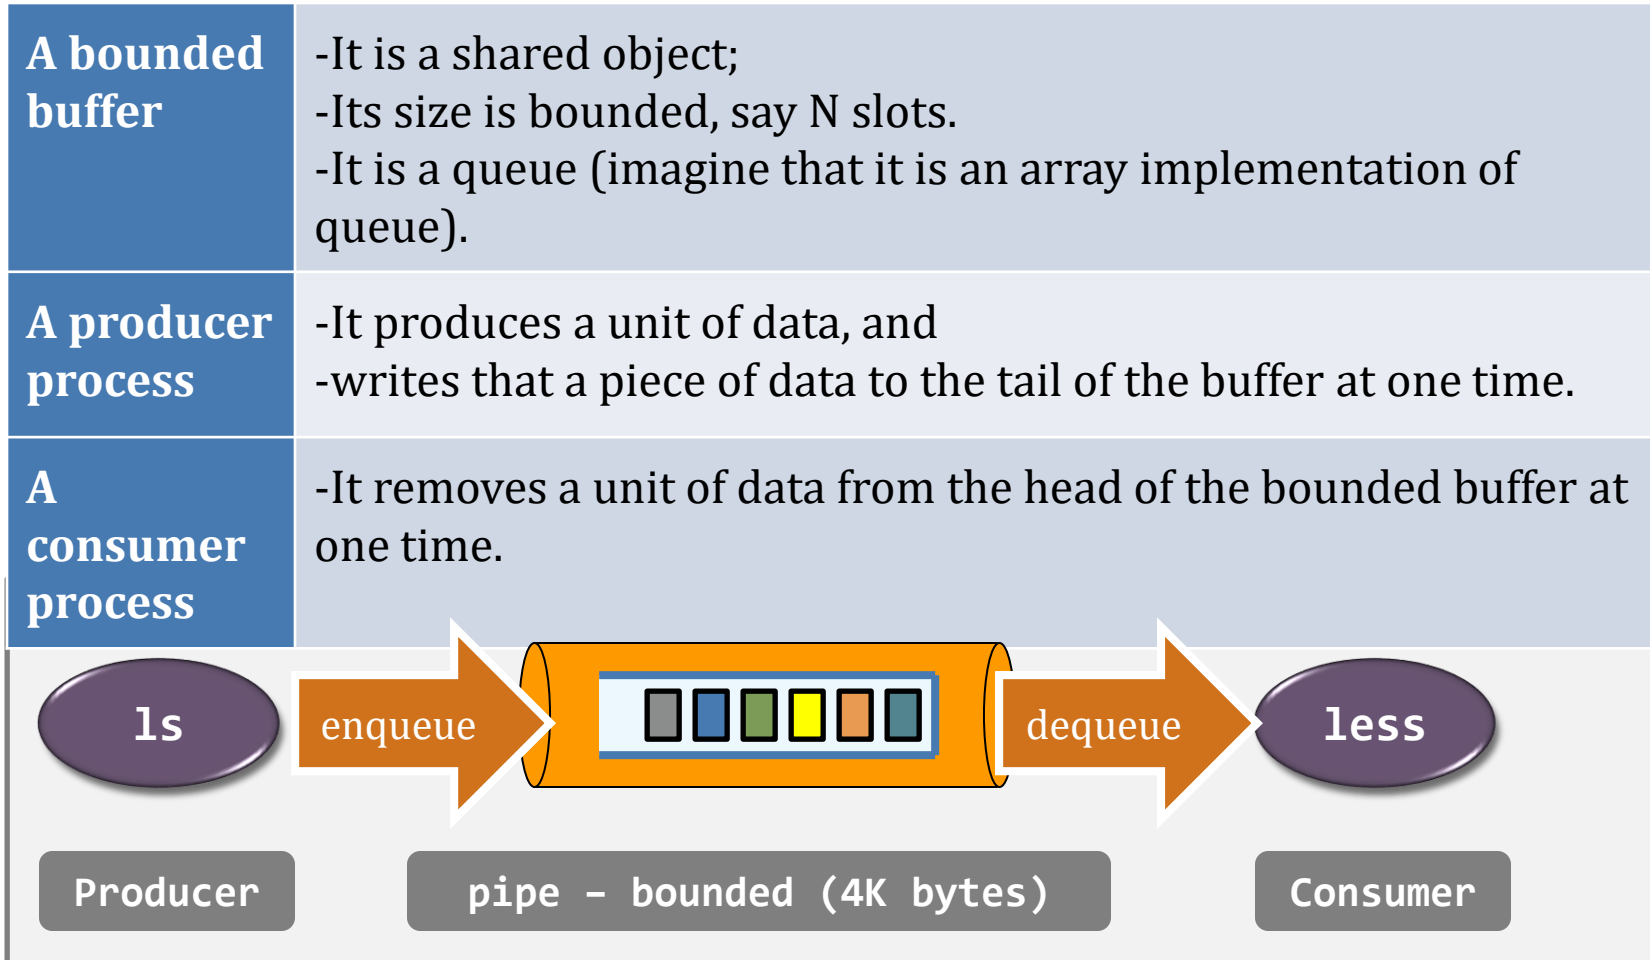
- Classic IPC problems.
- Producer-consumer problem.

In the following, we demonstrate how to use semaphore to solve some IPC problems.

Semaphore is not the only way. You might use (i) lock + condition variable, (ii) use X more shared variables (like Peterson's solution) directly,

Producer-consumer problem – introduction

- ❖ Also known as the **bounded-buffer problem**.
- ❖ Single-object synchronization



Producer-consumer problem – introduction

Requirement #1

When the **producer** wants to
(a) put a new item in the buffer, but
(b) **the buffer is already full...**

Then, **the producer should wait.**

The consumer should notify the producer after she has dequeued an item.

Requirement #2

When the **consumer** wants to
(a) consumes an item from the buffer, but
(b) **the buffer is empty...**

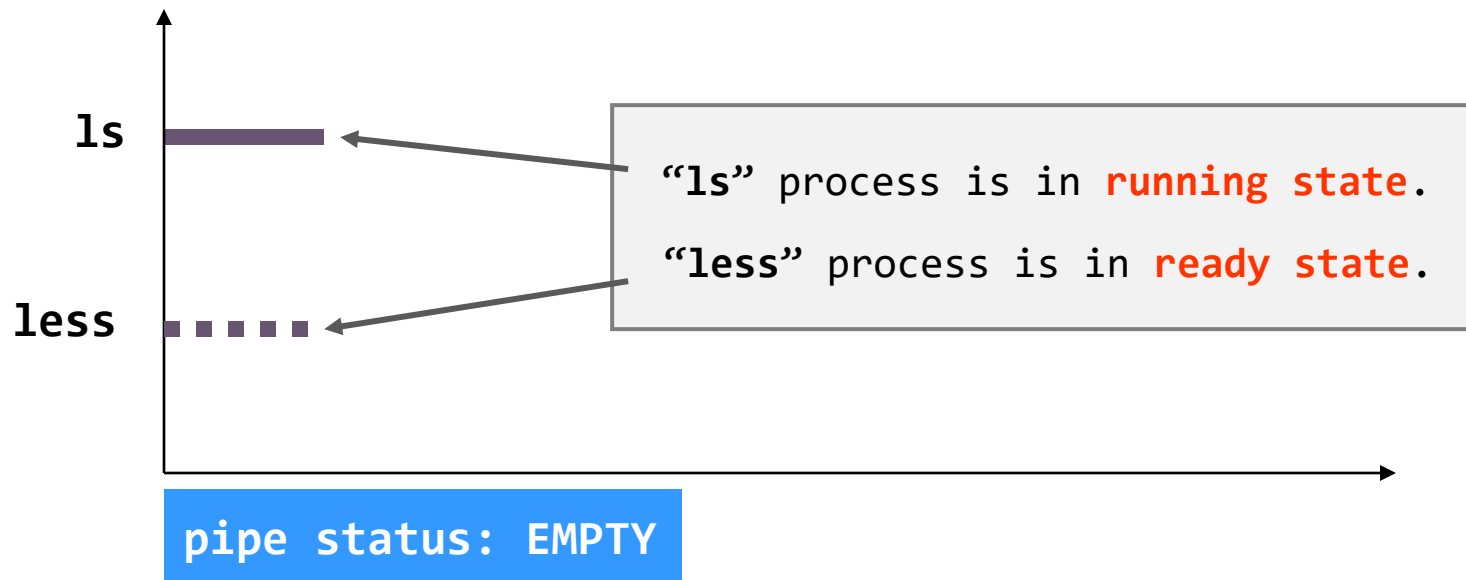
Then, **the consumer should wait.**

The producer should notify the consumer after she has enqueued an item.

Producer-consumer problem – pipe

Assumptions

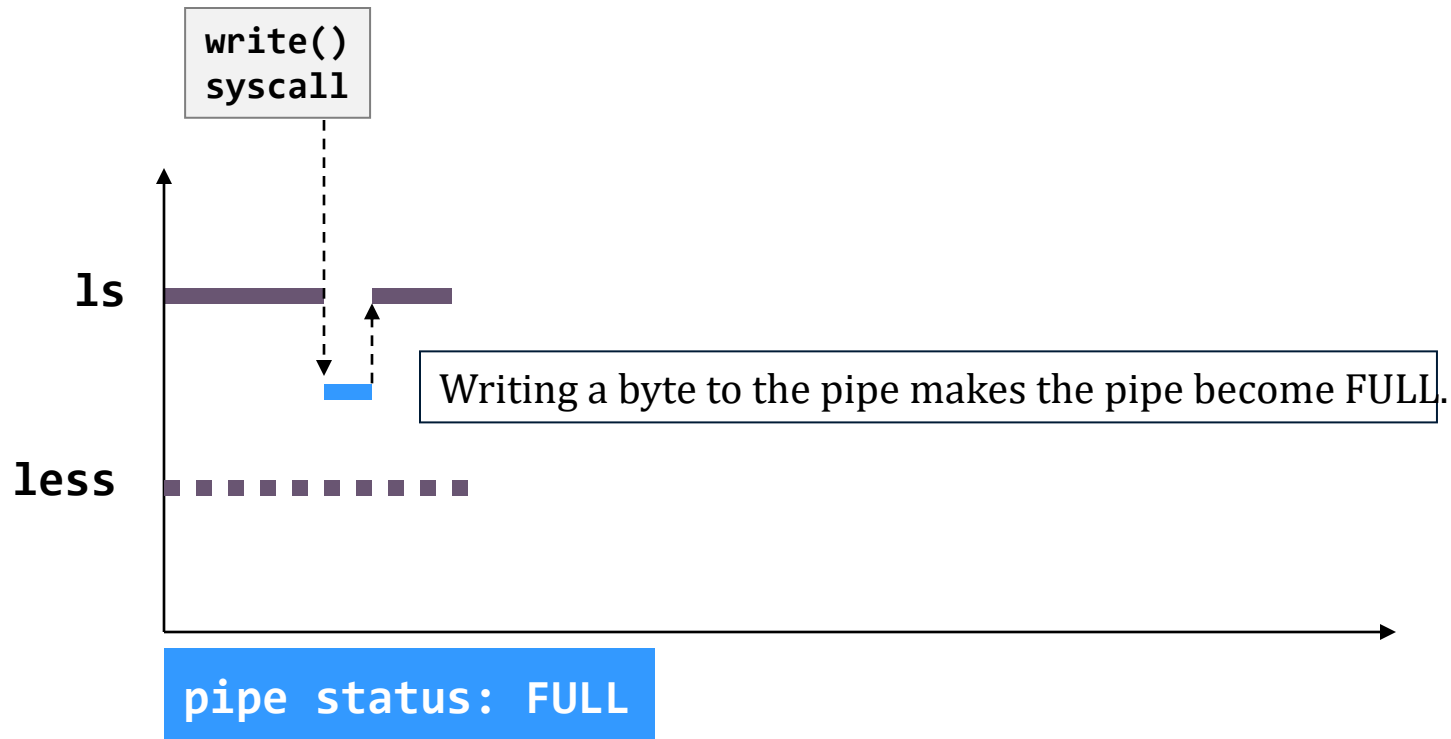
- The pipe is a queue of **1 byte** only!
- Each **write()** system call writes **1 byte** to the pipe.
- Each **read()** system call reads **1 byte** from the pipe.



Producer-consumer problem – pipe

Assumptions

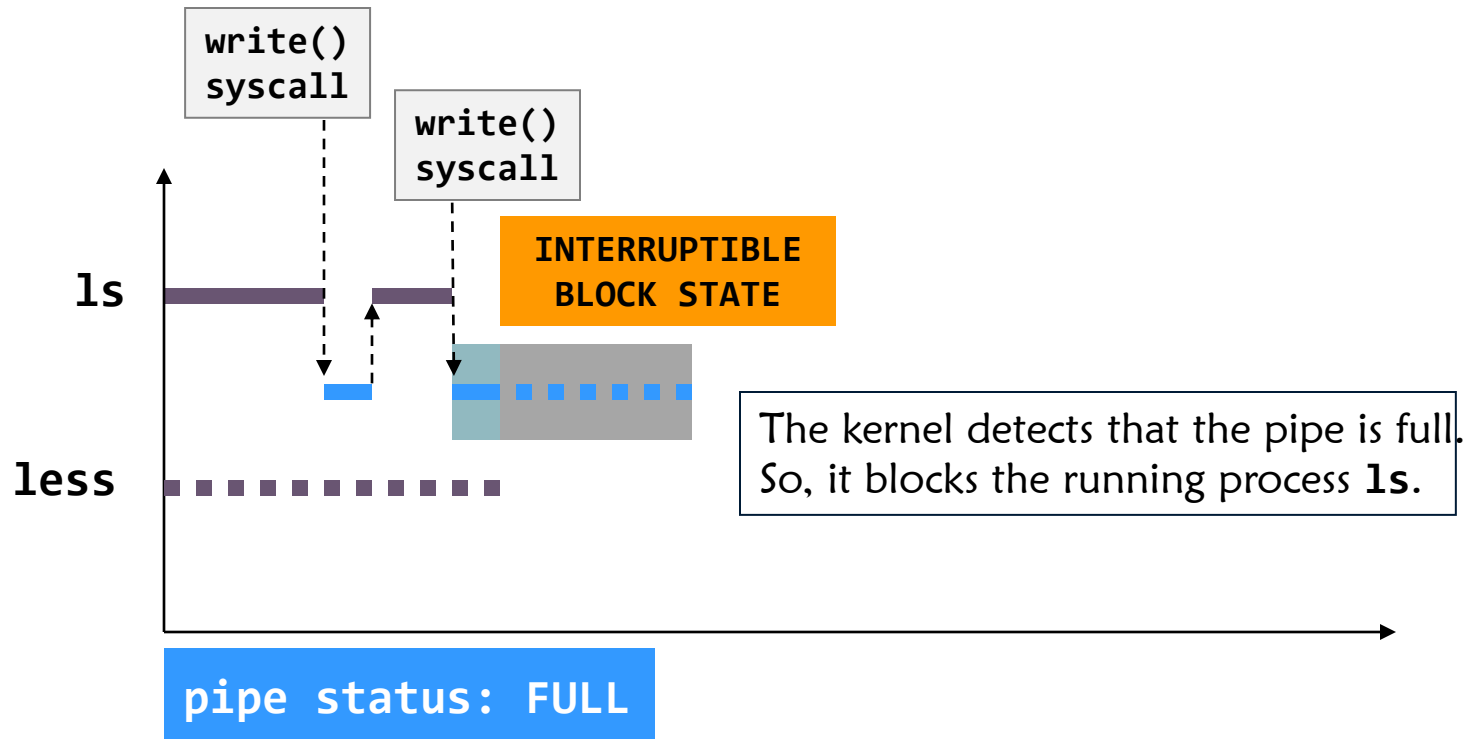
- The pipe is a queue of **1 byte** only!
- Each **write()** system call writes **1 byte** to the pipe.
- Each **read()** system call reads **1 byte** from the pipe.



Producer-consumer problem – pipe

Assumptions

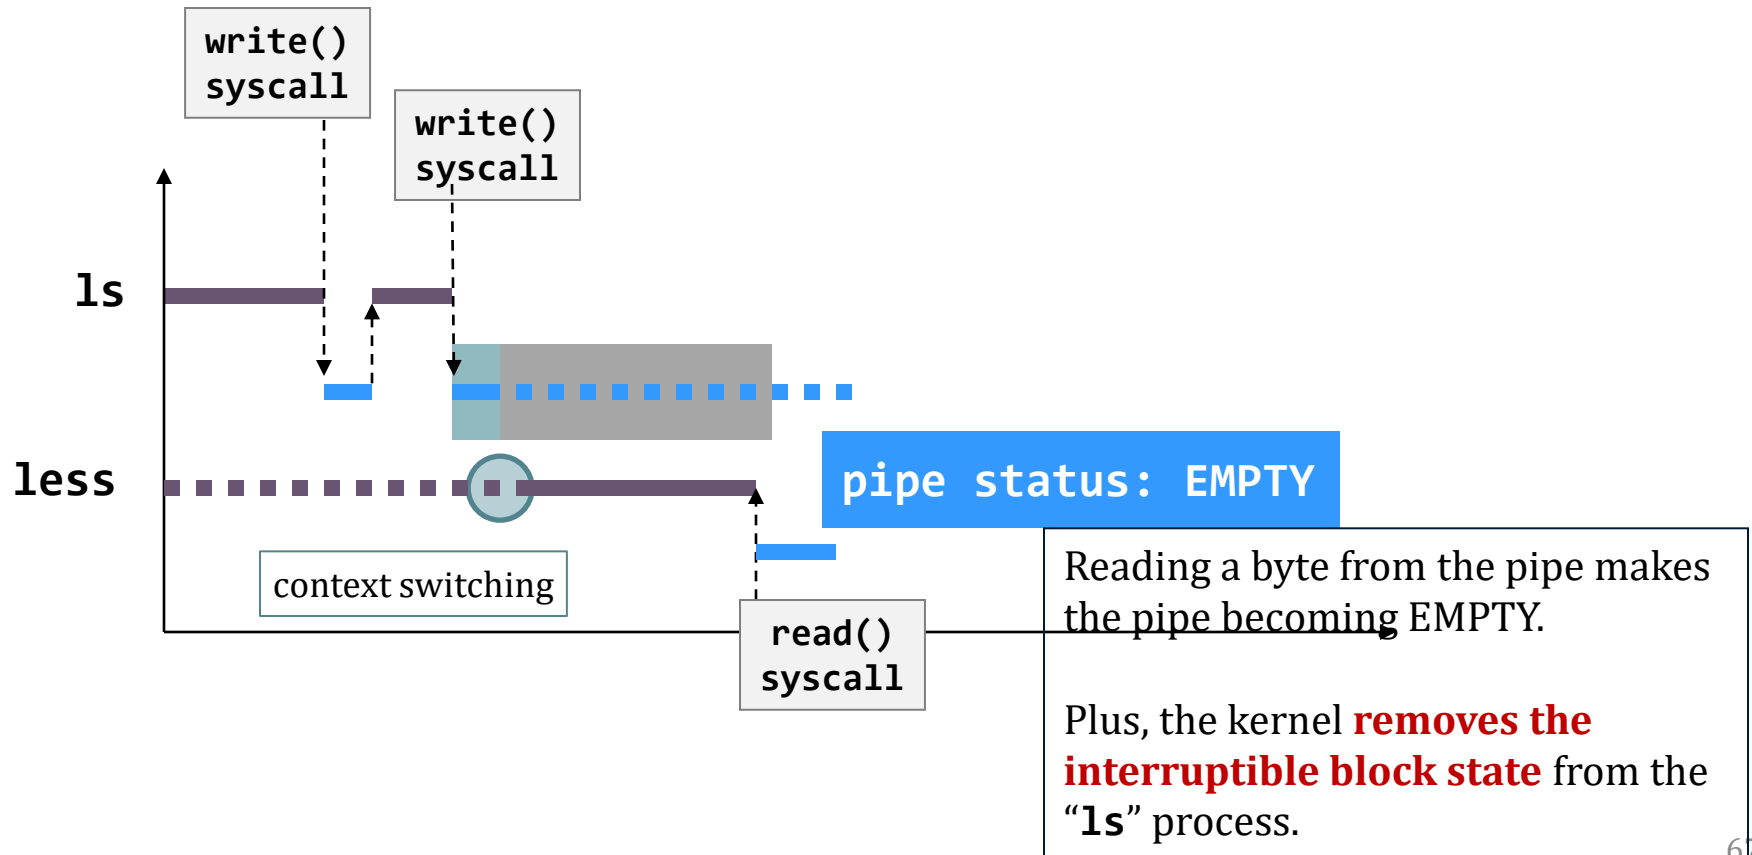
- The pipe is a queue of **1 byte** only!
- Each **write()** system call writes **1 byte** to the pipe.
- Each **read()** system call reads **1 byte** from the pipe.



Producer-consumer problem – pipe

Assumptions

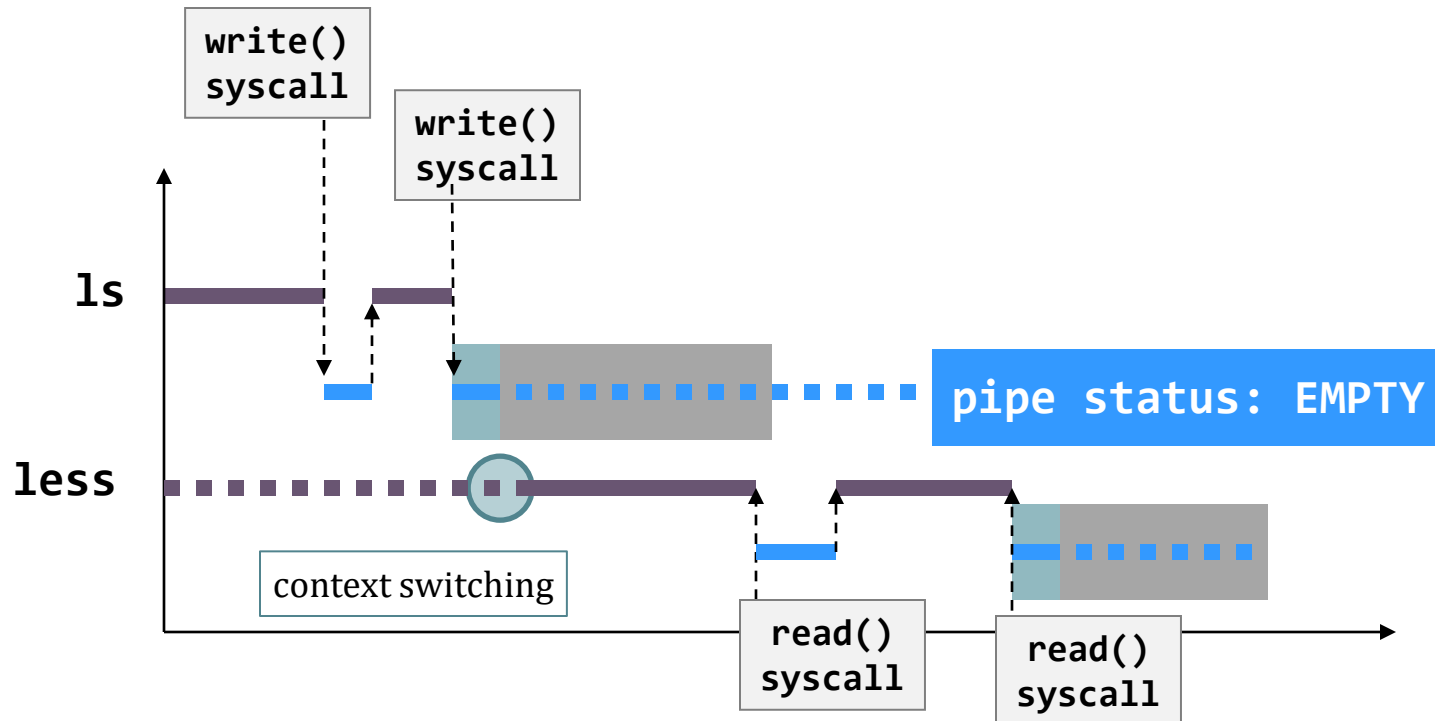
- The pipe is a queue of **1 byte** only!
- Each `write()` system call writes **1 byte** to the pipe.
- Each `read()` system call reads **1 byte** from the pipe.



Producer-consumer problem – pipe

Assumptions

- The pipe is a queue of **1 byte** only!
- Each **write()** system call writes **1 byte** to the pipe.
- Each **read()** system call reads **1 byte** from the pipe.

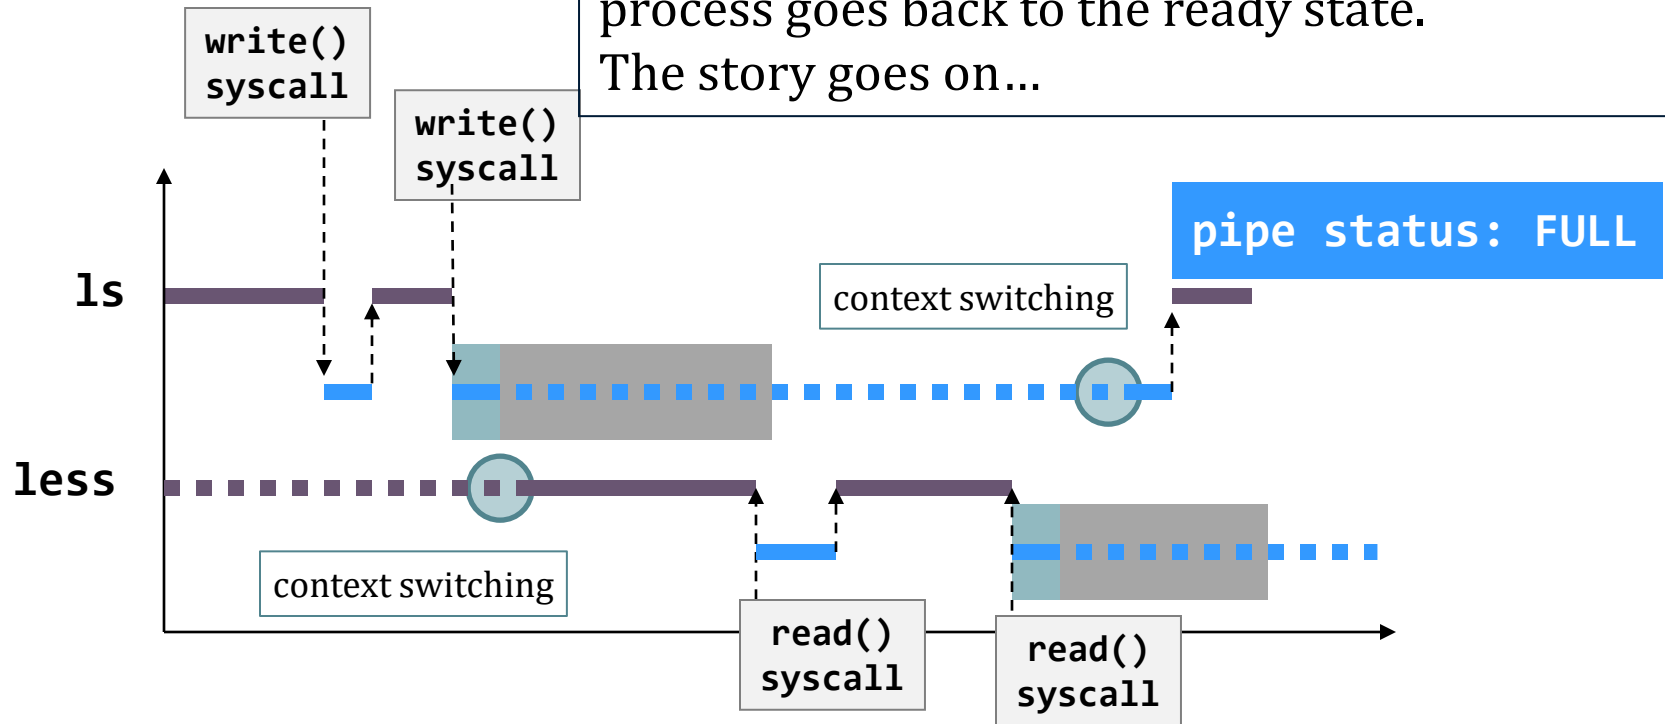


The kernel detects that the pipe is empty.
So, it blocks the running process **less**.

Producer-consumer problem – pipe


When “**ls**” process goes back to the running state, it **immediately writes a byte to the pipe so as to complete the execution of the write() system call**.

Then, the pipe is no longer empty, and the “**less**” process goes back to the ready state.
The story goes on...



Producer-consumer problem: semaphore

- ◆ The Producer-consumer problem is **more general** than the pipe story
 - ◆ Pipe cannot work with >1 producers/consumers
- ◆ The problem can be divided into two sub-problems.
 - ◆ Mutual exclusion.
 - ◆ The buffer is a shared object. **Mutual exclusion** is needed. Done by one binary semaphore
 - ◆ Synchronization.
 - ◆ Because the buffer's size is bounded, **coordination** is needed. Done by two semaphores
 - ◆ **Notify** the producer to stop producing when the buffer is **full**
 - ◆ In other words, **notify** the producer to produce when the buffer is NOT full
 - ◆ **Notify** the consumer to stop eating when the buffer is **empty**
 - ◆ In other words, **notify** the consumer to consume when the buffer is NOT empty



Producer-consumer problem: semaphore

Shared object

```
#define N 100
semaphore mutex = 1;
semaphore avail = N;
semaphore fill = 0;
```

Note

The size of the bounded buffer is “N”.

fill : number of occupied slots in buffer
avail: number of empty slots in buffer

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         wait(&avail);
7         wait(&mutex);
8         insert_item(item);
9         post(&mutex);
10        post(&fill);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         wait(&fill);
6         wait(&mutex);
7         item = remove_item();
8         post(&mutex);
9         post(&avail);
10        //consume the item;
11    }
12 }
```

Producer-consumer problem: semaphore

Note

6: (Producer) I wait for an **available** slot and acquire it if I can

10: (Producer) I **notify** the others that I have **filled** the buffer

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         wait(&avail);
7         wait(&mutex);
8         insert_item(item);
9         post(&mutex);
10        post(&fill);
11    }
12 }
```


Producer-consumer problem: semaphore

Note

6: (Producer) I wait for an **available** slot and acquire it if I can

10: (Producer) I **notify** the others that I have **filled** the buffer

Note

5: (Consumer) I wait for someone to **fill** up the buffer and proceed if I can

9: (Consumer) I **notify** the others that I have made the buffer with a new **available** slot

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         wait(&avail);
7         wait(&mutex);
8         insert_item(item);
9         post(&mutex);
10        post(&fill);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         wait(&fill);
6         wait(&mutex);
7         item = remove_item();
8         post(&mutex);
9         post(&avail);
10        //consume the item;
11    }
12 }
```

Producer-consumer problem – question #1

Necessary to use both “avail” and “fill”?

Let us try to remove semaphore fill?

Shared object

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore avail = N;
semaphore fill = 0;
```

Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6         wait(&avail);
7         wait(&mutex);
8         insert_item(item);
9         post(&mutex);
10        post(&fill);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         wait(&fill);
6         wait(&mutex);
7         item = remove_item();
8         post(&mutex);
9         post(&avail);
10        //consume the item;
11    }
12 }
```

Producer-consumer problem – question #1

Just view wait(avail) as -- resource?
Just view post(avail) as ++ resource?

wait s--
post s++

So,

- producer s-- by wait
 - consumer s++ by post
- Problem solved?

Producer function

```
1 void producer(void) {  
2     int item;  
3  
4     while(TRUE) {  
5         item = produce_item();  
6         wait(&avail);  
7         wait(&mutex);  
8         insert_item(item);  
9         post(&mutex);  
10        post(&fill);  
11    }  
12 }
```

Consumer Function

```
1 void consumer(void) {  
2     int item;  
3  
4     while(TRUE) {  
5         wait(&fill);  
6         wait(&mutex);  
7         item = remove_item();  
8         post(&mutex);  
9         post(&avail);  
10        //consume the item;  
11    }  
12 }
```

Producer-consumer problem – question #1

Just view wait(avail) as -- resource?
Just view post(avail) as ++ resource?

wait s--
post s++
So,

If consumer gets CPU first, it removes item from NULL

E R R O R

Producer function

```
1 void producer(void) {  
2     int item;  
3  
4     while(TRUE) {  
5         item = produce_item();  
6         wait(&avail);  
7         wait(&mutex);  
8         insert_item(item);  
9         post(&mutex);  
10        post(&fill);  
11    }  
12 }
```

Consumer Function

```
1 void consumer(void) {  
2     int item;  
3  
4     while(TRUE) {  
5         wait(&fill);  
6         wait(&mutex);  
7         item = remove_item();  
8         post(&mutex);  
9         post(&avail);  
10        //consume the item;  
11    }  
12 }
```

Producer-consumer problem – question #2

Question #2.

Can we swap Lines 6 & 7 of the producer?

Let us simulate what will happen with the modified code!

Shared object

```
#define N 100
semaphore mutex = 1;
semaphore avail = N;
semaphore fill = 0;
```

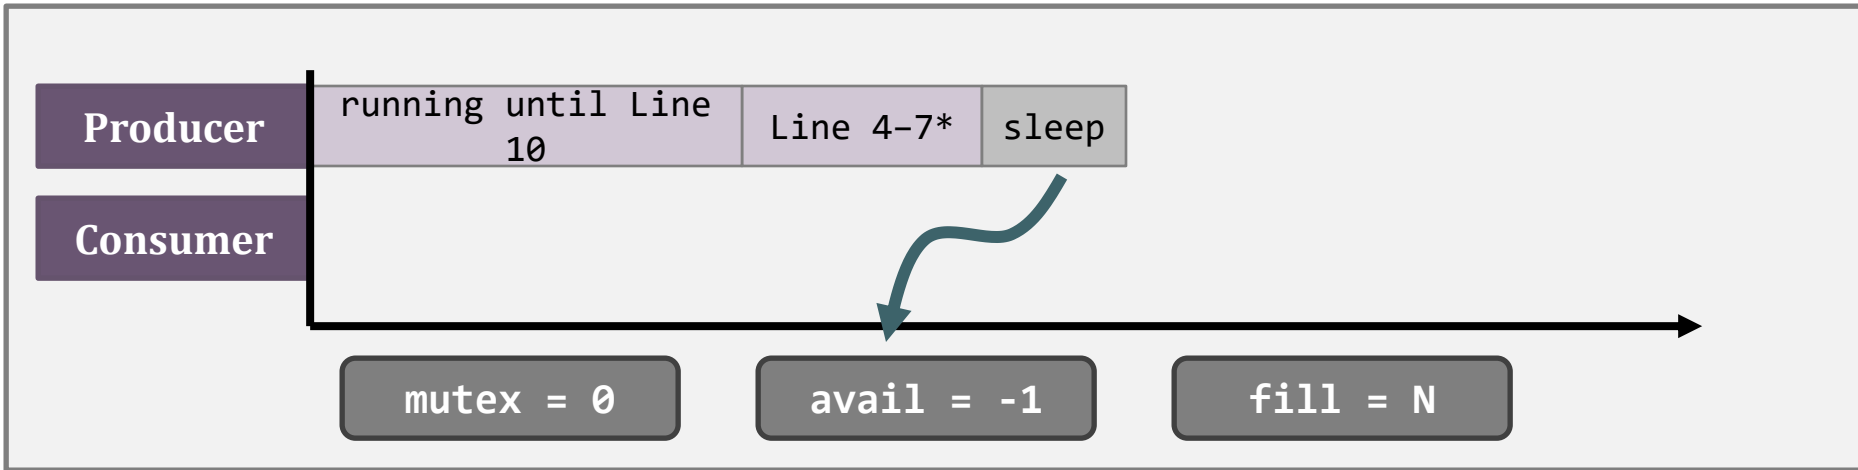
Producer function

```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6*      wait(&mutex);
7*      wait(&avail);
8         insert_item(item);
9         post(&mutex);
10        post(&fill);
11    }
12 }
```

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         wait(&fill);
6         wait(&mutex);
7         item = remove_item();
8         post(&mutex);
9         post(&avail);
10        //consume the item
11    }
12 }
```

Producer-consumer problem – question #2



Producer function

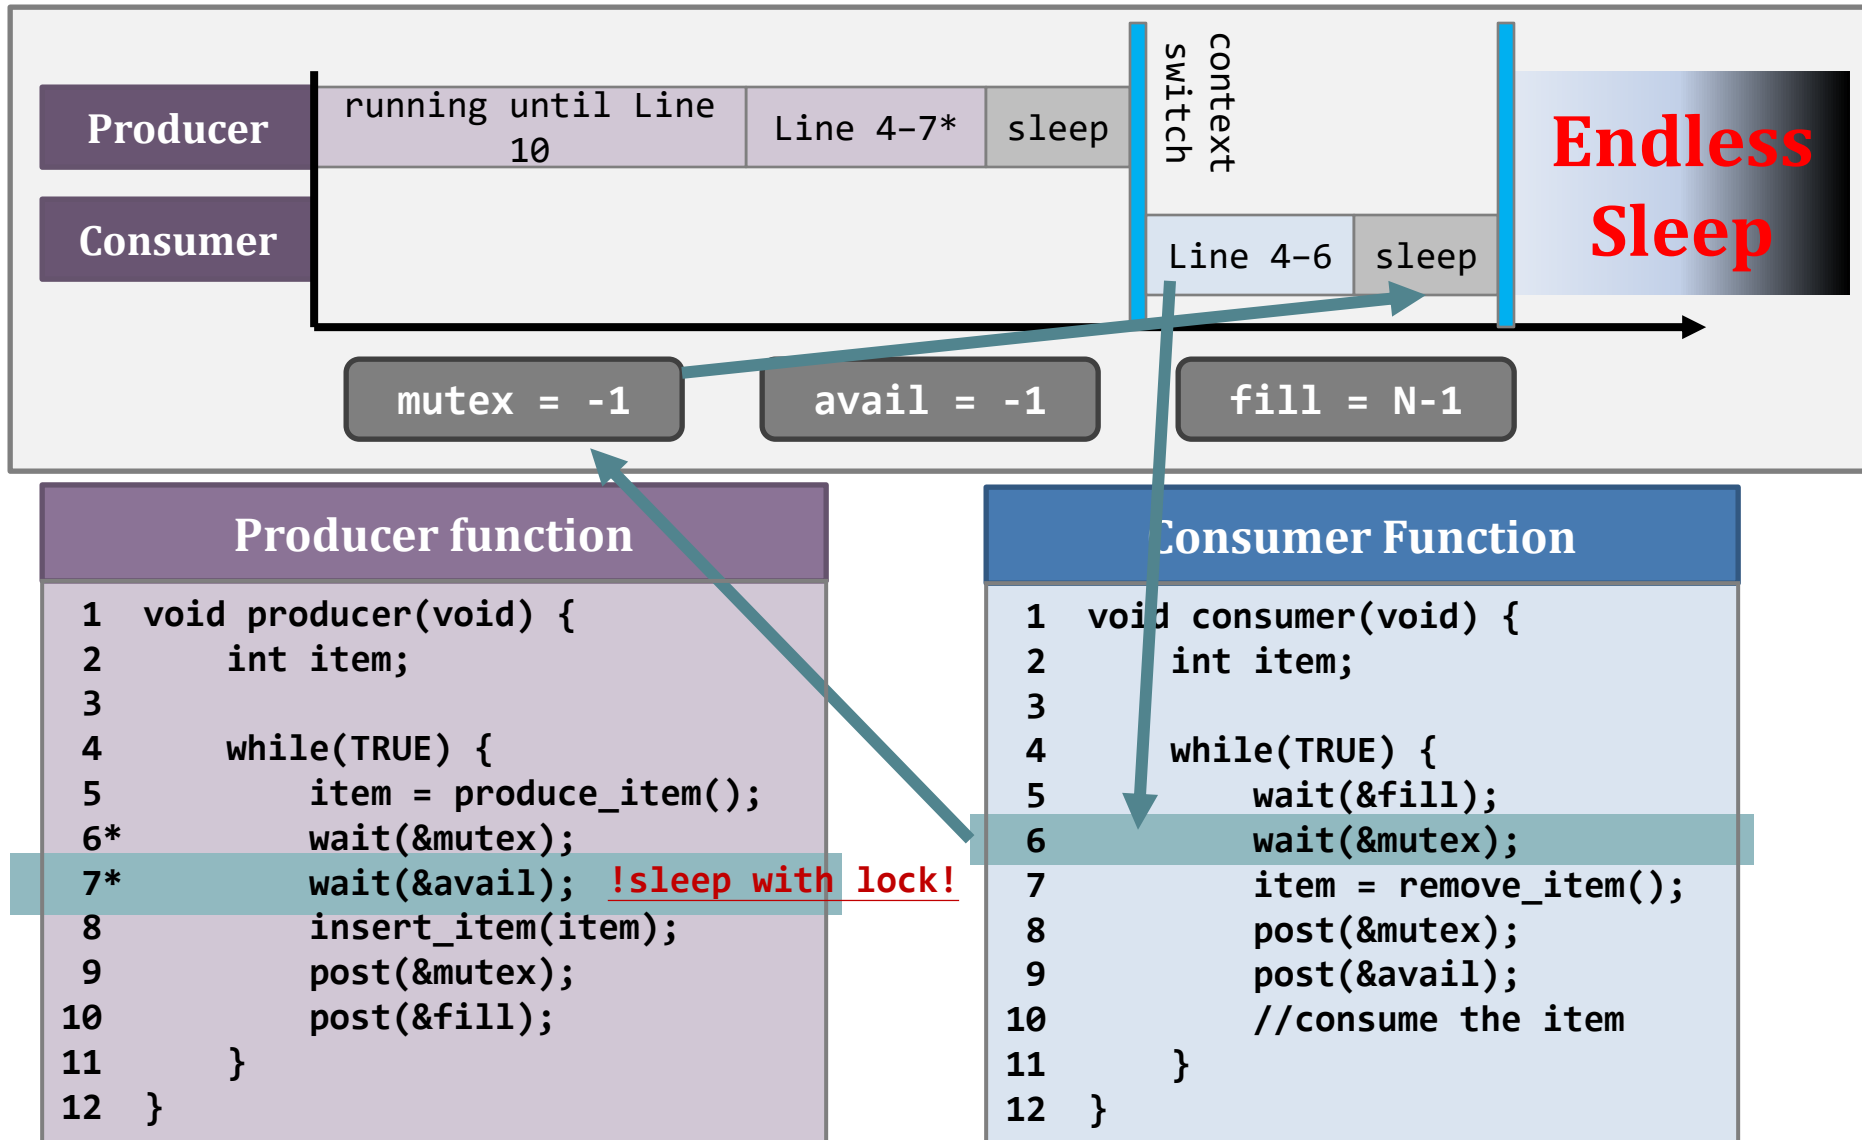
```
1 void producer(void) {
2     int item;
3
4     while(TRUE) {
5         item = produce_item();
6*        wait(&mutex);
7*        wait(&avail);
8        insert_item(item);
9        post(&mutex);
10       post(&fill);
11     }
12 }
```

Consider: producer gets the CPU to keep producing until the buffer is full

Consumer Function

```
1 void consumer(void) {
2     int item;
3
4     while(TRUE) {
5         wait(&fill);
6         wait(&mutex);
7         item = remove_item();
8         post(&mutex);
9         post(&avail);
10        //consume the item
11    }
12 }
```

Producer-consumer problem – question #2




Producer-consumer problem – question #2

- ◆ This scenario is called a **deadlock**
 - ◆ Consumer waits for Producer's **mutex** at line 6
 - ◆ i.e., it waits for Producer (line 9) to unlock the **mutex**
 - ◆ Producer waits for Consumer's **avail** at line 7
 - ◆ i.e., it waits for Consumer (line 9) to release **avail**
- ◆ **Implication:** careless implementation of the producer-consumer solution can be disastrous.

Deadlock

I won't give you the paper unless you come out

I won't come out unless you give me the paper



waiting

holding

waiting



Summary on producer-consumer problem


- ◆ How to avoid race condition on the shared buffer?
 - ◆ E.g., Use a **binary semaphore**.
- ◆ How to achieve synchronization?
 - ◆ E.g., Use two semaphores: fill and avail

Inter-process communication (IPC)

- Classic IPC problems.**
 - Producer-consumer problem.**
 - Dining philosopher problem.**


Dining philosopher – introduction

- ◆ 5 philosophers, 5 plates of spaghetti, and 5 chopsticks.
- ◆ The jobs of each philosopher are to think and to eat
- ◆ They **need exactly two chopsticks** in order to eat the spaghetti.
- ◆ Question: how to construct a synchronization protocol such that they
 - ◆ will not **starve to death**, and
 - ◆ will not result in any **deadlock scenarios**?
 - ◆ A waits for B's chopstick
 - ◆ B waits for C's chopstick
 - ◆ C waits for A's chopstick




It's a multi-object synchronization problem



Dining philosopher – introduction



Philosophers



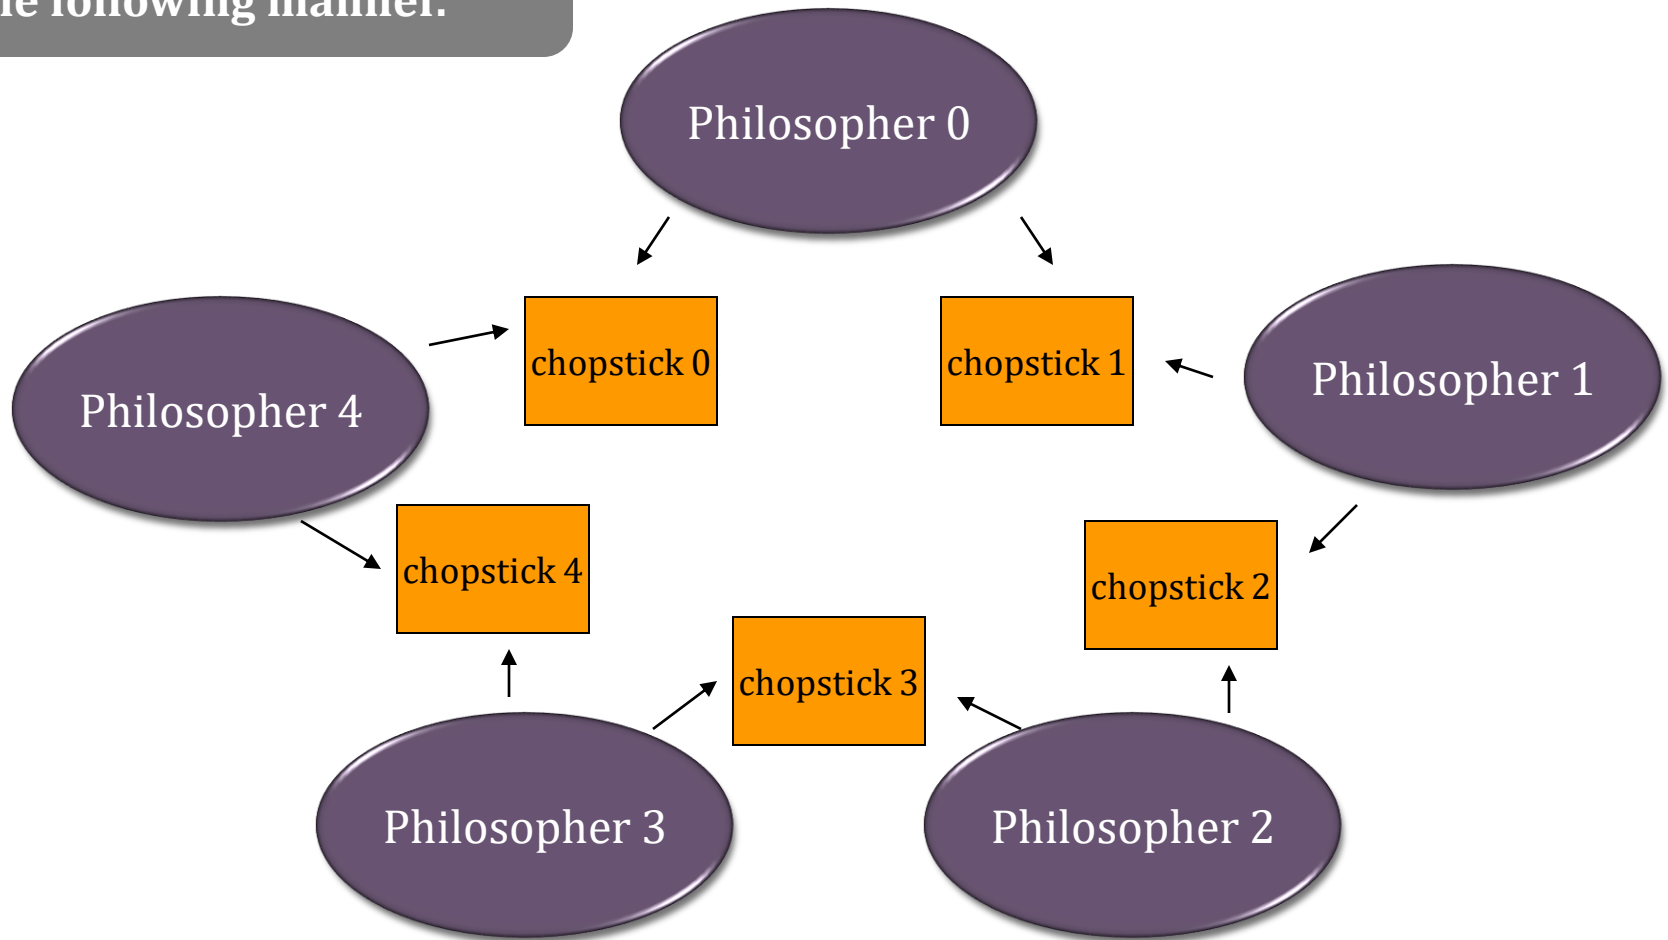
Chopsticks



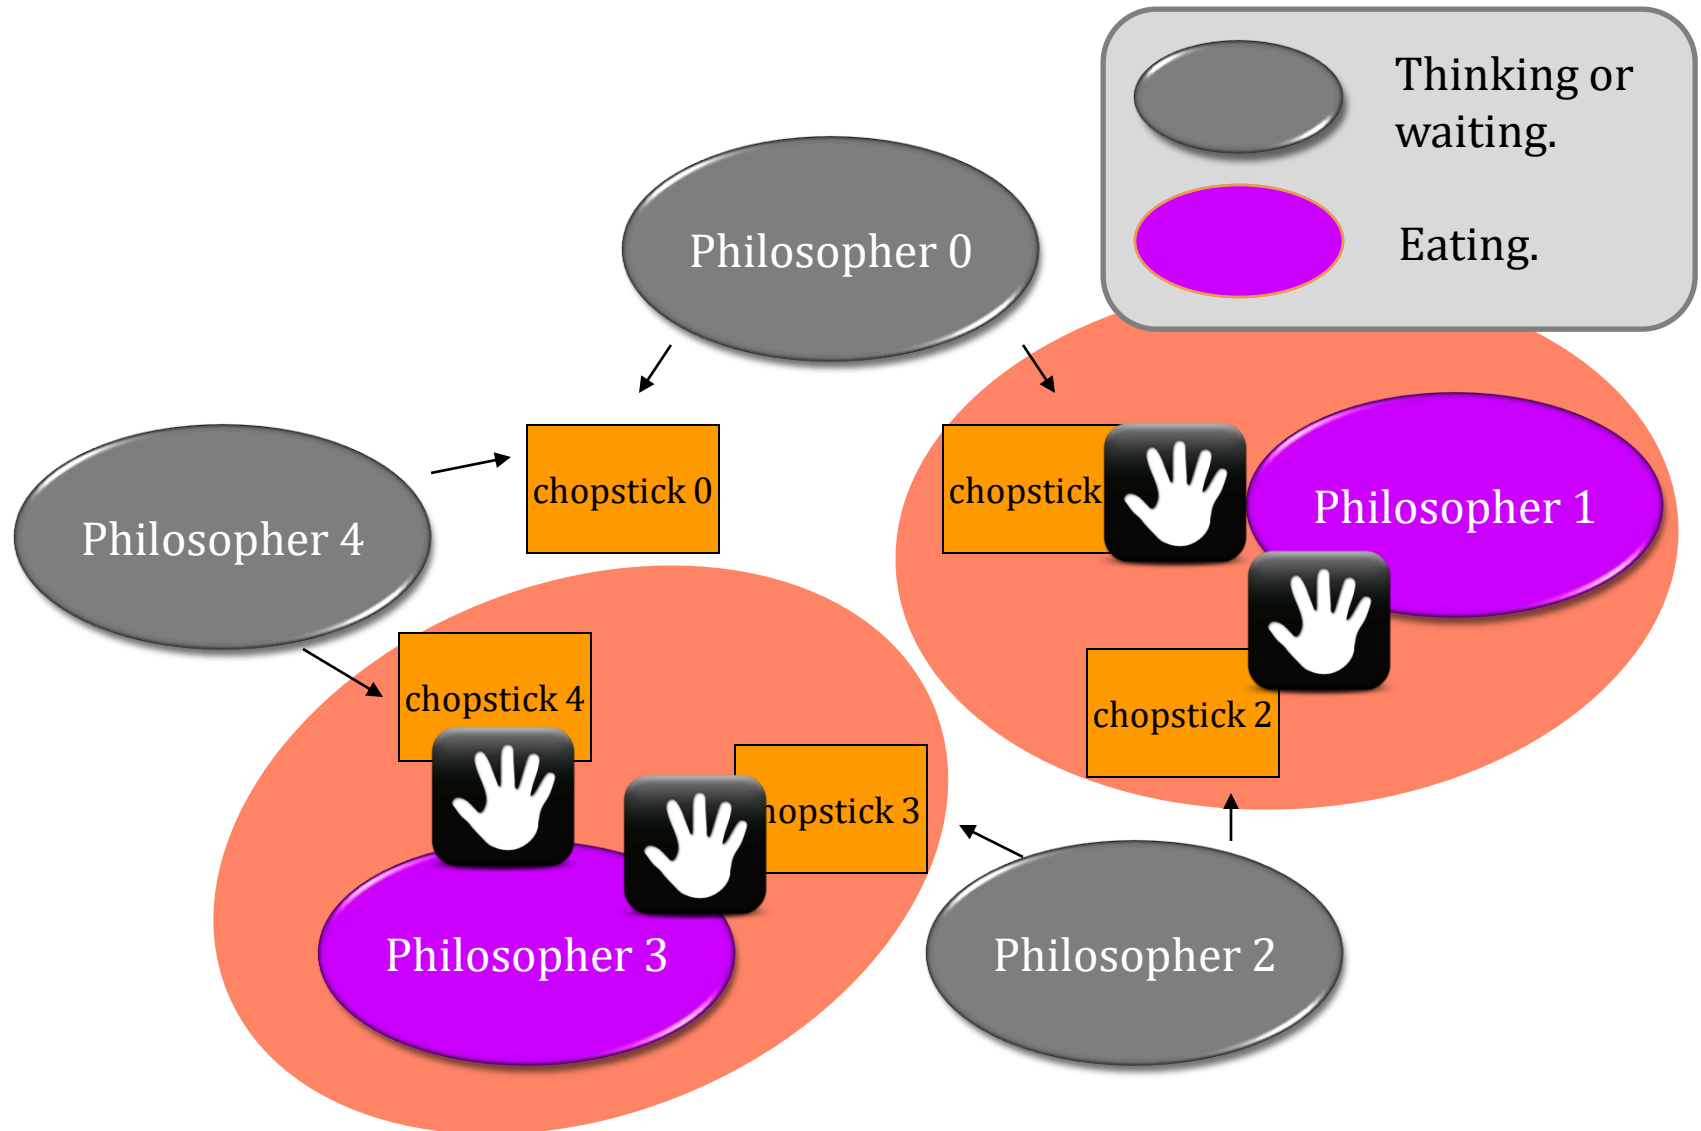
A process needs two shared resources in order to do some work

Dining philosopher – introduction

The chopsticks are arranged in the following manner.



Dining philosopher – introduction



Dining philosopher – requirement #1

◆ Mutual exclusion

- ◆ While you are eating, people cannot steal your chopstick
- ◆ Two persons cannot hold the same chopstick

◆ Let's propose the following solution:

- ◆ When you are hungry, you have to check if anyone is using the chopsticks that you need.
- ◆ If yes, you wait.
- ◆ If no, **seize both chopsticks**.
- ◆ After eating, put down all your chopsticks.

Dining philosopher – meeting requirement #1?

Shared object

```
#define N 5  
semaphore chopstick[N];
```

Five binary semaphores

Helper Functions

```
void take_chopstick(int i)  
{  
    wait(&chopstick[i]);  
}
```

```
void put_chopstick(int i) {  
    post(&chopstick[i]);  
}
```

Section
Entry

Critical
Section

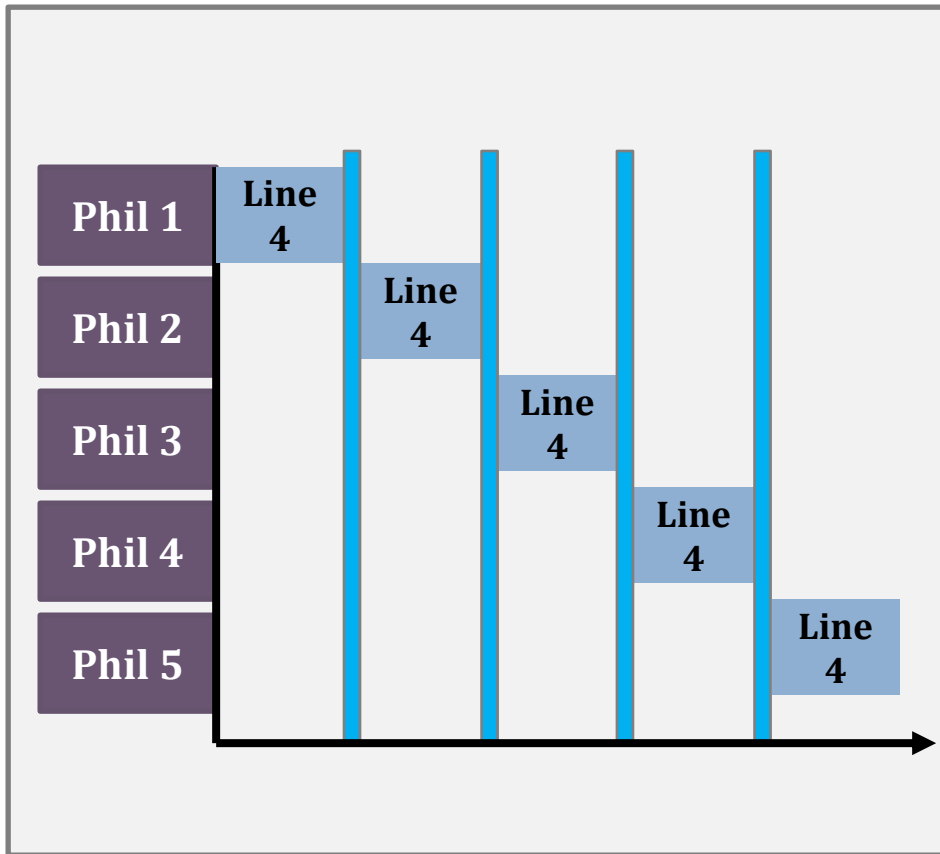
Section
Exit

Main Function

```
1 void philosopher(int i) {  
2     while (TRUE) {  
3         think();  
4         take_chopstick(i);  
5         take_chopstick((i+1) % N);  
6         eat();  
7         put_chopstick(i);  
8         put_chopstick((i+1) % N);  
9     }  
10 }
```

Dining philosopher – deadlock

- Each philosopher finishes thinking at the same time and each first grabs her left chopstick
- All chopsticks[i]=0
- When executing line 5, all are waiting



Main Function

```
1 void philosopher(int i) {  
2     while (TRUE) {  
3         think();  
4         take_chopstick(i);  
5         take_chopstick((i+1) % N);  
6         eat();  
7         put_chopstick(i);  
8         put_chopstick((i+1) % N);  
9     }  
10 }
```

Dining philosopher – requirement #2

◆ Synchronization

- ◆ Should avoid **deadlock**.

- ◆ How about the following suggestions:

- ◆ First, a philosopher takes a chopstick.

- ◆ If a philosopher finds that she cannot take the second chopstick, then she should put it down.

- ◆ Then, the philosopher goes to sleep for a while.

- ◆ When wake up, she retries

- ◆ Loop until both chopsticks are seized.

Dining philosopher – meeting requirement #2?

Potential Problem: Philosophers are all busy (no deadlock), but no progress (starvation)

Imagine:

- all pick up their left chopsticks,
- seeing their right chopsticks unavailable (because P1's right chopstick is taken by P2 as her left chopstick) and then putting down their left chopsticks,
- all sleep for a while
- all pick up their left chopsticks,

Dining philosopher – before the final solution

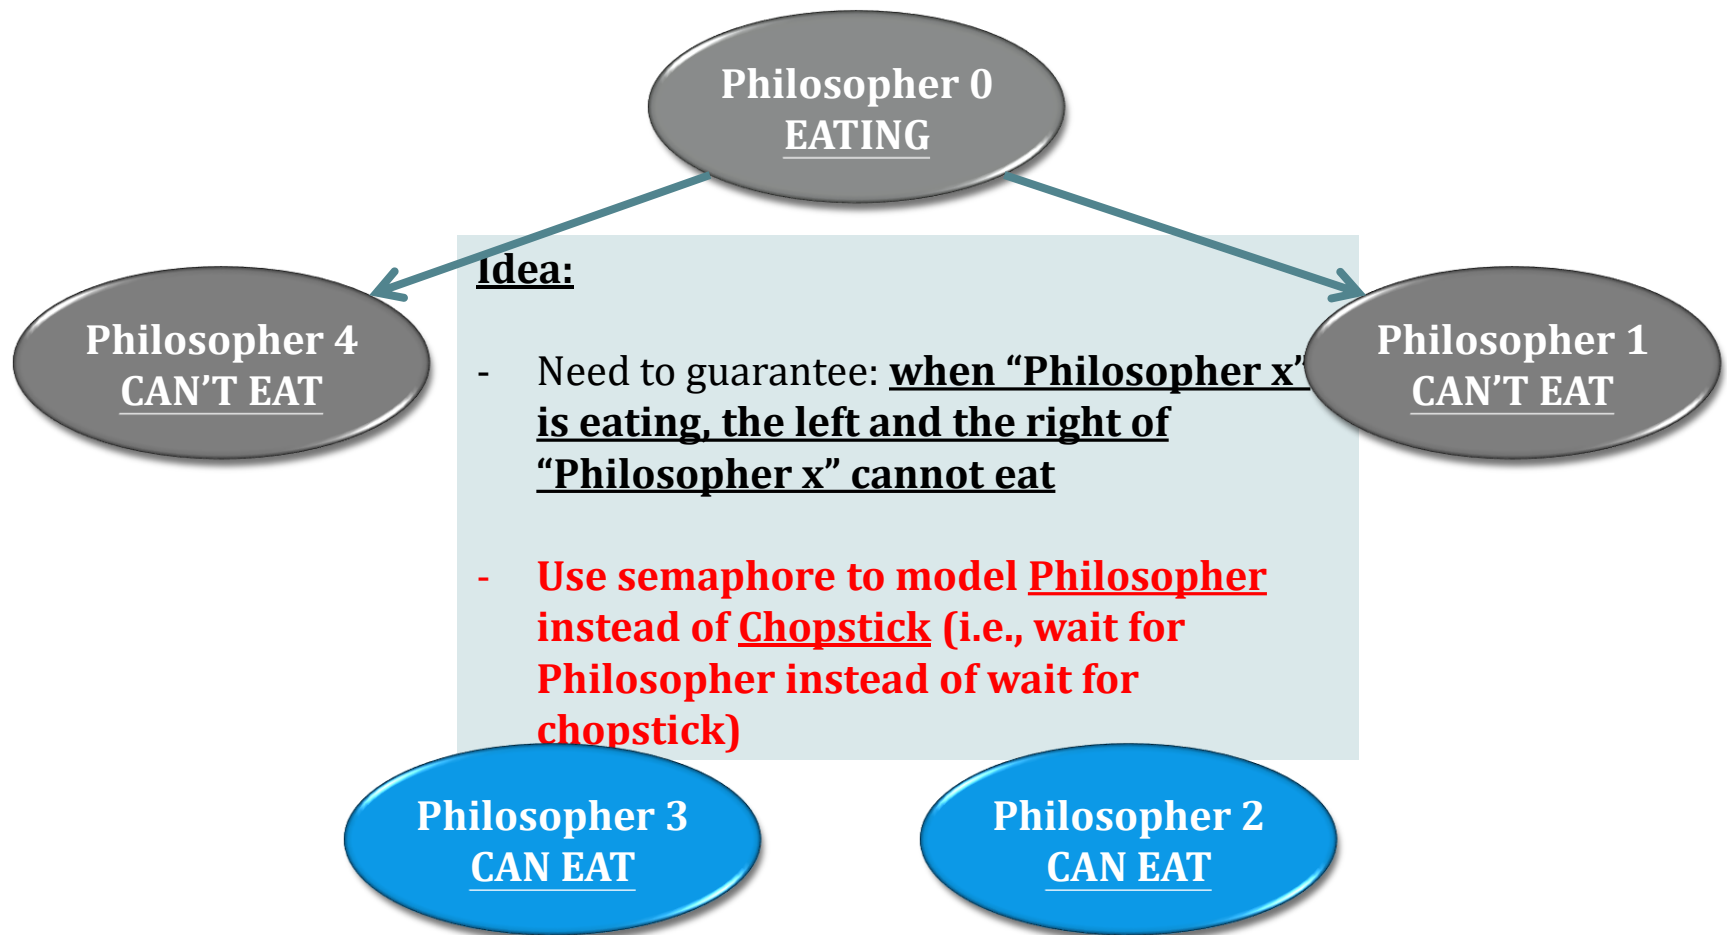
- ◆ Before we present the final solution, let us see what problems we have.

Problems

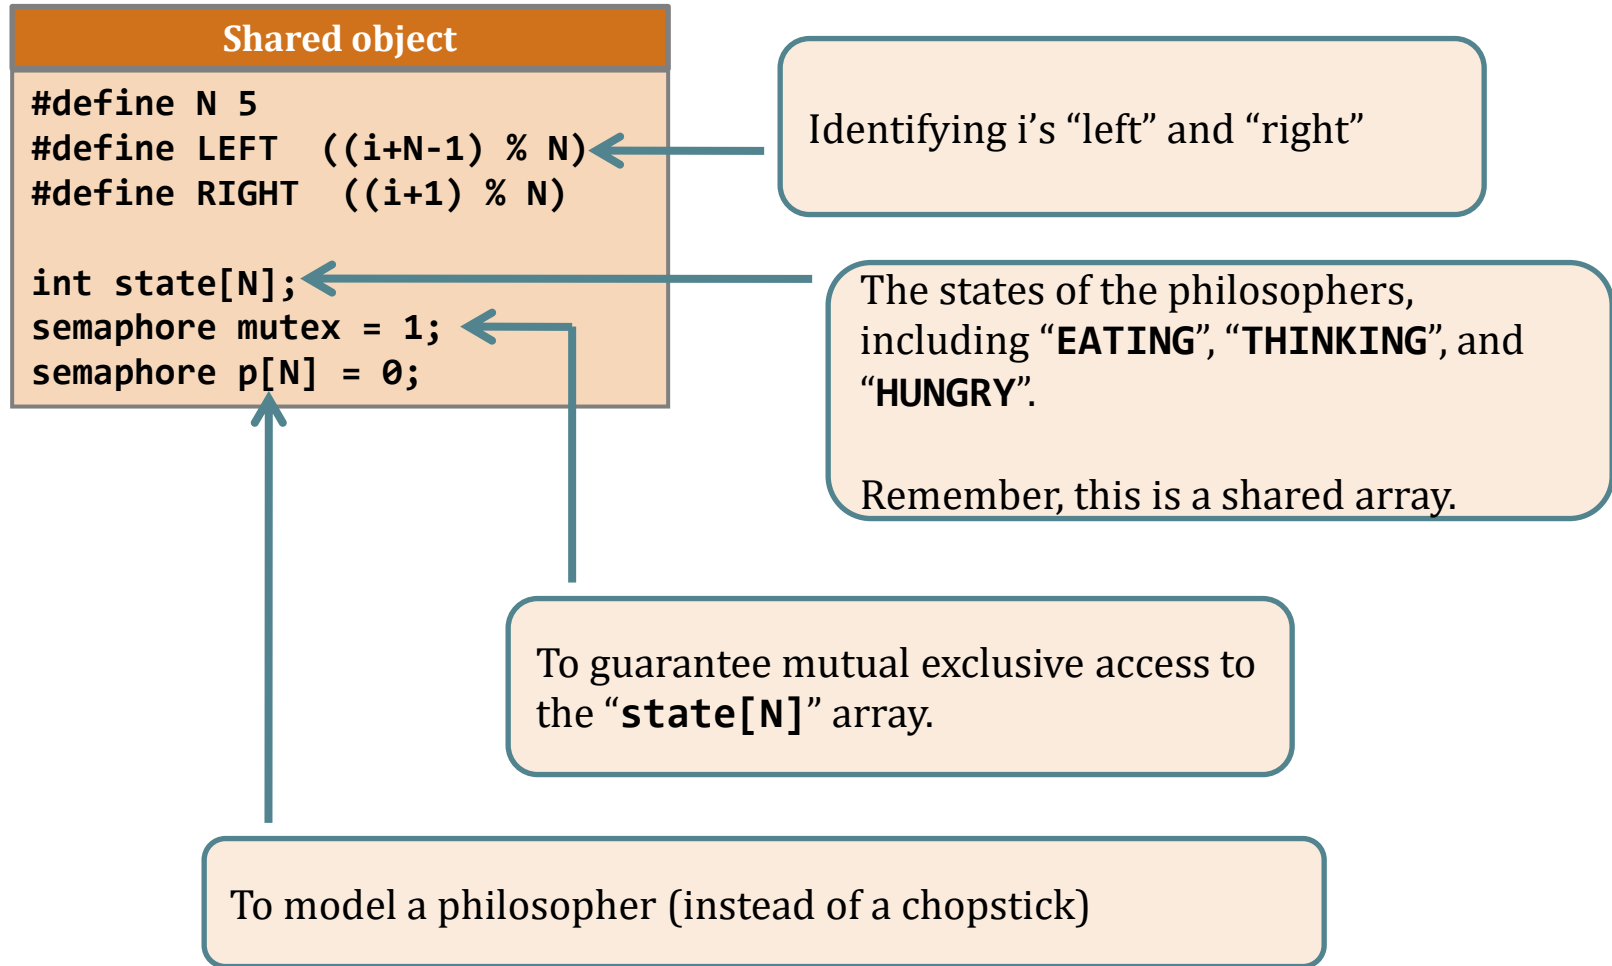
Model each chopstick as a semaphore is intuitive, but may cause deadlock

Using sleep() to avoid deadlock is effective, yet creating starvation.

Dining philosopher – before the final solution.



Dining philosopher – the final solution.



Dining philosopher – the final solution.

Shared object

```
#define N 5
#define LEFT  ((i+N-1) % N)
#define RIGHT  ((i+1) % N)

int state[N];
semaphore mutex = 1;
semaphore p[N] = 0;
```

Main function

```
1 void philosopher(int i) {
2     think();
3     take_chopsticks(i);
4     eat();
5     put_chopsticks(i);
6 }
```

```
void wait(semaphore *s) {
    disable_interrupt();
    *s = *s - 1;
    if ( *s < 0 ) {
        enable_interrupt();
        sleep();
        disable_interrupt();
    }
    enable_interrupt();
}
```

Section entry

```
1 void take_chopsticks(int i) {
2     wait(&mutex);
3     state[i] = HUNGRY;
4     captain(i);
5     post(&mutex);
6     wait(&p[i]);
7 }
```

Section exit

```
1 void put_chopsticks(int i) {
2     wait(&mutex);
3     state[i] = THINKING;
4     captain(LEFT);
5     captain(RIGHT);
6     post(&mutex);
7 }
```

```
void post(semaphore *s) {
    disable_interrupt();
    *s = *s + 1;
    if ( *s <= 0 )
        wakeup();
    enable_interrupt();
}
```

Extremely important helper function

```
1 void captain(int i) {
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
3         state[i] = EATING;
4         post(&p[i]);
5     }
6 }
```


Dining philosopher – Hungry

Tell the captain that you are hungry

If one of your neighbors is eating, the captain just does nothing for you and returns

Then, you wait for your chopsticks (later, the captain will notify you when chopsticks are available)

Section entry

```
1 void take_chopsticks(int i) {  
2     wait(&mutex);  
3     state[i] = HUNGRY;  
4     captain(i);  
5     post(&mutex);  
6     wait(&p[i]);  
7 }
```

Critical Section

The captain is “indivisible”

Extremely important helper function

```
1 void captain(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         post(&p[i]);  
5     }  
6 }
```

Dining philosopher – Finish eating

Tell the captain
Try to let your **left neighbor**
to eat.

Tell the captain
Try to let your right **neighbor**
to eat.

Section exit

```
1 void put_chopsticks(int i)
{
2     wait(&mutex);
3     state[i] = THINKING;
4     captain(LEFT);
5     captain(RIGHT);
6     post(&mutex);
7 }
```

Extremely important helper function

```
1 void captain(int i) {
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {
3         state[i] = EATING;
4         post(&p[i]);
5     }
6 }
```

Wake up the one who is sleeping

Dining philosopher – the final solution

Don't
print

An illustration: How can
Philosopher 1 start
eating?

Philosopher 0
THINKING

Philosopher 4
THINKING

Philosopher 1
THINKING

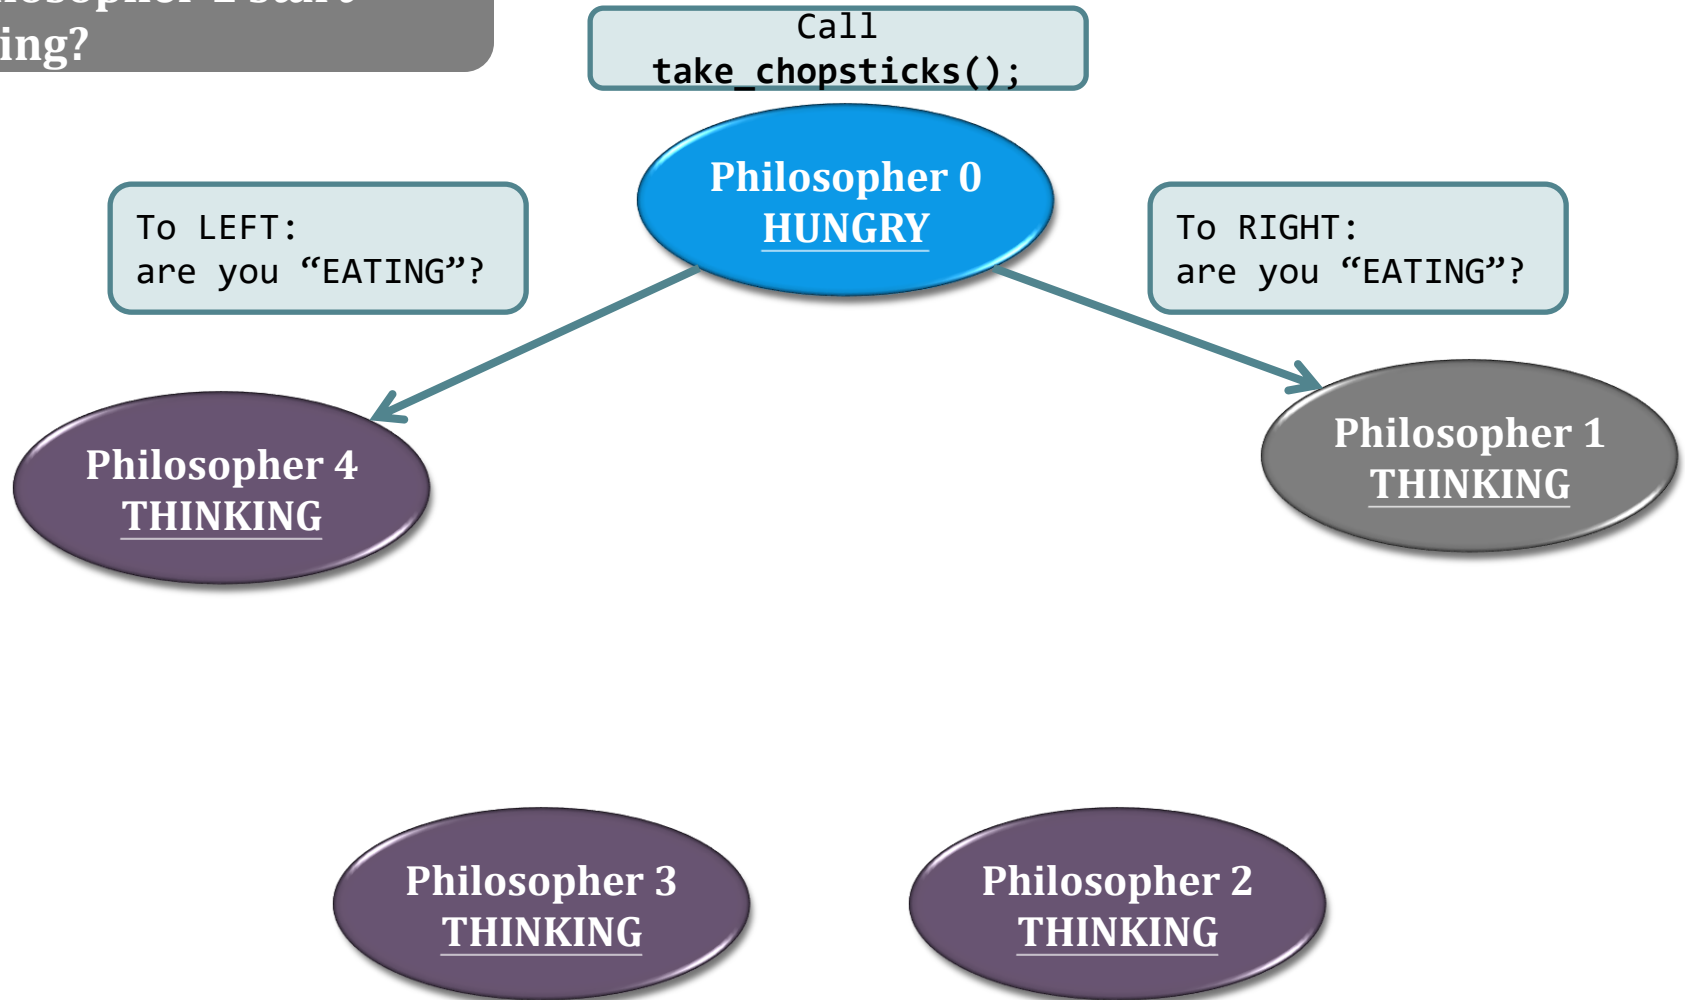
Philosopher 3
THINKING

Philosopher 2
THINKING

Dining philosopher – the final solution

Don't
print

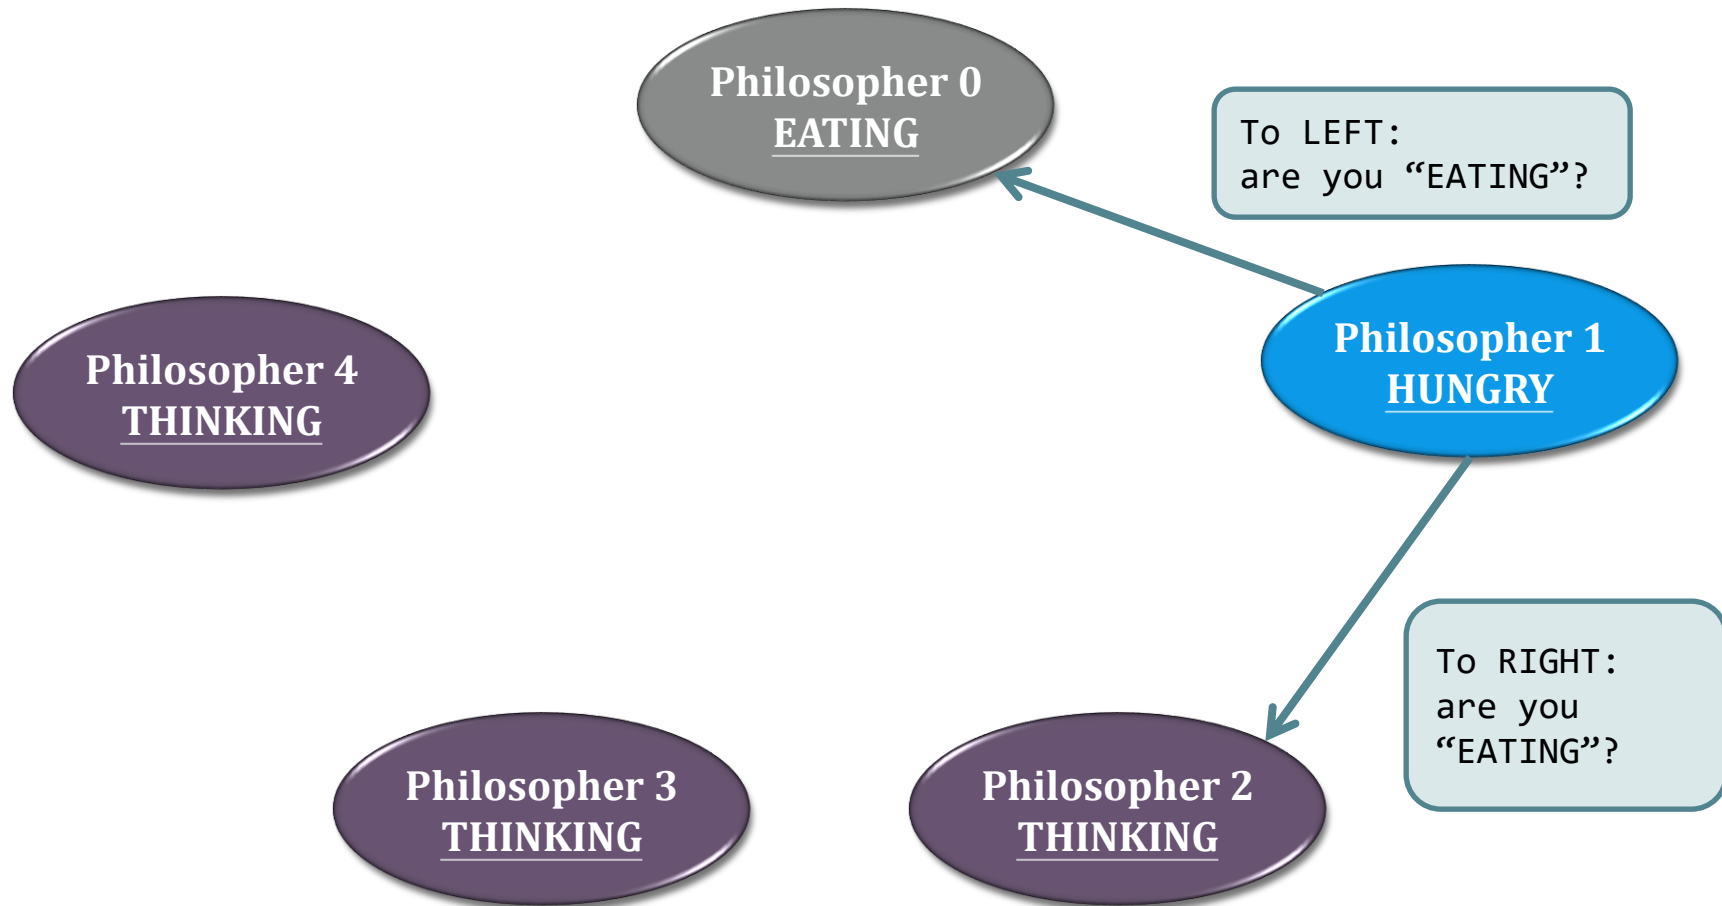
An illustration: How can
Philosopher 1 start
eating?



Dining philosopher – the final solution

Don't
print

An illustration: How can
Philosopher 1 start
eating?



Dining philosopher – the final solution

Don't
print

An illustration: How can
Philosopher 1 start
eating?

Philosopher 0
EATING

```
Section entry
1 void take_chopsticks(int i) {
2     wait(&mutex);
3     state[i] = HUNGRY;
4     captain(i);
5     post(&mutex);
6     wait(&p[i]);
7 }
```

//as P0 is eating, captain(i) returns
w/o doing anything;
wait(&p[1]);

Philosopher 1
HUNGRY

Philosopher 4
THINKING

To LEFT:
are you
"EATING"?

Philosopher 3
HUNGRY

To RIGHT:
are you
"EATING"?

Philosopher 2
THINKING

Dining philosopher – the final solution

Don't
print

An illustration: How can
Philosopher 1 start
eating?

Philosopher 0
EATING

Philosopher 4
THINKING

Philosopher 1
HUNGRY

Blocked;
because of
`wait(&p[1]);`

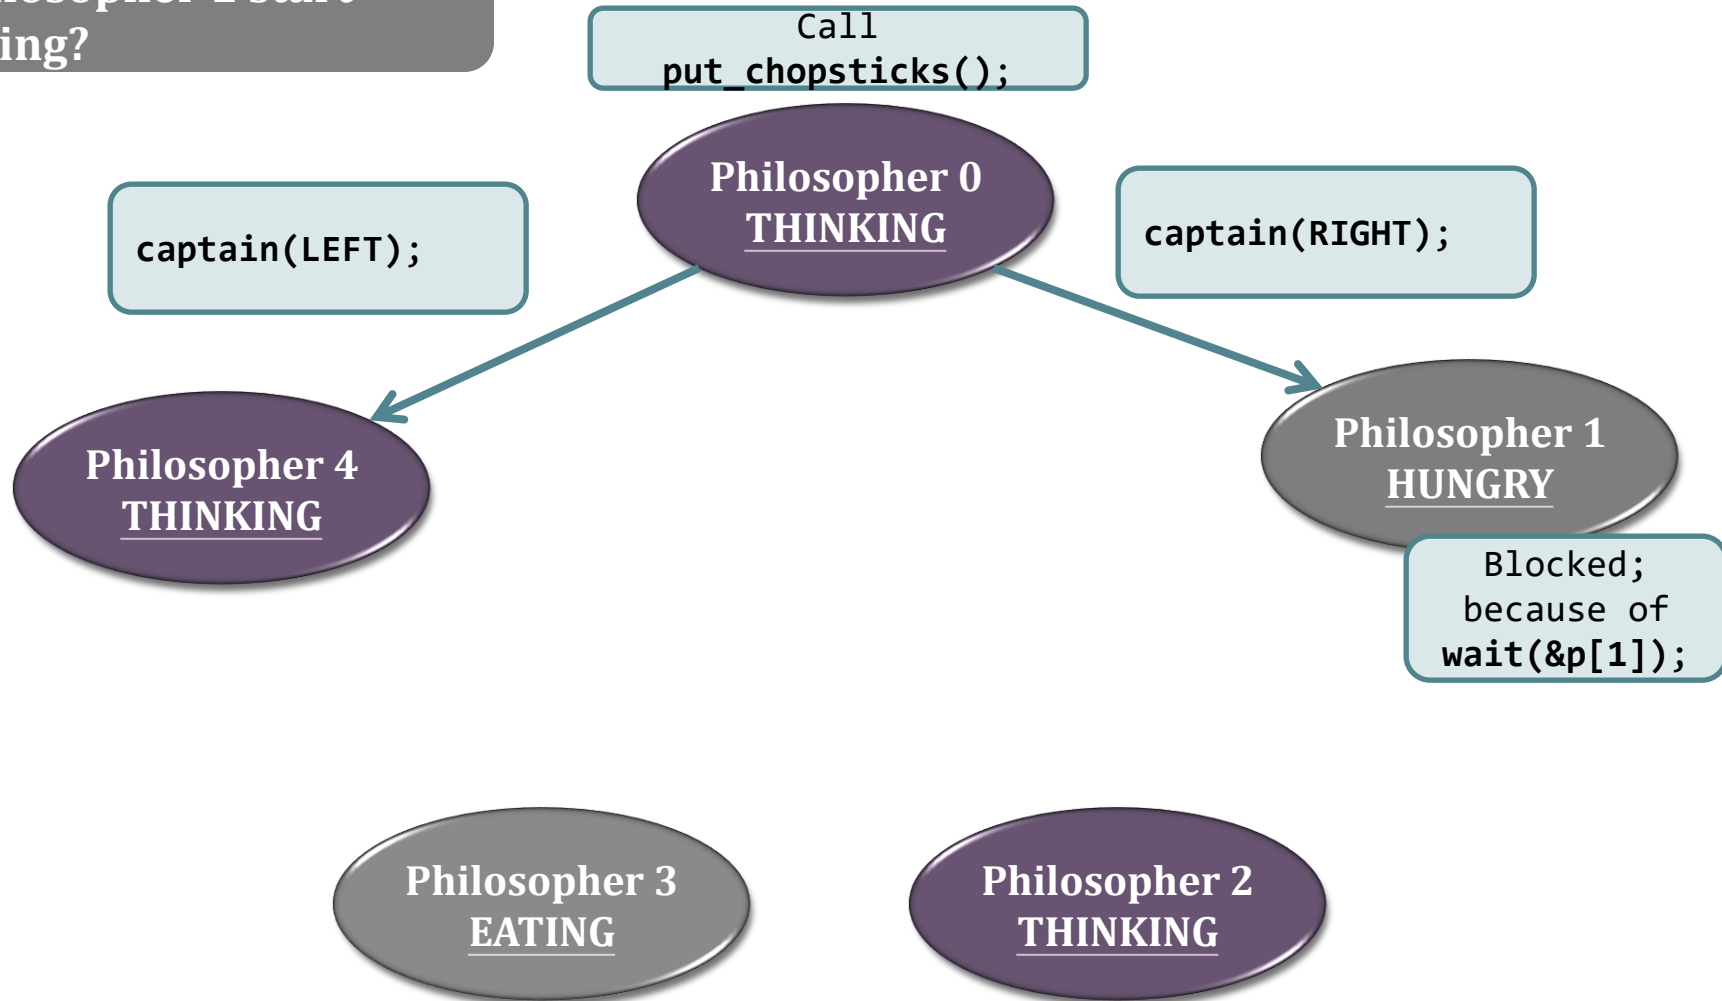
Philosopher 3
EATING

Philosopher 2
THINKING

Dining philosopher – the final solution

Don't
print

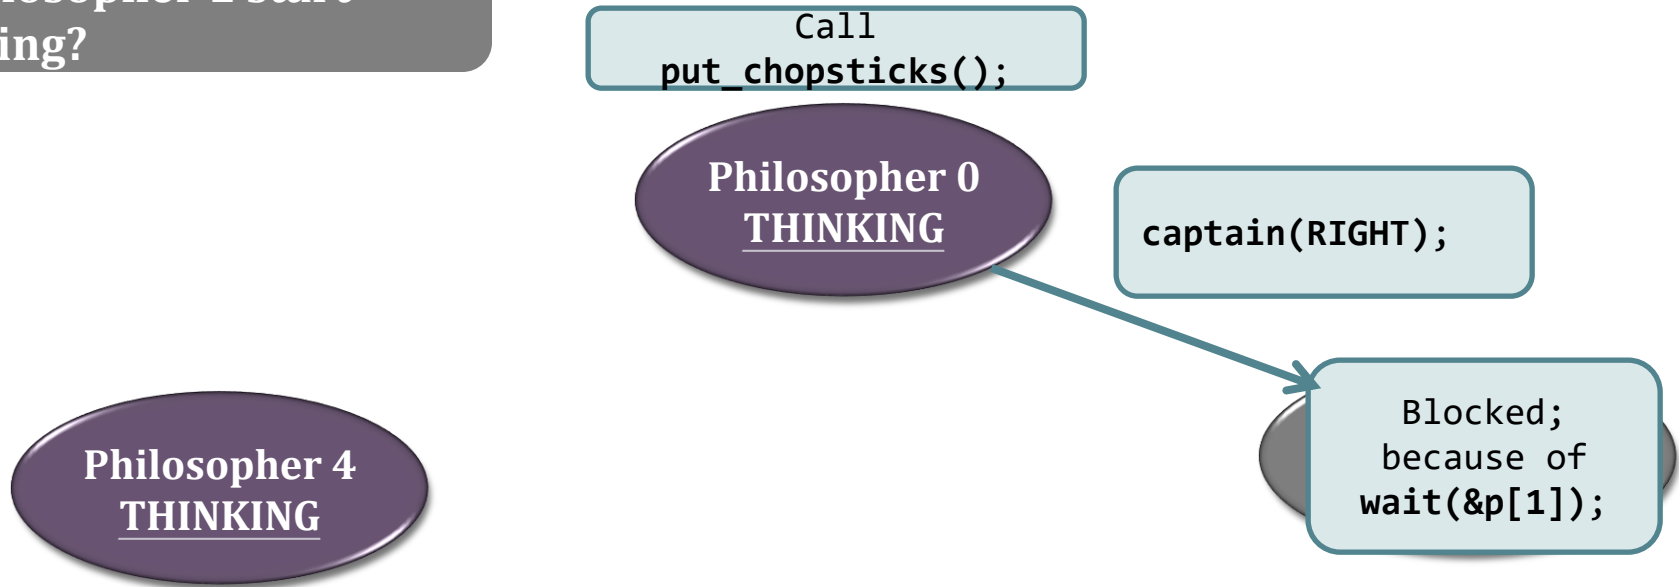
An illustration: How can
Philosopher 1 start
eating?



Dining philosopher – the final solution

Don't
print

An illustration: How can
Philosopher 1 start
eating?



```
1 void captain(int i) {  
2     if(state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING) {  
3         state[i] = EATING;  
4         post(&p[i]); ← Wake up !  
5     }  
6 }
```

Dining philosopher – the final solution

Don't
print

An illustration: How can
Philosopher 1 start
eating?

Philosopher 0
THINKING

Philosopher 4
THINKING

Philosopher 3
EATING

Philosopher 2
THINKING

Wake up

Philosopher 1
EATING

Section entry

```
1 void take_chopsticks(int i) {  
2     wait(&mutex);  
3     state[i] = HUNGRY;  
4     captain(i);  
5     post(&mutex);  
6     wait(&p[i]);  
7 }
```

Dining philosopher – the core

5 philosophers → ideally how many chopsticks

how many chopsticks do we have now?

Very common in today's cloud computing multi-tenancy model

Summary on IPC problems

- ◆ The problems have the following properties in common:
 - ◆ Multiple number of processes;
 - ◆ Processes have to be synchronized in order to generate useful output;
 - ◆ Each resource may be shared as well as limited, and there may be more than one shared processes.
- ◆ The synchronization algorithms have the following requirements in common:
 - ◆ Guarantee mutual exclusion;
 - ◆ Uphold the correct synchronization among processes; and
 - ◆ (must be) Deadlock-free.

Heisenbugs

- ◆ Jim Gray, 1998 ACM Turing Award winner, coined that term
- ◆ You find your program P has a concurrency bug
- ◆ You insert 'printf' statements or GDB to debug P
- ◆ Then because of those debugging things added, P behaves normally when you are in debug mode

Heisenbugs

THE BUG FAIRY
PRESENTS:

FUN, FUN
BUGS!



THE HEISENBUG: A BUG THAT NEVER SHOWS UP WHILE
BEING OBSERVED IN A DEBUGGER, YET ALWAYS APPEARS
IN THE RELEASE BUILD!



THE REBUG: WHEN SPECIAL DEBUG CODE INTRODUCES A BUG!

```
while (count<10)
{
    count++;
    #ifdef _DEBUG
        count--;
    #endif
}
```



THE WTF BUG: AFTER DISCOVERING THE CAUSE OF THE BUG,
THE DEVELOPER WILL WONDER HOW THE ROUTINE EVER
WORKED AT ALL!



HALLEY'S BUG: A BUG THAT DEFIES BEING REPRODUCED AND
SHOWS UP INCREDIBLY RARELY.



WISH HER LUCK! SHE'S OFF TO VISIT YOUR CODE!



Thank You!