


Lecture 7: Deadlock

Bo Tang @ 2020, Spring

Deadlock

I won't give you the paper unless you come out

I won't come out unless you give me the paper




waiting

holding

waiting




Starvation vs Deadlock



◆ Starvation vs. Deadlock

- ◆ Starvation: thread waits indefinitely
 - ◆ Low-priority thread waiting for resources constantly in use by high-priority threads
- ◆ Deadlock: circular waiting for resources
 - ◆ Thread A owns Res 1 and is waiting for Res 2
Thread B owns Res 2 and is waiting for Res 1



- ◆ Deadlock \Rightarrow Starvation but not vice versa
 - ◆ Starvation can end (but does not have to)
 - ◆ Deadlock cannot end without external intervention

Conditions for Deadlock

◆ Deadlock not always deterministic


<u>Thread A</u>	<u>Thread B</u>
<code>sem_wait(x);</code>	<code>sem_wait(y);</code>
<code>sem_wait(y);</code>	<code>sem_wait(x);</code>
<code>sem_post(y);</code>	<code>sem_wait(x);</code>
<code>sem_post(x);</code>	<code>sem_wait(y);</code>

- ◆ Deadlock will not always happen with this code
 - ◆ Have to have exactly the right timing
 - ◆ So you release a piece of software, and you tested it, and there it is, controlling a nuclear power plant...

◆ Deadlocks occur with multiple resources

- ◆ Means you cannot decompose the problem
- ◆ Cannot solve deadlock for each resource independently
- ◆ System with 2 disk drives and two threads
 - ◆ Each thread needs 2 disk drives to function
 - ◆ Each thread gets one disk and waits for another one


Bridge Crossing Example



- ◆ Each segment of road can be viewed as a resource
 - ◆ Car must own the segment under them
 - ◆ Must acquire segment that they are moving into
- ◆ For bridge: must acquire both halves
 - ◆ Traffic only in one direction at a time
 - ◆ Problem occurs when two cars in opposite directions on bridge: each acquires one segment and needs next
- ◆ If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
 - ◆ Several cars may have to be backed up
- ◆ Starvation is possible
 - ◆ East-going traffic really fast \Rightarrow no one goes west


Train Example

- ◇ Circular dependency (Deadlock!)
- ◇ Each train wants to turn right
- ◇ Blocked by other trains
- ◇ Similar problem to multiprocessor networks



Train Example

- ◆ Fix? Imagine grid extends in all four directions
 - ◆ **Force ordering of channels** (tracks)
 - ◆ Protocol: Always go east-west first, then north-south
 - ◆ Called “dimension ordering” (X then Y)



Four requirements for Deadlock

◆ Mutual exclusion

- ◆ Only one thread at a time can use a resource.

◆ Hold and wait

- ◆ Thread holding at least one resource is waiting to acquire additional resources held by other threads

◆ No preemption

- ◆ Resources are released only voluntarily by the thread holding the resource, after thread is finished with it


◆ Circular wait

- ◆ There exists a set $\{T_1, \dots, T_n\}$ of waiting threads
 - ◆ T_1 is waiting for a resource that is held by T_2
 - ◆ T_2 is waiting for a resource that is held by T_3
 - ◆ ...
 - ◆ T_n is waiting for a resource that is held by T_1

Resource-Allocation Graph

◆ System Model

- ◆ A set of Threads T_1, T_2, \dots, T_n
- ◆ Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- ◆ Each resource type R_i has W_i instances
- ◆ Each thread utilizes a resource as follows:
 - ◆ Request () / Use () / Release ()




◆ Resource-Allocation Graph:

- ◆ V is partitioned into two types:
 - ◆ $T = \{T_1, T_2, \dots, T_n\}$, the set threads in the system.
 - ◆ $R = \{R_1, R_2, \dots, R_m\}$, the set of resource types in system
- ◆ request edge – directed edge $T_1 \rightarrow R_j$
- ◆ assignment edge – directed edge $R_j \rightarrow T_i$


Resource Allocation Graph Examples

◆ Recall:


- ◆ request edge – directed edge $T_1 \rightarrow R_j$
- ◆ assignment edge – directed edge $R_j \rightarrow T_i$



Simple Resource
Allocation Graph




Allocation Graph
With Deadlock



Allocation Graph with
Cycle, but No Deadlock

Methods for Handling Deadlocks




- ◆ Allow system to enter deadlock and then recover
 - ◆ Requires deadlock detection algorithm
 - ◆ Some technique for forcibly preempting resources and/or terminating tasks
- ◆ Ensure that system will *never* enter a deadlock
 - ◆ Need to monitor all lock acquisitions
 - ◆ Selectively deny those that *might* lead to deadlock
- ◆ Ignore the problem and pretend that deadlocks never occur in the system
 - ◆ Used by most operating systems, including UNIX

Deadlock Detection Algorithm

- ◆ Only one of each type of resource \Rightarrow look for loops
- ◆ More general deadlock detection algorithm
 - ◆ Let $[X]$ represent an m-ary vector of non-negative integers (quantities of resources of each type):
 - ◆ $[FreeResources]$: current free resources each type
 - ◆ $[Request_x]$: current requests from thread X
 - ◆ $[Alloc_x]$: current resources held by thread X
 - ◆ See if tasks can eventually terminate on their own

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

- ◆ Nodes left in UNFINISHED \Rightarrow deadlocked



What to do when detect deadlock?

- ◆ Terminate thread, force it to give up resources
 - ◆ In Bridge example, Godzilla picks up a car, hurls it into the river. Deadlock solved!
 - ◆ Shoot a dining philosopher
 - ◆ But, not always possible – killing a thread holding a mutex leaves world inconsistent
- ◆ Preempt resources without killing off thread
 - ◆ Take away resources from thread temporarily
 - ◆ Does not always fit with semantics of computation
- ◆ Roll back actions of deadlocked threads
 - ◆ For bridge example, make one car roll backwards (may require others behind him)
 - ◆ Common technique in databases (transactions)
 - ◆ Of course, if you restart in exactly the same way, may reenter deadlock once again
- ◆ Many operating systems use other options

Techniques for Preventing Deadlock

◆ Infinite resources

- ◆ Include enough resources so that no one ever runs out of resources. Examples:
 - ◆ Bay bridge with 12,000 lanes. Never wait!
 - ◆ Infinite disk space (not realistic yet?)

◆ No sharing of resources (totally independent threads)

- ◆ Not very realistic

◆ Do not allow waiting

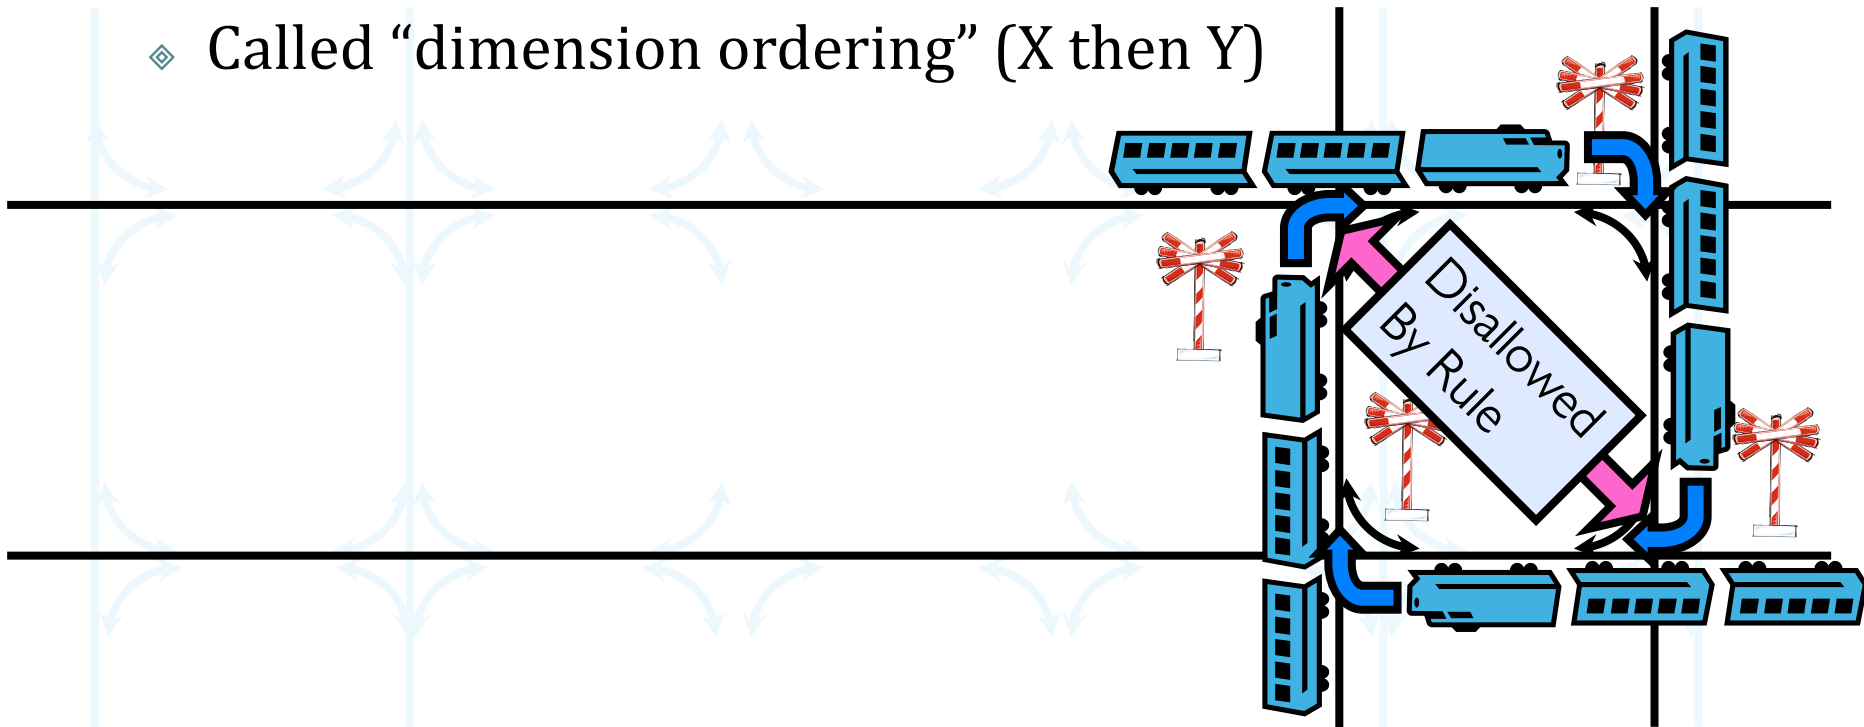
- ◆ Technique used in Ethernet/some multiprocessor nets
 - ◆ Everyone speaks at once. On collision, back off and retry
- ◆ Inefficient, since have to keep retrying
 - ◆ Consider: driving to SUSTech; when hit traffic jam, suddenly you are transported back home and told to retry!

Techniques for Preventing Deadlock

- ◆ Make all threads request everything they will need at the beginning.
 - ◆ Problem: Predicting future is hard, tend to over-estimate resources. Example:
 - ◆ If need 2 chopsticks, request both at same time
 - ◆ Don not leave home until we know no one is using any intersection between home and SUSTech; only one car on the Bay Bridge at a time
- ◆ Force all threads to request resources in a particular order preventing any cyclic use of resources
 - ◆ Thus, preventing deadlock
 - ◆ Example (x.P, y.P, z.P,...)
 - ◆ Make tasks request disk, then memory, then...
 - ◆ Keep from deadlock on freeways around SF by requiring everyone to go clockwise

Review: Train Example

- ◈ Circular dependency (Deadlock!)
 - ◈ Each train wants to turn right, blocked by other trains
 - ◈ Similar problem to multiprocessor networks
- ◈ Fix? Imagine grid extends in all four directions
 - ◈ Force ordering of channels (tracks)
 - ◆ Protocol: Always go east-west first, then north-south
 - ◈ Called “dimension ordering” (X then Y)



Banker's Algorithm for Preventing Deadlock

- ◆ Toward right idea:

- ◆ State maximum (max) resource needs in advance
- ◆ Allow particular thread to proceed if:
 - ◆ $(\text{available resources} - \text{\#requested}) \geq \text{max}$
remaining that might be needed by any thread

- ◆ Banker's algorithm (less conservative):

- ◆ Allocate resources dynamically
 - ◆ Evaluate each request and grant if some ordering of threads is still deadlock free afterward
 - ◆ Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([\text{Maxnode}] - [\text{Allocnode}] \leq [\text{Avail}])$ for $([\text{Requestnode}] \leq [\text{Avail}])$
Grant request if result is deadlock free (conservative!)



Banker's Algorithm for Preventing Deadlock

```
◆ [Avail] = [FreeResources]
  Add all nodes to UNFINISHED
  do {
    done = true
    Foreach node in UNFINISHED {
      if ([Requestnode] <= [Avail]) {
        remove node from UNFINISHED
        [Avail] = [Avail] + [Allocnode]
        done = false
      }
    }
  } until(done)
```

◆ Allocate resources dynamically

- ◆ Evaluate each request and grant if some ordering of threads is still deadlock free afterward
- ◆ **Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([Max_{node}] - [Alloc_{node}] \leq [Avail])$ for $([Request_{node}] \leq [Avail])$**
Grant request if result is deadlock free (conservative!)



Banker's Algorithm for Preventing Deadlock

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

ordering of threads is still deadlock free afterward

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Maxnode] - [Allocnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```




Banker's Algorithm for Preventing Deadlock

- ◆ Toward right idea:
 - ◆ State maximum resource needs in advance
 - ◆ Allow particular thread to proceed if:
 $(\text{available resources} - \text{\#requested}) \geq \text{max remaining that might be needed by any thread}$
- ◆ Banker's algorithm (less conservative):
 - ◆ Allocate resources dynamically
 - ◆ Evaluate each request and grant if some ordering of threads is still deadlock free afterward
 - ◆ **Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([\text{Max}_{\text{node}}] - [\text{Alloc}_{\text{node}}] \leq [\text{Avail}])$ for $([\text{Request}_{\text{node}}] \leq [\text{Avail}])$**
Grant request if result is deadlock free (conservative!)
 - ◆ Keeps system in a "SAFE" state, i.e. there exists a sequence $\{T_1, T_2, \dots, T_n\}$ with T_1 requesting all remaining resources, finishing, then T_2 requesting all remaining resources, etc..
 - ◆ Algorithm allows the sum of maximum resource needs of all current threads to be greater than total resources



Banker's Algorithm Example



- ◆ Banker's algorithm with dining philosophers
 - ◆ “Safe” (will not cause deadlock) if when try to grab chopstick either:
 - ◆ Not last chopstick
 - ◆ Is last chopstick but someone will have two afterwards
 - ◆ What if k-handed philosophers? Do not allow if:
 - ◆ It is the last one, no one would have k
 - ◆ It is 2nd to last, and no one would have k-1
 - ◆ It is 3rd to last, and no one would have k-2
 - ◆ ...



Thank You!