


Lecture 8:

Address Translation

Yinqian @ 2021, Spring


copyright@Bo Tang

Want Processes to Co-exist



- ✧ Consider multiprogramming on physical memory
 - ✧ What happens when pintos needs to expand?
 - ✧ If emacs needs more memory than is on the machine?
 - ✧ If pintos has an error and writes to address 0x7100?
 - ✧ When does gcc have to know it will run at 0x4000?
 - ✧ What if emacs is not using its memory?

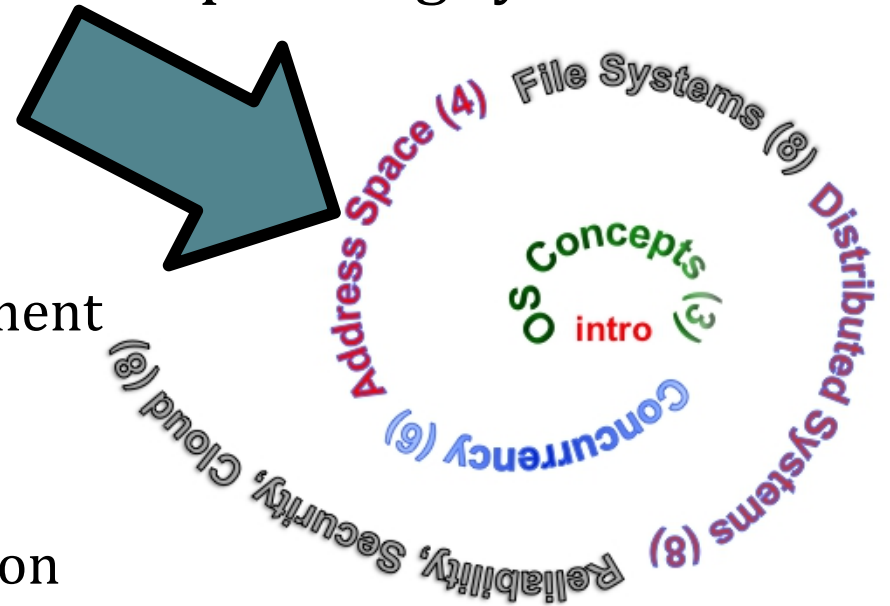
Virtualizing Resources



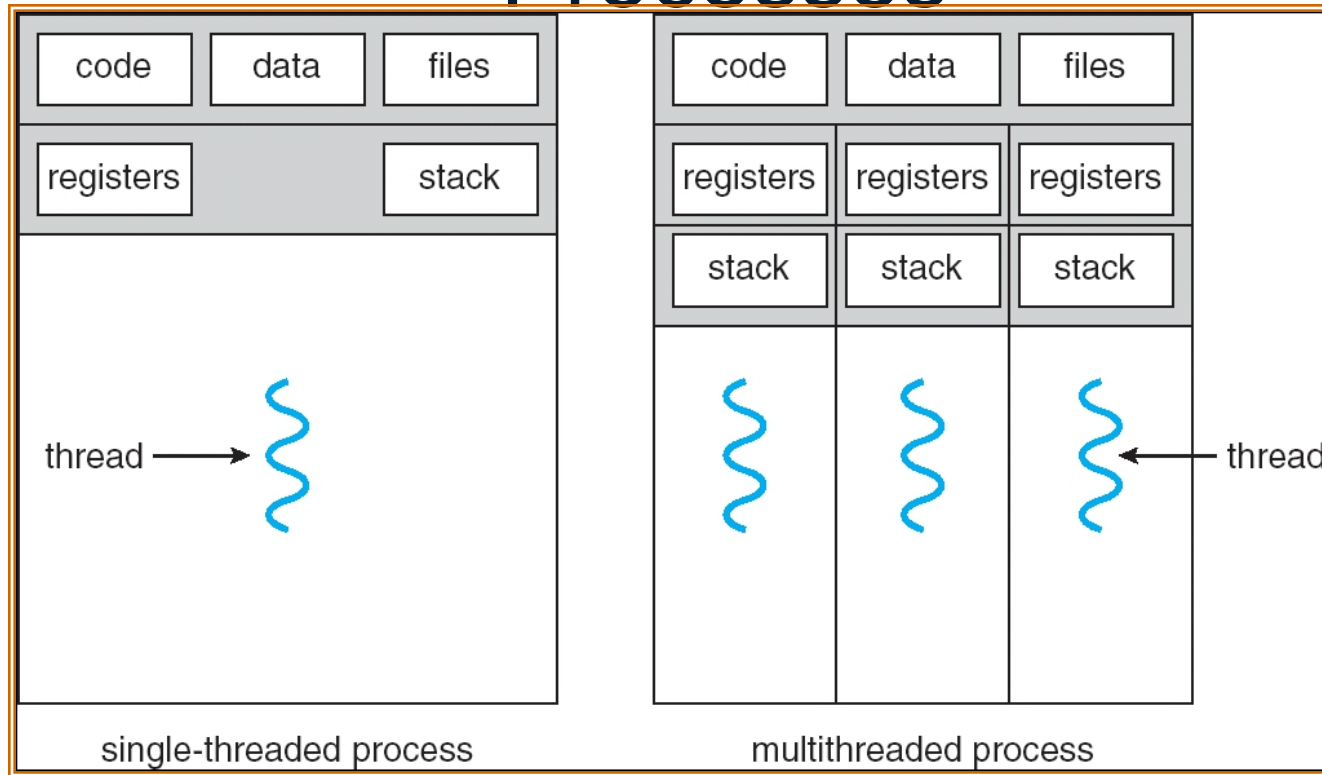
- ✧ **Physical Reality:** different processes/threads share the same hardware
 - ✧ Need to multiplex CPU (done)
 - ✧ Need to multiplex use of Memory (today)
 - ✧ Need to multiplex disk and devices (later in term)
- ✧ **Why worry about memory sharing?**
 - ✧ The complete working state of a process is defined by its data in memory (and registers)
 - ✧ Consequently, two different processes cannot use the same memory
 - ✧ Two different data cannot occupy same locations in memory
 - ✧ May not want different threads to have access to each other's memory

Next Objective

- ✧ Dive deeper into the concepts and mechanisms of memory sharing and address translation
- ✧ Enabler of many key aspects of operating systems
 - ✧ Protection
 - ✧ Multi-programming
 - ✧ Isolation
 - ✧ Memory resource management
 - ✧ I/O efficiency
 - ✧ Sharing
 - ✧ Inter-process communication
 - ✧ Demand paging
- ✧ Today: Linking, Segmentation



Recall. Single and Multithreaded Processes



- ✧ Threads encapsulate concurrency
 - ✧ “Active” component of a process
- ✧ Address spaces encapsulate protection
 - ✧ Keeps buggy program from trashing the system
 - ✧ “Passive” component of a process

Important Aspects of Memory Multiplexing

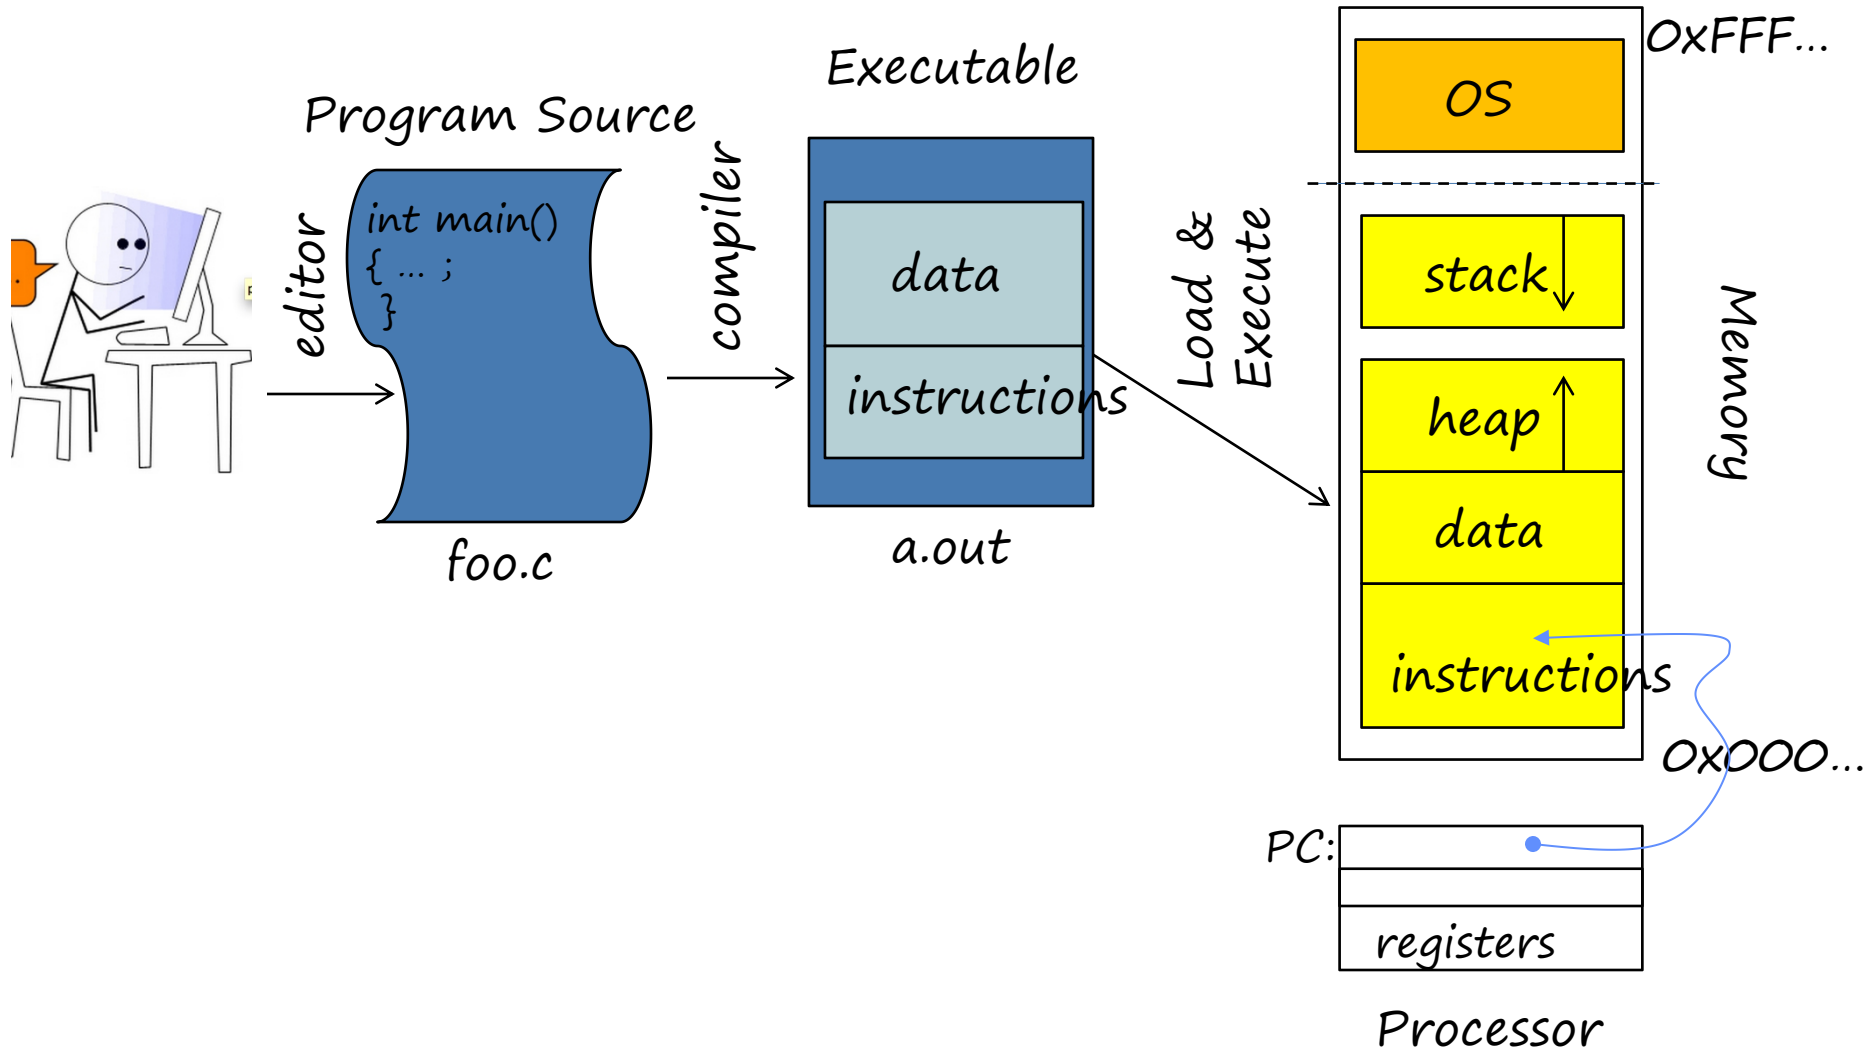
- ✧ **Protection:** prevent access to private memory of other processes
 - ✧ Kernel data protected from User programs
 - ✧ Programs protected from themselves
 - ✧ May want to give special behavior to different memory regions (Read Only, Invisible to user programs, etc)
- ✧ **Controlled overlap:** sometimes we want to share memory across processes.
 - ✧ E.g., communication across processes, share code
 - ✧ Need to control such overlap

Important Aspects of Memory Multiplexing

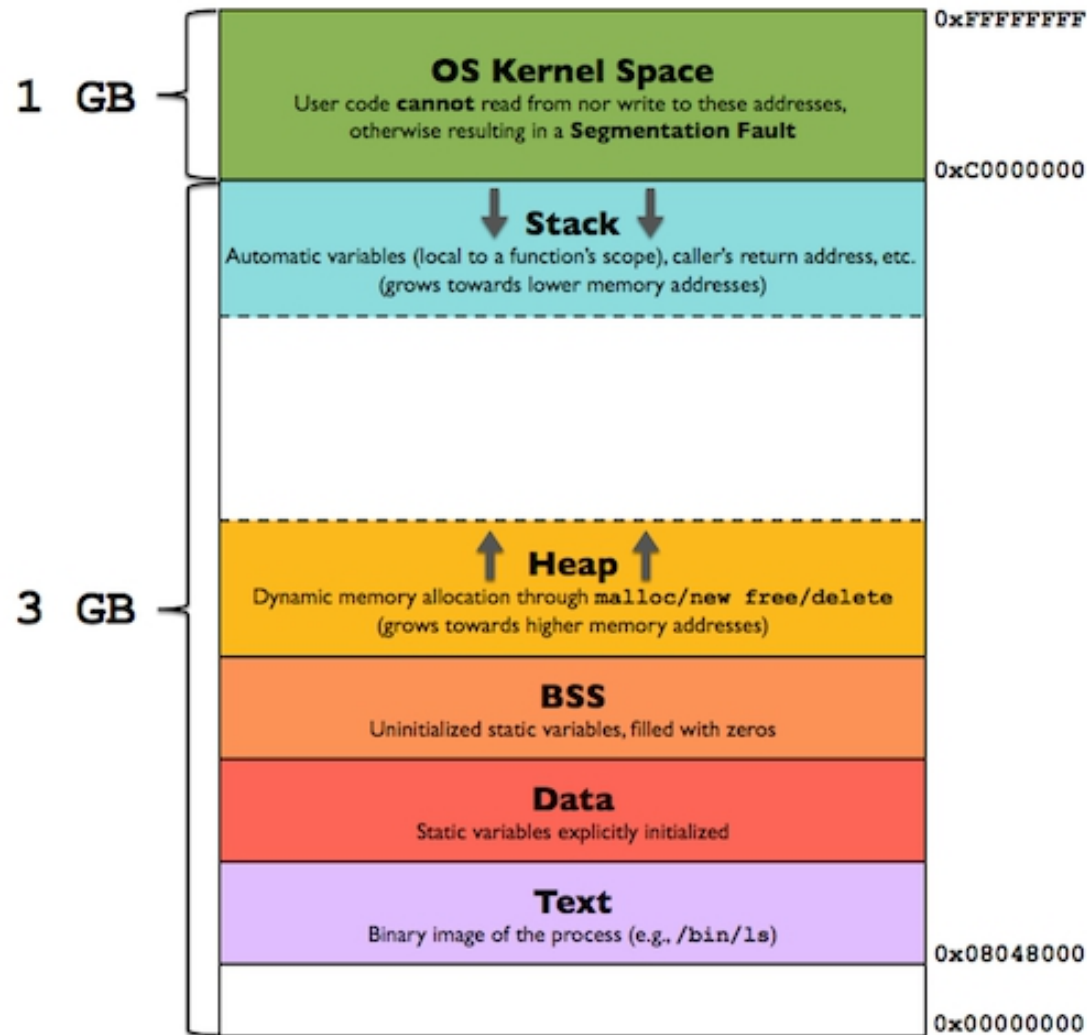
✧ Translation:

- ✧ Ability to translate accesses from one address space (virtual) to a different one (physical)
- ✧ When translation exists, processor uses virtual addresses, physical memory uses physical addresses
- ✧ Side effects:
 - ✳ Can be used to give uniform view of memory to programs
 - ✳ Can be used to provide protection (e.g., avoid overlap)
 - ✳ Can be used to control overlap

Recall: OS Bottom Line: Run Programs



Recall: Address Space

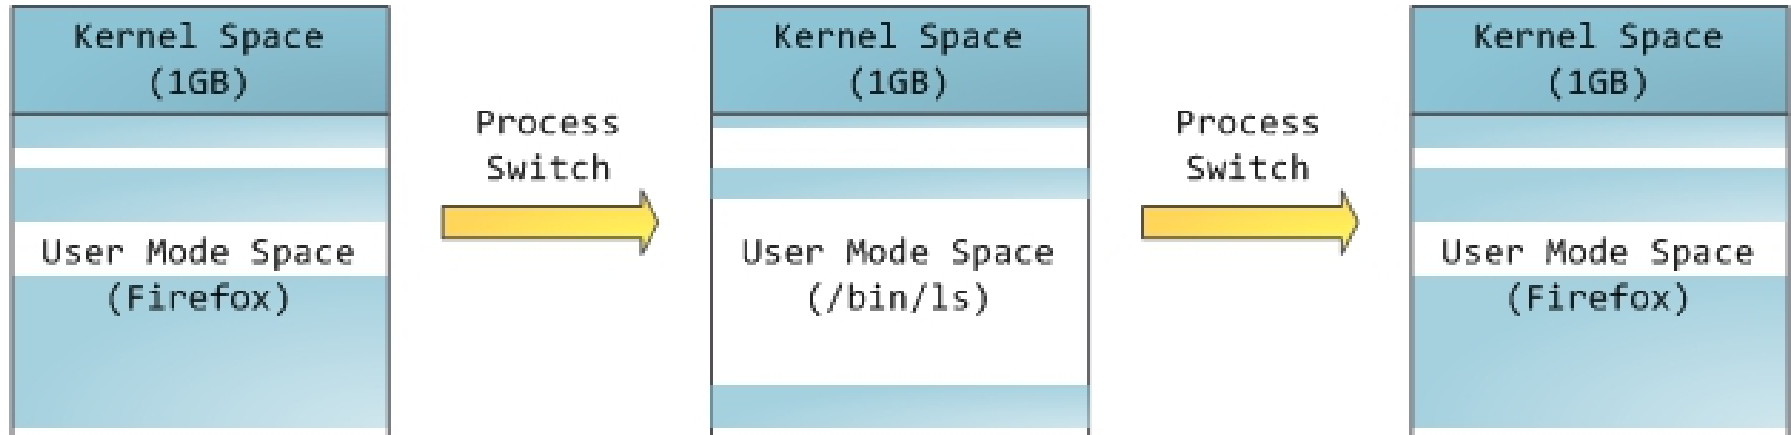


Each process has
its own user
(address) space

Each process
thinks it has 2^{32}
(2^{64}) byte
memory

<https://gabrieletolomei.wordpress.com/miscellanea/operating-systems/in-memory-layout/> (text → code segment, read only / execute), BSS: block started by symbol, “static int i” (data segment: read-write)

Recall: Context Switch



One address maps to one byte.

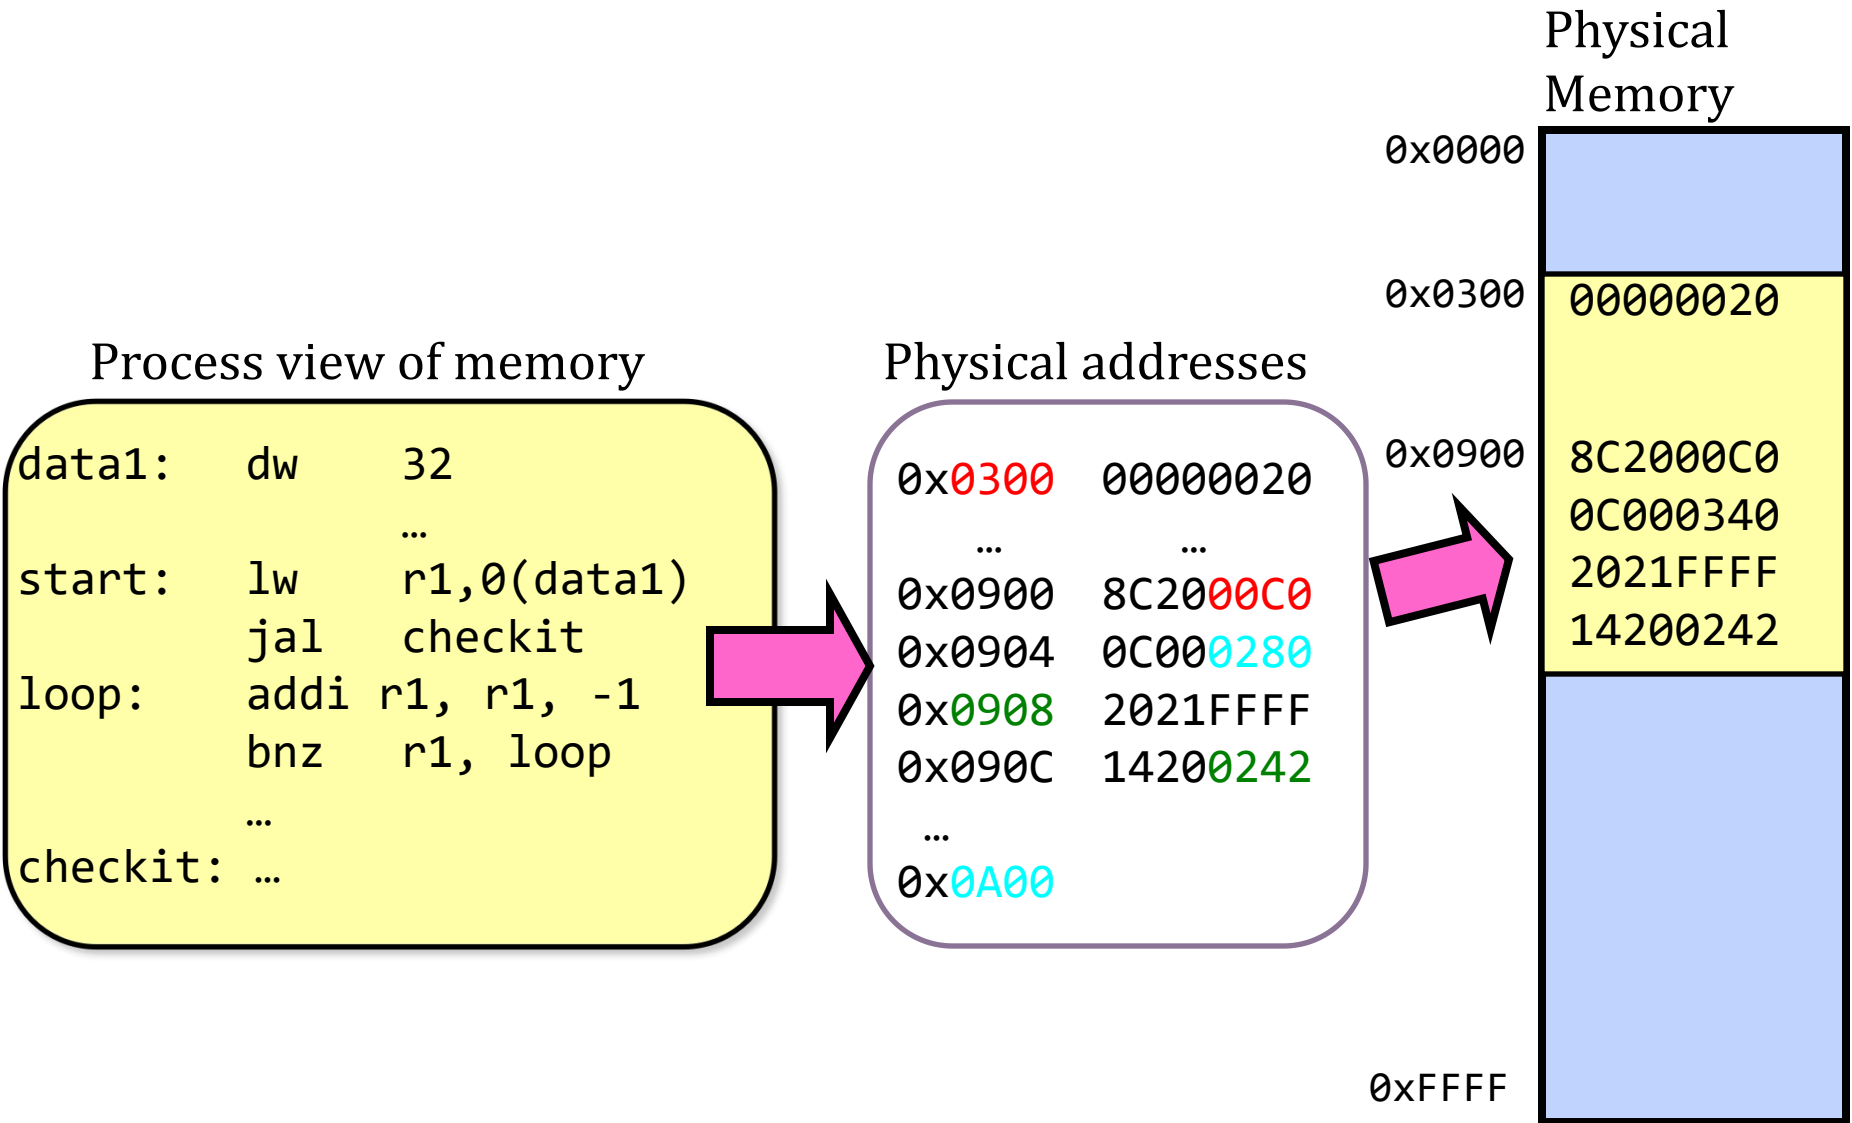
On a 32-bit system,

- The maximum amount of memory in a process is 2^{32} bytes = **4GB**.

Then, how about a 64-bit system? = **16EB**

<http://duartes.org/gustavo/blog/post/anatomy-of-a-program-in-memory/>

Binding of Instructions and Data to Memory



2nd copy of program from previous example

- Call `fork()`

Process view of memory

```
data1:    dw    32
          ...
start:    lw     r1,0(data1)
          jal    checkit
loop:     addi   r1, r1, -1
          bnz    r1, r0, loop
          ...
checkit:  ...
```

Physical addresses

```
0x0300    00000020
          ...
0x0900    8C2000C0
0x0904    0C000280
0x0908    2021FFFF
0x090C    14200242
          ...
0x0A00
```

Physical
Memory

0x0000

0x0300

0x0900

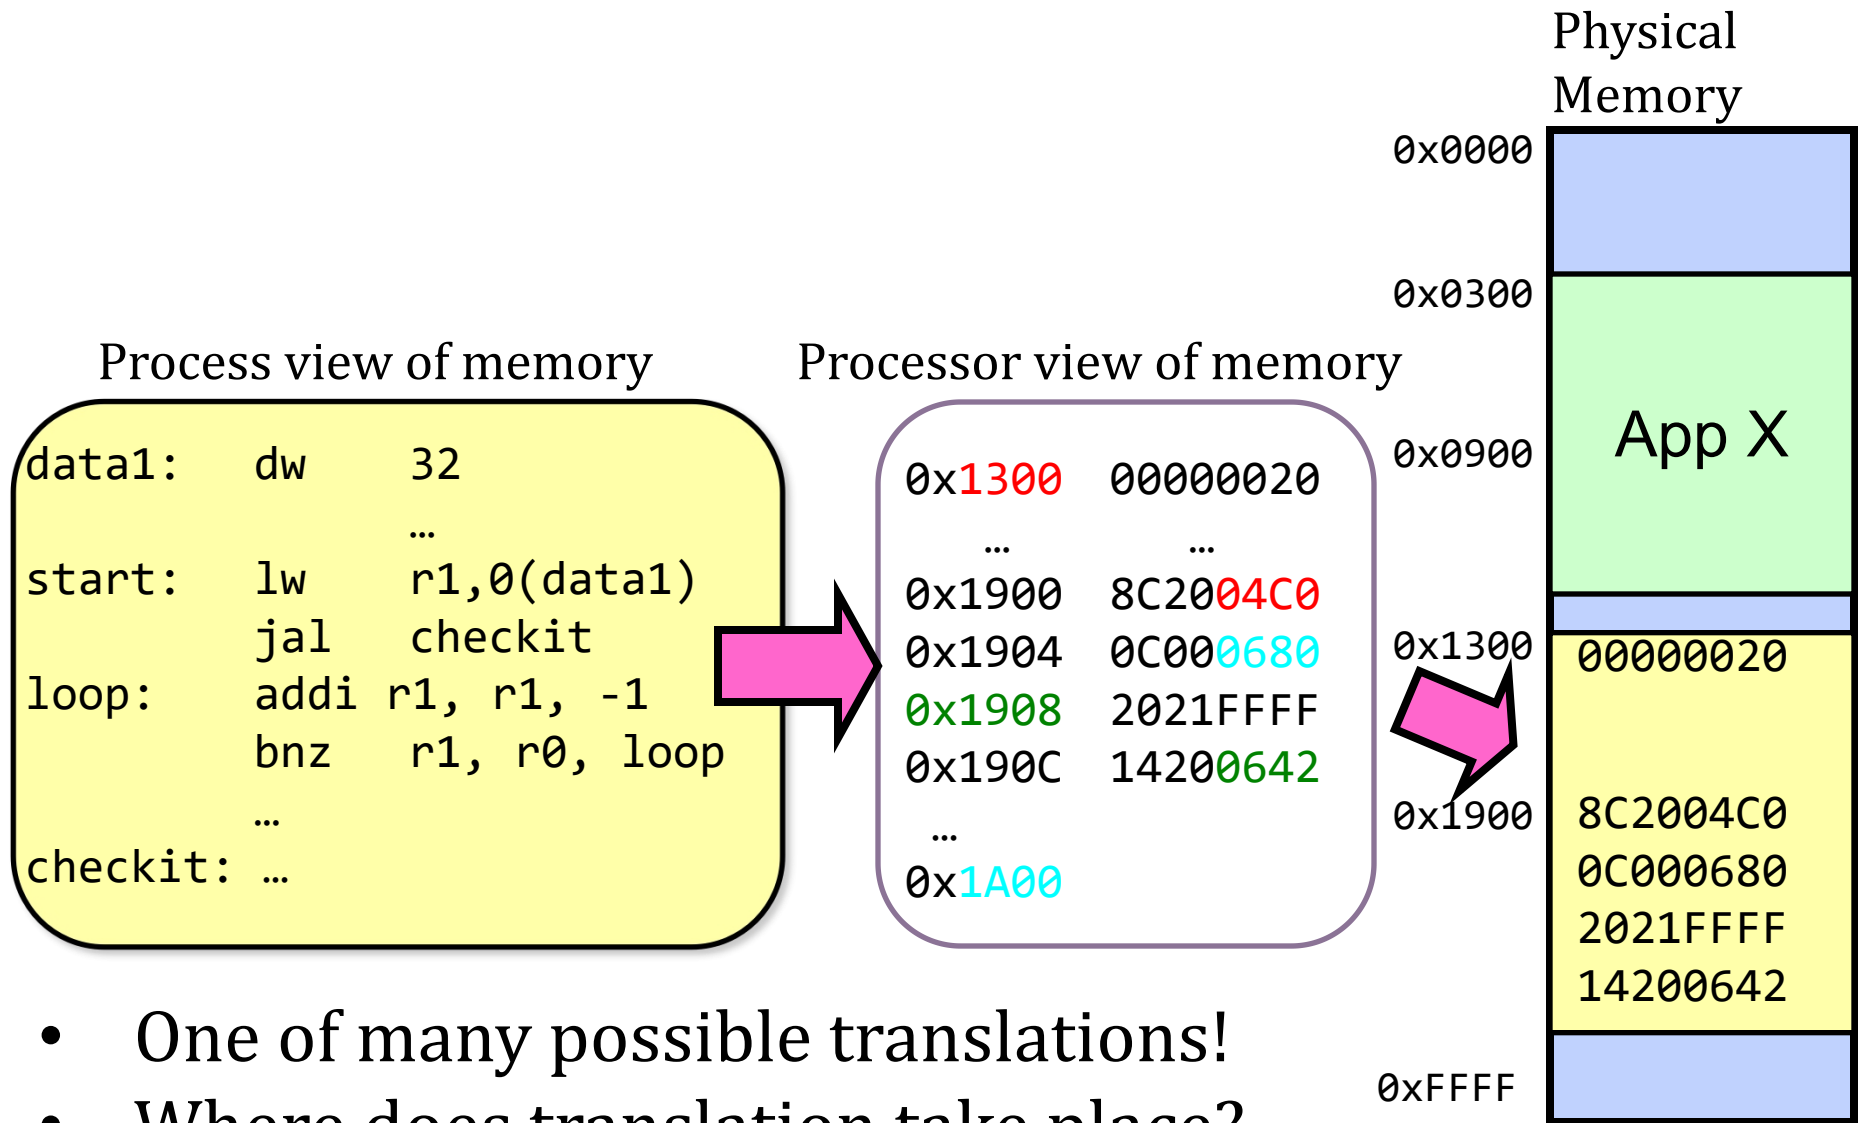
?

App X

0xFFFF

Need address translation!

2nd copy of program from previous example

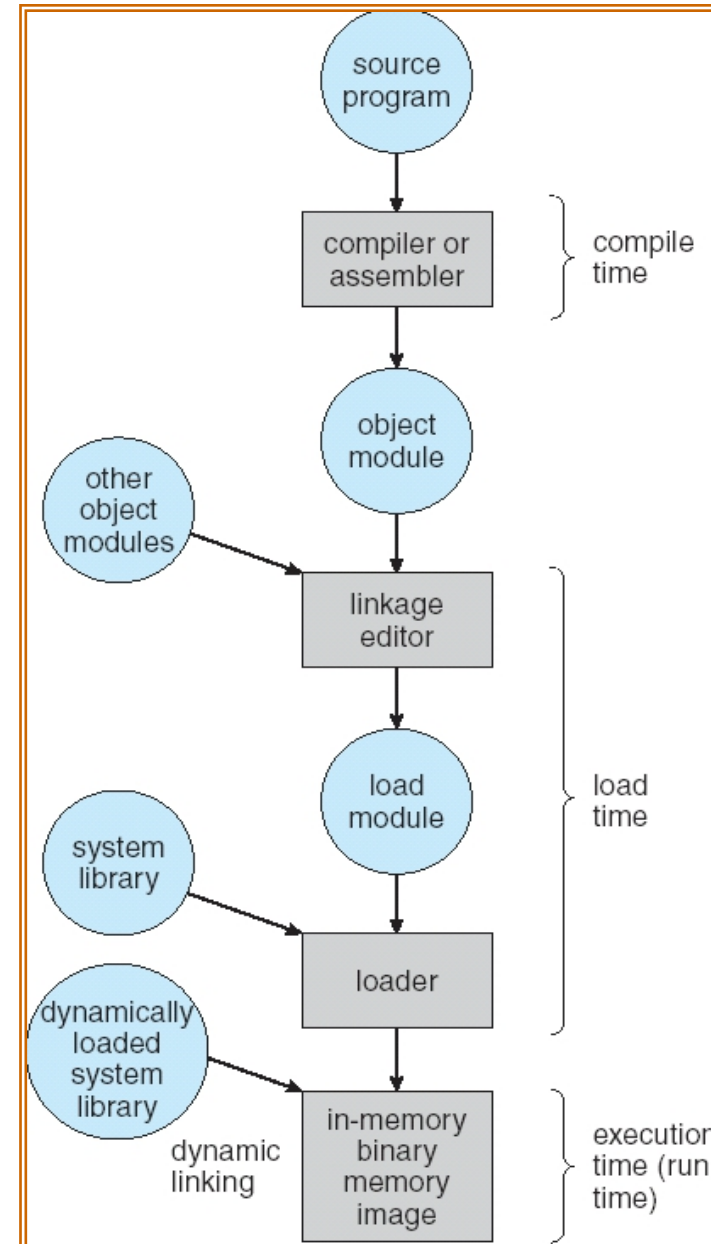


- One of many possible translations!
- Where does translation take place?

Compile time, Link/Load time, or Execution time?

Multi-step Processing of a Program for Execution

- ❖ Preparation of a program for execution involves components at:
 - (1) Compile time (i.e., “gcc”)
 - (2) Link time (UNIX “ld” does link)
 - (3) Load time
 - (4) Execution time (e.g., dynamic libs)
- ❖ Addresses can be bound to final values anywhere in this path
 - ❏ Depends on hardware support
 - ❏ Also depends on operating system
- ❖ Dynamic Libraries
 - ❏ Linking postponed until execution
 - ❏ Small piece of code, *stub*, used to locate appropriate memory-resident library routine
 - ❏ Stub replaces itself with the address of the routine, and executes routine

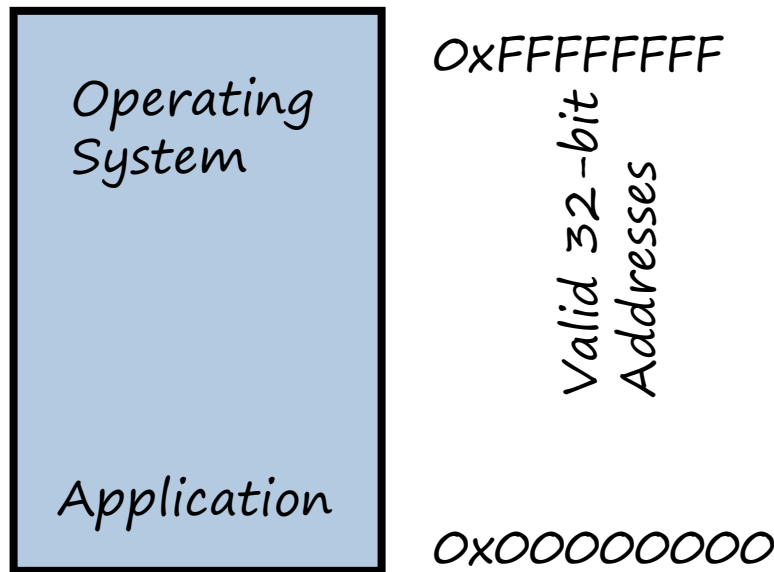


Multiplexing Memory Approaches

- ✧ Uniprogramming
- ✧ Multiprogramming
 - ✧ Without protection
 - ✧ With protection (base+bound)
- ✧ Virtual memory
 - ✧ Base & Bound
 - ✧ Segmentation
 - ✧ Paging
 - ✧ Paging + Segmentation

Uniprogramming


- ✧ Uniprogramming (no Translation or Protection)
 - ✧ Application always runs at same place in physical memory since only one application at a time
 - ✧ Application can access any physical address



- ✧ Application given illusion of dedicated machine by giving it reality of a dedicated machine

Multiprogramming (primitive stage)


- ✧ Multiprogramming without Translation or Protection
 - ✧ Must somehow prevent address overlap between threads



- ✧ Use Loader/Linker: Adjust addresses while program loaded into memory (loads, stores, jumps)
 - ✧ Everything adjusted to memory location of program
 - ✧ Translation done by a linker-loader (relocation)
 - ✧ Common in early days (... till Windows 3.x, 95?)
- ✧ With this solution, no protection: bugs in any program can cause other programs to crash or even the OS

Multiprogramming (Version with Protection)


- ✧ Can we protect programs from each other without translation?



- ✧ Yes: use two special registers *BaseAddr* and *LimitAddr* to prevent user from straying outside designated area
 - ✱ If user tries to access an illegal address, cause an error
- ✧ During switch, kernel loads new base/limit from PCB (Process Control Block)
 - ✱ User not allowed to change base/limit registers

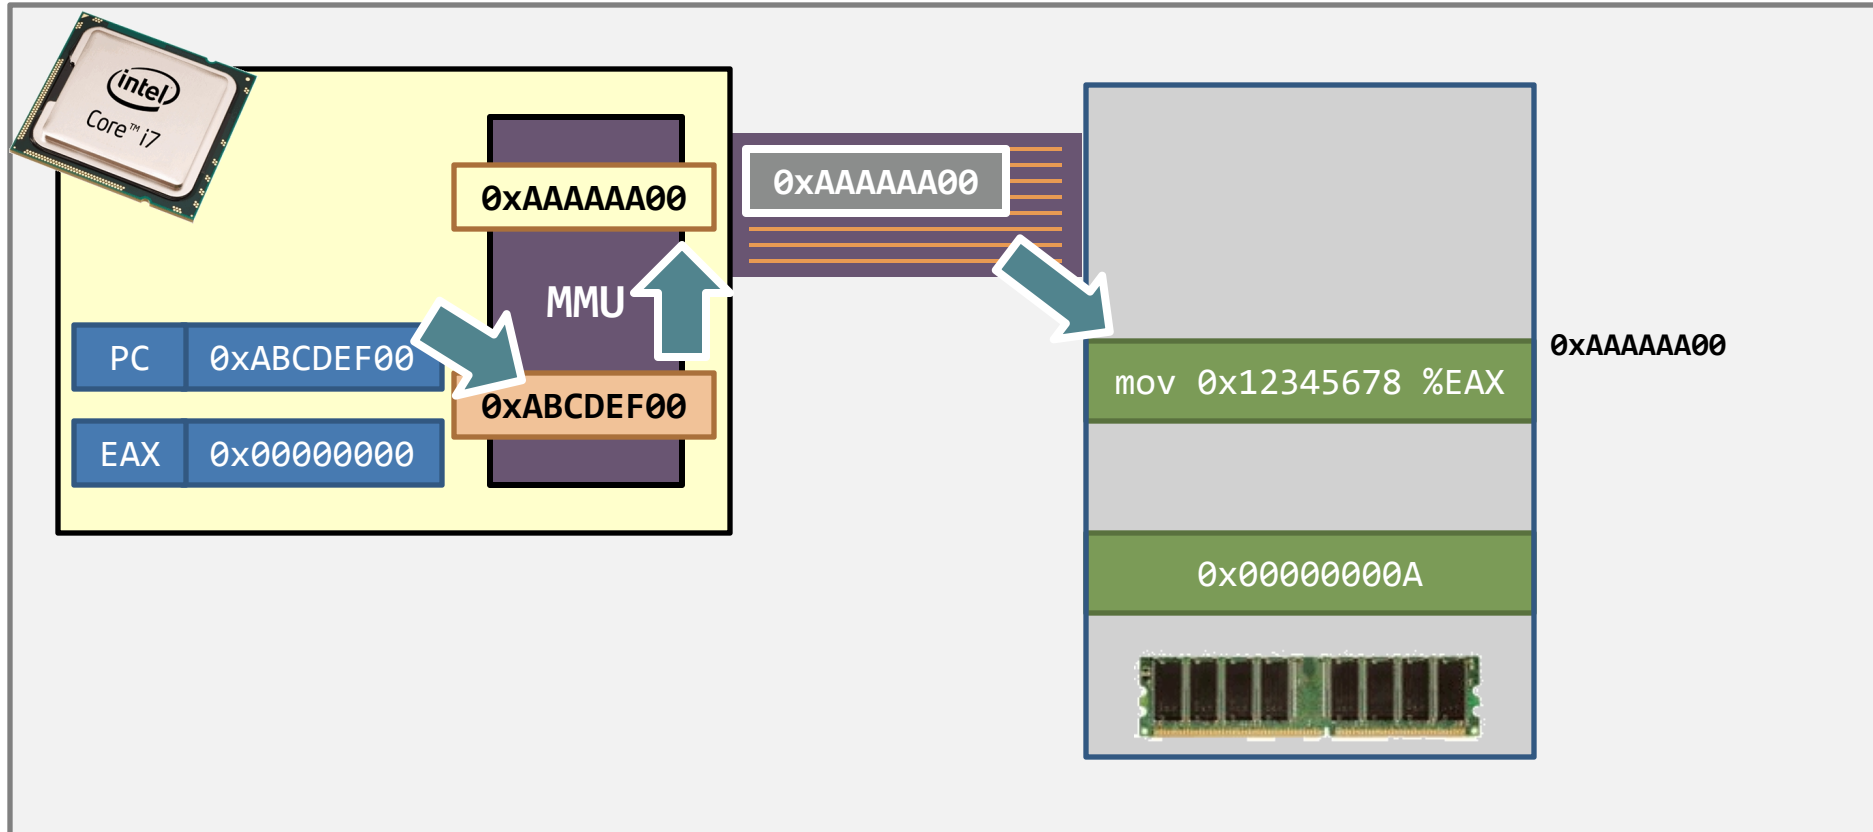
Virtual memory support in modern CPUs

- ✧ The **MMU – memory management unit**
 - ✧ Usually on-chip (but some architecture may off-chip or no hardware MMU)



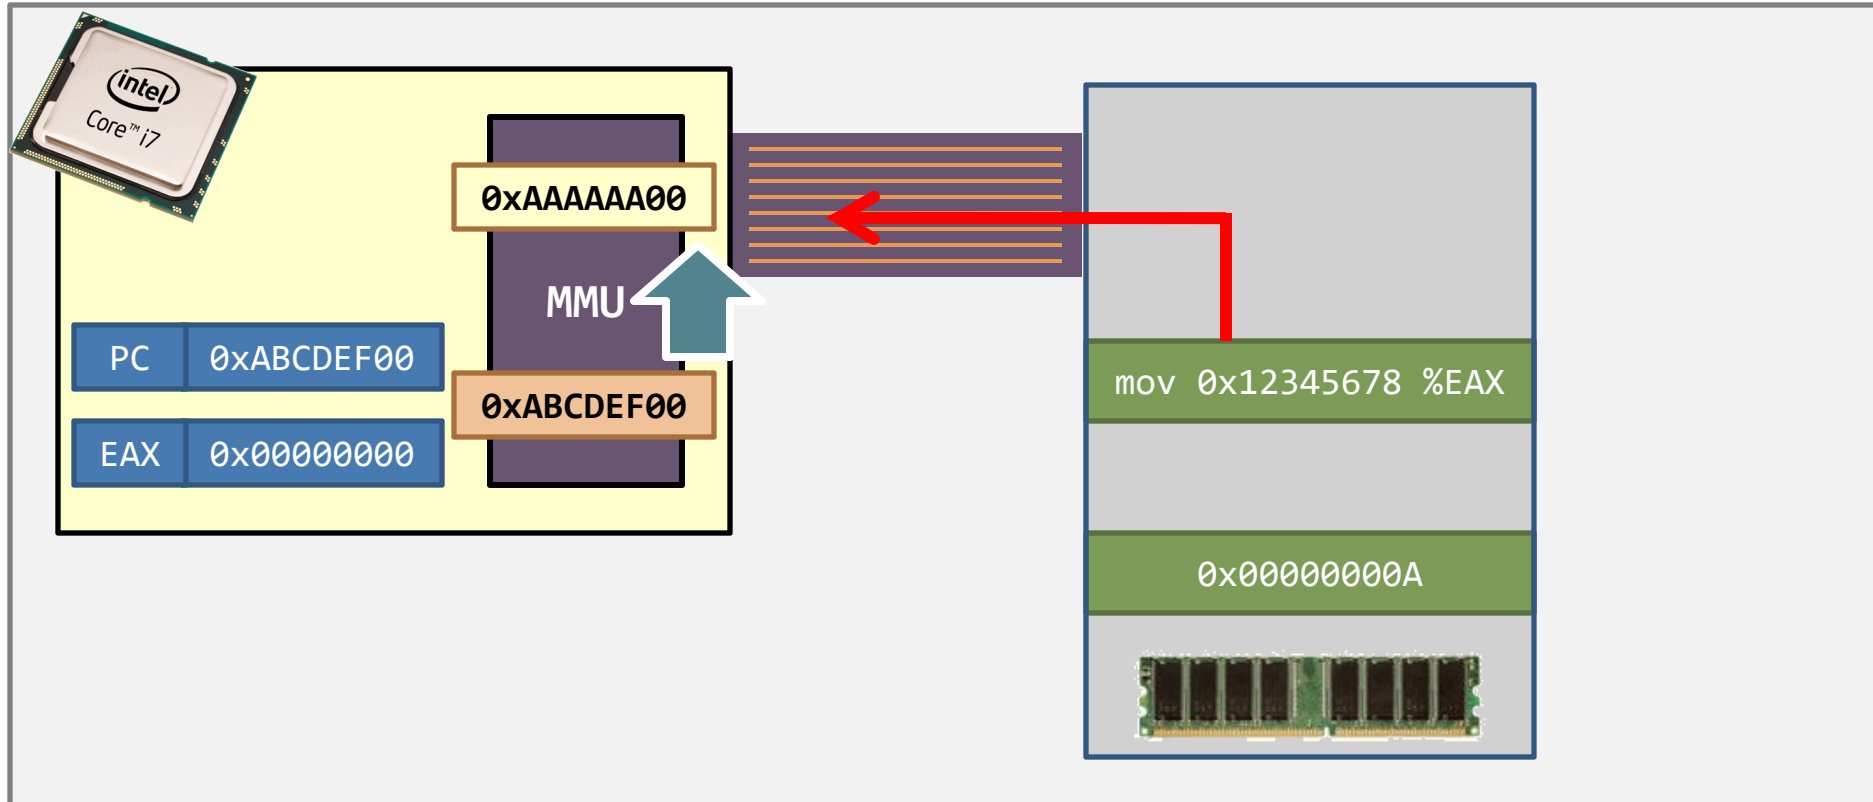
Virtual memory – how does it work?

- ❖ Step 1. When CPU wants to fetch an instruction
 - ❖ the **virtual address** is sent to MMU and
 - ❖ is translated into a **physical address**.



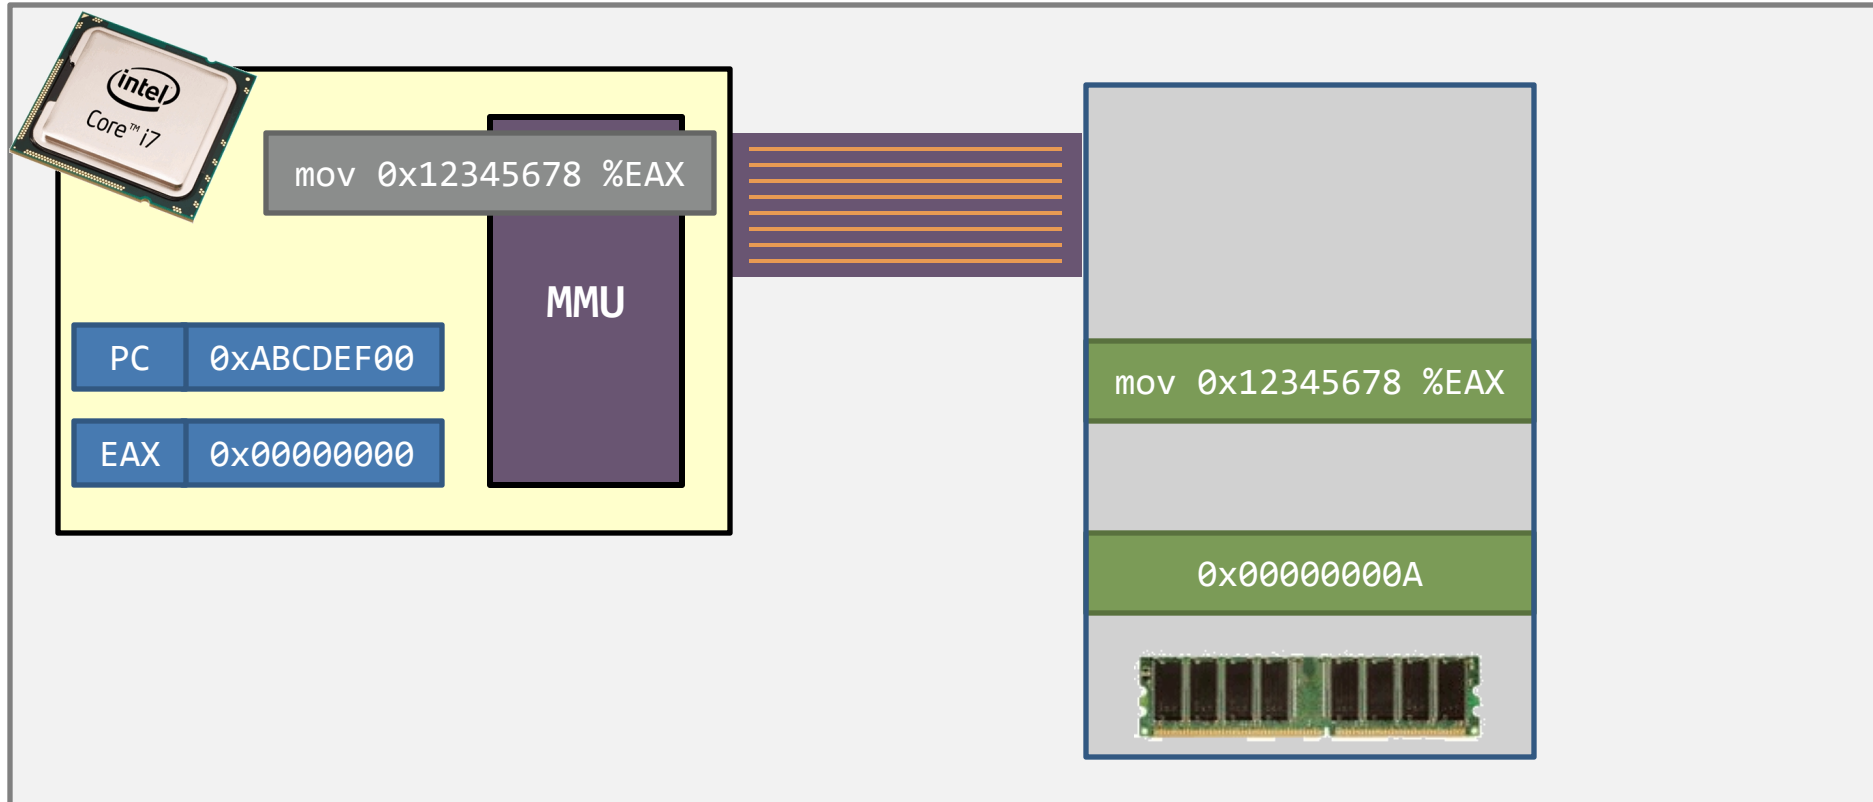
Virtual memory – how does it work?

- ✧ Step 2. The memory returns the instruction



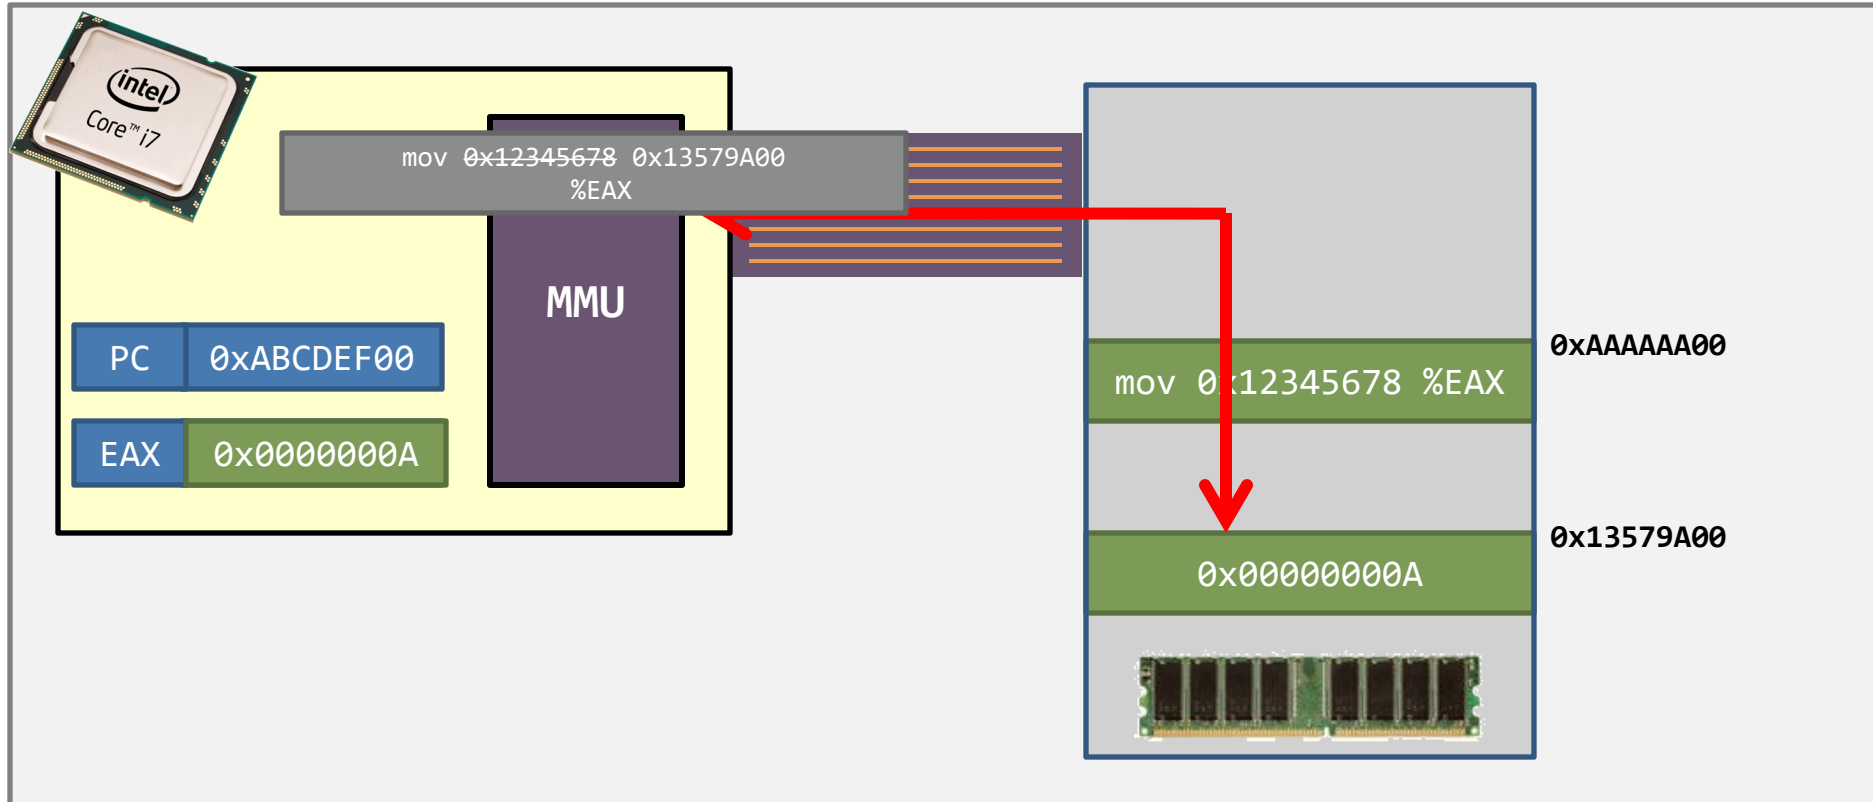
Virtual memory – how does it work?

- ✧ Step 3. The CPU decodes the instruction.
 - ✧ An instruction **uses virtual addresses**
 - ✧ but not physical addresses.

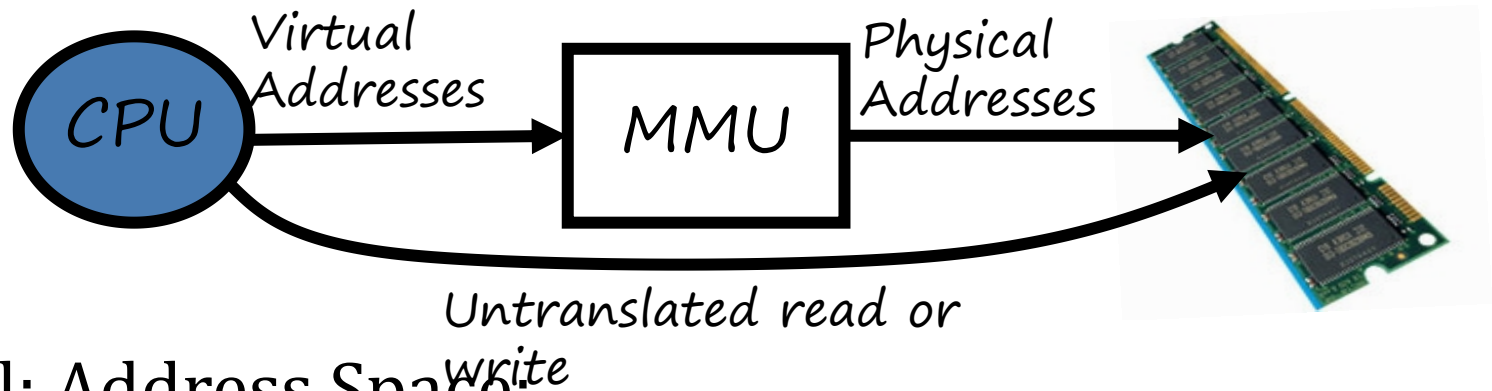


Virtual memory – how does it work?

- ✧ Step 4. With the help of the MMU, the target memory is retrieved.




General Address translation



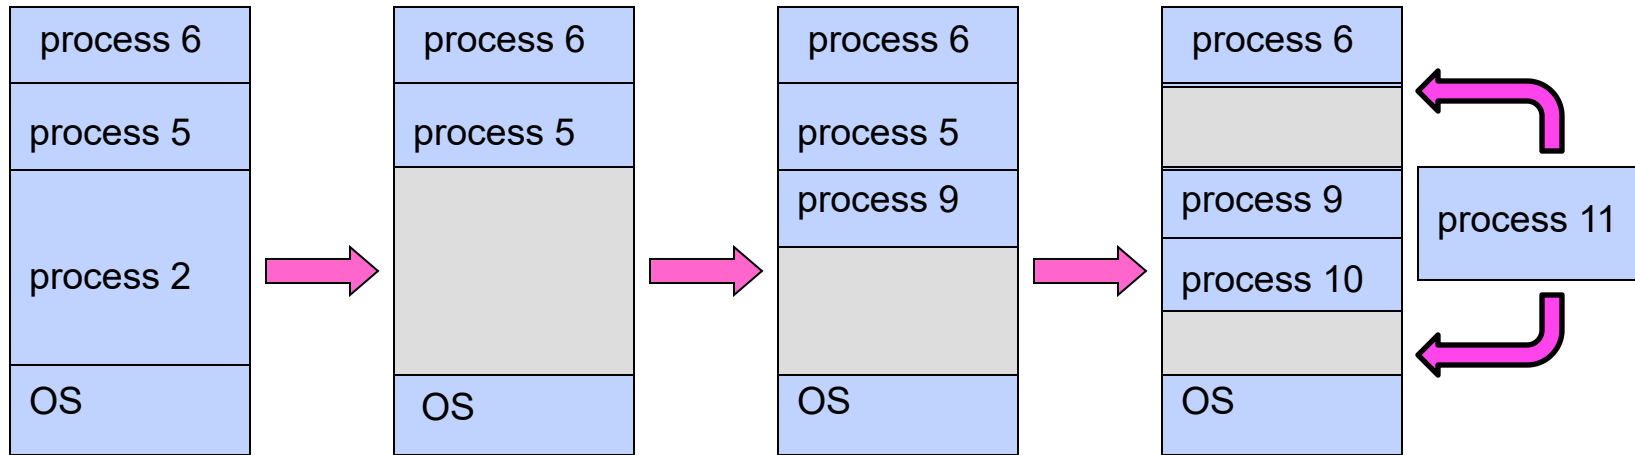
- ✧ Recall: Address Space:
 - ✧ All the addresses and state a process can touch
 - ✧ Each process has different address space
- ✧ Consequently, two views of memory:
 - ✧ View from the CPU (what program sees, virtual memory)
 - ✧ View from memory (physical memory)
 - ✧ Translation box (MMU) converts between the two views
- ✧ Translation makes it much easier to implement protection
 - ✧ If task A cannot even gain access to task B's data, no way for A to adversely affect B
- ✧ With translation, every program can be linked/loaded into same region of user address space

Simple Example: Base and Bounds




- ✧ Could use base/bounds for **dynamic address translation** – translation happens at execution:
 - ✧ Alter address of every load/store by adding “base”
 - ✧ Generate error if address bigger than limit
- ✧ This gives program the illusion that it is running on its own dedicated machine, with memory starting at 0
 - ✧ Program gets continuous region of memory
 - ✧ Addresses within program do not have to be relocated when program placed in different region of DRAM

Issues with Simple B&B Method



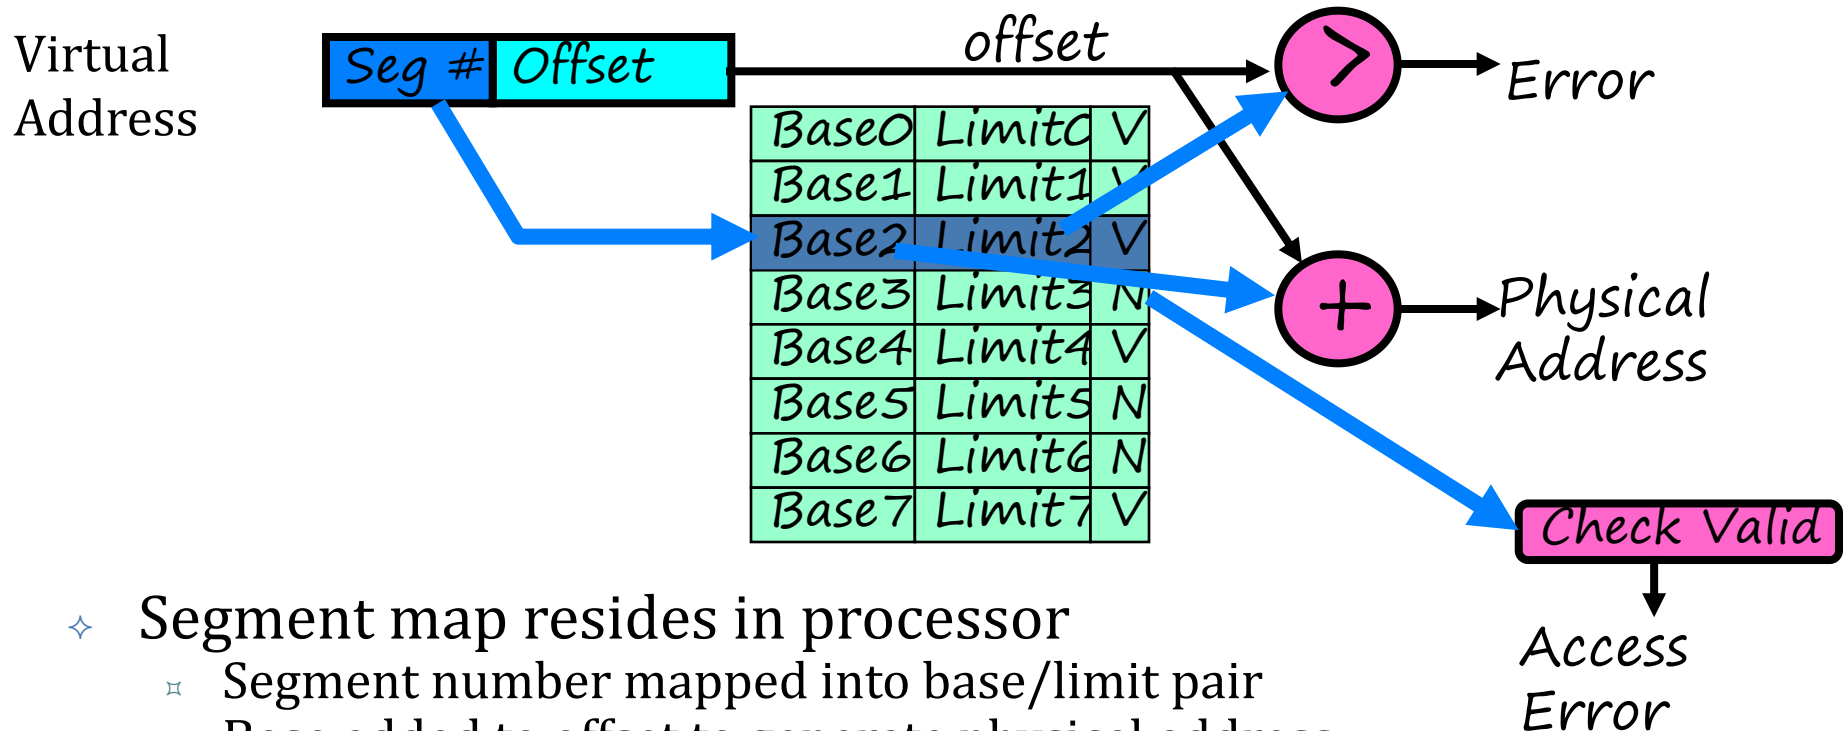
- ✧ Fragmentation problem over time
 - ✧ Not every process is same size → memory becomes fragmented
- ✧ Missing support for sparse address space
 - ✧ Would like to have multiple chunks/program (Code, Data, Stack)
- ✧ Hard to do inter-process sharing
 - ✧ Want to share code segments when possible
 - ✧ Want to share memory between processes
 - ✧ Helped by providing multiple segments per process

More Flexible Segmentation




- ✧ **Logical View:** multiple separate segments
 - ✧ Typical: Code, Data, Stack
 - ✧ Others: memory sharing, etc
- ✧ **Each segment is given region of contiguous memory**
 - ✧ Has a base and limit
 - ✧ Can reside anywhere in physical memory

Implementation of Multi-Segment Model




- ❖ Segment map resides in processor
 - ❏ Segment number mapped into base/limit pair
 - ❏ Base added to offset to generate physical address
 - ❏ Error check catches offset out of range
- ❖ As many chunks of physical memory as entries
 - ❏ Segment addressed by portion of virtual address
 - ❏ However, could be included in instruction instead:
 - ✱ x86 Example: `mov [es:bx],ax.`
- ❖ What is “V/N” (valid / not valid)?
 - ❏ Can mark segments as invalid; requires check as well

Example: Four Segments (16 bit addresses)




Virtual Address Format




Virtual Address Space

Seg ID #	Base	Limit
0 (code)	0x400 0	0x080 0
1 (data)	0x480 0	0x140 0
2 (shared)	0xF00 0	0x100 0
3 (stack)	0x000 0	0x300 0



Physical Address Space

Example: Four Segments (16 bit addresses)



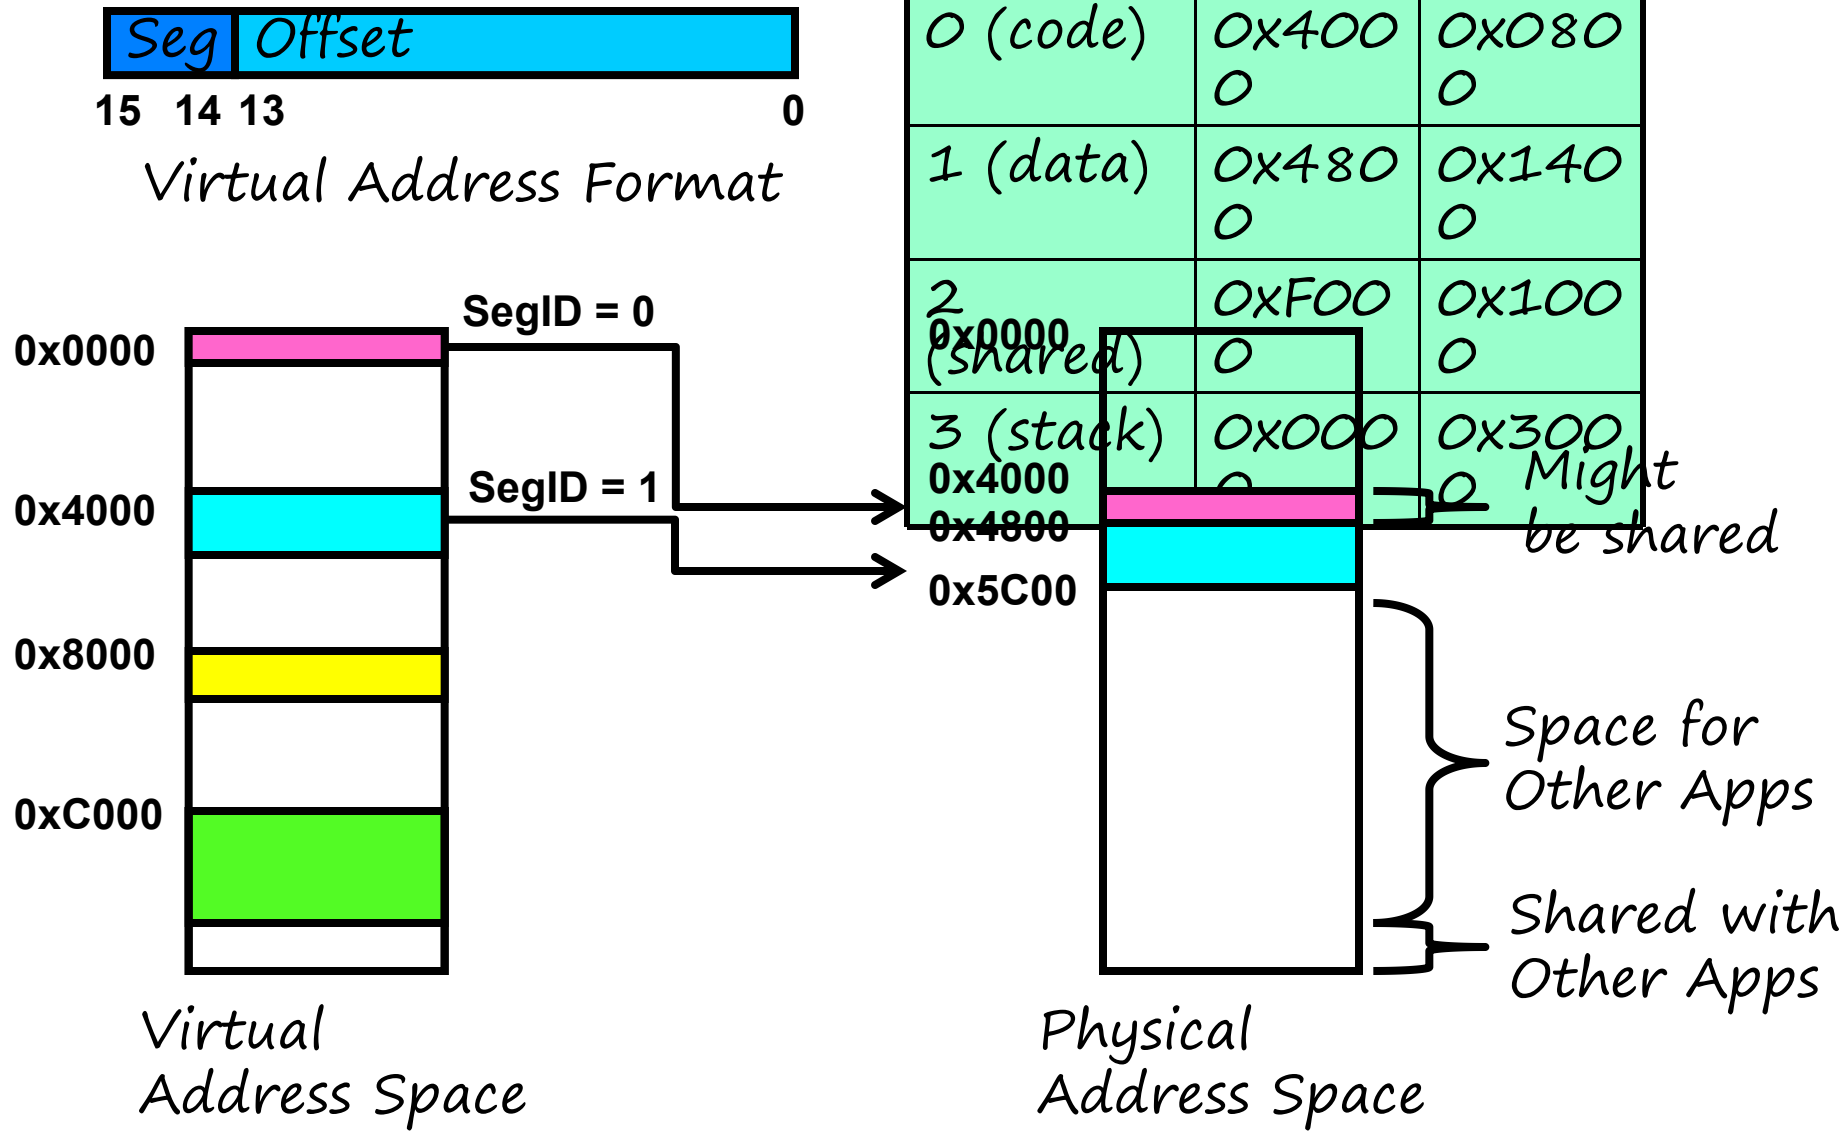
SegID = 0



Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

Physical Address Space

Example: Four Segments (16 bit addresses)




Example: Four Segments (16 bit addresses)

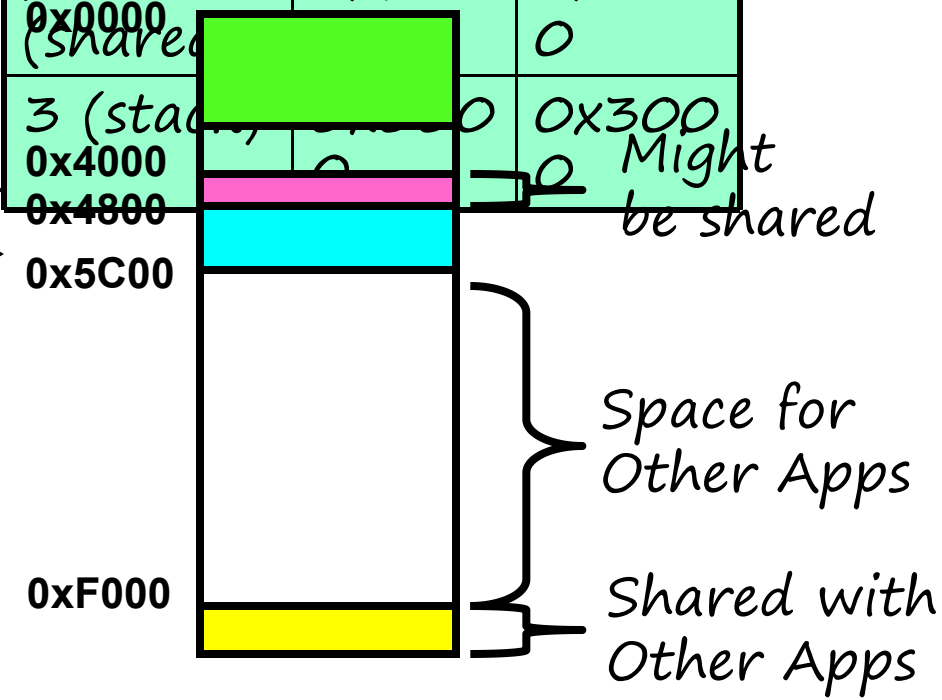


Virtual Address Format

Seg ID #	Base	Limit
0 (code)	0x400 0	0x080 0
1 (data)	0x480 0	0x140 0
2	0xF00 0	0x100 0
3 (stack)	0x4000 0x4800	0x300 0



Virtual Address Space



Physical Address Space

Example of Segment Translation (16b address)

```
0x240  main:   la $a0, varx
0x244                jal strlen
...
0x360  strlen: li $v0, 0 ;count
0x364  loop:   lb $t0, ($a0)
0x368                beq $r0,$t0, done
...
0x4050  varx    dw 0x314159
```

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

Let us simulate a bit of this code to see what happens

1. Fetch 0x240. Virtual segment #? 0; Offset? 0x240

Physical address? Base=0x4000, so physical addr=0x4240

Fetch instruction at 0x4240. Get "la \$a0, varx"

Move 0x4050 → \$a0, Move PC+4→PC

Example of Segment Translation (16b address)

0x240	main:	la	\$a0, varx
0x244		jal	strlen
...			...
0x360	strlen:	li	\$v0, 0 ;count
0x364	loop:	lb	\$t0, (\$a0)
0x368		beq	\$r0,\$t0, done
...			...
0x4050	varx	dw	0x314159

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

Let us simulate a bit of this code to see what happens

1. Fetch 0x240. Virtual segment #? 0; Offset? 0x240
Physical address? Base=0x4000, so physical addr=0x4240
Fetch instruction at 0x4240. Get "la \$a0, varx"
Move 0x4050 → \$a0, Move PC+4 → PC
2. Fetch 0x244. Translated to Physical=0x4244. Get "jal strlen"
Move 0x0248 → \$ra (return address!), Move 0x0360 → PC

Example of Segment Translation (16b address)

0x240	main:	la \$a0, varx
0x244		jal strlen
...		...
0x360	strlen:	li \$v0, 0 ;count
0x364	loop:	lb \$t0, (\$a0)
0x368		beq \$r0,\$t0, done
...		...
0x4050	varx	dw 0x314159

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

Let us simulate a bit of this code to see what happens

1. Fetch 0x240. Virtual segment #? 0; Offset? 0x240
Physical address? Base=0x4000, so physical addr=0x4240
Fetch instruction at 0x4240. Get "la \$a0, varx"
Move 0x4050 → \$a0, Move PC+4→PC
2. Fetch 0x244. Translated to Physical=0x4244. Get "jal strlen"
Move 0x0248 → \$ra (return address!), Move 0x0360 → PC
3. Fetch 0x360. Translated to Physical=0x4360. Get "li \$v0, 0"
Move 0x0000 → \$v0, Move PC+4→PC

Example of Segment Translation (16b address)

0x240	main:	la \$a0, varx
0x244		jal strlen
...		...
0x360	strlen:	li \$v0, 0 ;count
0x364	loop:	lb \$t0, (\$a0)
0x368		beq \$r0,\$t0, done
...		...
0x4050	varx	dw 0x314159

Seg ID #	Base	Limit
0 (code)	0x4000	0x0800
1 (data)	0x4800	0x1400
2 (shared)	0xF000	0x1000
3 (stack)	0x0000	0x3000

Let us simulate a bit of this code to see what happens

1. Fetch 0x240. Virtual segment #? 0; Offset? 0x40
Physical address? Base=0x4000, so physical address = 0x4000 + 0x40 = 0x4040
Fetch instruction at 0x4040. Get "la \$a0, varx"
Move 0x4050 → \$a0, Move PC+4 → PC
2. Fetch 0x244. Translated to Physical=0x4044. Get "jal strlen"
Move 0x0248 → \$ra (return address!), Move 0x0360 → PC
3. Fetch 0x360. Translated to Physical=0x4360. Get "li \$v0, 0"
Move 0x0000 → \$v0, Move PC+4 → PC
4. Fetch 0x0364. Translated to Physical=0x4364. Get "lb \$t0, (\$a0)" Since \$a0 is 0x4050, try to load byte from 0x4050, Translate 0x4050 (0100 0000 0101 000).
Virtual segment #? 1; Offset? 0x50 Physical address? Base=0x4800, Physical address = 0x4800 + 0x50 = 0x4850, Load Byte from 0x4850 → \$t0, Move PC+4 → PC

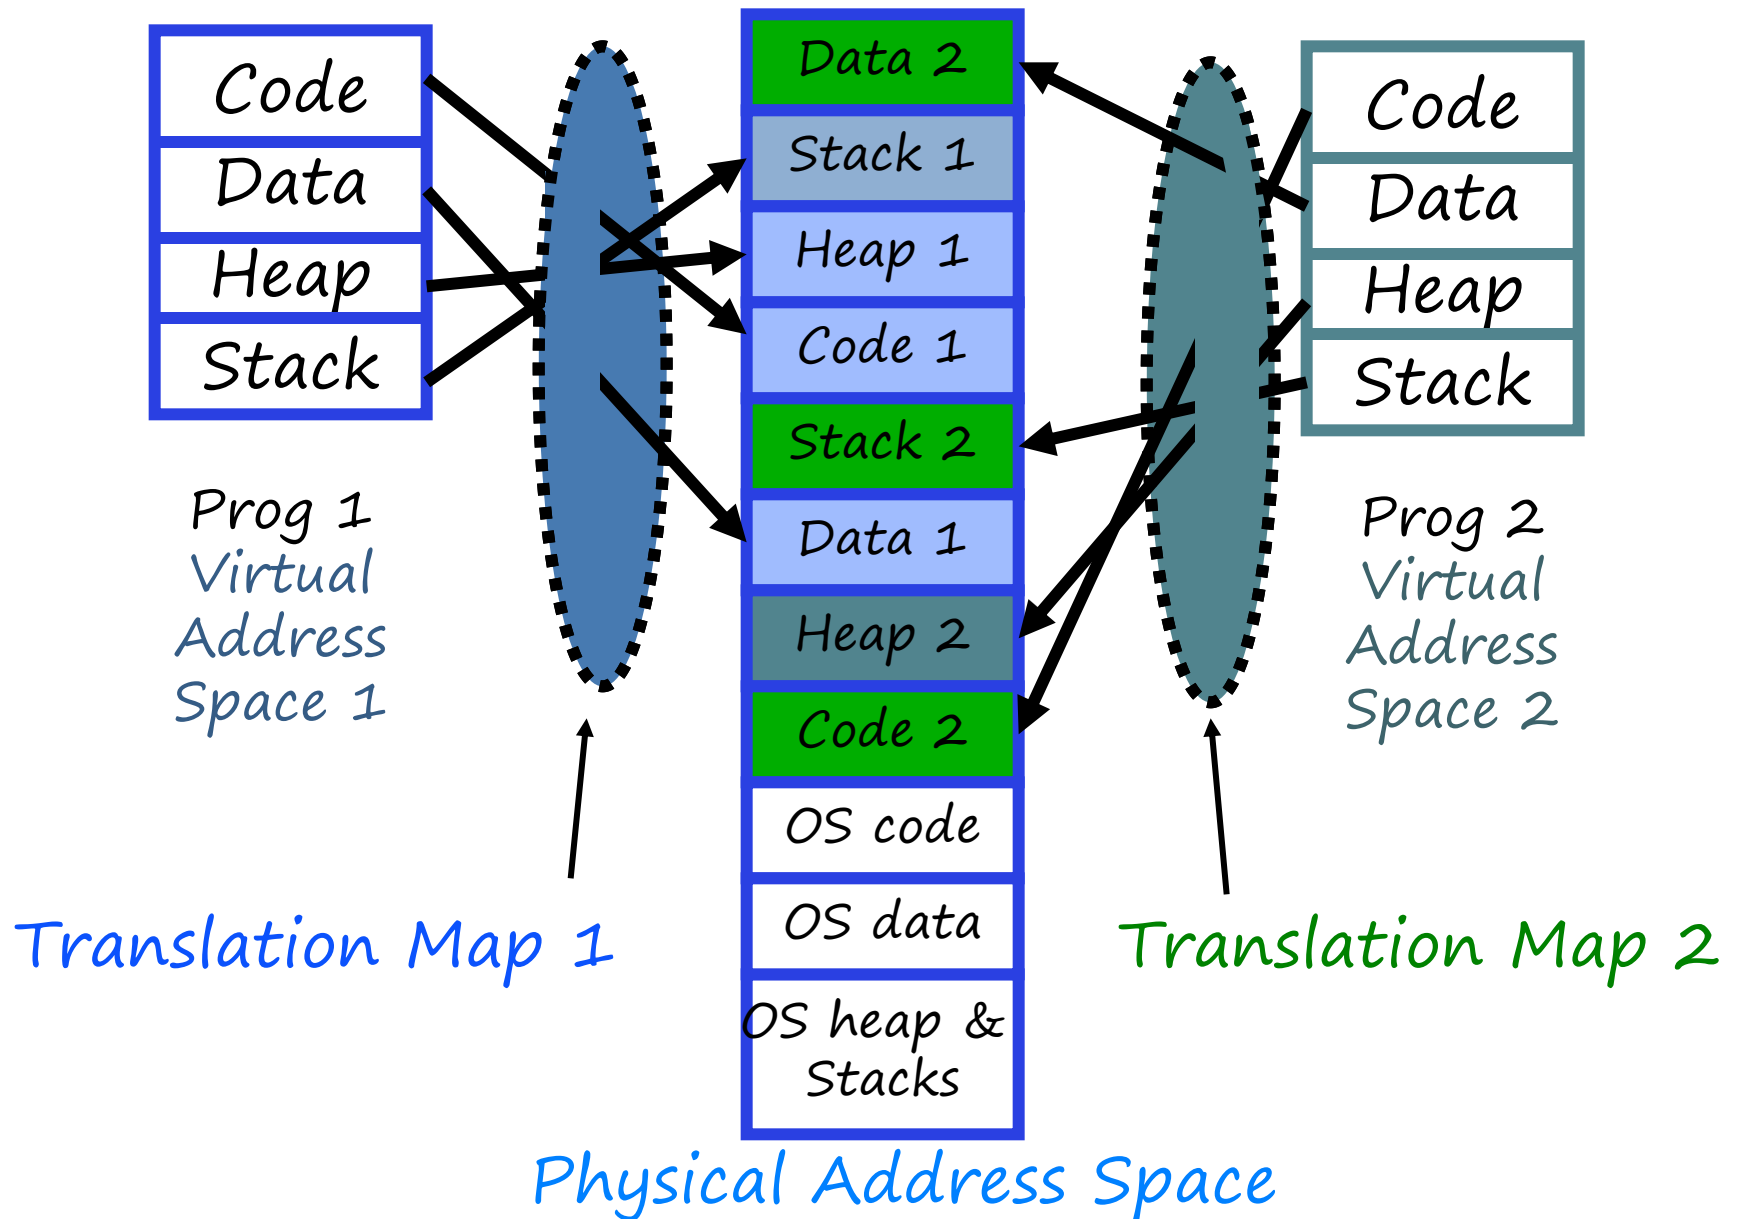
Observations about Segmentation

- ✧ Virtual address space has holes
 - ✧ Segmentation efficient for sparse address spaces
 - ✧ A correct program should never address gaps
 - ✱ If it does, trap to kernel and dump core
- ✧ When is it OK to address outside valid range?
 - ✧ This is how the stack and heap are allowed to grow
 - ✧ For instance, stack takes fault, system automatically increases size of stack
- ✧ Need protection mode in segment table
 - ✧ For example, code segment would be read-only
 - ✧ Data and stack would be read-write (stores allowed)
 - ✧ Shared segment could be read-only or read-write
- ✧ What must be saved/restored on context switch?
 - ✧ Segment table stored in CPU, not in memory (small)
 - ✧ Might store all of processes memory onto disk when switched (called “swapping”)

Problems with Segmentation

- ✧ Must fit variable-sized chunks into physical memory
- ✧ May move processes multiple times to fit everything
- ✧ **Fragmentation**: wasted space
 - ✧ **External**: free gaps between allocated chunks
 - ✧ **Internal**: do not need all memory within allocated chunks

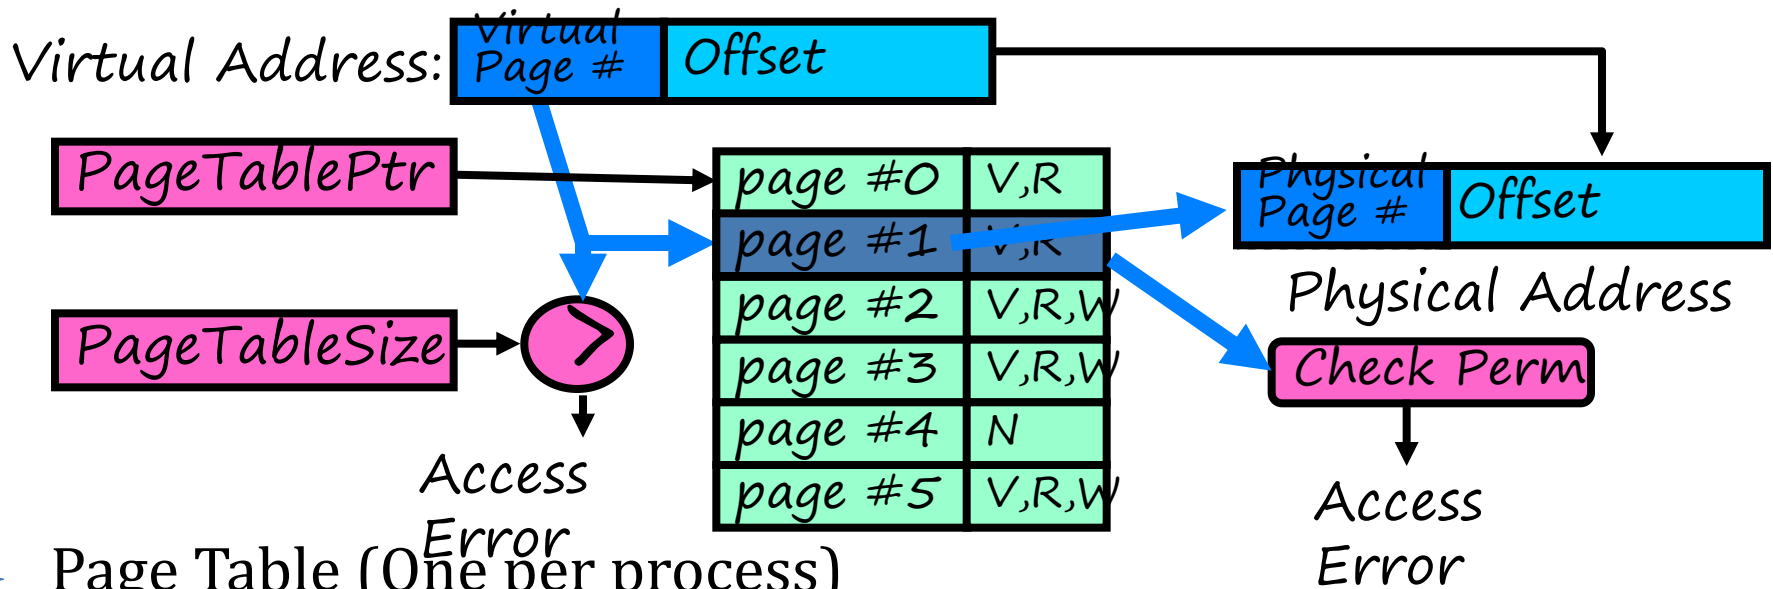
General Address Translation



Paging: Physical Memory in Fixed Size Chunks

- ✧ Solution to fragmentation from segments?
 - ✧ Allocate physical memory in fixed size chunks (“pages”)
 - ✧ Every chunk of physical memory is equivalent
 - ✱ Can use simple vector of bits to handle allocation:
00110001110001101 ... 110010
 - ✱ Each bit represents page of physical memory
1 \Rightarrow allocated, 0 \Rightarrow free
- ✧ Should pages be as big as our previous segments?
 - ✧ No: Can lead to lots of internal fragmentation
 - ✱ Typically have small pages (1K-16K)
 - ✧ Consequently: need multiple pages per segment

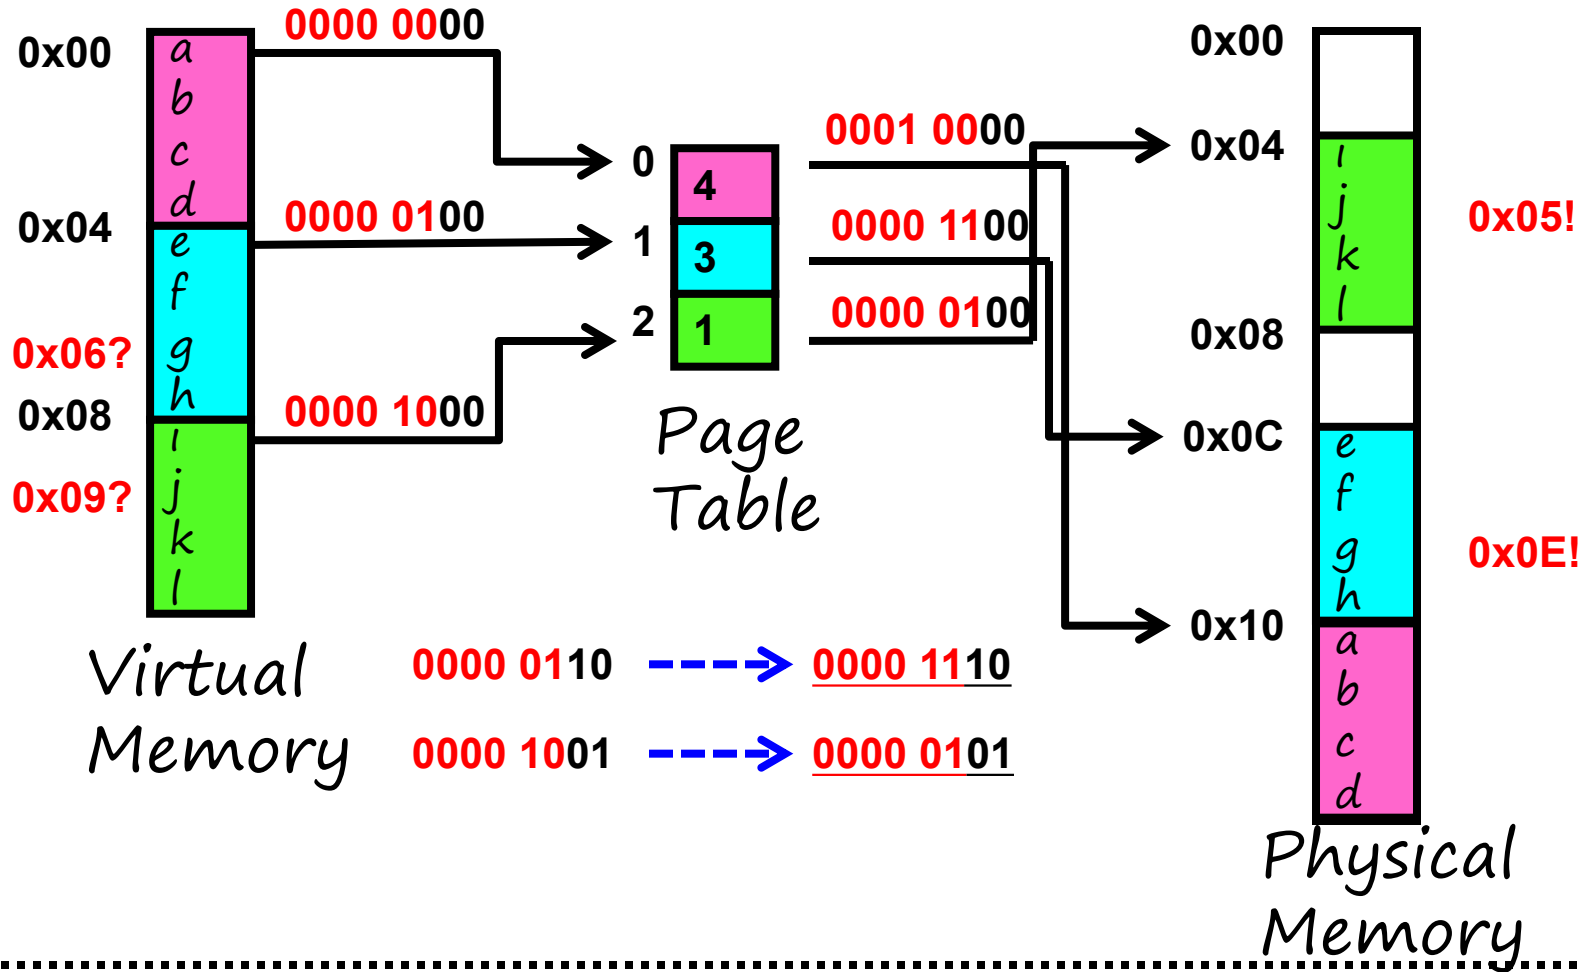
How to Implement Paging?



- ❖ Page Table (One per process)
 - ❏ Resides in physical memory
 - ❏ Contains physical page and permission for each virtual page
 - ✱ Permissions include: Valid bits, Read, Write, etc
- ❖ Virtual address mapping
 - ❏ Offset from Virtual address copied to Physical Address
 - ✱ Example: 10 bit offset \Rightarrow 1024-byte pages
 - ❏ Virtual page # is all remaining bits
 - ✱ Example for 32-bits: $32 - 10 = 22$ bits, i.e. 4 million entries
 - ✱ Physical page # copied from table into physical address
 - ❏ Check Page Table bounds and permissions

Simple Page Table Example

Example (4 byte pages)



What about Sharing?

Virtual Address
(Process A):

Virtual Page #	Offset
-------------------	--------

PageTablePtrA


page #0	V,R
page #1	V,R
page #2	V,R,W
page #3	V,R,W
page #4	N
page #5	V,R,W

PageTablePtrB

page #0	V,R
page #1	N
page #2	V,R,W
page #3	N
page #4	V,R
page #5	V,R,W

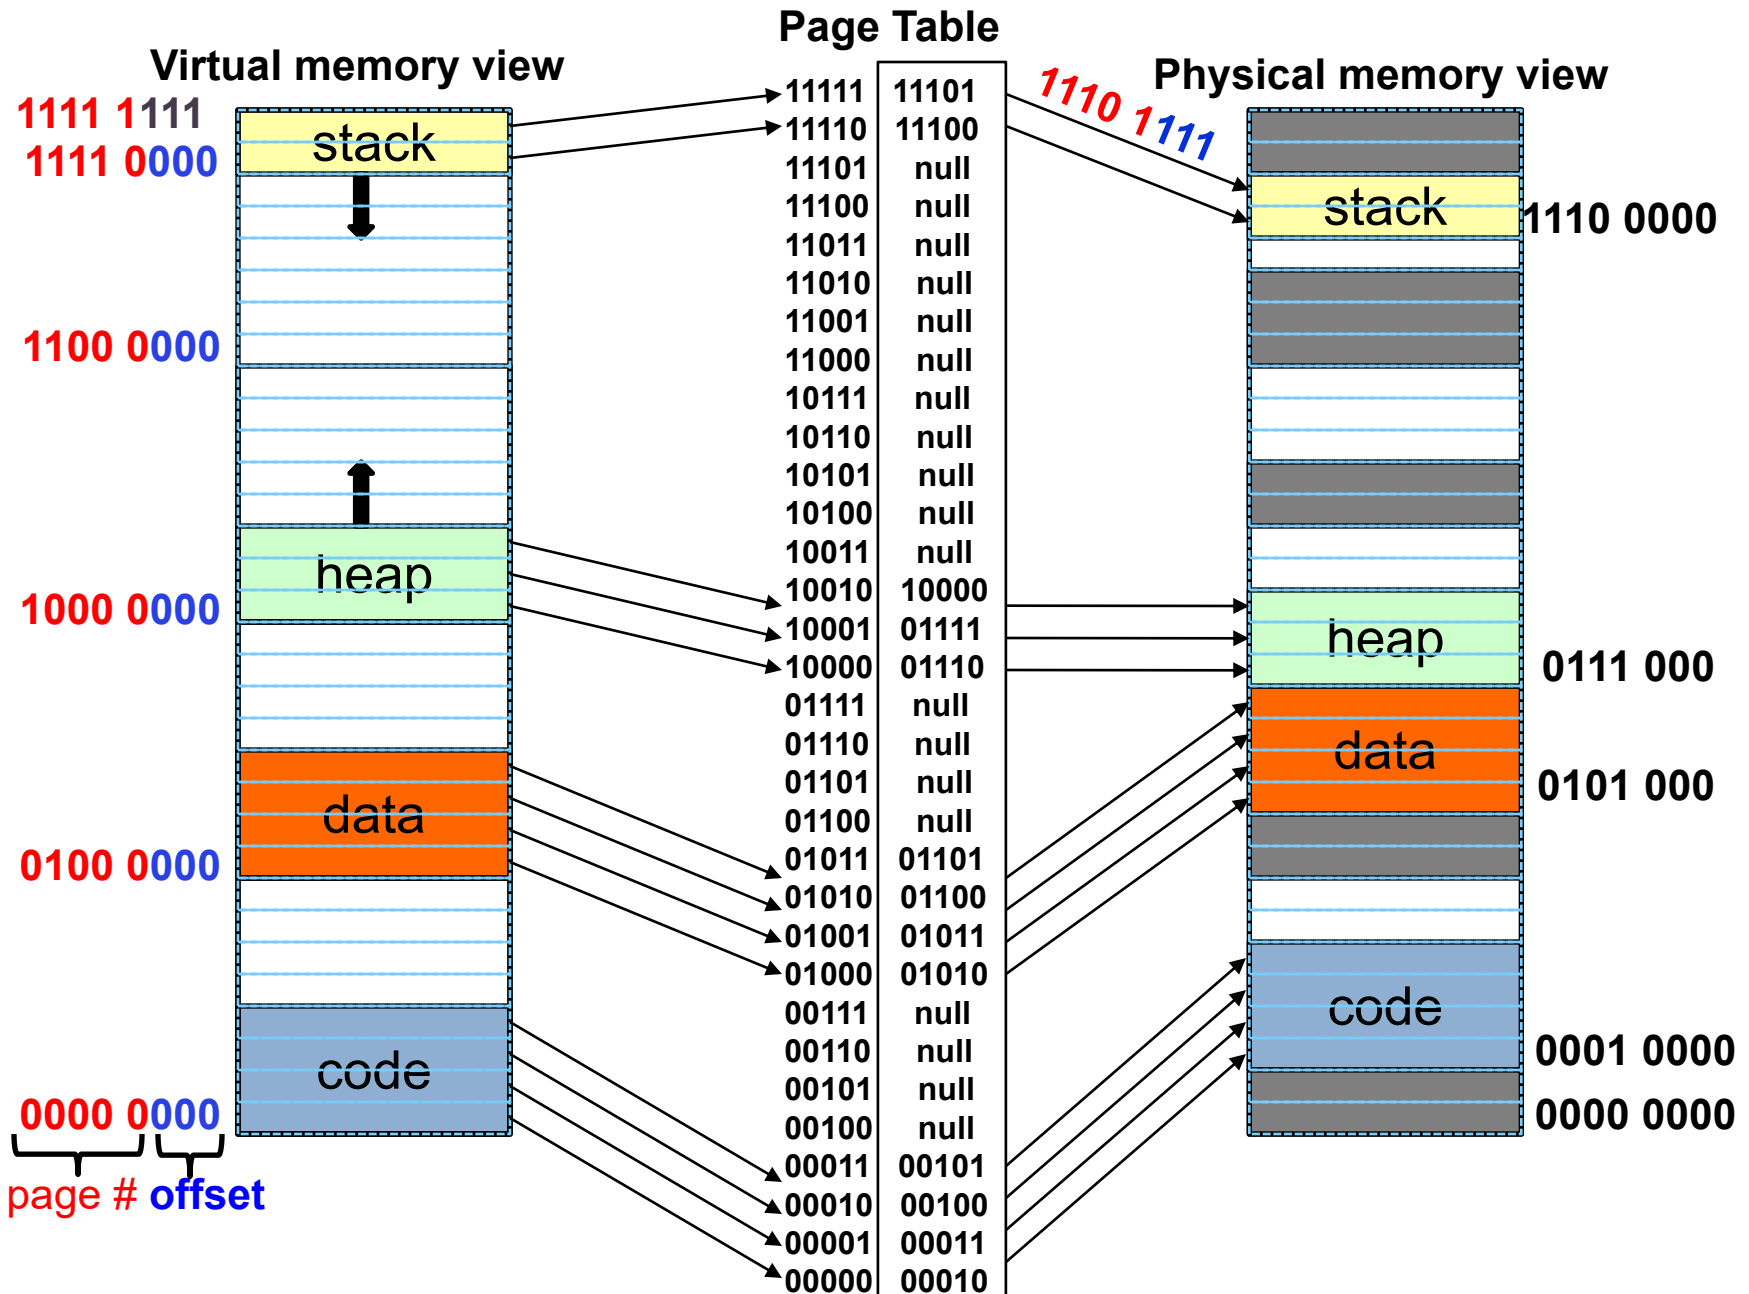
Virtual Address
(Process B):

Virtual Page #	Offset
-------------------	--------

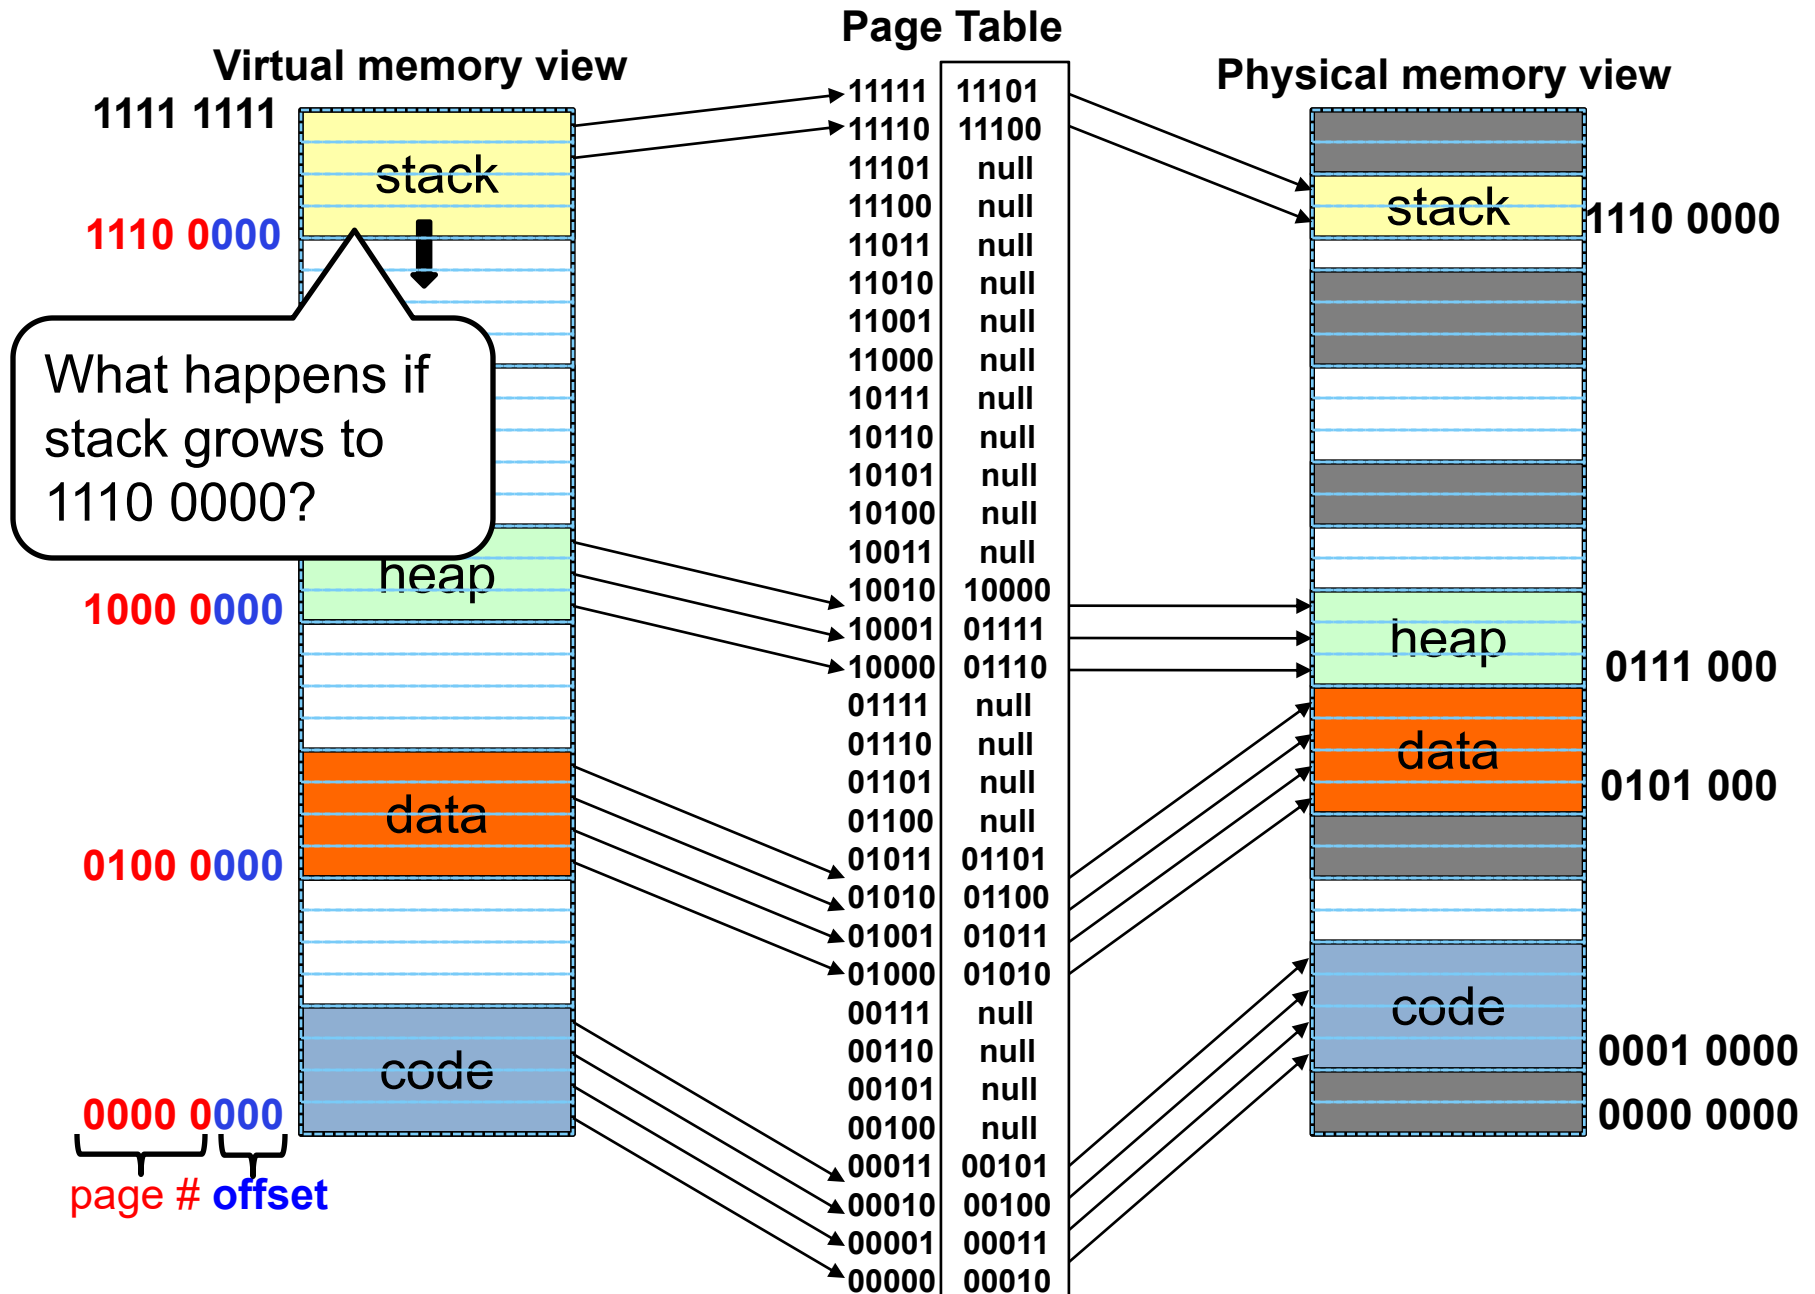


This physical page
appears in address
space of both processes

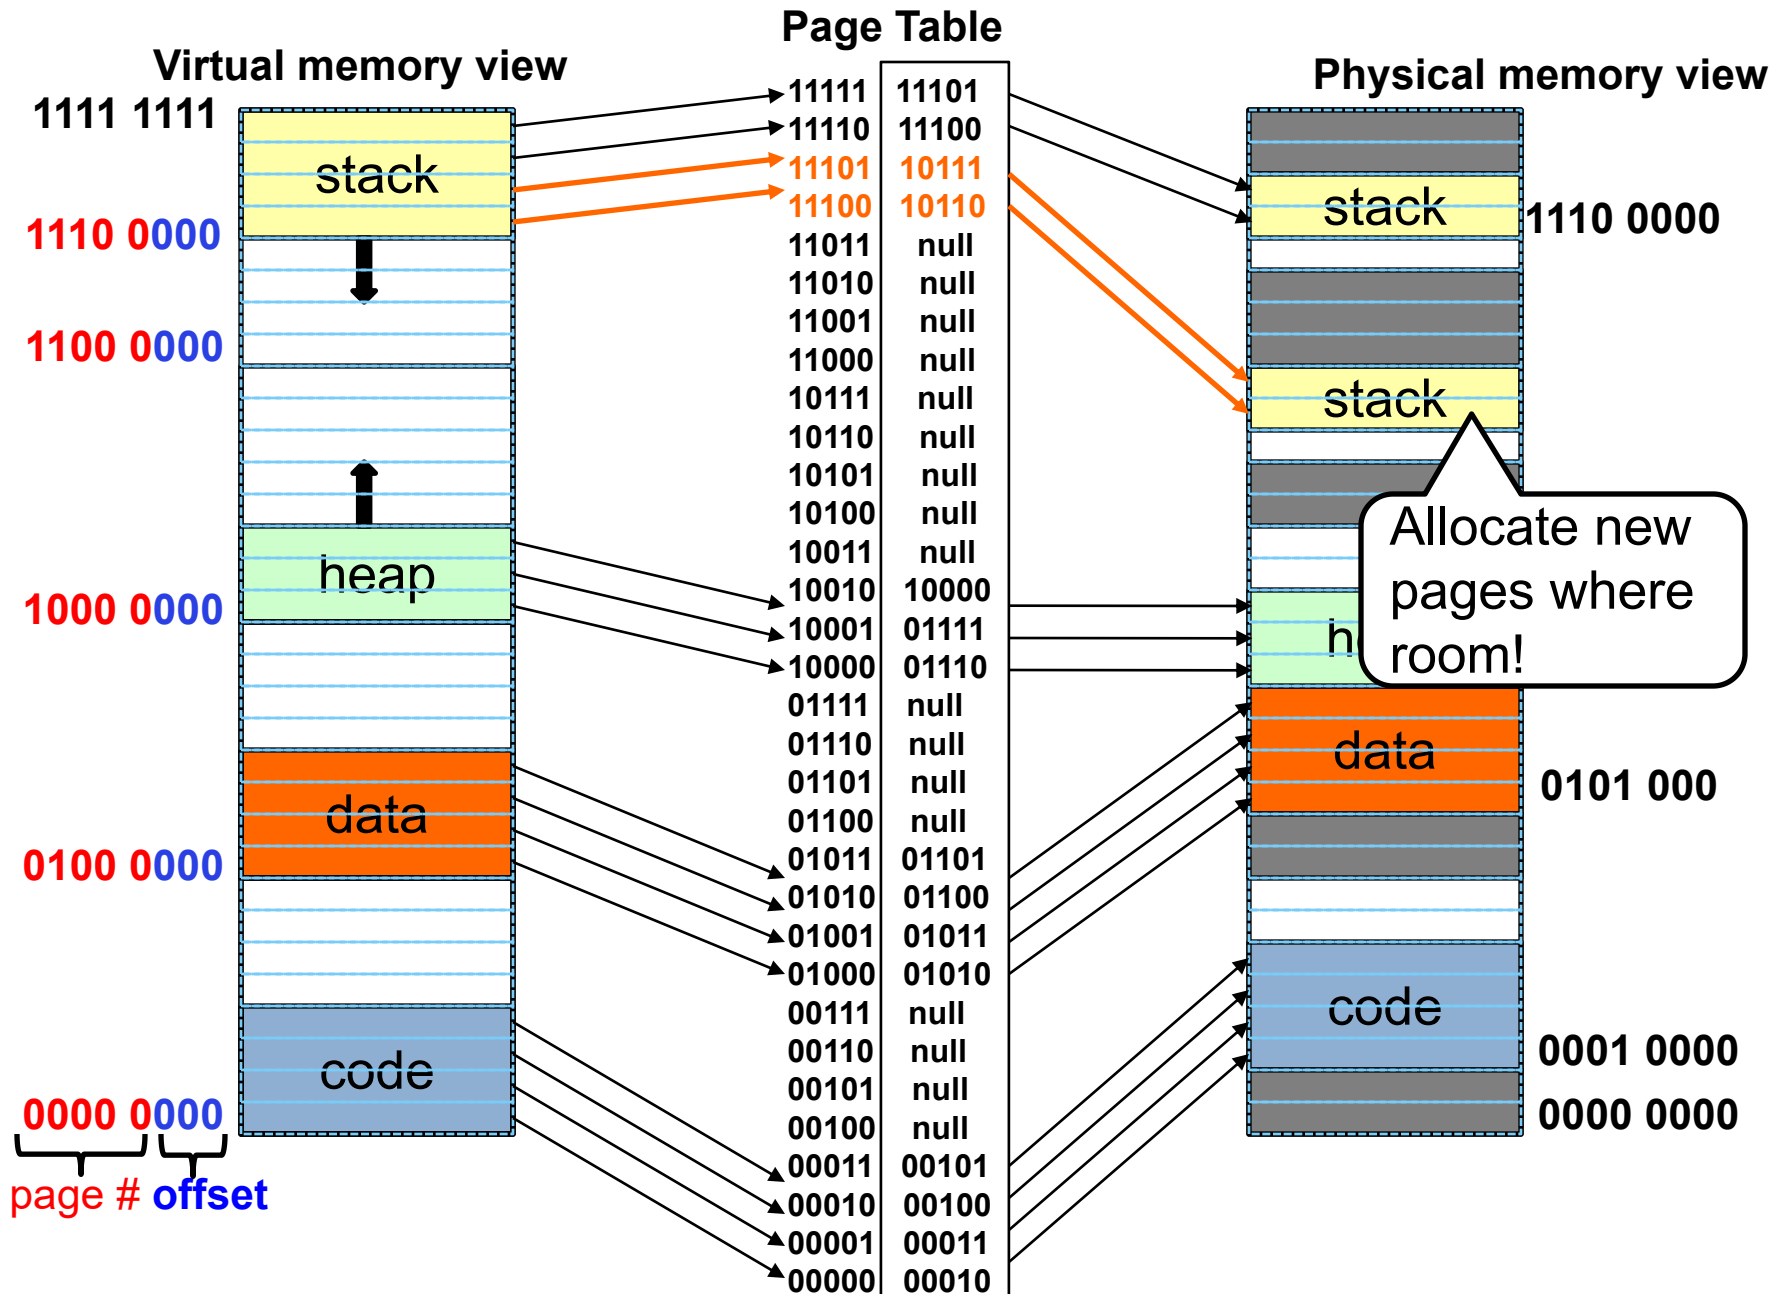
Summary: Paging



Summary: Paging



Summary: Paging

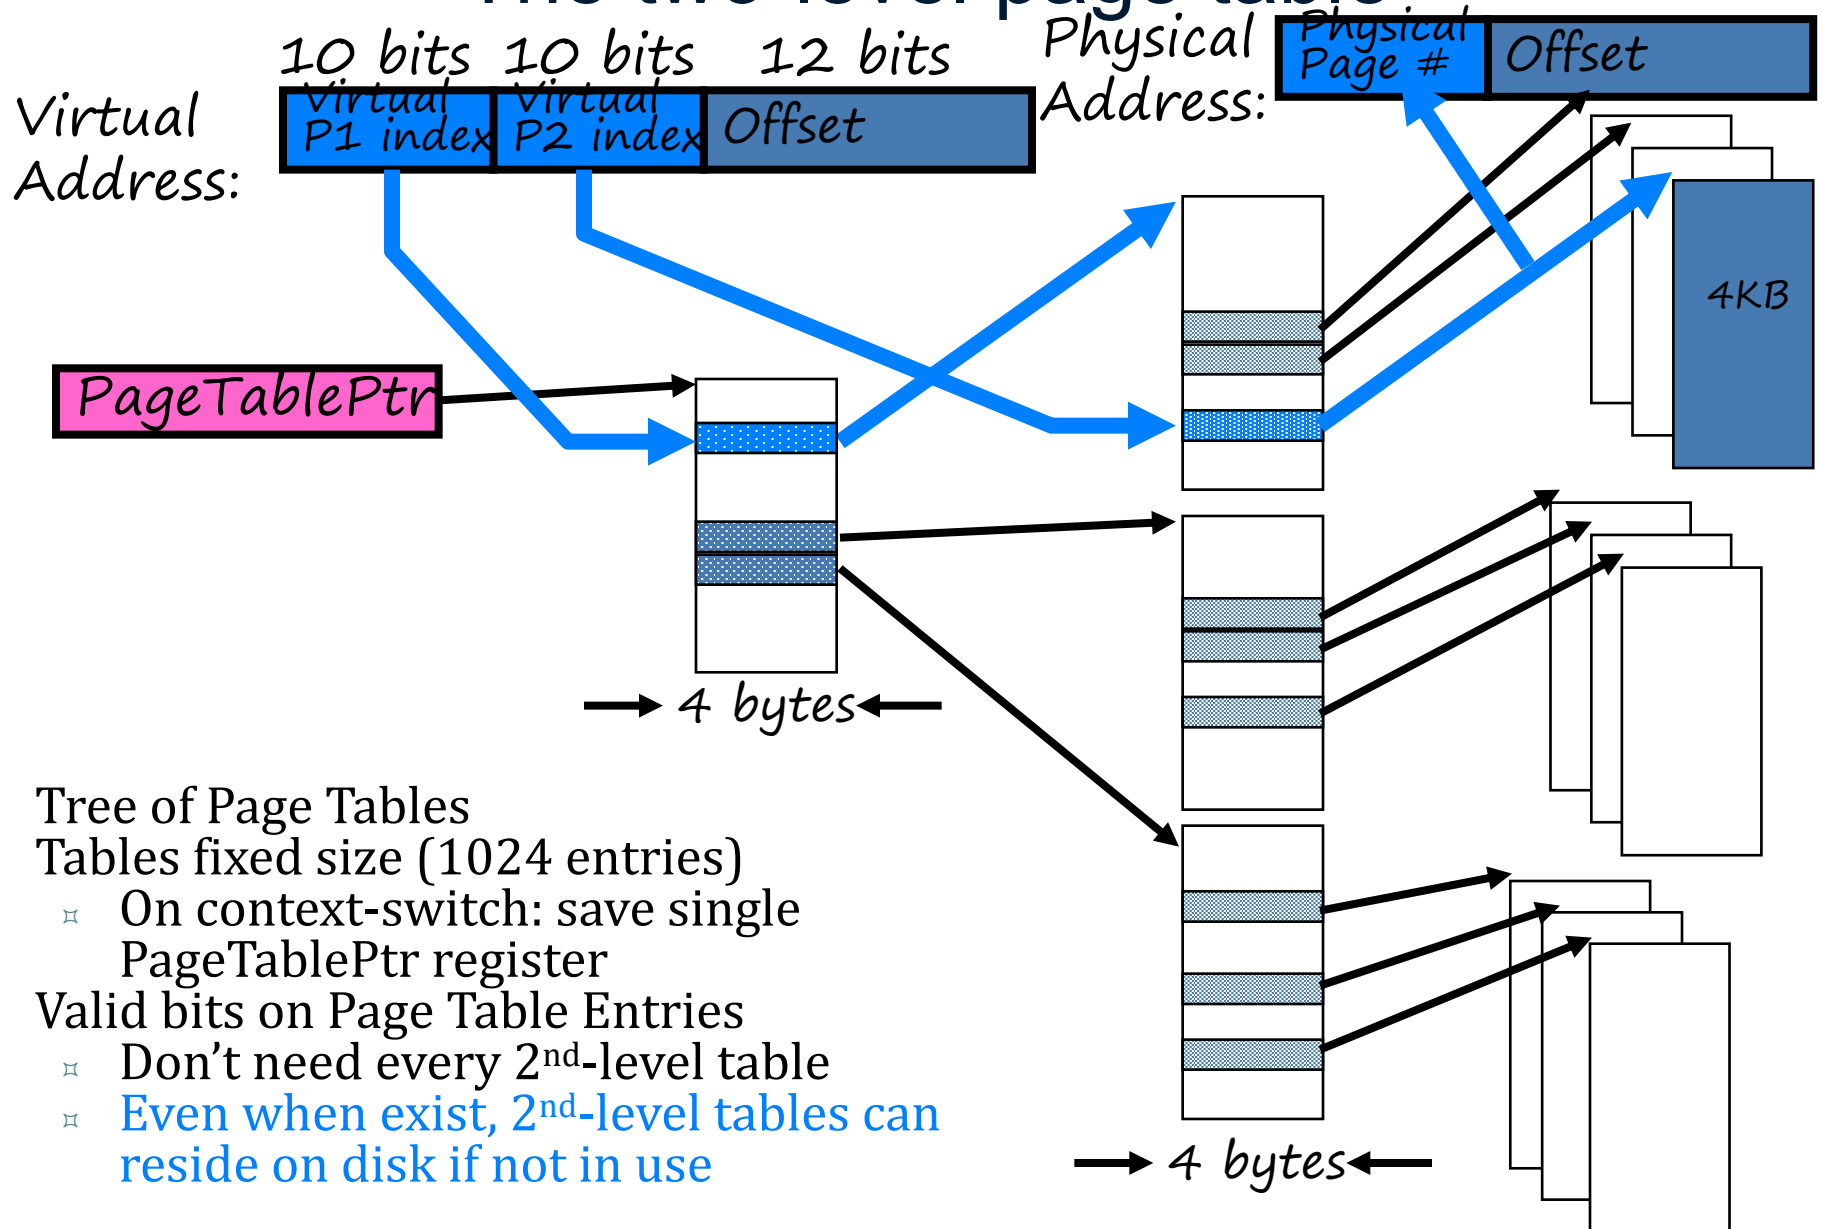


Page Table Discussion

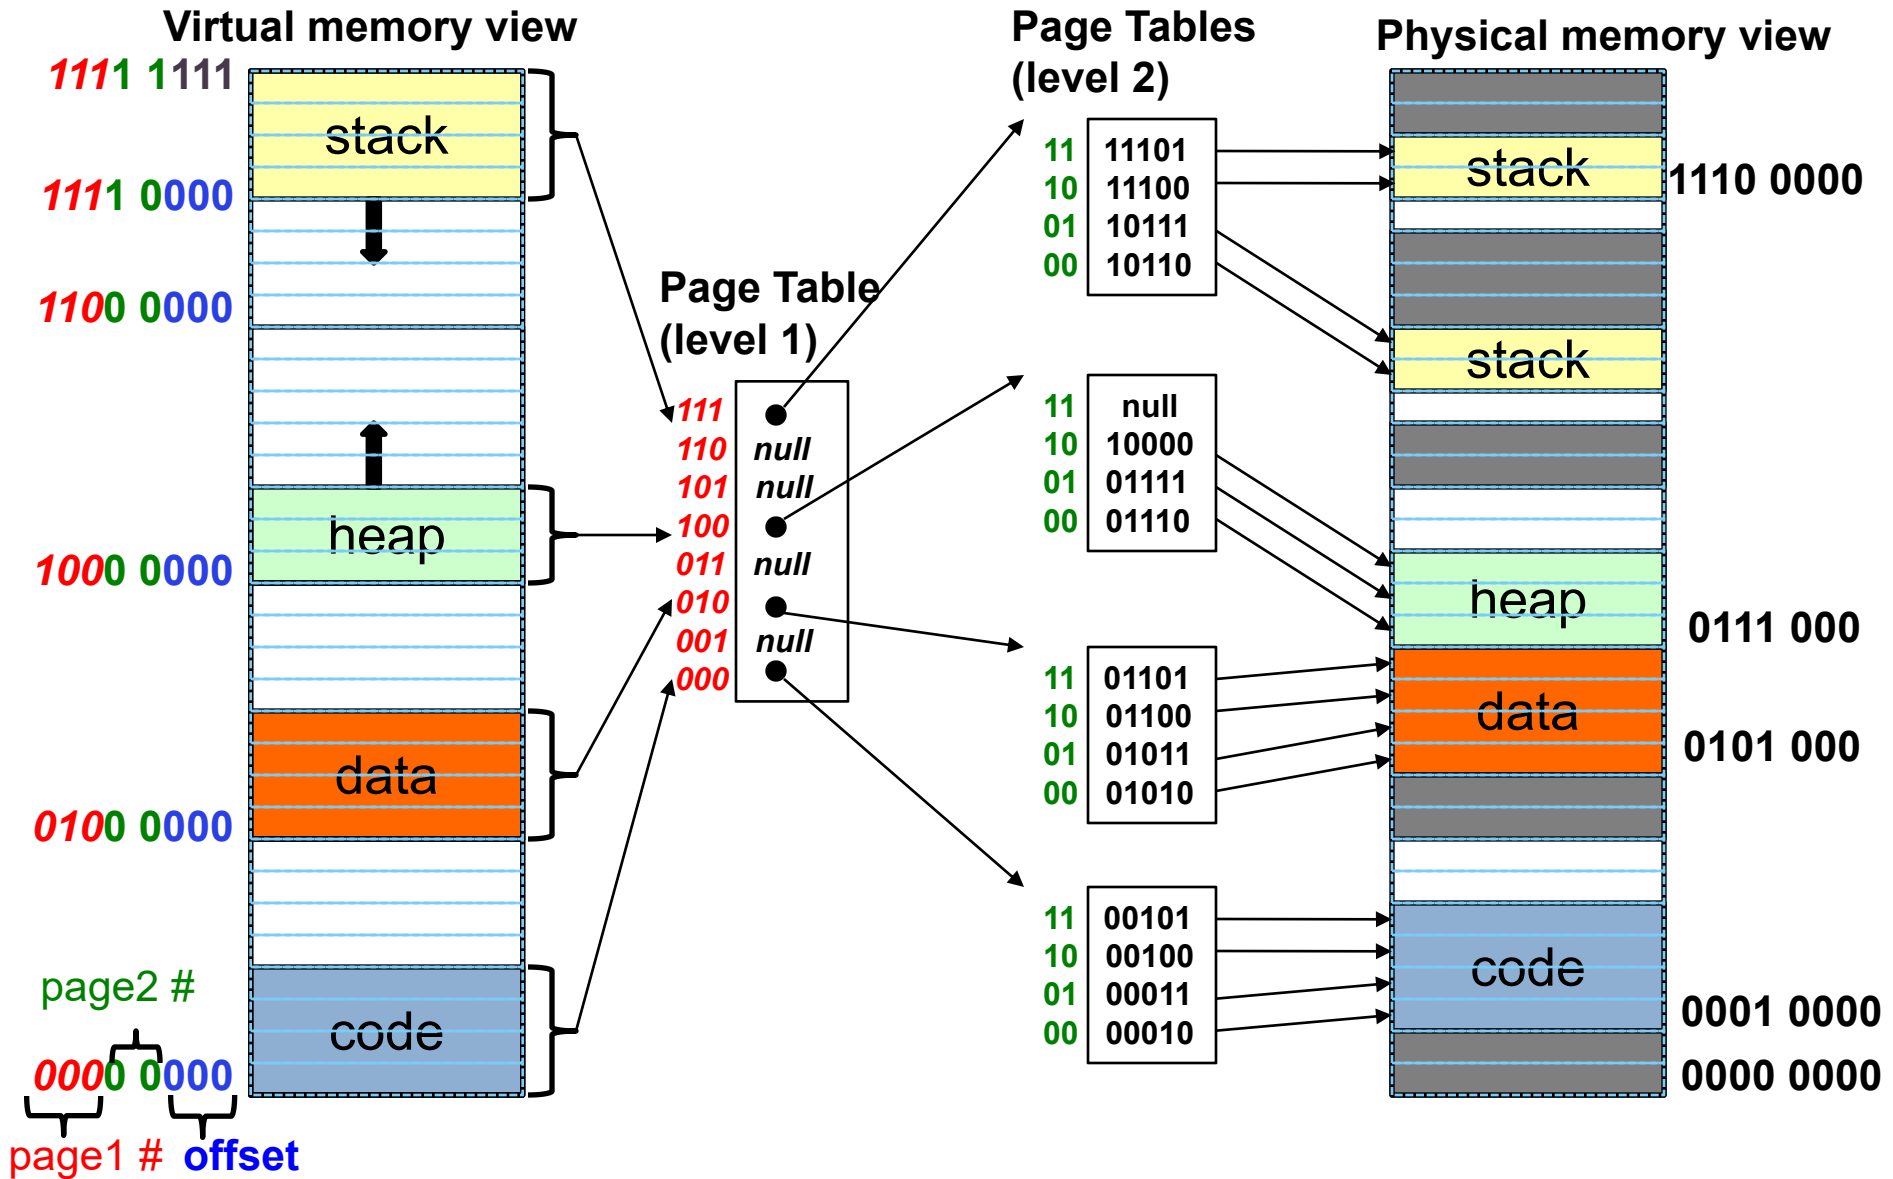
- ✧ What needs to be switched on a context switch?
 - ✧ Page table pointer and limit
- ✧ Analysis
 - ✧ Pros
 - ✱ Simple memory allocation
 - ✱ Easy to share
 - ✧ Con: What if address space is sparse?
 - ✱ E.g., on UNIX, code starts at 0, stack starts at $(2^{31}-1)$
 - ✱ With 1K pages, need 2 million page table entries!
 - ✧ Con: What if table really big?
 - ✱ Not all pages used all the time \Rightarrow would be nice to have working set of page table in memory
- ✧ How about multi-level paging or combining paging and segmentation?

Fix for sparse address space:

The two-level page table




Summary: Two-Level Paging



Summary: Two-Level Paging

Virtual memory view



1001 0000
(0x90)

Page Table
(level 1)

111	●
110	null
101	null
100	●
011	null
010	●
001	null
000	●

Page Tables
(level 2)

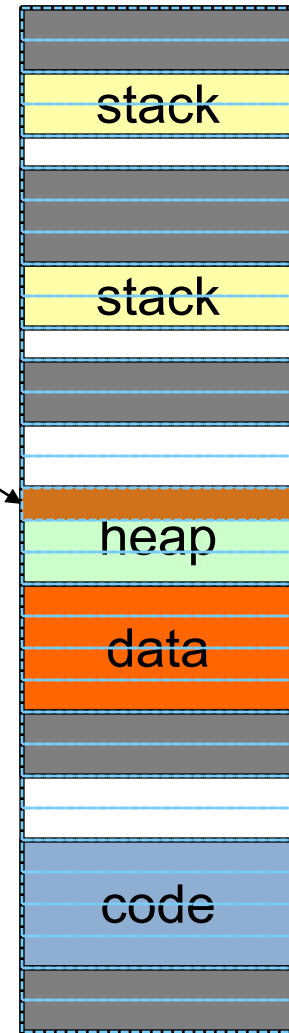
11	11101
10	11100
01	10111
00	10110

11	null
10	10000
01	01111
00	01110

11	01101
10	01100
01	01011
00	01010

11	00101
10	00100
01	00011
00	00010

Physical memory view



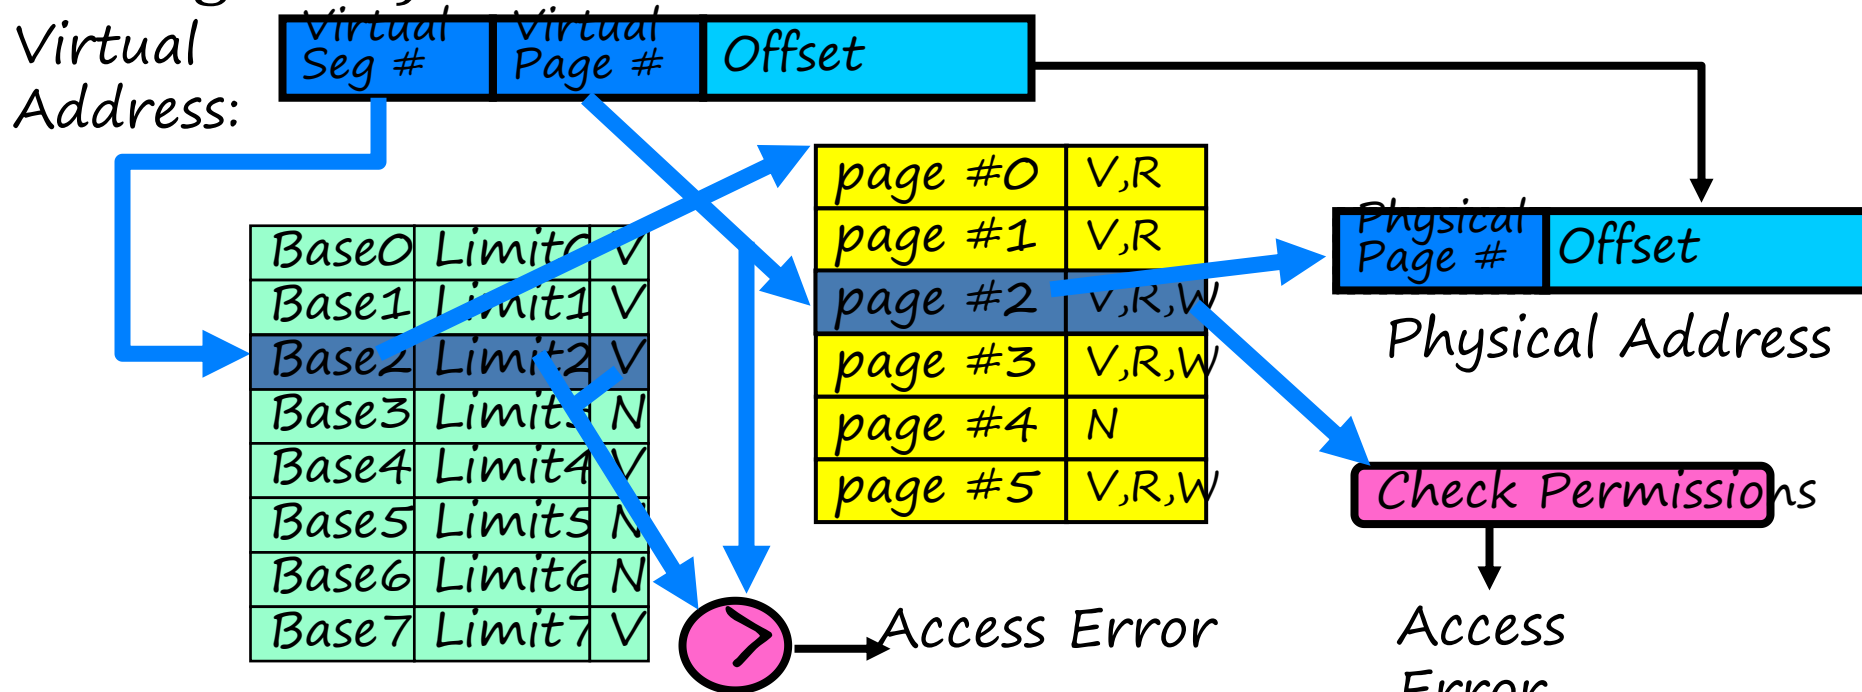
1110 0000

1000 0000
(0x80)

0001 0000
0000 0000

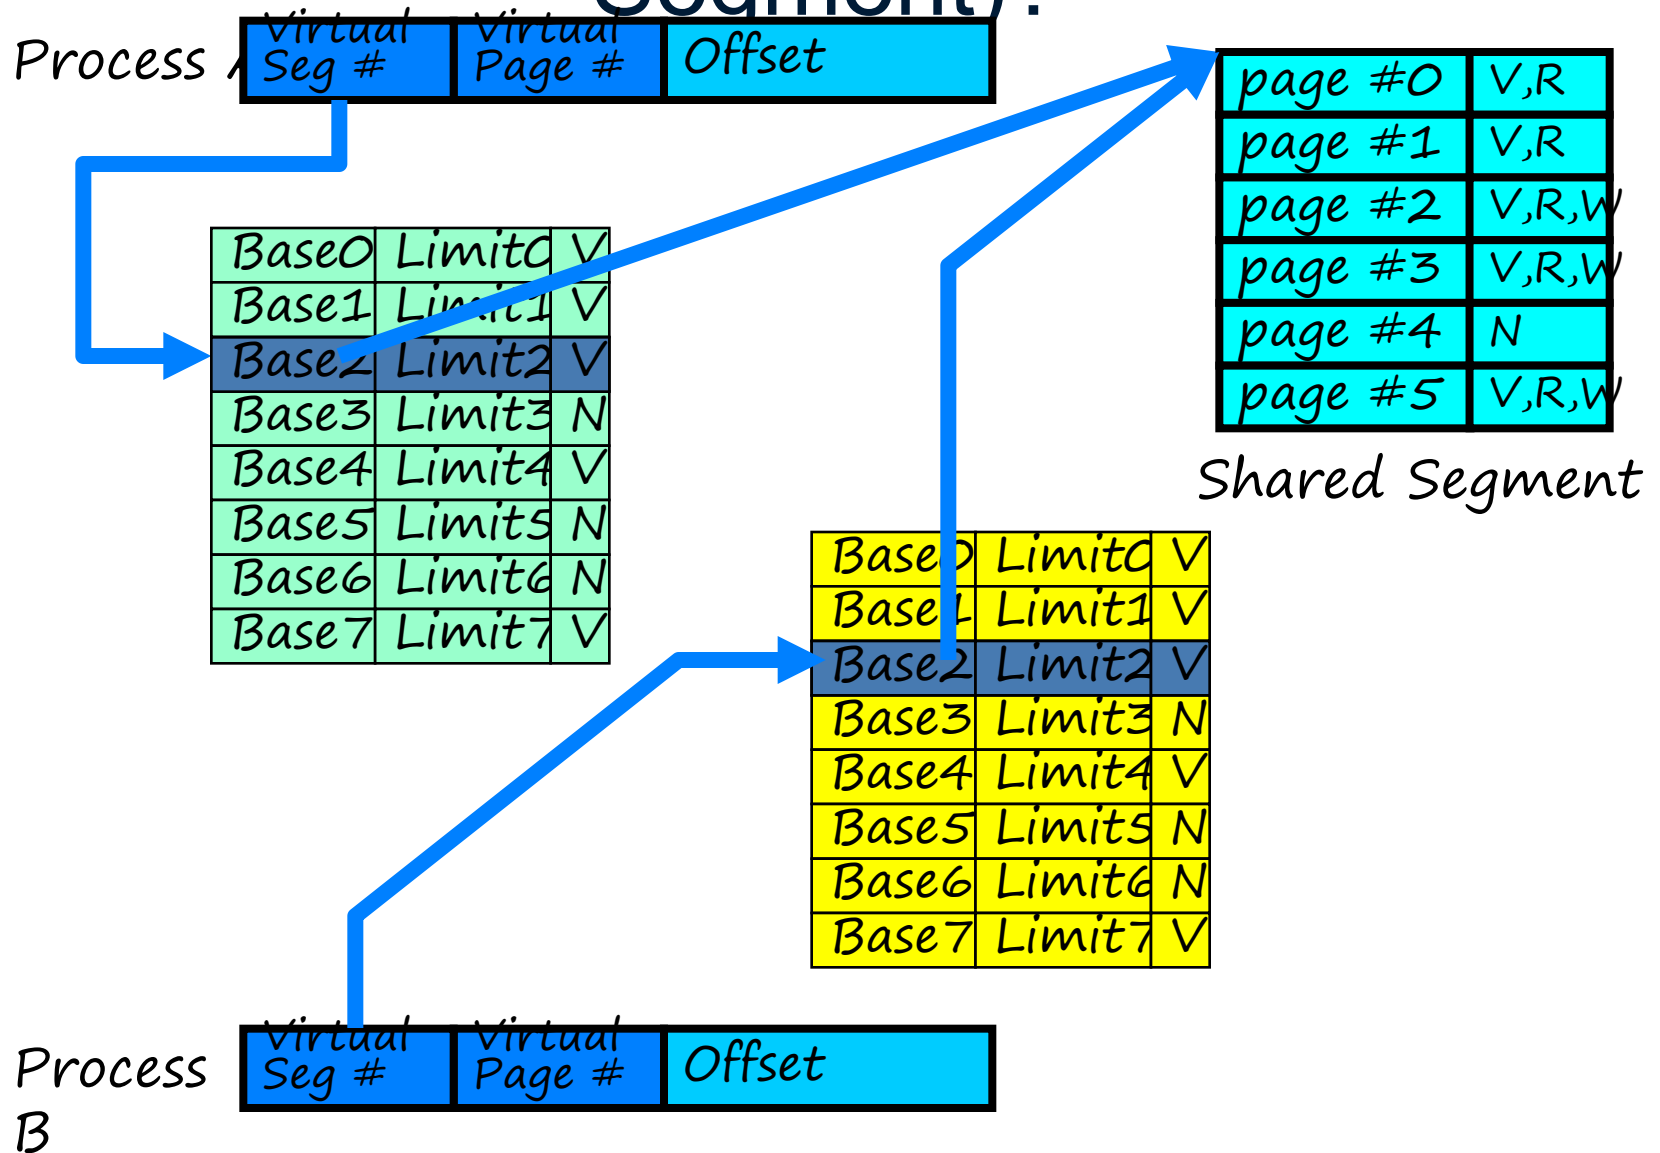
Multi-level Translation. Segments + Pages

- What about a tree of tables?
 - Lowest level page table \Rightarrow memory still allocated with bitmap
 - Higher levels often segmented
- Could have any number of levels. Example (top segment):



- What must be saved/restored on context switch?
 - Contents of top-level segment registers (for this example)
 - Pointer to top-level table (page table)

What about Sharing (Complete Segment)?



Thank You!